

# Agent Architectures



# Outline

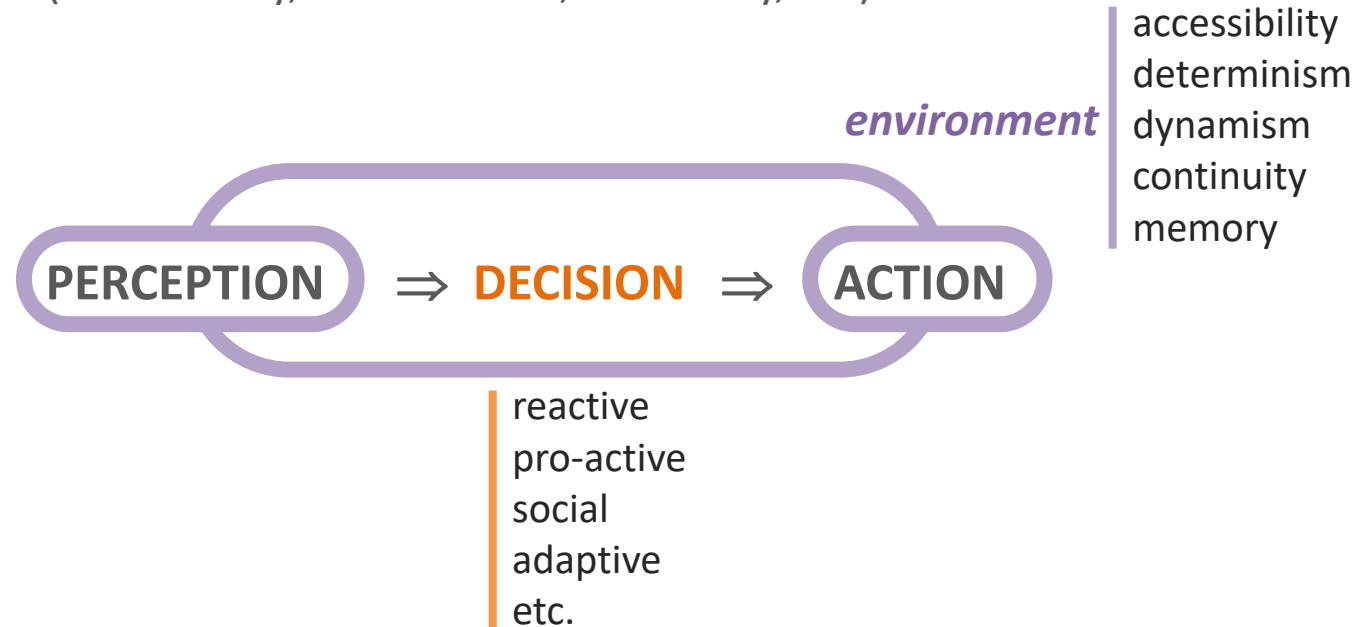
- **Introduction to agent architectures**
- Abstract architectures for agents
- Deductive reasoning agents
- Agents as intentional systems
- Reactive agents
- Hybrid architectures



# Recalling concepts

## AGENT & ENVIRONMENT

- Sense – Decide - Act
- **Agent properties** (autonomous, reactive, pro-active, social, adaptive, etc.)
- **Social Ability** (cooperation, coordination, negotiation)
- **Environment properties** (accessibility, determinism, continuity, etc)



# The architectural stance

## Modular thinking in Computer Science

- System components

- perception
- decision making
  - planning
  - learning
- action


⇒ *an architecture, please!*

- Component interactions

- between internal components
- between agent and environment
- between agents


# Agent architectures

- **What is an agent architecture?**
  - Principles describing agent behavior with
    - formal/abstract view of agents
    - map of the internals (control-flow)
- **What is the goal of an architecture?**
  - Guide agent design and engineering



# Agent architectures

- **What is an abstract architecture?**
  - Common principles describing agent behavior independently of their specificities
  - A template for the construction of agents




# Outline

- Introduction to agent architectures
- **Abstract architectures for agents**
- Deductive reasoning agents
- Agents as intentional systems
- Reactive agents
- Hybrid architectures





# Abstract architectures for agents



Let us make formal the abstract view of agents:

- **Environment states:** (*finite*) set of (*discrete*) states

$$E = \{e_0, e_1, \dots\}$$

- Agents have a set of **actions** (which transforms the environment's state) :


$$Ac = \{\alpha_0, \alpha_1, \dots\}$$

- **Run:** finite sequence of interleaved *states* and *actions*

$$r: e_0 \xrightarrow{\alpha_0} e_1 \xrightarrow{\alpha_1} e_2 \xrightarrow{\alpha_2} \dots \xrightarrow{\alpha_{n-1}} e_n$$



# Abstract architectures for agents



Let:

- $R$  be the set of **all such possible runs** (over  $E$  and  $Ac$ )  
where  $r, r'$  are members of  $R$
- $R^{Ac}$  is the subset of these that *end with an action*
- $R^E$  is the subset of these that *end with an environment state*

# Abstract architectures for agents

- A **state transformer** (environment changes):

$$\tau: R^{Ac} \rightarrow \wp(E)$$

Maps a run (ending in an action) to a set of possible environment states

Important points about this definition:

- History dependent
- Non-determinism

if  $\tau(r)=\emptyset$ , there are no possible successor states to  $r$  (system has *ended* its run)

# Abstract architectures for agents

- An **environment** can now be fully defined as a triple

$Env = \langle E, e_0, \tau \rangle$  where:

- $E$  is a set of environment states
- $e_0 \in E$  is the initial state
- $\tau$  is a state transformer function



# Abstract architectures for agents

- **Agent** is a function which maps runs to actions:

$$Ag: R^E \rightarrow Ac$$

An **agent makes a decision** (i.e., action to perform) based on the **history** of system that it has witnessed to date (i.e.,  $R^E$ )

- Let  **$AG$**  be the **set of all agents**,  $Ag \in AG$



# Abstract architectures for agents

- A **system** is a **pair with**:
  - **one (or more) agent(s)** and
  - **an environment**
- Any **system** has a **set of possible runs**

# Abstract architectures for agents

- We denote the **set of runs of agent  $Ag$  in Environment  $Env$**  by

$$R(Ag, Env)$$

- We assume  $R(Ag, Env)$  contains only **runs that have ended**

# Outline

- Introduction to agent architectures
- Abstract architectures for agents
- **Deductive reasoning agents**
- Agents as intentional systems
- Reactive agents
- Hybrid architectures





# The oldest agent architecture

- **1956–1985:** pretty much all agents designed within AI were ***symbolic reasoning*** agents (mostly *deductive reasoning*)
  - agents with *explicit logical reasoning* to decide what to do
- **1985–present:** problems with symbolic reasoning led to the so-called ***reactive agents*** movement (1985–present)
- **1990–present:** diverse alternatives, including ***hybrid*** architectures, combining the best of *deliberative* reasoning and *reactiveness*

# Deductive reasoning agents


- Classical approach for creating deductive agents:
  - **Agent as a knowledge-based system**
- This paradigm is known as **symbolic AI**

# Deductive reasoning agents

Two key problems to be solved:

## 1. *Transduction problem:*

- translating real-world environment into an accurate, adequate symbolic description.



*Fields:* computer vision, speech understanding, learning...

# Deductive reasoning agents

Two key problems to be solved:

2. *Representation/reasoning problem:*

- symbolically representing information
- how to get agents to reason with this information



*Fields: knowledge representation, automated reasoning, planning*

# Deductive reasoning agents

- *Deductive reasoning agent* (architecture) is one that:
  - contains an explicit **symbolic representation of the world**
  - **internal state** given by formulae (**predicate logic**)
- Example with first-order predicate logic:

```
isopen(valve221)  
temperature(reactor4726, 321)  
pressure(tank776, 28)
```


# Deductive reasoning agents

- ***Deductive reasoning agent*** (architecture):
  - **analogous to beliefs in humans:**
    - internal state may include incorrect/outdated info
  - **makes decisions via *symbolic reasoning***
    - proving theorems without breaking axioms on what is possible

# Agents with state

Agent decision:

- $D$  (internal state = set of formulae or database)
- $see: S \rightarrow Per$  (observe the environment)
- $next: D \times Per \rightarrow D$  (update internal state)
- $action: D \rightarrow Ac$  (decision making with deduction rules)





# Deductive reasoning agents

How can an **agent decide what to do** using theorem proving?

- Use logic to encode a theory stating the *best* action to perform in a given situation

# Deductive reasoning agents

- Let:
  - $\rho$  be this theory (typically **deduction rules**)
  - $DB$  be the logical data (database) describing **current state** of the world
  - $Ac$  be the **set of actions** the agent can perform
  - $DB \vdash_{\rho} \phi$  means that **formula  $\phi$  can be proved from database  $DB$  using deduction rules  $\rho$**

# Deductive reasoning agents




- Agent's action selection function (i.e.,  $action: D \rightarrow Ac$ ) is defined in terms of its deduction rules

```
function action(DB:D) returns an action Ac
begin
  /* for each action, attempts to prove Do(a) from its database using deduction rules */
  for each  $a \in Ac$  do
    if  $DB \vdash_{\rho} Do(a)$  then return  $a$ 
  end for
  /* attempts to find an action such that  $\neg Do(a)$  cannot be derived (i.e., not explicitly forbidden) */
  for each  $a \in Ac$  do
    if  $DB \vdash_{\rho} \neg Do(a) = false$  then return  $a$ 
  end for
  return null /* no action found */ /
end function
```


# Vacuum world



# Vacuum world

(0,2) 	(1,2) 	(2,2)
(0,1)	(1,1)	(2,1)
(0,0) 	(1,0)	(2,0)

Agent starts here (facing north)



# Vacuum world

- **Environment state**

- $S = \{(0,0,d_{0,0}), (0,1,d_{0,1}), (0,2,d_{0,2}), (1,0,d_{1,0}), (1,1,d_{1,1}), (1,2,d_{1,2}), (2,0,d_{2,0}), (2,1,d_{2,1}), (2,2,d_{2,2})\}$

- Agent can receive a **percept** *dirt* or *null*

- i.e., either **dirt under the agent** or not
- $Per = \{dirt, null\}$

- **Actions:** *forward*, *suck*, or *turn* (right 90°)

- $Ac = \{forward, suck, turn\}$



# Vacuum world

- Internal state  $DB$ : three domain predicates
  - $In(x, y)$  – agent is at  $(x, y)$
  - $Dirt(x, y)$  – there is dirt under the agent at  $(x, y)$
  - $Facing(d)$  – the agent is facing direction  $d$



# Vacuum world

**We need deduction rules for agent's behavior!**



# Vacuum world

- Deduction rules (agent's behavior):

- $In(x, y) \wedge Dirt(x, y) \rightarrow Do(suck)$

- $In(0,0) \wedge Facing(north) \wedge \neg Dirt(0,0) \rightarrow Do(forward)$

- $In(0,0) \wedge \neg Facing(north) \wedge \neg Dirt(0,0) \rightarrow Do(turn)$

- $In(0,1) \wedge Facing(north) \wedge \neg Dirt(0,1) \rightarrow Do(forward)$

- $In(0,1) \wedge \neg Facing(north) \wedge \neg Dirt(0,1) \rightarrow Do(turn)$

- ...

# Final remarks: deductive reasoning agents

- **Agent's decision making strategy** – encoded as logical theory
- **Agent's action** – reduces to a problem of proof
- Logic-based approach are **elegant** and have **(clean) logical semantics**

# Final remarks: deductive reasoning agents

- **Disadvantages:**
  - Inherent **computational complexity** of theorem proving
  - Cannot **operate effectively in time-constrained environments**
  - The **environment cannot change** while the agent is making a decision
  - Not easy to **represent and reason about complex and dynamic environments**

# Outline

- Introduction to agent architectures
- Abstract architectures for agents
- Deductive reasoning agents
- **Agents as intentional systems**
- Reactive agents
- Hybrid architectures



# Agents as Intentional Systems

## *Intentional stance*

Develop **agent behaviors** in terms of ***mental states*** (beliefs, desires, wishes, hopes, ...)



# Agents as Intentional Systems

## Examples:

- “Michael took his umbrella because he *believed* it was going to rain.”
- “John worked hard because he *wanted* to obtain a PhD.”

...such attitudes are called *intentional notions*



# Agents as Intentional Systems

This approach can be useful to model:

- **complex systems**
- systems whose **structure is incompletely known**



Is it legitimate or useful to attribute beliefs, desires, and so on, to artificial agents?

# How to implement agents using the Intentional Stance?



**Intentional Systems** are the base for deliberative agents,  
which follow the *intentional stance* through practical reasoning

# How to implement agents using the Intentional Stance?

- **Intentional systems** are the base for **deliberative agents**
  - This **agent architecture** has its origins in the **philosophical work** of Bratman:
    - Michael E. Bratman. *Intention, Plans and Practical Reason*. Harvard University Press, 1987.

# How to implement agents using the Intentional Stance?

“Practical reasoning is a matter of ***weighing conflicting considerations*** for and against **competing options**, where the relevant considerations are provided by what the **agent *desires/values/cares*** about and what the **agent *believes***.” [Bratman, 1990]

# Practical reasoning

Practical Reasoning = Deliberation + Means-Ends Reasoning

- ***Deliberation***: deciding what state of affairs an agent wants to achieve from (possibly conflicting) desires
- ***Means-Ends Reasoning***: deciding how an agent wants to achieve these states of affairs



# The B.D.I. model

- Bratman's philosophical work was used as an **inspiration for creating an architecture for intelligent agents** based on the mental attitudes of:
  - *beliefs*
  - *desires*
  - *intentions*



# The B.D.I. model

- **Beliefs**

- Information about the environment, other agents, and itself

- **Desires**

- Desires/goals are state of affairs to achieve

- **Intentions**

- Commitments to achieving particular goals

# Intentions in practical reasoning

## Intentions are stronger than mere desires:

“My desire to play basketball this afternoon is merely a potential influencer of my conduct.

It must be viewed with my other relevant desires...

Once I intend to play basketball this afternoon, the matter is settled: I normally need not continue to weigh the pros and cons.

When the afternoon arrives, I will normally just proceed with my intentions.” (Bratman, 1990)





# Intentions in practical reasoning

And Intentions drive means-ends reasoning

- They *lead to action* because I *attempt to achieve* them
- I try to decide *how* to achieve them
- If one course of action *fails*, I usually *attempt others*

# Intentions in practical reasoning

*For example:*

- *You might consider a career as an academic or a career in industry (deliberation)*
- *You have to decide the career (deliberation)*
- *Your decision is to be an academic (intention/state of affairs)*
- *You make a plan: apply for PhD program, get a PhD, etc.*
  - *You decide how to achieve the state of affairs (means-end reasoning)*

# Intentions in practical reasoning

Property: **Intentions persist**

- *I do not give up without good reason*

*E.g., If your intention is to become an academic, then you should persist with this intention*

# Intentions in practical reasoning

Property: Intentions constrain future deliberation

- *I will not entertain options that are inconsistent*

*filter of admissibility*

*E.g., If I have an intention to write a book, so I cannot consider the option of partying every night*

# Intentions in practical reasoning

Property: **Intentions influence beliefs**

- *Intentions are closely related to beliefs about the future*

*E.g., If you intend to become an academic, then you should believe that, assuming some background conditions, you will indeed become an academic*

# Deliberation and belief revision

So how do we model *deliberation in agents*?

- **Belief revision function**
  - Update beliefs with sensory input and previous belief
- **Function to generate options**
  - Use beliefs and existing intentions to generate options (=desires)
- **Filtering function**
  - Choose between competing alternatives and commit to their achievement

# Deliberation and belief revision

So how do we model *deliberation in agents*?

- revise the agent's beliefs (belief revision function):

$$brf: 2^{Bel} \times Per \rightarrow 2^{Bel}$$

- produce the agent's desires/options (option generation function):

$$options: 2^{Bel} \times 2^{Int} \rightarrow 2^{Des}$$

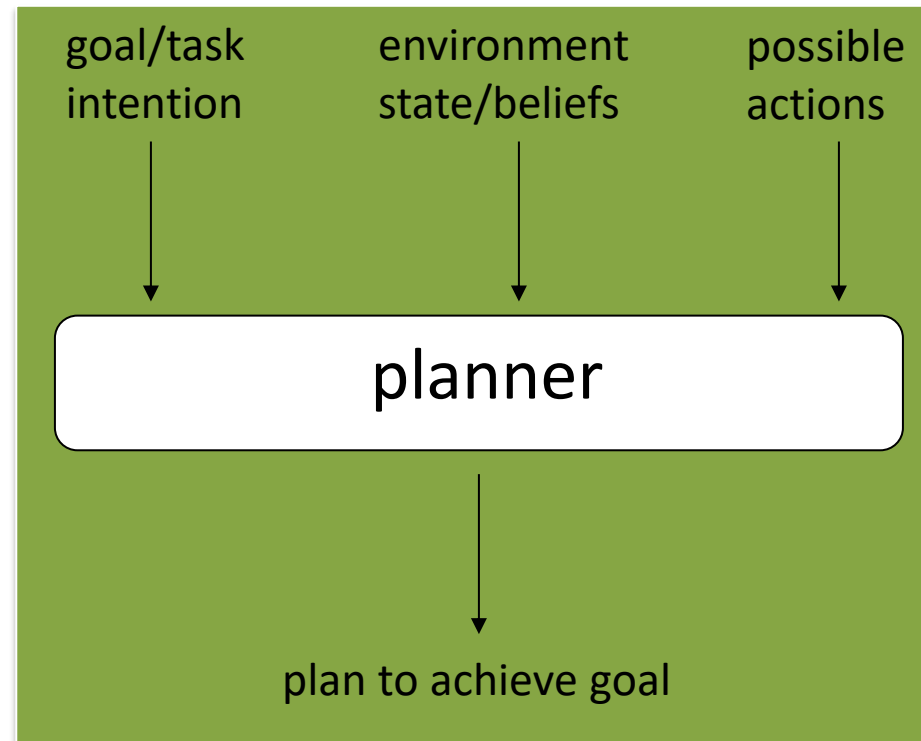
- select the *best* option(s) for the agent to commit to (filter function):

$$filter: 2^{Bel} \times 2^{Des} \times 2^{Int} \rightarrow 2^{Int}$$

# Means-Ends reasoning

An agent's means-ends function

$$plan: 2^{Bel} \times 2^{Int} \times 2^{Ac} \rightarrow Plan$$





# Implementing a practical reasoning agent

Decision-making is a *loop*:

1. *Observe* the world and *update beliefs*
2. *Deliberate* to decide the *intention(s)*
  - determine available *options*
  - *filter*
3. Use *means-ends reasoning* to find a *plan* for the intention(s)
4. *Execute* the plan
5. Return to 1

# Commitments

How committed an agent should be to its intention?

How long should an intention persist?

- A commitment implies *temporal persistence*.
- But to what extent?



*When do I give up  
pursuing an Intention?*

# Commitment strategies

- Blind commitment
- Single-minded commitment
- Open-minded commitment



# Blind commitment

Maintains an intention *until* it believes the intention has been *achieved*



Aka *fanatical* commitment

# Single-minded commitment

Maintains an intention *until* it believes that either:

- the intention has been *achieved*
- is *no longer possible* to achieve



# Open-minded commitment

Maintains an intention *as long as*

- it has not been *achieved*
- it is still *desired*



# Outline

- Introduction to agent architectures
- Abstract architectures for agents
- Deductive reasoning agents
- Agents as intentional systems
- **Reactive agents**
- Hybrid architectures



# Reactive agents

- Many **problems** with **symbolic/logical approaches**, for example:
  - Inherent **computational complexity** of theorem proving
  - Cannot **operate effectively** in **time-constrained environments**



# Reactive agents

- Many **problems** with **symbolic/logical approaches**, for example:
  - The **environment cannot change** while the **agent is making a decision**
  - **Not easy to represent and reason** about complex and dynamic environments

# Reactive agents

- In the mid to late 1980s, researchers started to **investigate alternatives to symbolic AI paradigm**
- These **new approaches** had a few themes in common:
  - **Rejection of symbolic representation** and syntactic manipulation (e.g., logic programming)

# Reactive agents

- These **new approaches** had a few themes in common:
  - The idea that **intelligent behavior** is **linked to the environment**
  - **Intelligent behavior can emerge** from the interaction of various simpler behaviors

# Reactive agents

What are reactive agents?

- **Agents equipped with simple processing units** that perceive and quickly **react to changes** in the environment
- **Do not use complex symbolic reasoning**

# Reactive agents

What are reactive agents?

- In reactive agent systems, **intelligence** is **not a property of a single component**
- The **intelligence is distributed** in the system and **emerges from the interaction** among agent components and the environment

# Purely reactive agents

- **Purely reactive agents** make **no reference to their history**
  - no internal state
  - decision making entirely on the present
- Formally:

$$Ag: E \rightarrow Ac$$

```
function decide(perception)
  current_state <- INTERPRET-INPUT(perception)
  rule <- RULE-MATCH(current_state, rules)
  action <- RULE-ACTION(rule)
  return action
```

(Russell and Norvig, 1994)

# Reactive architectures

Inspiration: **intelligent behavior** of animals in the world

- **simple behaviors** of each **individual agent**
- **complex behaviors** comes from **combining individual behaviors**



# Brooks: subsumption architecture

Dr. Rodney Brooks' short bio:

- 1981: **PhD Stanford**
- 1984 - 2010: Professor at **MIT**
- 1997 – 2007: Director of **MIT Artificial Intelligence Lab**
- 1990 - 2011: Founder, Board Member, and CTO of **iRobot Corp**
- 2008 – 2018: Founder, Board Member, and CTO of **Rethink Robotics**
- 2019 - Present: Founder and CTO of **Robust.AI**

**Dr. Brooks was one of the most vocal and influential critics of the symbolic approach**





# Brooks: subsumption architecture

Decision-making:


- A set of *task-accomplishing behaviors*
- Each **behavior module** can be seen as an **action selection function**

# Brooks: subsumption architecture

- **Behavior module:**
  - **Perceptual inputs** are mapped into **actions**
  - Each **behavior module** is intended to **perform a task**


# Brooks: subsumption architecture

- Behavior modules are finite-state machines



# Subsumption architecture: layered control

- However, many behaviors can *fire simultaneously*
- Subsumption hierarchy:
  - behaviors arranged into layers
  - lower layers inhibit higher layers



# Agent development using Brooks' architecture

## 1. Create a module for a particular task

- should **link perception and action**
- should work by itself

## 2. Add more modules

- the **priorities** for the behaviors need to be **re-adjusted** every time **one module is added**

# Agent development using Brooks' architecture

## Common requirements:

1. Deal with **multiple goals**
2. Deal with **multiple sensors**
  - may provide conflicting data
3. Be **robust** in dealing with changes in the environment
4. Deal with **time constraints**

# Advantages of reactive agents

- Simplicity
- Economy
- Computational tractability



# Advantages of reactive agents

- Robustness against failure
- Elegance
- Extensibility





# Outline

- Introduction to agent architectures
- Abstract architectures for agents
- Deductive reasoning agents
- Agents as intentional systems
- Reactive agents
- **Hybrid architectures**



# Limitations of reactive agents

- Decisions are only based on local information: agents have a ***short-term view***
- Agents ***need sufficient local information*** to make decisions
- **No learning**: how to guarantee reactive rules evolve?
- Not trivial to ***engineer*** agents with complex behavior
- **Not easy to predict complex behavior** when agents have a *high number of layers*

# Criticisms to deliberative agents

- **Deliberation** and **Planning** with **incomplete info** can be a problem
- **Speed of decisions can be slow** to deal with the real world
- Many architectures rely on a **symbolic approach**
- The need to be **grounded to the real world**



# Hybrid agents

Many researchers argue that **neither a completely deliberative nor completely reactive approach is suitable**

⇒ *hybrid* systems to marry classical and alternative approaches



# Hybrid architectures

- Requirement: **agent must have both reactive and deliberative behaviors**
- Often, the **reactive subsystem is given some kind of precedence over the deliberative subsystem**
- This kind of structuring leads naturally to the **idea of a *layered* architecture**



# Hybrid architectures


A key problem in hybrid architectures: **what kind of control flow should we consider between the agent's subsystems?**

- *Horizontal layering*

**Layers are directly connected to the sensory input and action output.** Each layer itself acts like an agent, producing suggestions on what action to perform


- *Vertical layering*

**Sensory input and action output are each dealt with by at most one layer each**



# Hybrid architectures: horizontal layering

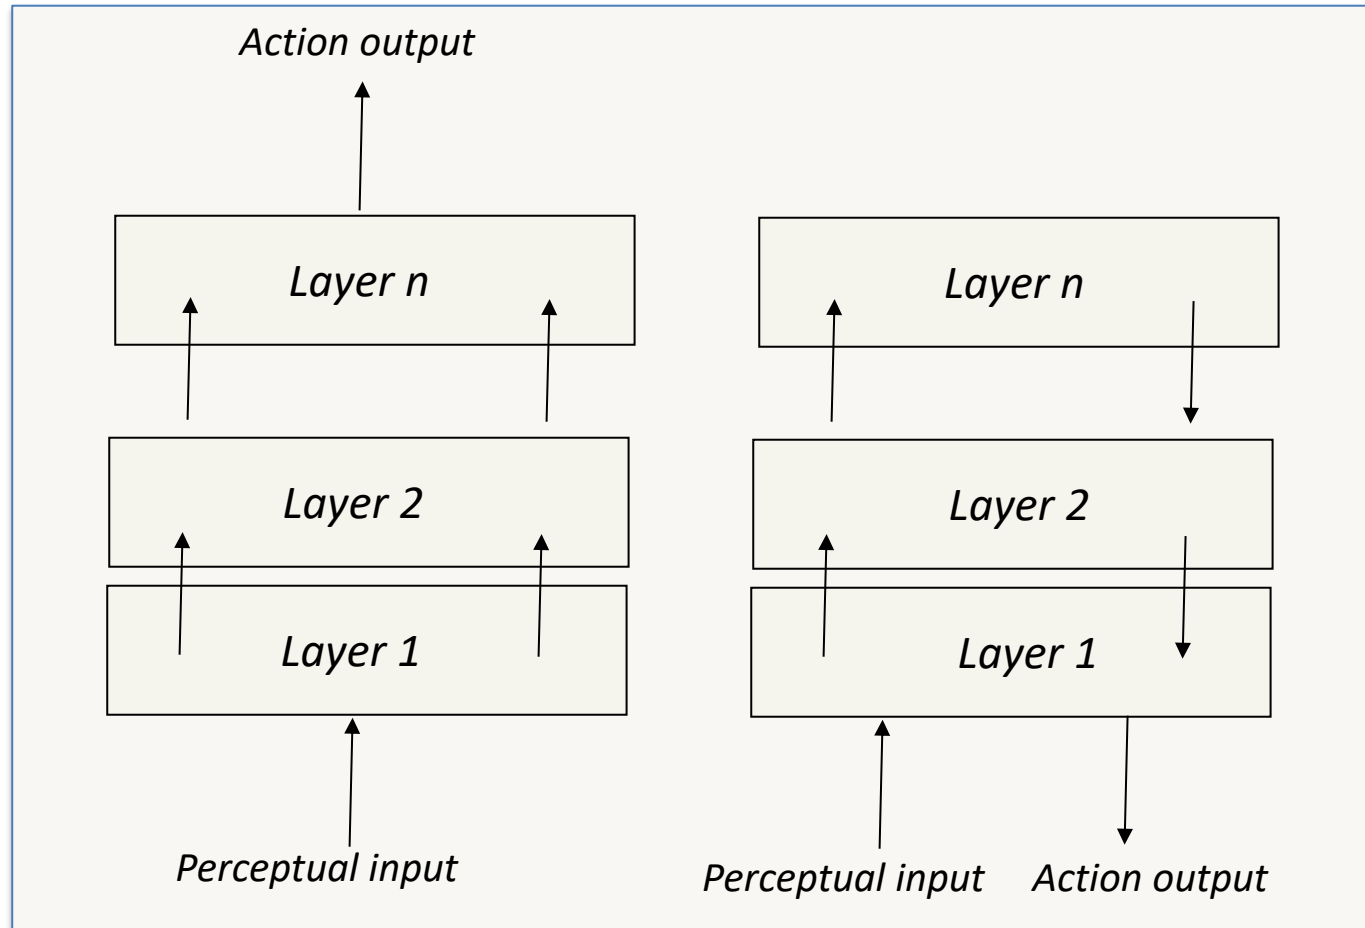
- **Advantages:** simple, distributed, fault-tolerant
- **Disadvantages:**
  - global behavior may not be coherent
  - difficult to avoid conflict between layers



# Hybrid architectures: vertical layering

**Vertical** (1 pass control)


**Vertical** (2 passes control)






# Comparing hybrid architectures


$m$  possible actions suggested by each layer,  $n$  layers



(a) Horizontal layering



(b) Vertical layering  
(One pass control)



(c) Vertical layering  
(Two pass control)

**complex decision making  
regarding the action output**

**not tolerant to  
layer failure**

# DARPA grand challenge<sup>1</sup>

Main goal of the challenge:

- develop an autonomous car (i.e., self-driving car):
  - capable of **traversing unrehearsed off-road terrain**
  - travel a **175 mile** long course through the **Mojave desert**
  - **take no more than 10 hours**

<sup>1</sup>[https://en.wikipedia.org/wiki/DARPA\\_Grand\\_Challenge](https://en.wikipedia.org/wiki/DARPA_Grand_Challenge)

# DARPA grand challenge

- Teams
  - 2004 (1M\$ prize): 107 teams (15 finalists, 0 finished)
  - 2005 (2M\$ prize): 195 teams (23 finalists, 5 finished)
- **The route is kept secret** until 2h before the race
  - At this time they received a route description in RDDF format

# Stanley – 2005 Winner

Volkswagen R5 (4WD)

- treats autonomous navigation as a software problem



# Stanley: sensory equipment

**Perception:** roof rack that houses

- 5 SICK laser range finders
- camera for long range perception
- 2 RADAR sensors
- antennae: 1 for GPS, 1 for GPS compass, and 1 radio antennae



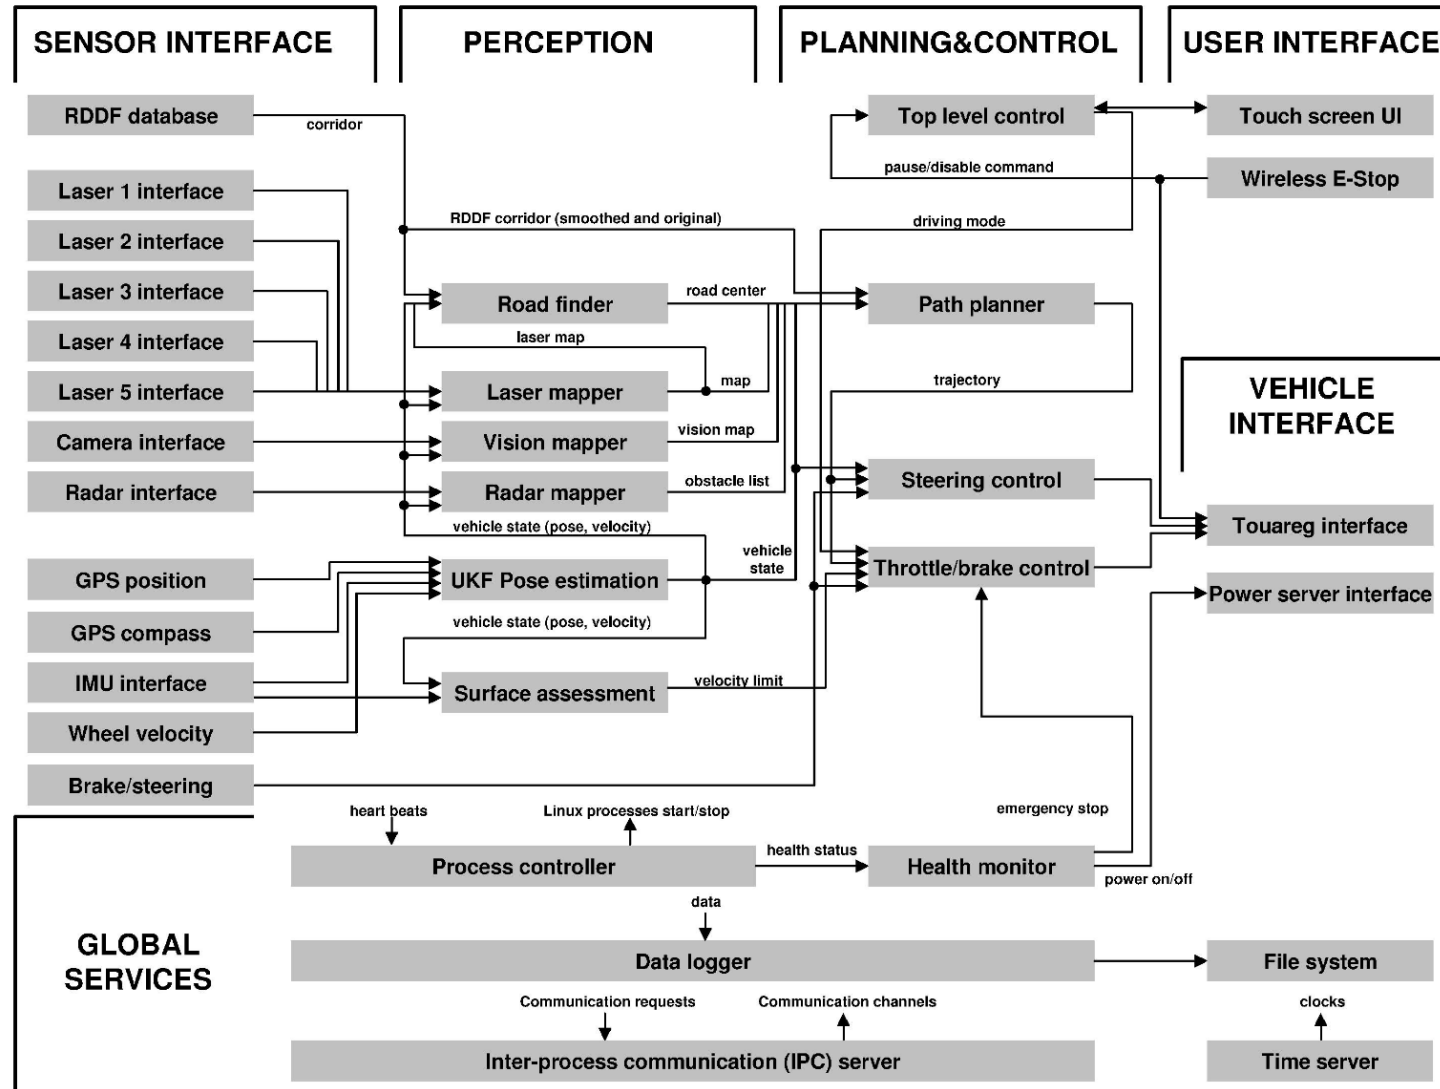
# Stanley control

3 main actuators

- brakes
- throttle
- steering



# Stanley architecture



# Stanley in action



<http://www.youtube.com/watch?v=M2AcMnfzpNg>



Since then...



# Advantages of hybrid architectures

- One of the **most used architectures** currently
- Allows for a **real-time response** combined with **goal oriented-behavior**
- Reactivity can be privileged in relation to deliberation
- **Knowledge** about the world can be subdivided into layers
  - different **levels of abstraction**



# Disadvantages of hybrid architectures

- **Vertical layering:** bottlenecks and fault intolerance
- **Horizontal layering:** complexity of decision making
- **Interactions between layers** are difficult to program and to test
  - one will need to analyze all the possible interactions between these layers
  - **an integration problem**



# Thank You



[rui.prada@tecnico.ulisboa.pt](mailto:rui.prada@tecnico.ulisboa.pt)