

# Markov decision process (MDP)



# Outline

- **Single-agent learning**

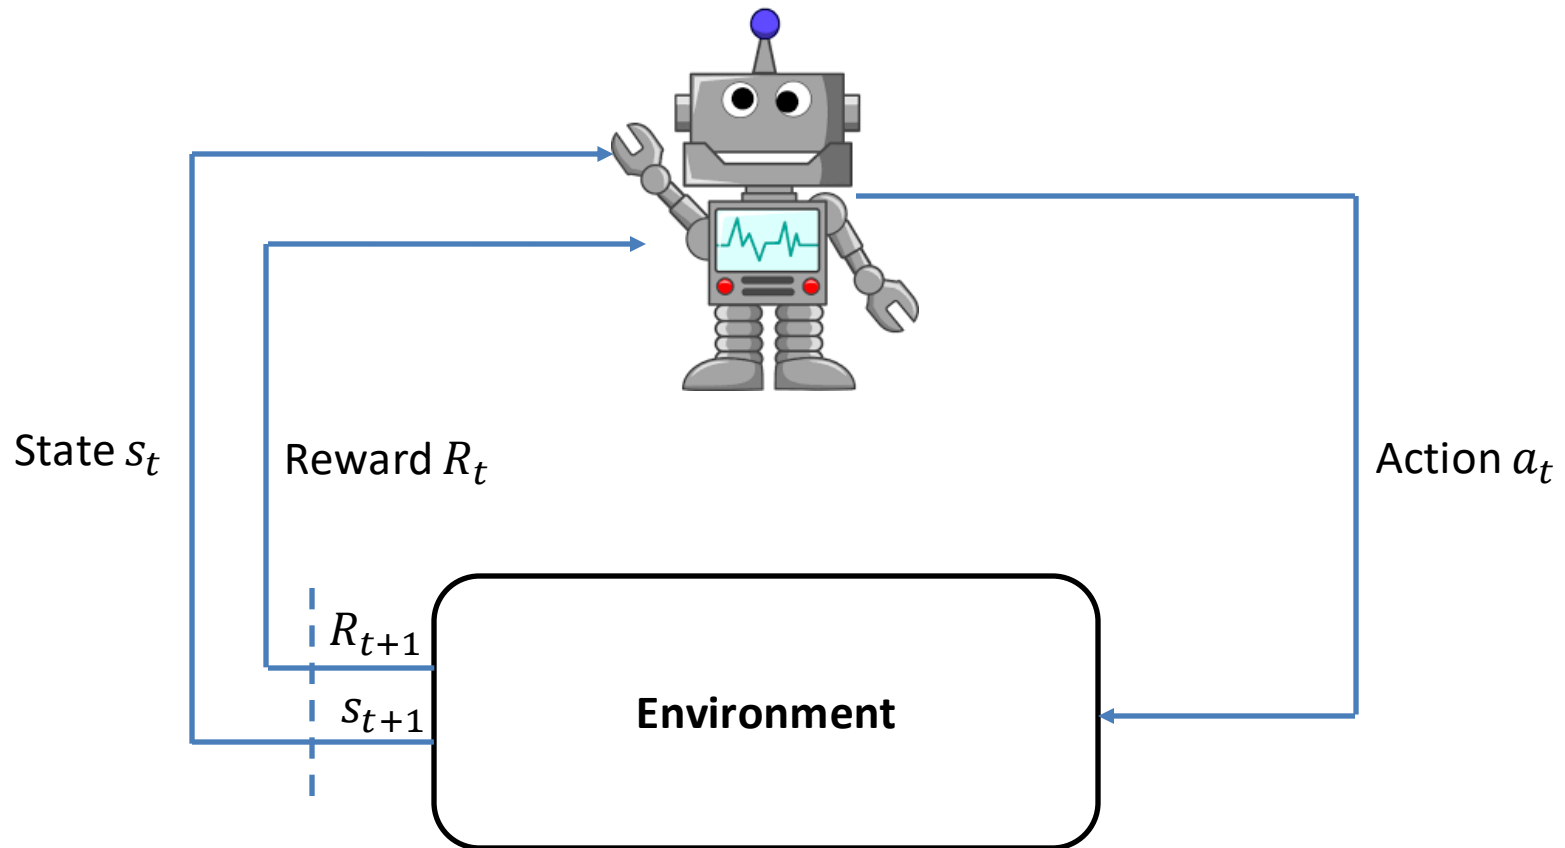


# Markov decision process (MDP)

- A framework for **sequential decision making** of a single agent
- **Markovian transition model** and **fully observable**
  - i.e., we assume it verifies the Markov property
- **Planning horizon** can be infinite



# Markov decision process (MDP)



# Markov decision process (MDP)

- We can formally define an MDP with following elements:
  - **Discrete time**  $t = 0, 1, 2, \dots$
  - **A discrete set of states**  $s \in S$
  - **A discrete set of actions**  $a \in A$
  - **A stochastic transition model**  $P(s'|s, a)$ 
    - the world transitions stochastically to state  $s'$  when the agent takes action  $a$  at state  $s$
  - **A reward function**  $R: S \times A \rightarrow \mathbb{R}$ 
    - An agent receives a reward  $R(s, a)$  when it takes action  $a$  at state  $s$

# Markov decision process (MDP)

- **Definition:** the **state-value function** of a state  $s$  under a policy  $\pi$  is the expected return the agent can receive when starting in state  $s$  and then following policy  $\pi$  :

$$V^{\pi}(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) | s_0 = s, a_t = \pi(s_t) \right]$$

**Definition:** the **action-value function (Q-values)** of taking an action  $a$  in state  $s$  under a policy  $\pi$  is the expected return the agent can receive when starting in state  $s$ , taking action  $a$ , and then following policy  $\pi$  :

$$Q^{\pi}(s, a) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) | s_0 = s, a_0 = a, a_{t>0} = \pi(s_t) \right]$$

# Markov decision process (MDP)

- A **policy  $\pi$**  is defined to be **better than or equal to a policy  $\pi'$**  if the expected return of  $\pi$  is greater or equal to expected return of  $\pi'$ :

$$\pi \geq \pi' \text{ if and only if } V^\pi(s) \geq V^{\pi'}(s), \text{ for all } s \in S$$

- **There is always at least one policy** that is better than or equal to all other policies, which is the **optimal policy**
  - Note that an MDP might have more than one optimal policy
  - We denote all the optimal policies by  $\pi^*$

# Markov decision process (MDP)

- These policies  $\pi^*$  share the same state-value function, called **optimal state-value function**, with the following definition:

$$V^*(s) = \max_{\pi} V^{\pi}(s), \text{ for all } s \in S$$

- These policies  $\pi^*$  also share the same action-value function, called **optimal action-value function**, with the following definition:

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a), \text{ for all } s \in S \text{ and } a \in A$$



# Value Iteration

- We initialize arbitrarily a state-value function (e.g., with zeros, ones, etc.)
- Then we iteratively apply the Bellman equation turned into an assignment operation:

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$V(s) := \max_{a \in A} Q(s, a), \quad \forall s$$

- Repeat the above two equations until  $V$  does not change significantly between two consecutive steps

# Value Iteration

- Value iteration converges to the optimal  $Q^*$  for any initialization
- After computing the optimal  $Q^*$ , we can extract the policy as follows:

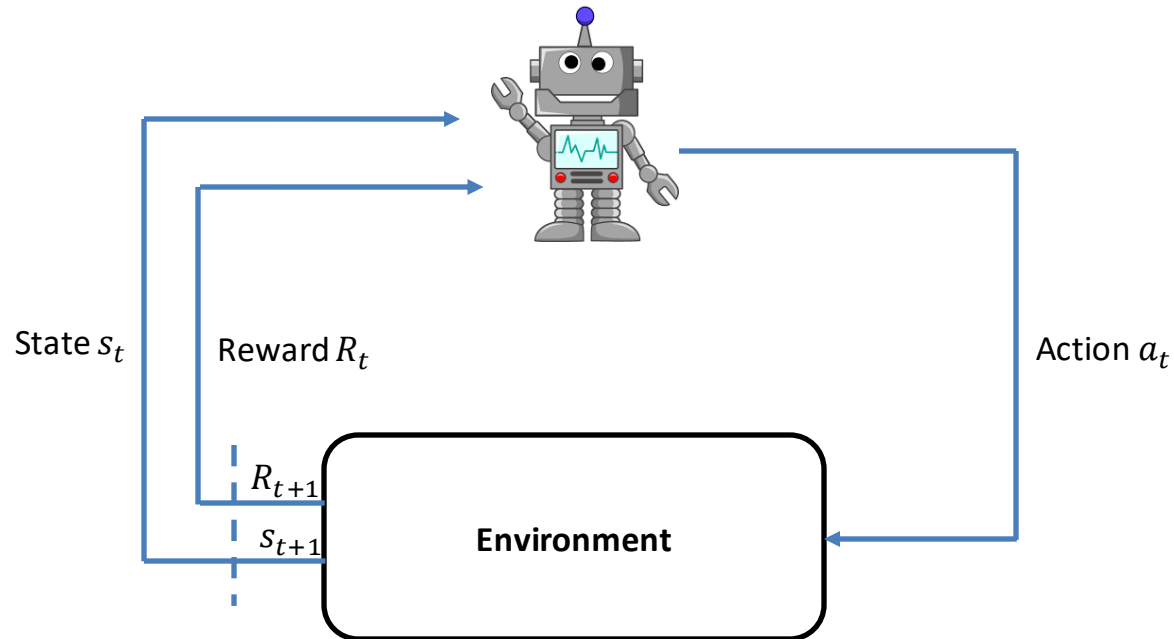
$$\pi^*(s) \in \operatorname{argmax}_{a \in A} Q^*(s, a)$$

# Markov decision process (MDP)

- What happens if we **do not know the stochastic transition model and reward function?**
- We can use **Reinforcement Learning**


# Reinforcement Learning

- Why not **interact** with the environment?



# Reinforcement Learning

- Why not **interact with the environment**?
  - At each time step  $t$ 
    - The agent observes the state  $s_t$
    - The agent takes action  $a_t$
    - The agent observes a reward  $R_t$  and the new state  $s_{t+1}$



# Reinforcement Learning

- Thus, my data point at each iteration is:

$$(s_t, a_t, R_t, s_{t+1})$$

- And what is the agent's goal?
  - Compute the **optimal policy of the MDP** with the data points
- Today, we focus on the most famous RL algorithm: **Q-learning**

# Reinforcement Learning

- We start with some estimate  $Q$
- Initialize current state  $s$
- Loop for each step:
  - choose some action  $a$  (e.g., using  $\epsilon$ -greedy)
  - Take action  $a$  and observe next state  $s'$  and reward  $r$
  - Update  $Q$  estimate according to
$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$
  - $s \leftarrow s'$

Q-learning

# Reinforcement Learning

- Important observations about Q-Learning:
  - Update depends on previous estimate (*bootstrap*)
  - Update rule **propagates** reward information
  - To compute Q, algorithm must visit **every** state-action



# Example

- We have the following MDP:

- $S = \{1, 2, 3\}$


- $A = \{left, right\}$

- $P(s'|s, a = left) = ?$

- $P(s'|s, a = right) = ?$

- $R(s, a) = ?$

- $\gamma = 0.9, \alpha = 0.3$



# Example

- We start with some estimate  $Q$

$$Q = \begin{bmatrix} 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{bmatrix}$$

- Initialize current state  $s$

$$s \rightarrow 1$$

# Example

- First iteration (current state  $s \rightarrow 1$ )
  - choose some action  $a$ 
    - $a \rightarrow left$
  - Take action  $a$  and observe next state  $s'$  and reward  $r$ 
    - $s' \rightarrow 1$
    - $r \rightarrow 0$
  - Update  $Q$  estimate according to
    - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
    - $Q(1, left) \leftarrow 1 + 0.3[0 + 0.9 \times 1 - 1]$
    - $Q(1, left) \leftarrow 0.97$

# Example

- Updated  $Q$

$$Q = \begin{bmatrix} 0.97 & 1 \\ 1 & 1 \\ 1 & 1 \end{bmatrix}$$

# Example

- Second iteration (current state  $s \rightarrow 1$ )
  - choose some action  $a$ 
    - $a \rightarrow right$
  - Take action  $a$  and observe next state  $s'$  and reward  $r$ 
    - $s' \rightarrow 2$
    - $r \rightarrow 0$
  - Update  $Q$  estimate according to
    - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
    - $Q(1, right) \leftarrow 1 + 0.3[0 + 0.9 \times 1 - 1]$
    - $Q(1, right) \leftarrow 0.97$

# Example

- Updated  $Q$

$$Q = \begin{bmatrix} 0.97 & 0.97 \\ 1 & 1 \\ 1 & 1 \end{bmatrix}$$

# Example

- Third iteration (current state  $s \rightarrow 2$ )
  - choose some action  $a$ 
    - $a \rightarrow left$
  - Take action  $a$  and observe next state  $s'$  and reward  $r$ 
    - $s' \rightarrow 1$
    - $r \rightarrow 0$
  - Update  $Q$  estimate according to
    - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
    - $Q(2, left) \leftarrow 1 + 0.3[0 + 0.9 \times 0.97 - 1]$
    - $Q(2, left) \leftarrow 0.9619$

# Example

- Updated  $Q$

$$Q = \begin{bmatrix} 0.97 & 0.97 \\ 0.9619 & 1 \\ 1 & 1 \end{bmatrix}$$



# Example

- Updated  $Q$  (after many iterations)

$$Q = \begin{bmatrix} 6.86 & 7.99 \\ 7.22 & 8.91 \\ 9.2 & 10 \end{bmatrix}$$

- This is  $Q^*$  with Value Iteration

$$Q^* = \begin{bmatrix} 6.89 & 7.66 \\ 7.08 & 8.73 \\ 9.07 & 9.95 \end{bmatrix}$$


# Example

- After computing  $Q$ , we can extract the policy as follows:

$$\pi^*(s) \in \operatorname{argmax}_{a \in A} Q(s, a)$$

$$\pi^* = \begin{bmatrix} \textit{right} \\ \textit{right} \\ \textit{right} \end{bmatrix}$$

$$\pi^* = \begin{bmatrix} 0 & 1 \\ 0 & 1 \\ 0 & 1 \end{bmatrix}$$



# Example

```
import numpy as np
np.set_printoptions(precision=2, suppress=True)

# States
S = ['1', '2', '3']

# Actions
A = ['L', 'R']

# Transition probabilities

L = np.array([[1.0, 0.0, 0.0],
              [0.8, 0.2, 0.0],
              [0.0, 0.8, 0.2]])

R = np.array([[0.2, 0.8, 0.0],
              [0.0, 0.2, 0.8],
              [0.0, 0.0, 1.0]])

P = [L, R]

# Reward function

R = np.array([[0.0, 0.0],
              [0.0, 0.0],
              [1.0, 1.0]])

gamma = 0.9
```

# Example

```
def egreedy(Q, state, eps):  
    p = np.random.random()  
  
    if p < eps:  
        action = np.random.choice(num_actions)  
    else:  
        action = np.argmax(Q[state,:])  
  
    return action
```

# Example

```
STEPS = 1000000
num_actions = len(A)
num_states = len(S)
ALPHA = 0.3

# Initialize Q-values
Q = np.ones((num_states, num_actions))

# Initialize current state
state = 0

for t in range(STEPS):

    # choose action
    action = egreedy(Q, state, 0.05)

    # choose next state
    next_state = np.random.choice(num_states, p=P[action][state, :])

    # obtain reward
    reward = R[state, action]

    # Update Q
    Q[state, action] = Q[state, action] + ALPHA*(reward + gamma*max(Q[next_state, :]) - Q[state, action])

    state = next_state

print(Q)
```