TÉCNICO LISBOA

Data Analysis and Integration

ETL Processes Data Warehousing

Introduction

- How to build a data warehouse
  - ETL process
    - **Extract** data from original database
    - **Transform** data to fit star schema
    - **Load** data onto data warehouse

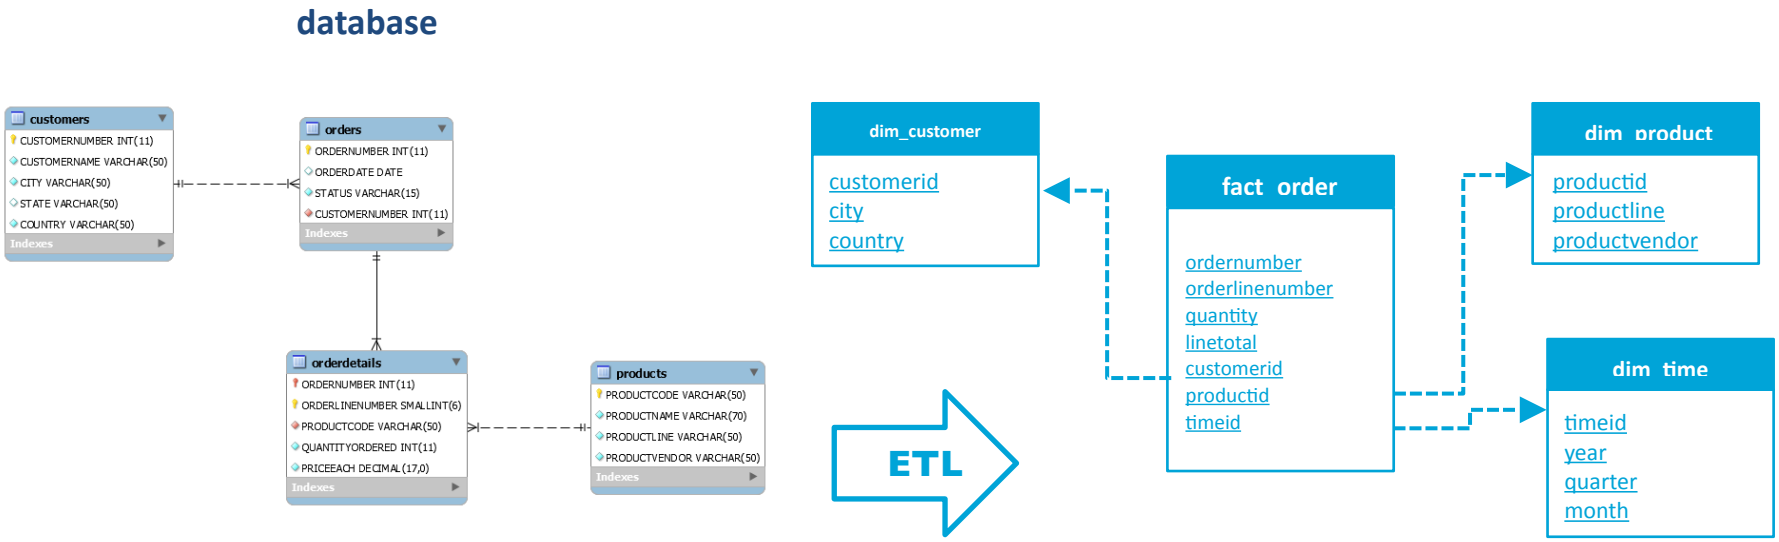

Introduction

- How to build a data warehouse
  - this usually involves
    - one transformation for each dimension table
    - one transformation for the fact table
    - a job that runs all transformations in the correct sequence

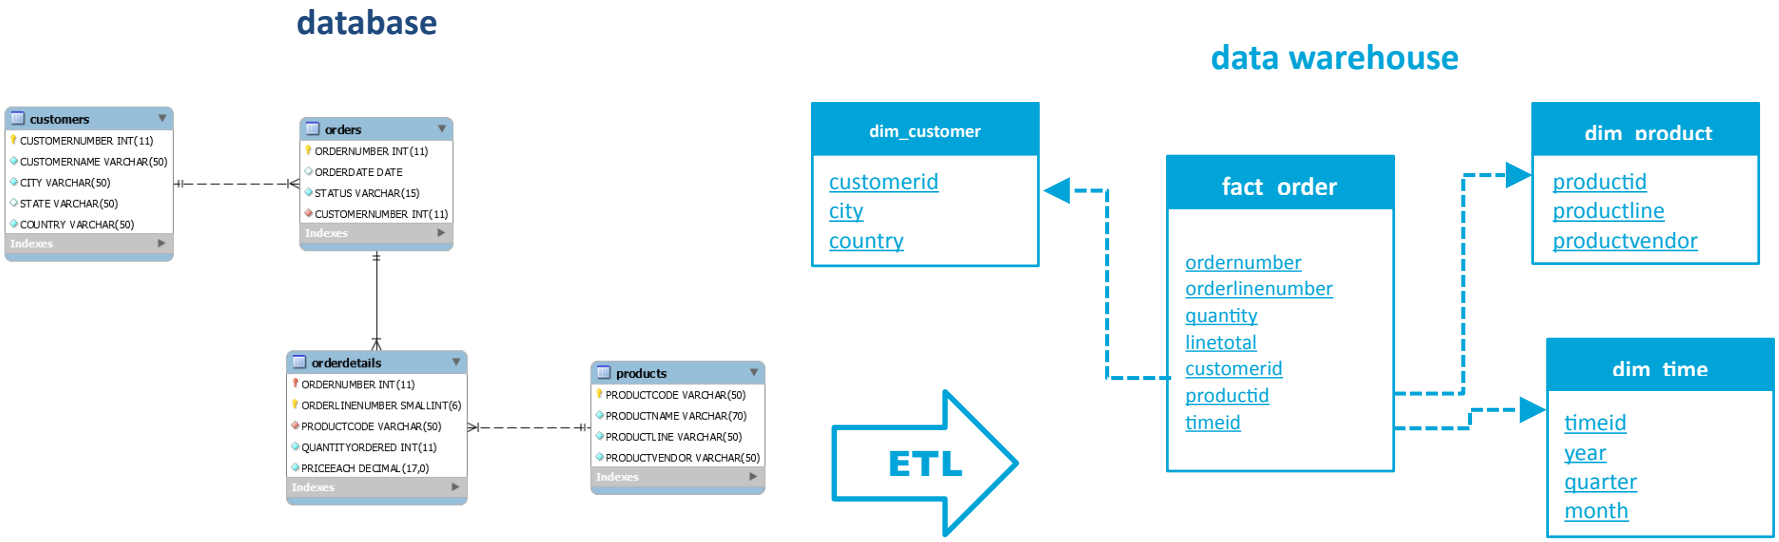

Database

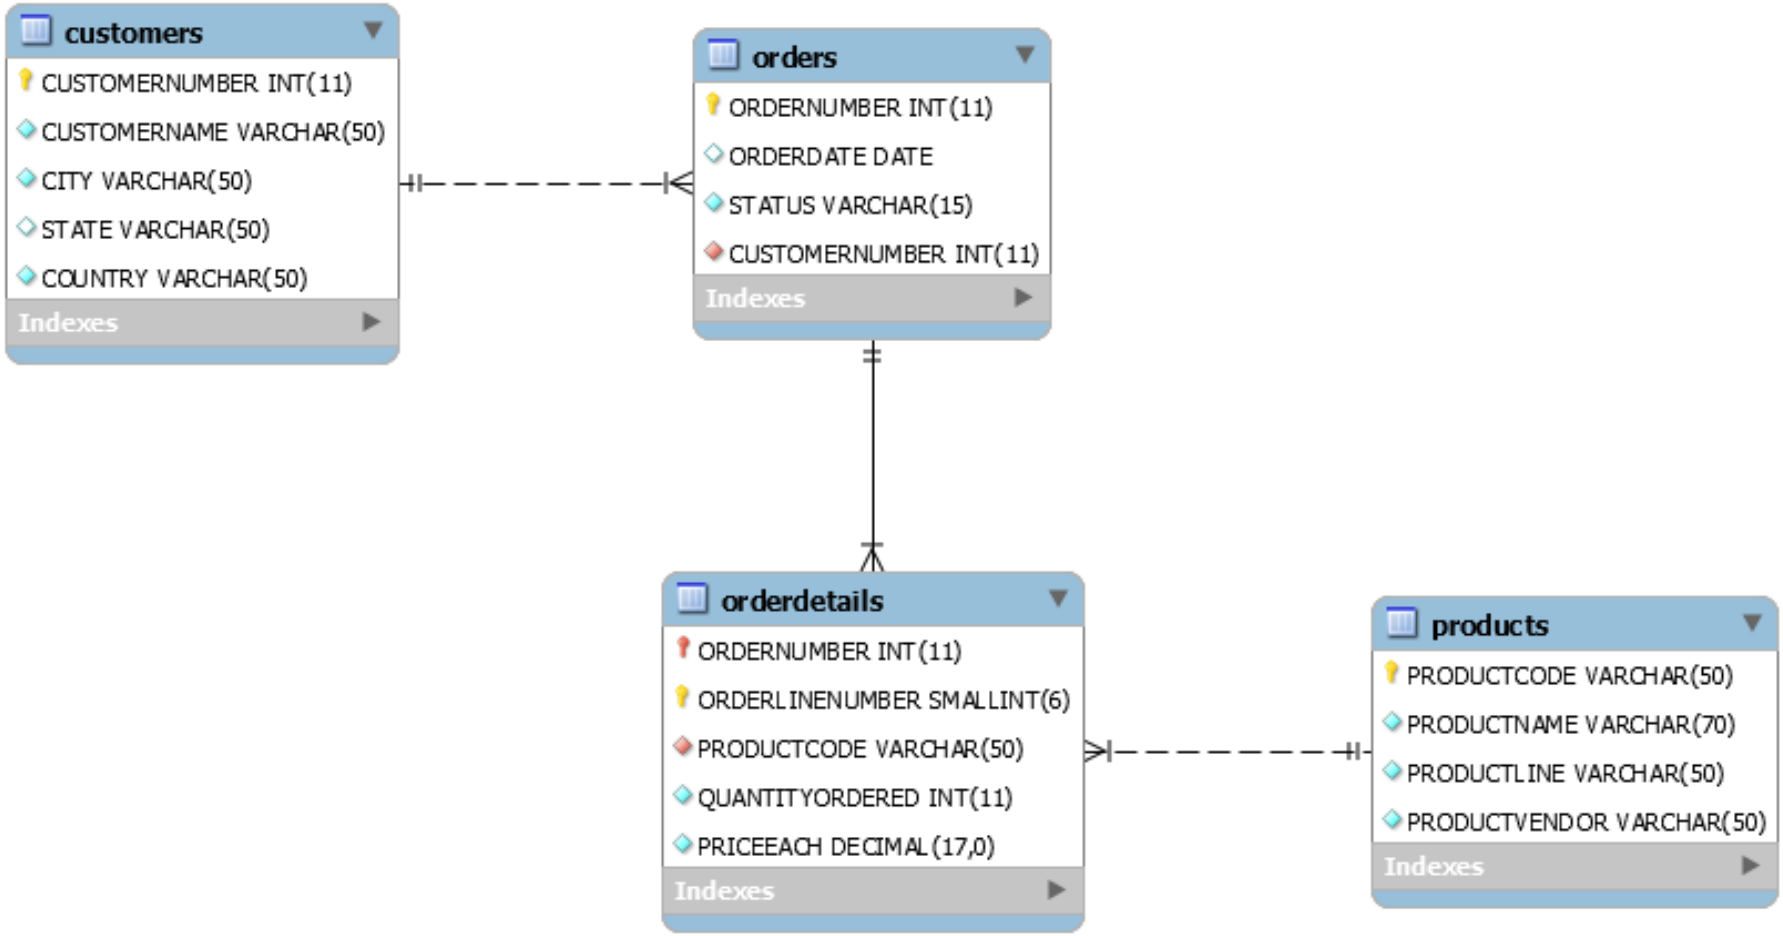
customers
CUSTOMERNUMBER INT(11)
CUSTOMERNAME VARCHAR(50)
CITY VARCHAR(50)
STATE VARCHAR(50)
COUNTRY VARCHAR(50)
Indexes
orders
ORDERNUMBER INT(11)
ORDERDATE DATE
STATUS VARCHAR(15)
CUSTOMERNUMBER INT(11)
Indexes
orderdetails
ORDERNUMBER INT(11)
ORDERLINENUMBER SMALLINT(6)
PRODUCTCODE VARCHAR(50)
QUANTITYORDERED INT(11)
PRICEEACH DECIMAL(17,0)
Indexes
products
PRODUCTCODE VARCHAR(50)
PRODUCTNAME VARCHAR(70)
PRODUCTLINE VARCHAR(50)
PRODUCTVENDOR VARCHAR(50)
Indexes

Data warehouse

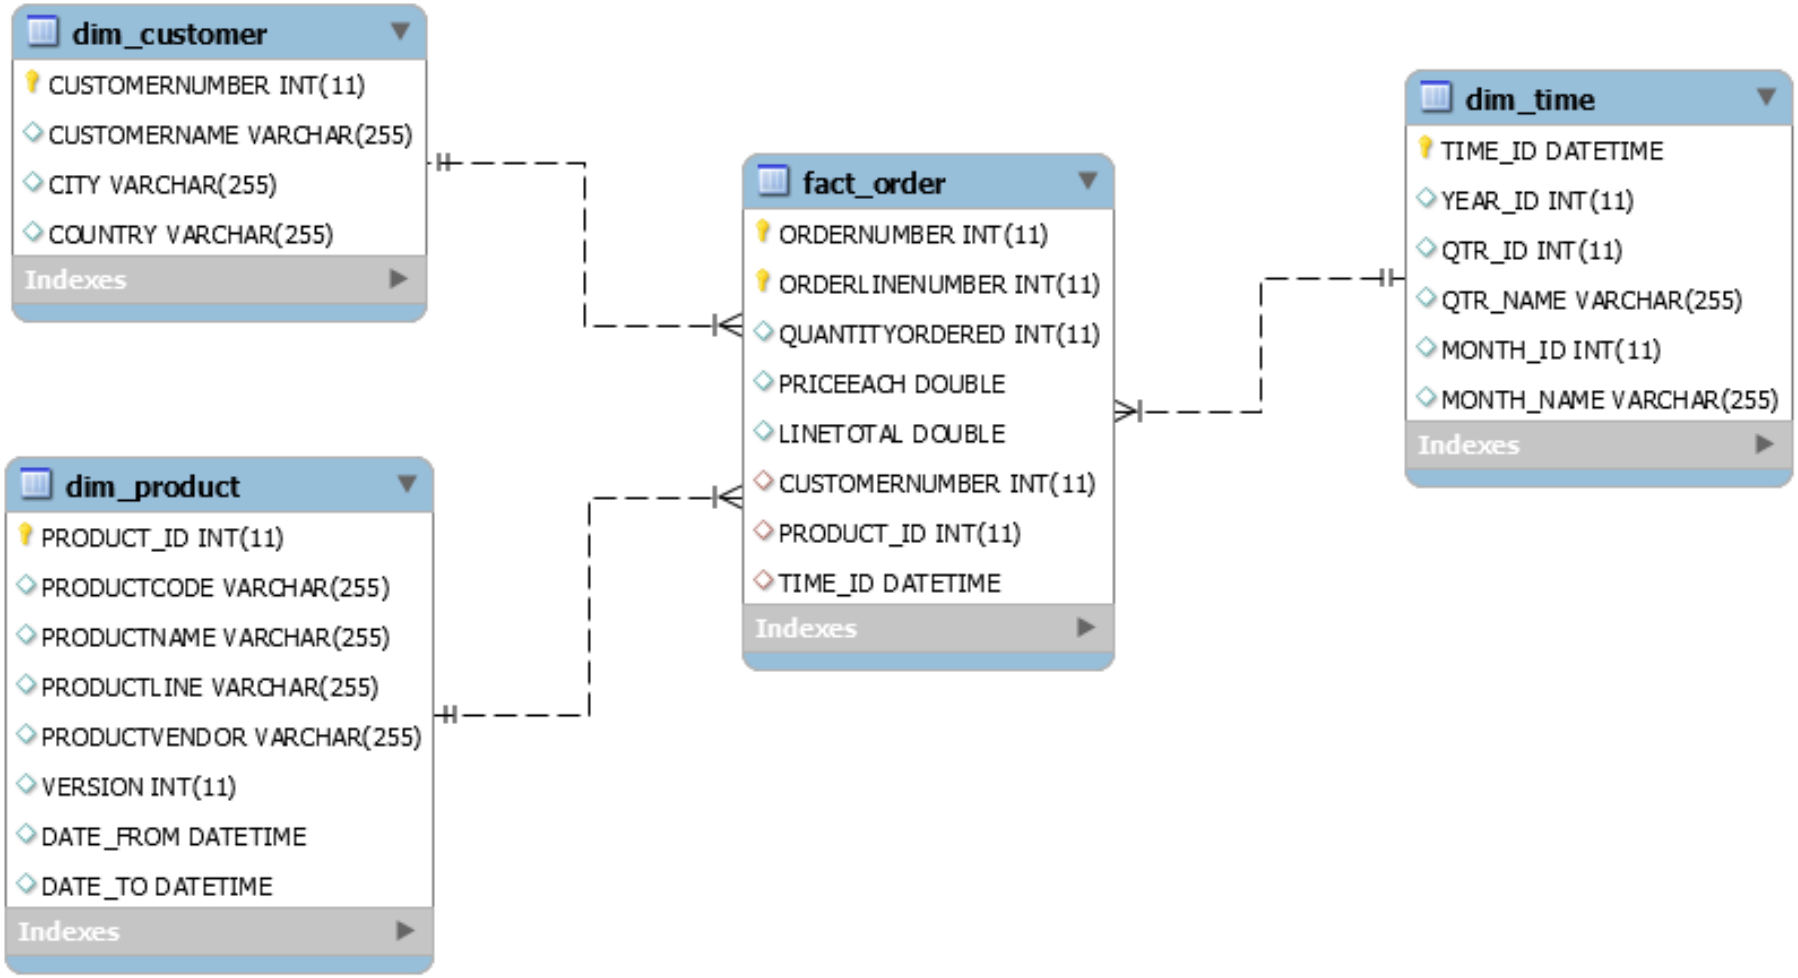

Data warehouse

- Some notes on this example
  - **dim_customer** is simplified, it should have a surrogate key
    - e.g. CUSTOMER_ID of type INT
    - in this case, it was simplified by reusing the natural key
  - **dim_time is simplified**, it should have a proper surrogate key
    - e.g. TIME_ID of type INT
    - in this case, it was simplified by using ORDERDATE as key
  - **dim_product** is a slowly-changing dimension of Type 2
    - surrogate key is PRODUCT_ID of type INT
    - includes version field and validity interval (date fields)

Fact and dimension tables

- Building the data warehouse
  - Star schema with fact table and dimension tables
  - Fact table has FKs to dimension tables
    - Therefore, dimension tables must be populated first

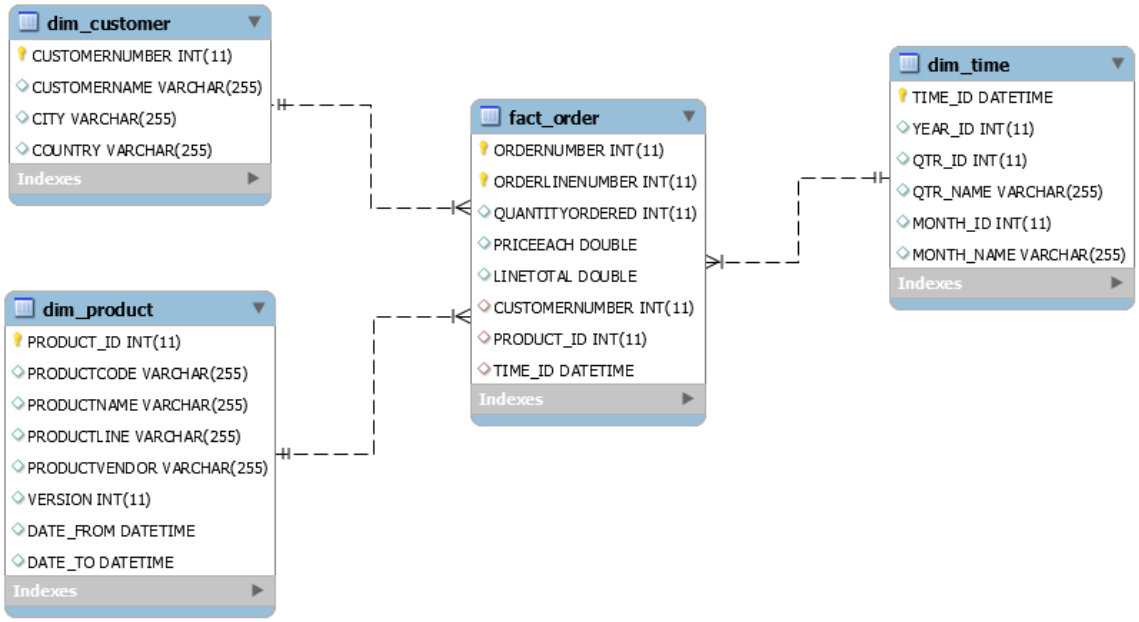

Loading Dimension Tables

Dimension tables

- The customer dimension
  - customer name, city and country (no state)
  - data comes from customers table
  - same key as customers table (natural key)

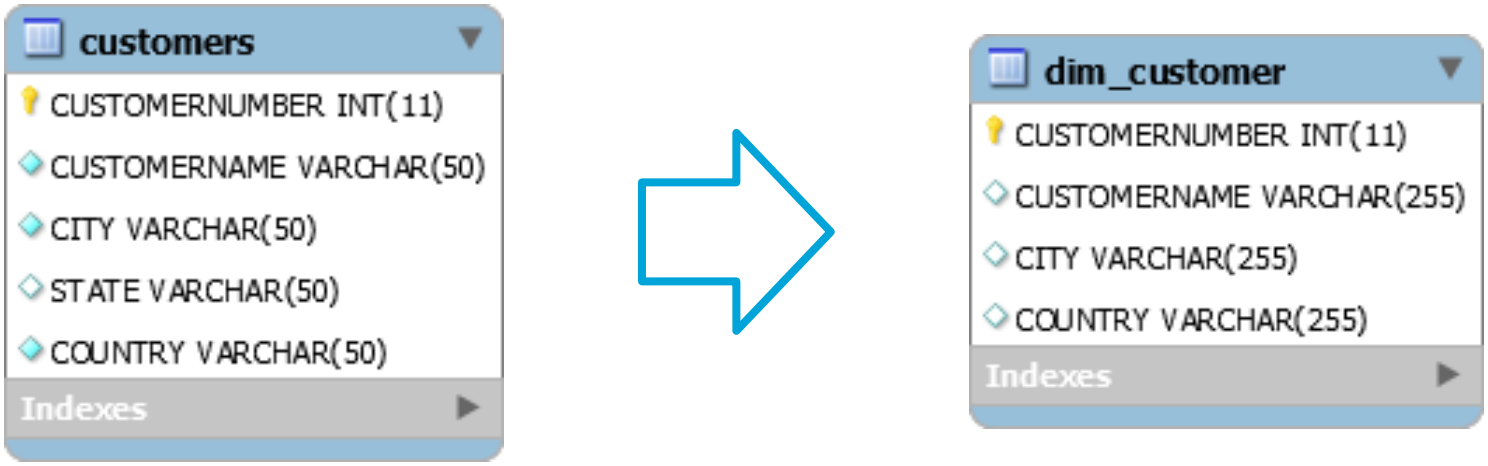

Dimension tables

- The customer dimension

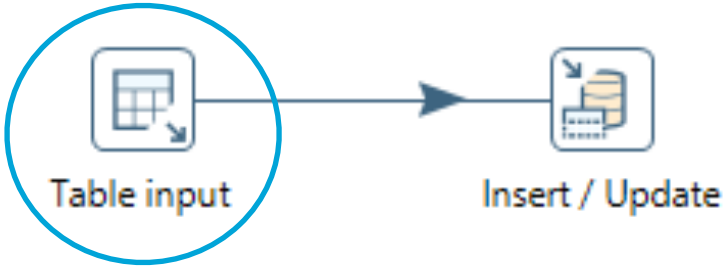

Table input

Step name: Table input

Connection: steelwheels | Edit... | New... | Wizard...

SQL | Get SQL select statement...

```
SELECT
  CUSTOMERNUMBER
, CUSTOMERNAME
, CITY
, COUNTRY
FROM customers
```

Dimension tables

- The customer dimension

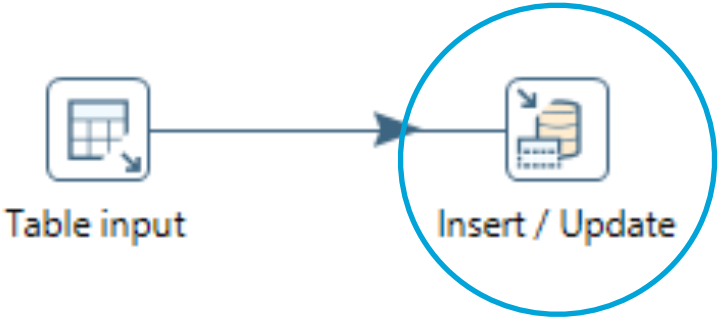

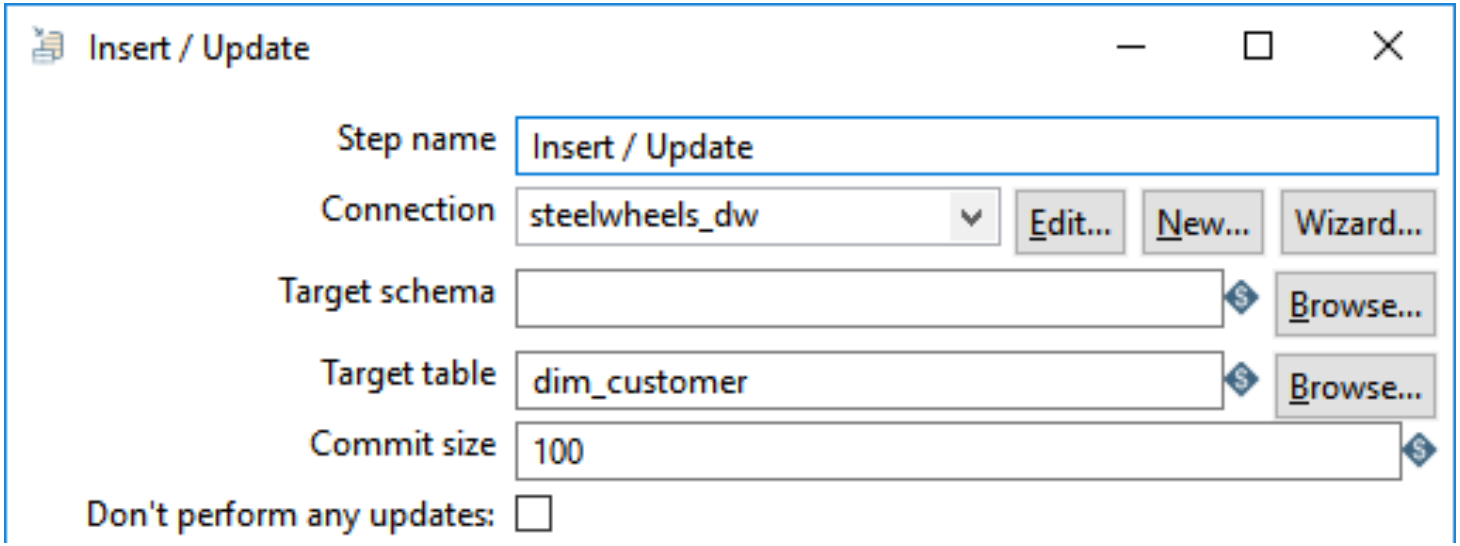

Dimension tables

- The customer dimension

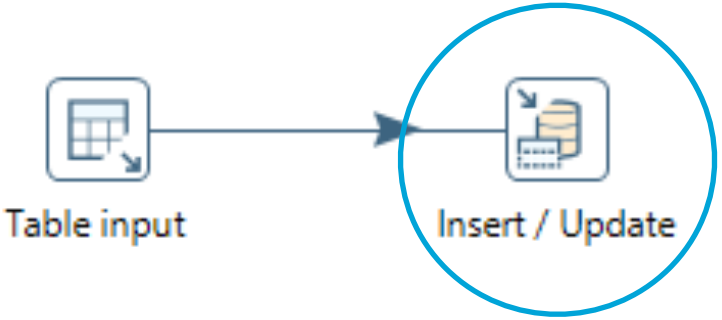

The key(s) to look up the value(s):

#	Table field	Comparator	Stream field1	Stre
1	CUSTOMERNUMBER	=	CUSTOMERNUMBER	

Get fields

Dimension tables

- The customer dimension

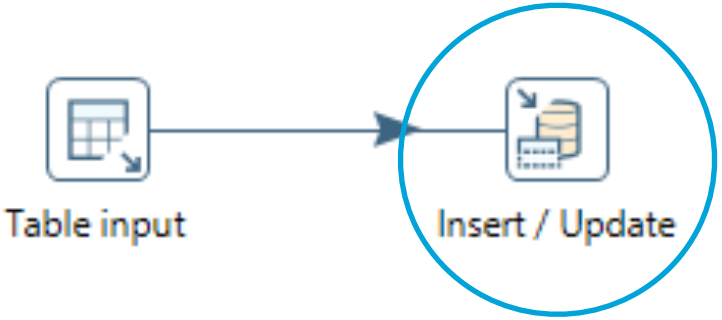

Update fields:

#	Table field	Stream field	Update
1	CUSTOMERNUMBER	CUSTOMERNUMBER	Y
2	CUSTOMERNAME	CUSTOMERNAME	Y
3	CITY	CITY	Y
4	COUNTRY	COUNTRY	Y

Get update fields

Edit mapping

Dimension tables

- The customer dimension

CUSTOMERNUMBER	CUSTOMERNAME	CITY	COUNTRY
97	Madison Inc	ST AUGUSTINE	USA
98	Johnson Inc	ST Cloud	USA
99	Tarallo Inc	Sanford	USA
100	Audio Video 'R' Us	Orlando	USA
103	Atelier graphique	Nantes	France
112	Signal Gift Stores	Las Vegas	USA
114	Australian Collectors, Co.	Melbourne	Australia
119	La Rochelle Gifts	Nantes	France
121	Baane Mini Imports	Stavern	Norway
124	Mini Gifts Distributors Ltd.	San Rafael	USA
125	Havel & Zbyszek Co	Warszawa	Poland
128	Blauer See Auto, Co.	Frankfurt	Germany
129	Mini Wheels Co.	San Francisco	USA
131	Land of Toys Inc.	NYC	USA
141	Euro+ Shopping Channel	Madrid	Spain
144	Volvo Model Replicas, Co	Luleå	Sweden
145	Danish Wholesale Imports	Kobenhavn	Denmark
146	Saveley & Henriot, Co.	Lyon	France
148	Dragon Souveniers, Ltd.	Singapore	Singapore
151	Muscle Machine Inc	NYC	USA

Dimension tables

- The product dimension
  - product name, line, vendor
  - data comes from products table
  - key is not product code but product id (surrogate key)
  - slowly-changing dimension (Type 2)

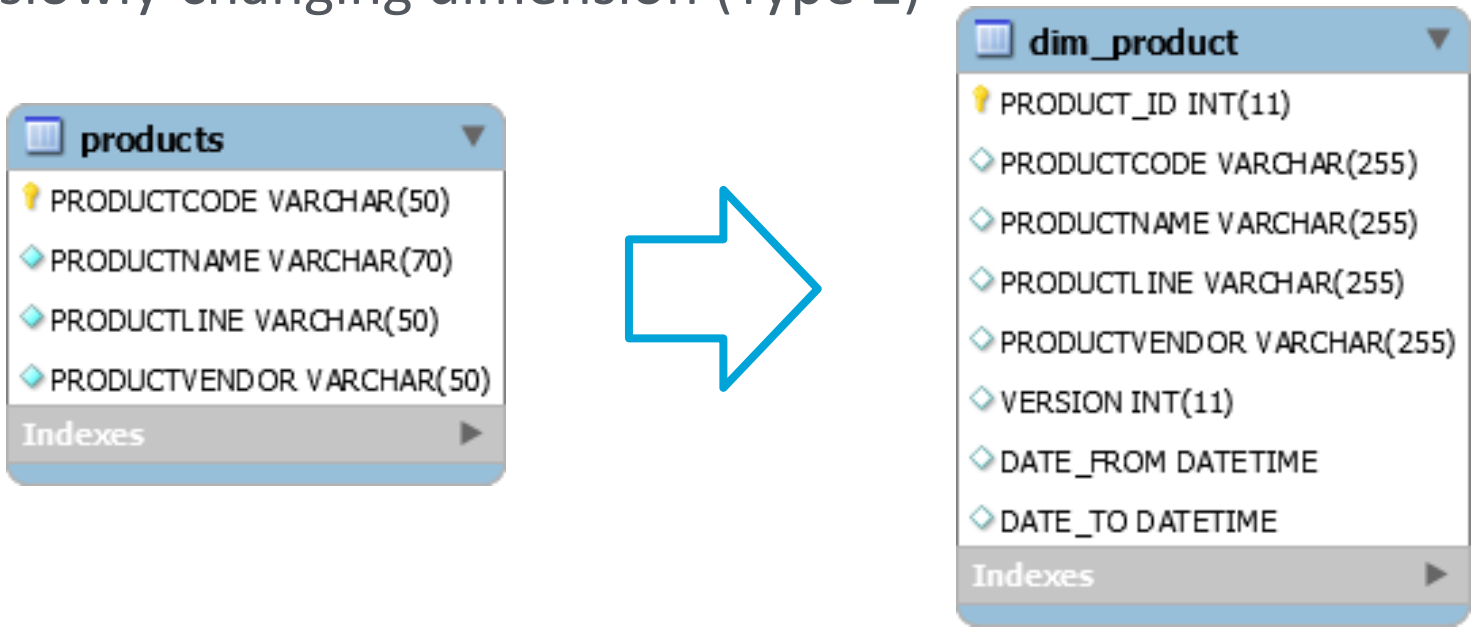

Dimension tables

- The product dimension

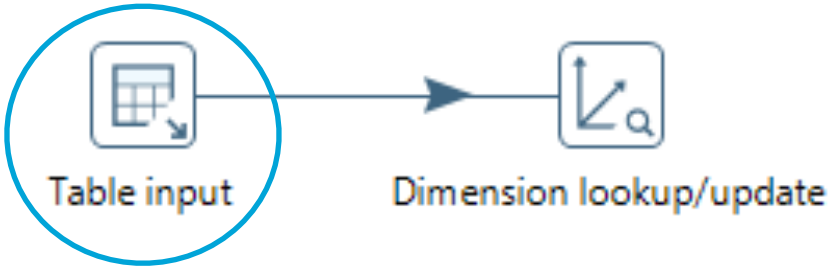

Table input

Step name: Table input

Connection: steelwheels | Edit... | New... | Wizard...

SQL | Get SQL select statement...

```
SELECT
  PRODUCTCODE
, PRODUCTNAME
, PRODUCTLINE
, PRODUCTVENDOR
FROM products
```

Dimension tables

- The product dimension

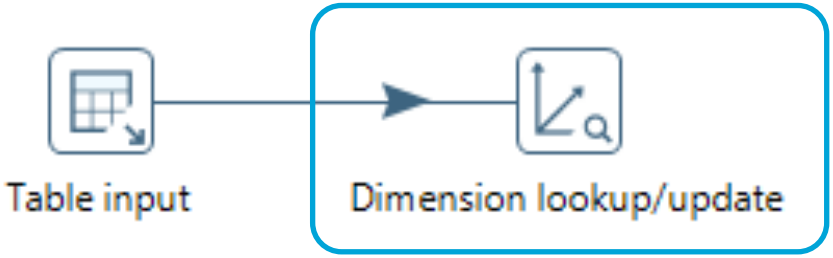

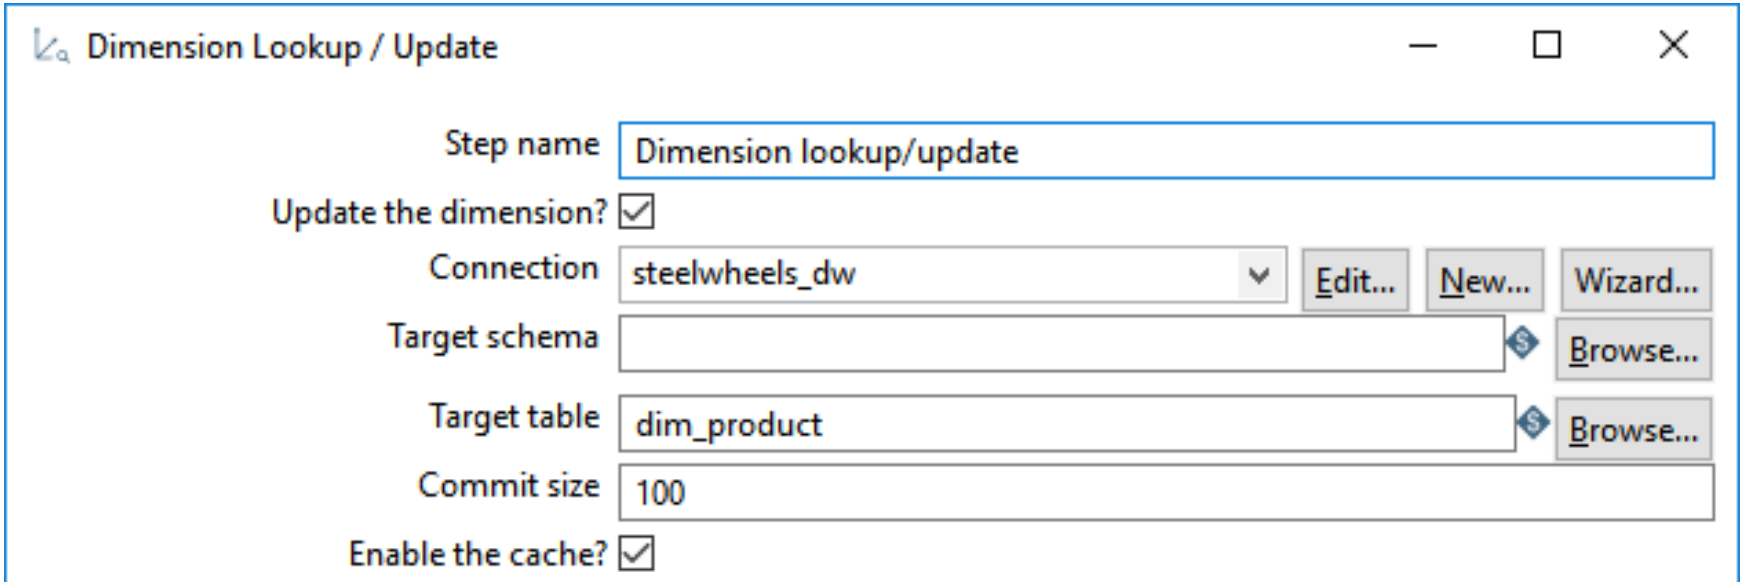

Dimension tables

- The product dimension

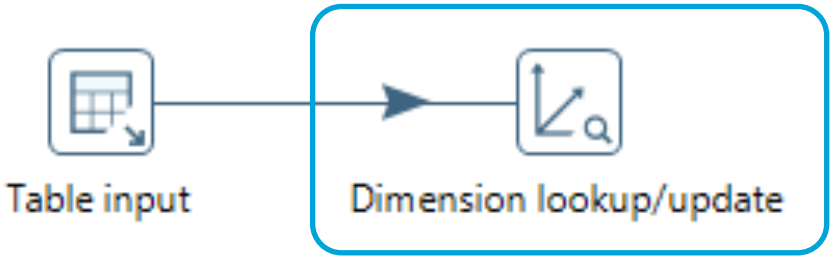

Keys | Fields

Key fields (to look up row in dimension):

#	Dimension field	Field in stream
1	PRODUCTCODE	PRODUCTCODE

Dimension tables

- The product dimension

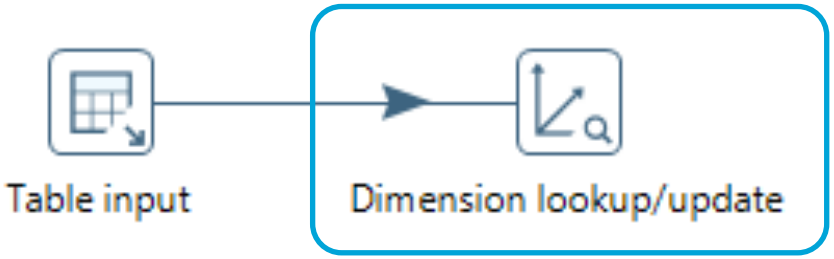

Keys | Fields

Lookup/Update fields

#	Dimension field	Stream field to compare with	Type of dimension update
1	PRODUCTNAME	PRODUCTNAME	Insert
2	PRODUCTLINE	PRODUCTLINE	Insert
3	PRODUCTVENDOR	PRODUCTVENDOR	Insert

Dimension tables

- The product dimension

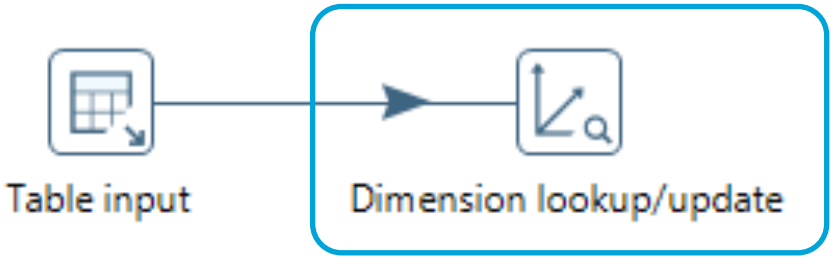

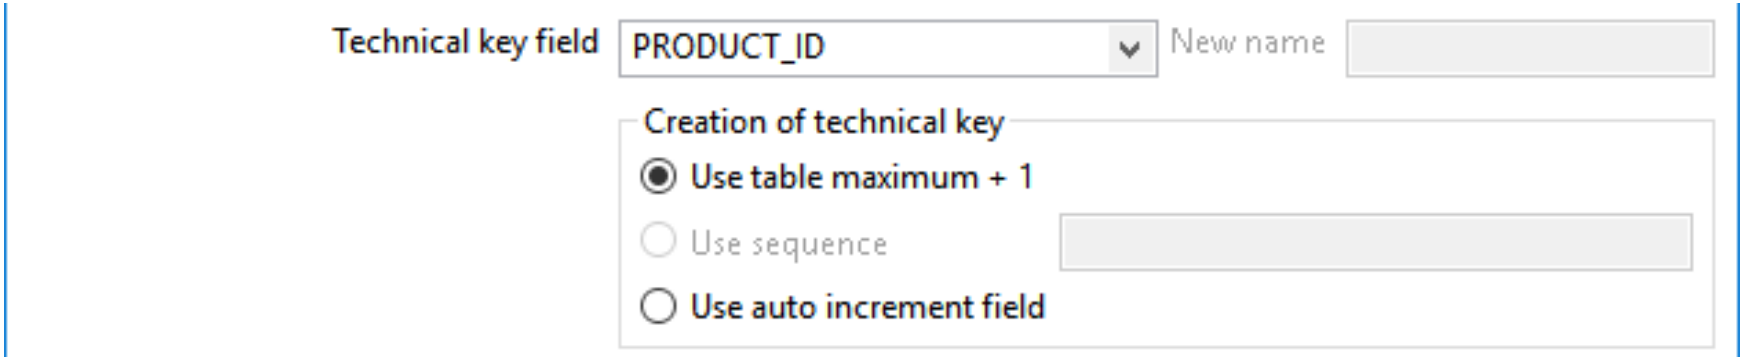

Dimension tables

- The product dimension

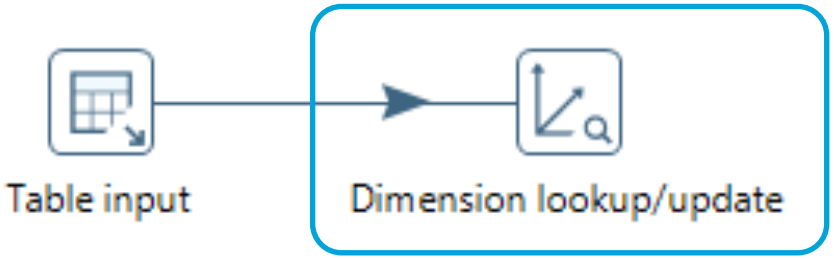

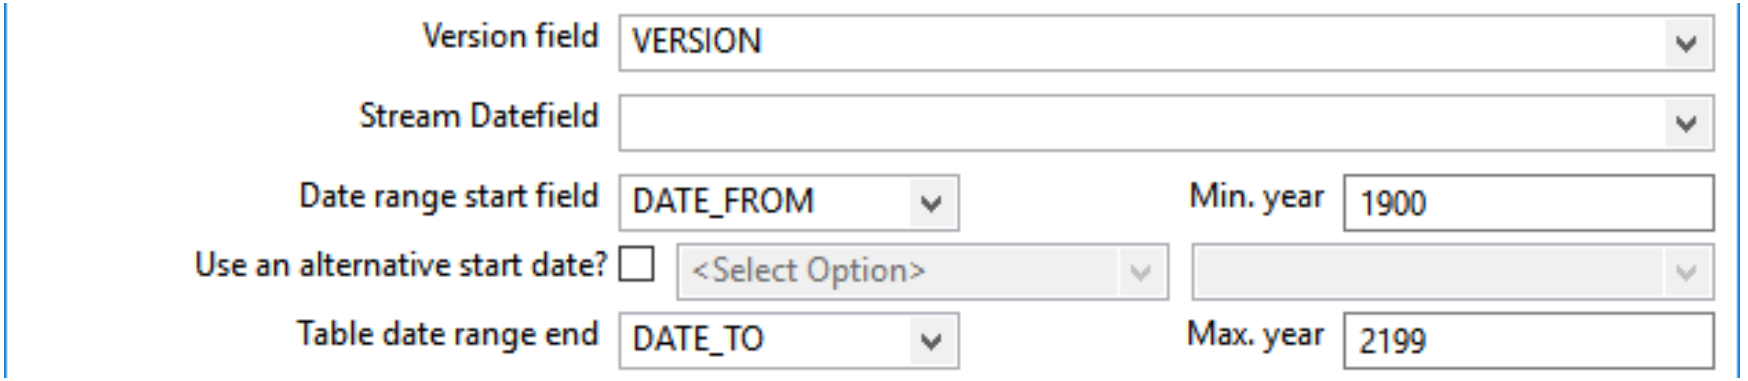

Dimension tables

- The product dimension

PRODUCT_ID	PRODUCTCODE	PRODUCTNAME	PRODUCTLINE	PRODUCTVENDOR	VERSION	DATE_FROM	DATE_TO
1	S10_1678	1969 Harley Davidson Ultimate Chopper	Motorcycles	Min Lin Diecast	1	1900-01-01 00:00:00	2200-01-01 00:00:00
2	S10_1949	1952 Alpine Renault 1300	Classic Cars	Classic Metal Creations	1	1900-01-01 00:00:00	2200-01-01 00:00:00
3	S10_2016	1996 Moto Guzzi 1100i	Motorcycles	Highway 66 Mini Classics	1	1900-01-01 00:00:00	2200-01-01 00:00:00
4	S10_4698	2003 Harley-Davidson Eagle Drag Bike	Motorcycles	Red Start Diecast	1	1900-01-01 00:00:00	2200-01-01 00:00:00
5	S10_4757	1972 Alfa Romeo GTA	Classic Cars	Motor City Art Classics	1	1900-01-01 00:00:00	2200-01-01 00:00:00
6	S10_4962	1962 LanciaA Delta 16V	Classic Cars	Second Gear Diecast	1	1900-01-01 00:00:00	2200-01-01 00:00:00
7	S12_1099	1968 Ford Mustang	Classic Cars	Autoart Studio Design	1	1900-01-01 00:00:00	2200-01-01 00:00:00
8	S12_1108	2001 Ferrari Enzo	Classic Cars	Second Gear Diecast	1	1900-01-01 00:00:00	2200-01-01 00:00:00
9	S12_1666	1958 Setra Bus	Trucks and Buses	Welly Diecast Productions	1	1900-01-01 00:00:00	2200-01-01 00:00:00
10	S12_2823	2002 Suzuki XREO	Motorcycles	Unimax Art Galleries	1	1900-01-01 00:00:00	2200-01-01 00:00:00
11	S12_3148	1969 Corvair Monza	Classic Cars	Welly Diecast Productions	1	1900-01-01 00:00:00	2200-01-01 00:00:00
12	S12_3380	1968 Dodge Charger	Classic Cars	Welly Diecast Productions	1	1900-01-01 00:00:00	2200-01-01 00:00:00
13	S12_3891	1969 Ford Falcon	Classic Cars	Second Gear Diecast	1	1900-01-01 00:00:00	2200-01-01 00:00:00
14	S12_3990	1970 Plymouth Hemi Cuda	Classic Cars	Studio M Art Models	1	1900-01-01 00:00:00	2200-01-01 00:00:00
15	S12_4473	1957 Chevy Pickup	Trucks and Buses	Exoto Designs	1	1900-01-01 00:00:00	2200-01-01 00:00:00
16	S12_4675	1969 Dodge Charger	Classic Cars	Welly Diecast Productions	1	1900-01-01 00:00:00	2200-01-01 00:00:00
17	S18_1097	1940 Ford Pickup Truck	Trucks and Buses	Studio M Art Models	1	1900-01-01 00:00:00	2200-01-01 00:00:00
18	S18_1129	1993 Mazda RX-7	Classic Cars	Highway 66 Mini Classics	1	1900-01-01 00:00:00	2200-01-01 00:00:00
19	S18_1342	1937 Lincoln Berline	Vintage Cars	Motor City Art Classics	1	1900-01-01 00:00:00	2200-01-01 00:00:00

Dimension tables

- The product dimension
  - **Testing the slowly changing dimension**
    - Change product line in the original database
    - Run transformation again
    - There should be now two rows for the same product

PRODUCT_ID	PRODUCTCODE	PRODUCTNAME	PRODUCTLINE	PRODUCTVENDOR	VERSION	DATE_FROM	DATE_TO
24	S18_1889	1948 Porsche 356-A Roadster	Classic Cars	Gearbox Collectibles	1	1900-01-01 00:00:00	2017-11-13 14:31:06
111	S18_1889	1948 Porsche 356-A Roadster	Vintage Cars	Gearbox Collectibles	2	2017-11-13 14:31:06	2200-01-01 00:00:00

Dimension tables

- The time dimension
  - Year, quarter, month (id and name for quarter and month)
  - Data comes from order date alone
  - Key is time_id

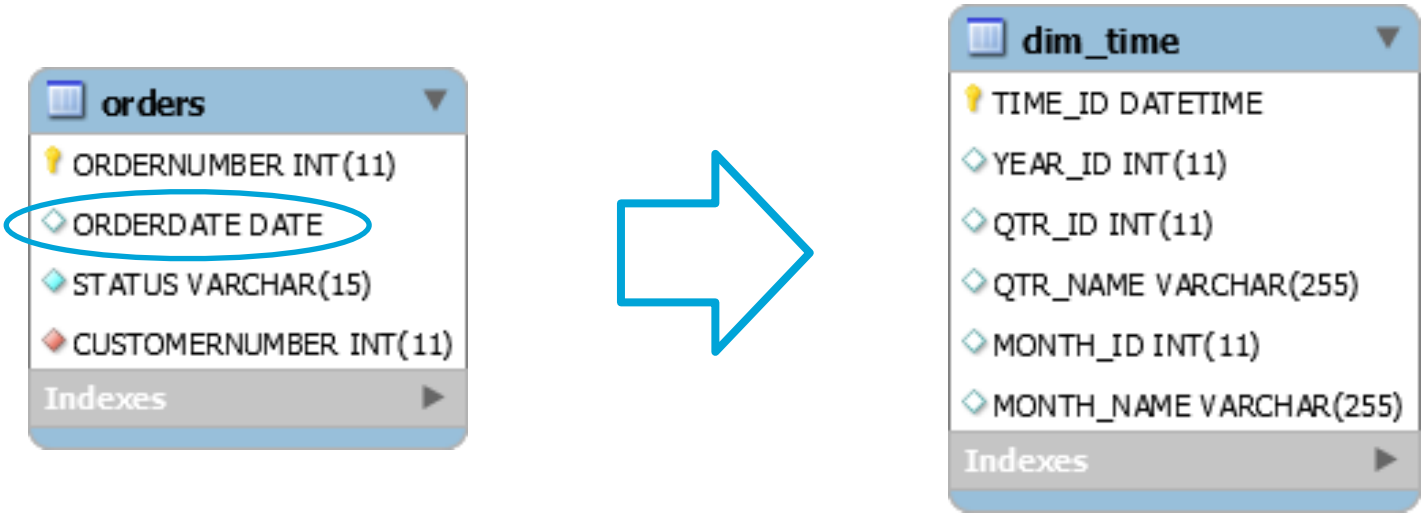

Dimension tables

Not the most robust way to do it! WHY?

- The time dimension

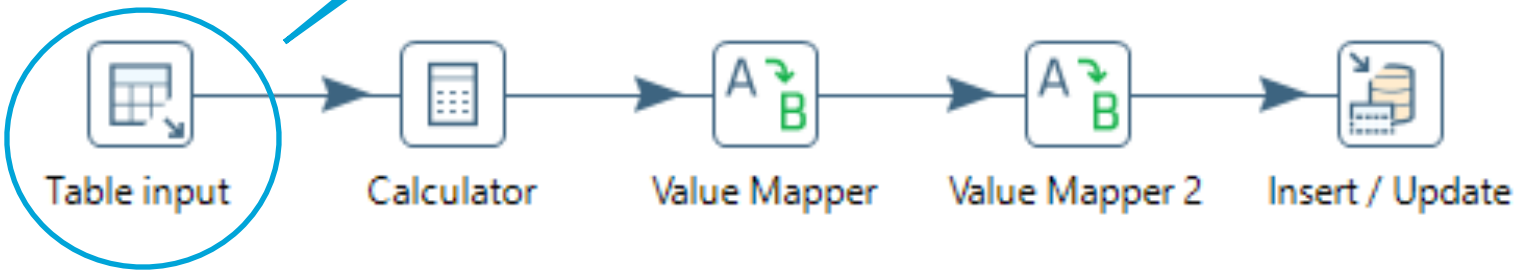

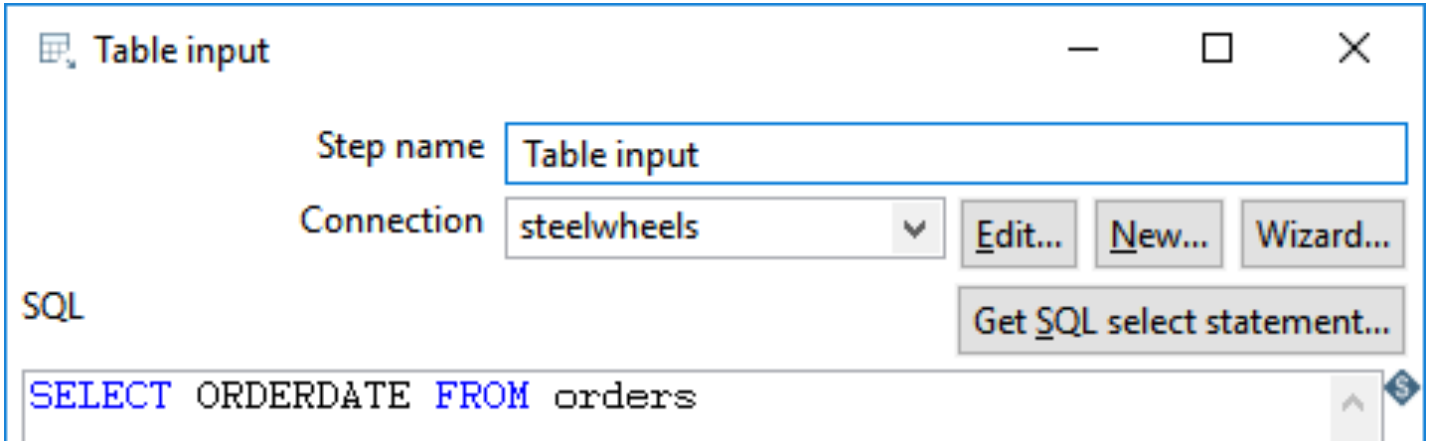

Dimension tables

- The time dimension

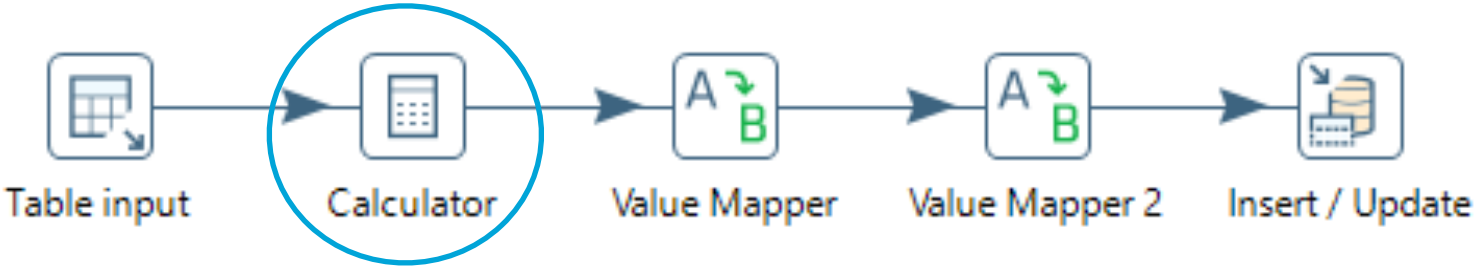

Calculator

Step name: Calculator

Fields:

#	New field	Calculation	Field A	Field B	Field C
1	year_id	Year of date A	ORDERDATE		
2	qtr_id	Quarter of date A	ORDERDATE		
3	month_id	Month of date A	ORDERDATE		

Dimension tables

- The time dimension

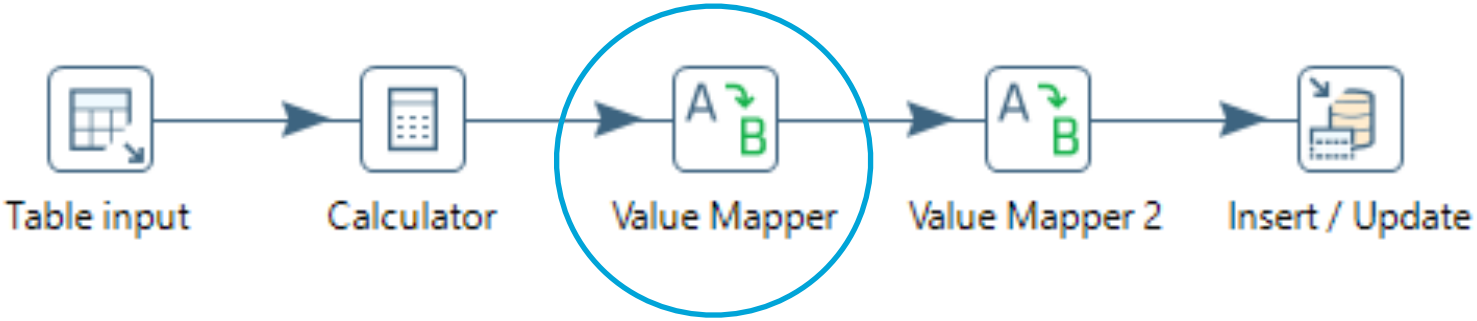

Value Mapper

Step name : Value Mapper

Fieldname to use : qtr_id

Target field name qtr_name

Default upon non-matching

Field values:

#	Source value	Target value
1	1	Q1
2	2	Q2

Dimension tables

- The time dimension

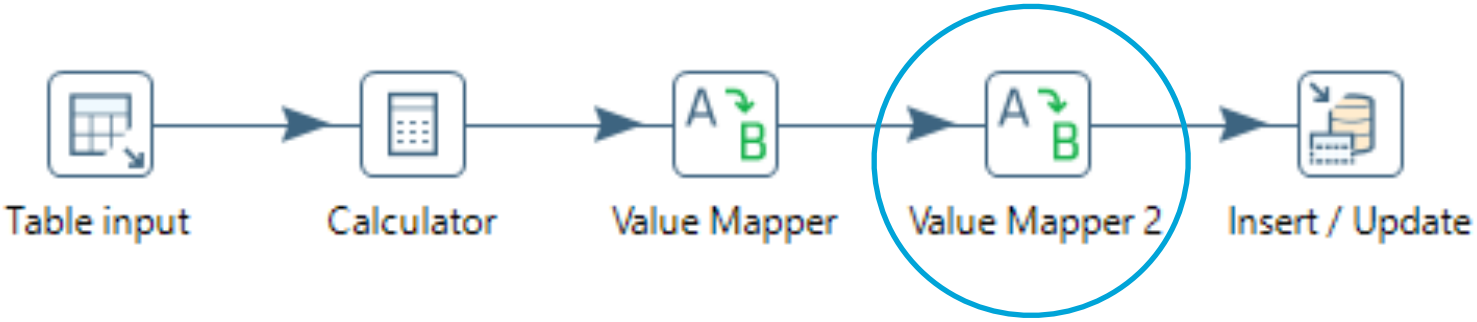

Value Mapper

Step name : Value Mapper 2

Fieldname to use : month_id

Target field name : month_name

Default upon non-matching :

Field values:

#	Source value	Target value
1	1	Jan
2	2	Feb

Dimension tables

- The time dimension

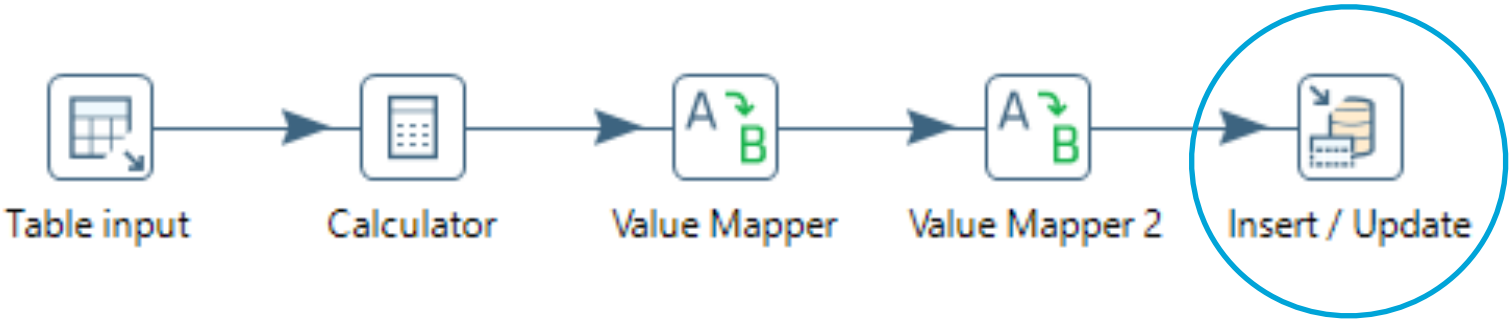

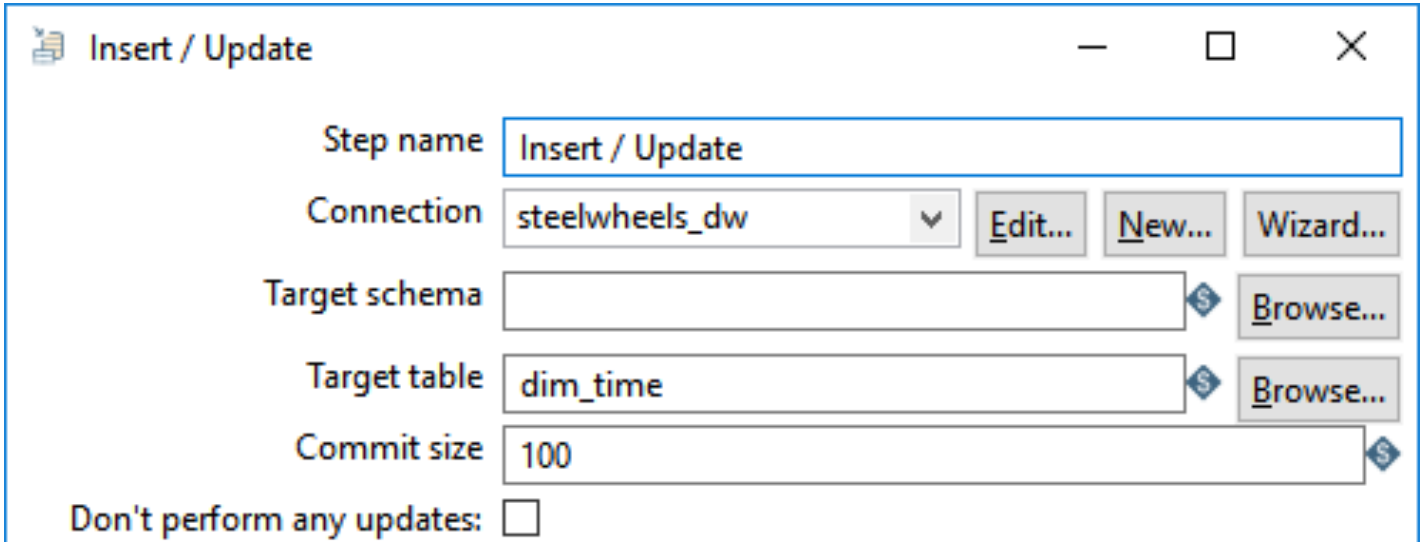

Dimension tables

- The time dimension

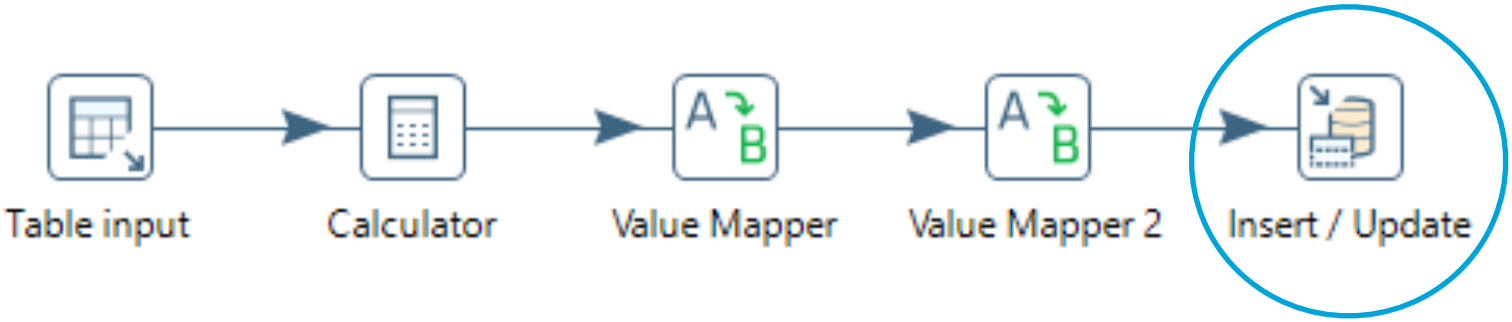

The key(s) to look up the value(s):

#	Table field	Comparator	Stream field1	Stream field2
1	TIME_ID	=	ORDERDATE	

Get fields

Dimension tables

- The time dimension

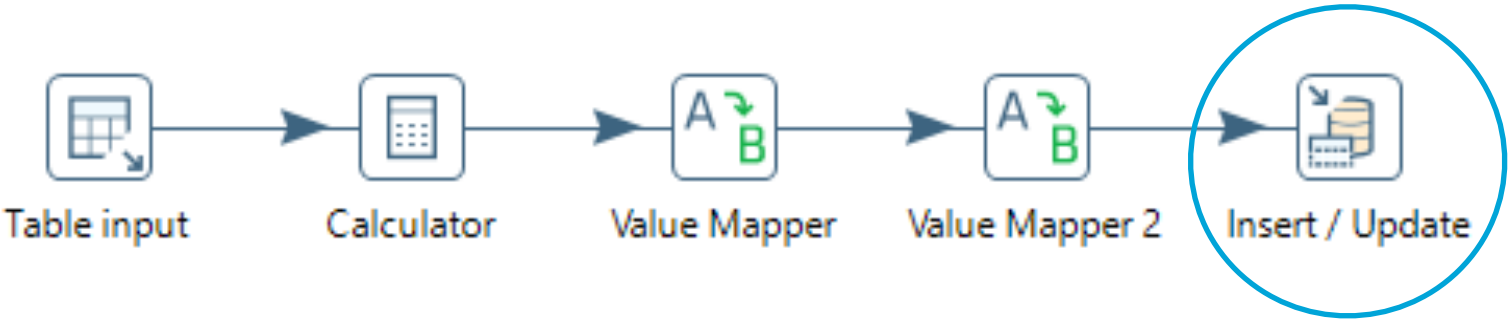

Update fields:

#	Table field	Stream field	Update
1	TIME_ID	ORDERDATE	Y
2	YEAR_ID	year_id	Y
3	QTR_ID	qtr_id	Y
4	QTR_NAME	qtr_name	Y
5	MONTH_ID	month_id	Y
6	MONTH_NAME	month_name	Y

Get update fields

Edit mapping

Dimension tables

- The time dimension

TIME_ID	YEAR_ID	QTR_ID	QTR_NAME	MONTH_ID	MONTH_NAME
2003-01-06 00:00:00	2003	1	Q1	1	Jan
2003-01-09 00:00:00	2003	1	Q1	1	Jan
2003-01-10 00:00:00	2003	1	Q1	1	Jan
2003-01-29 00:00:00	2003	1	Q1	1	Jan
2003-01-31 00:00:00	2003	1	Q1	1	Jan
2003-02-11 00:00:00	2003	1	Q1	2	Feb
2003-02-17 00:00:00	2003	1	Q1	2	Feb
2003-02-24 00:00:00	2003	1	Q1	2	Feb
2003-03-03 00:00:00	2003	1	Q1	3	Mar
2003-03-10 00:00:00	2003	1	Q1	3	Mar
2003-03-18 00:00:00	2003	1	Q1	3	Mar
2003-03-24 00:00:00	2003	1	Q1	3	Mar
2003-03-25 00:00:00	2003	1	Q1	3	Mar
2003-03-26 00:00:00	2003	1	Q1	3	Mar
2003-04-01 00:00:00	2003	2	Q2	4	Apr
2003-04-04 00:00:00	2003	2	Q2	4	Apr
2003-04-11 00:00:00	2003	2	Q2	4	Apr
2003-04-16 00:00:00	2003	2	Q2	4	Apr
2003-04-21 00:00:00	2003	2	Q2	4	Apr
2003-04-28 00:00:00	2003	2	Q2	4	Apr

Loading the Facts table

Data warehouse

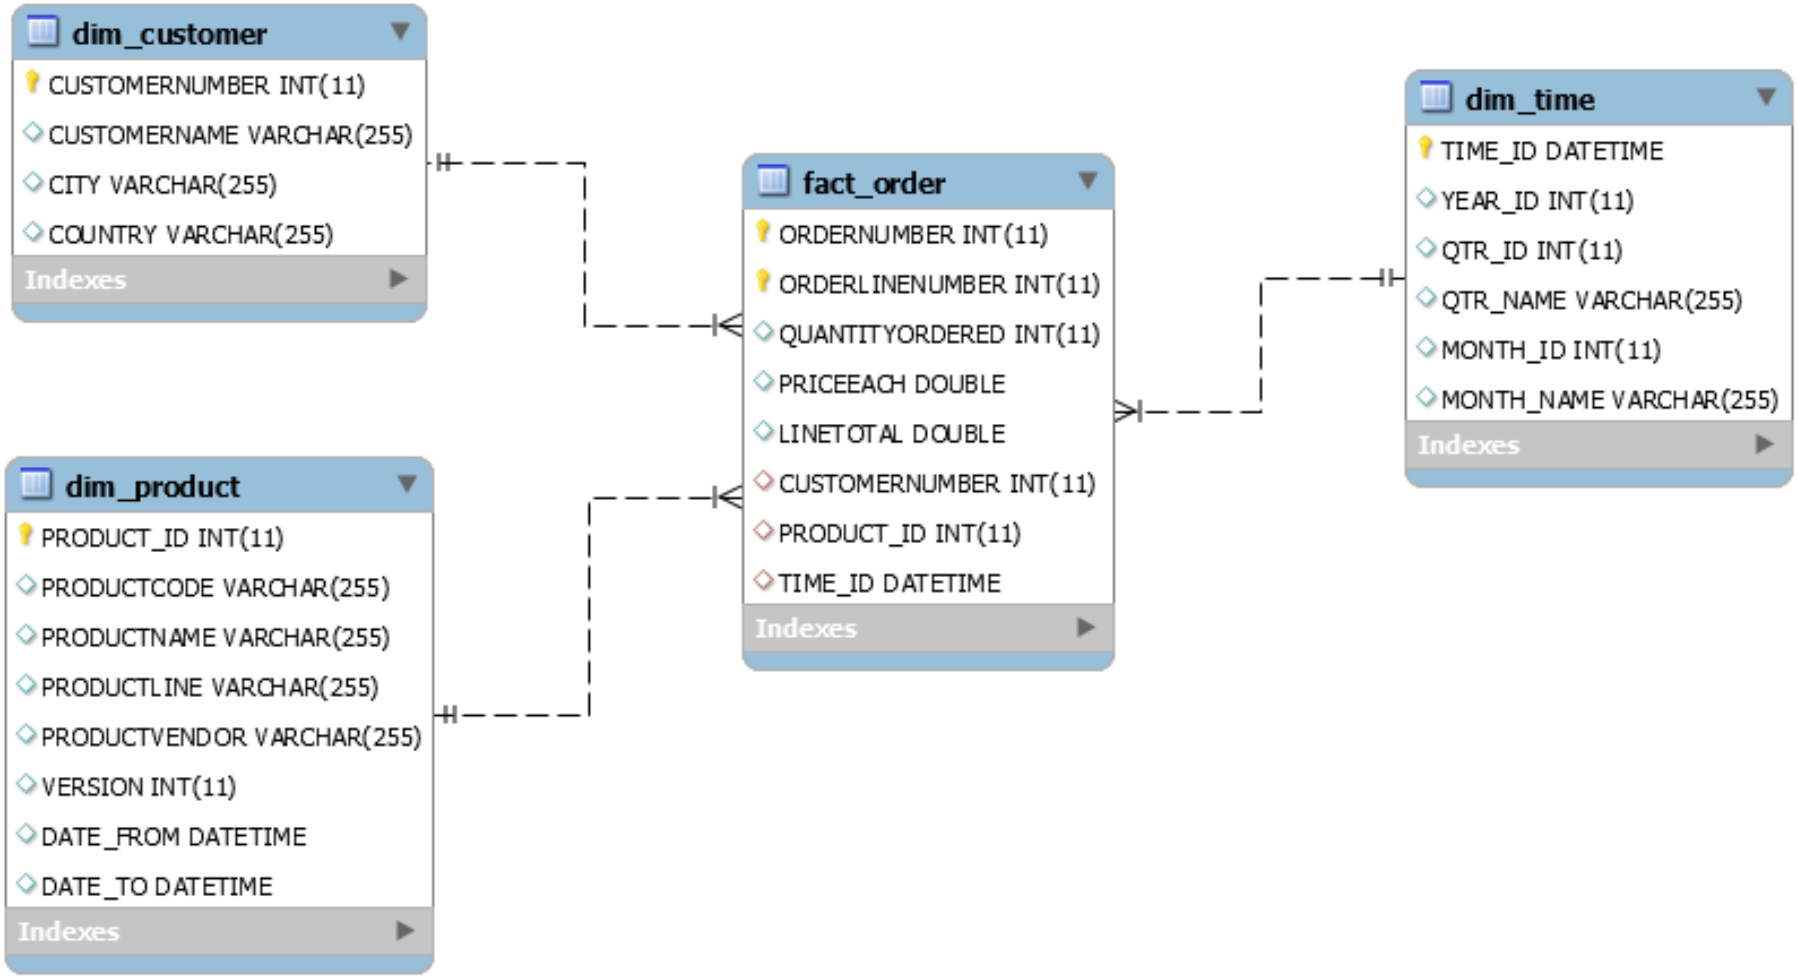

Fact table

- The fact table
  - quantity, unit price, and line total (must be calculated)
  - data comes from orderdetails and orders
    - but product id must come from product dimension (!)

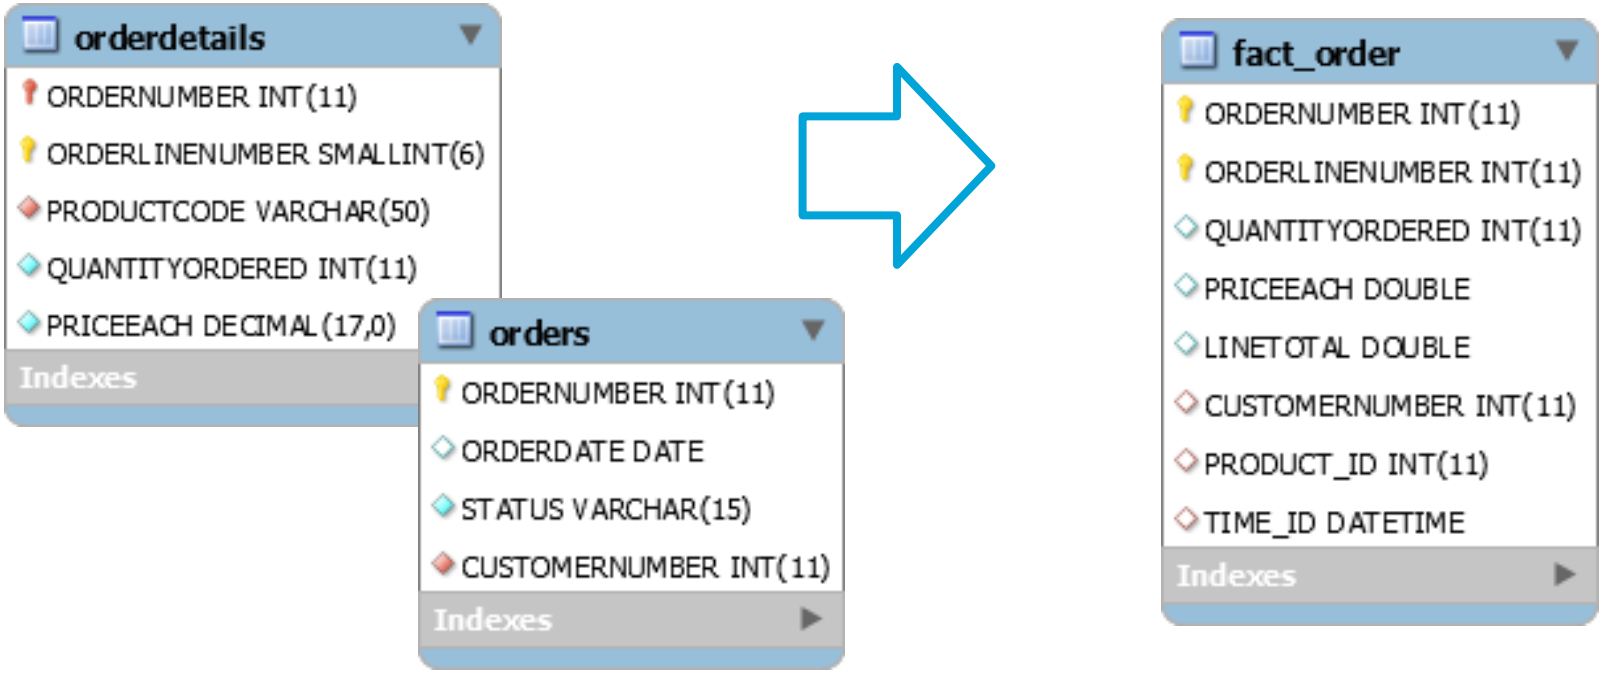

Fact table

- The fact table

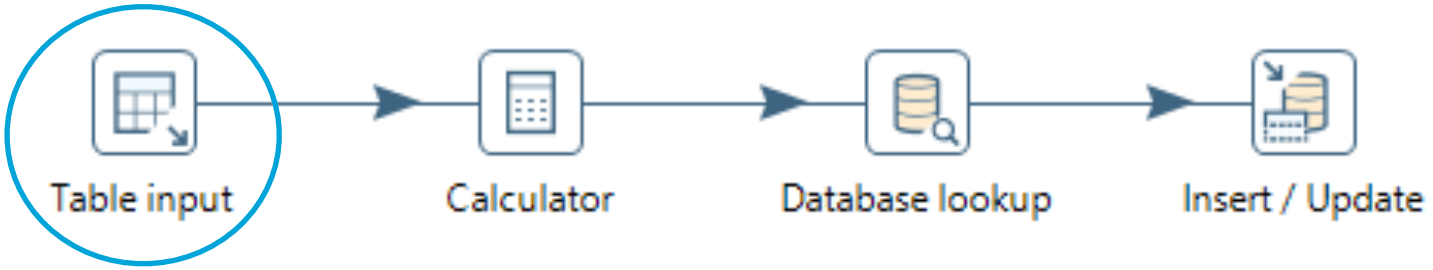

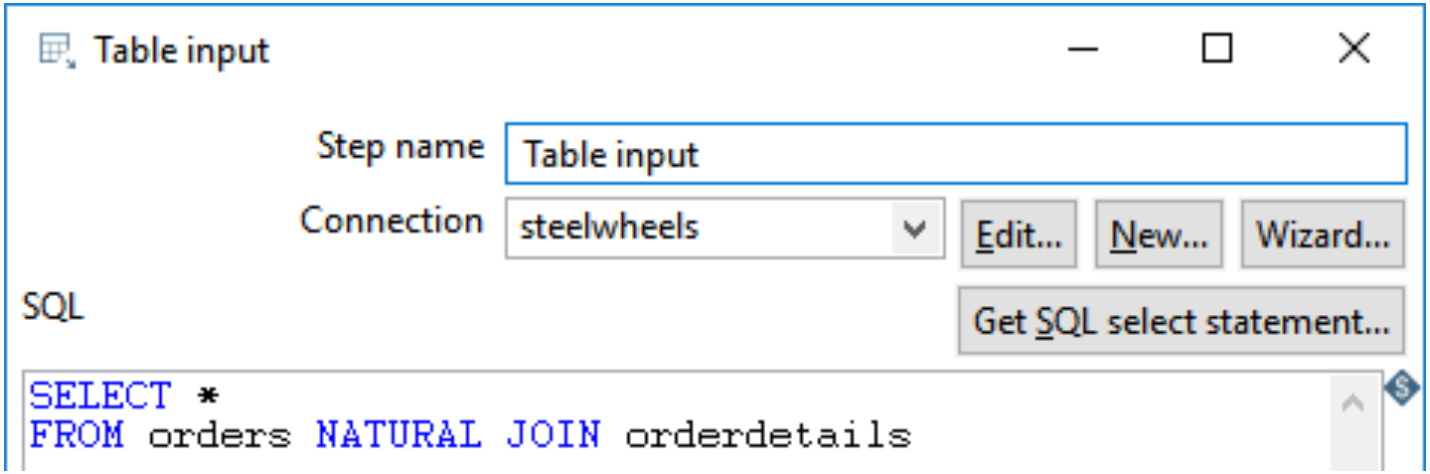

Fact table

- The fact table

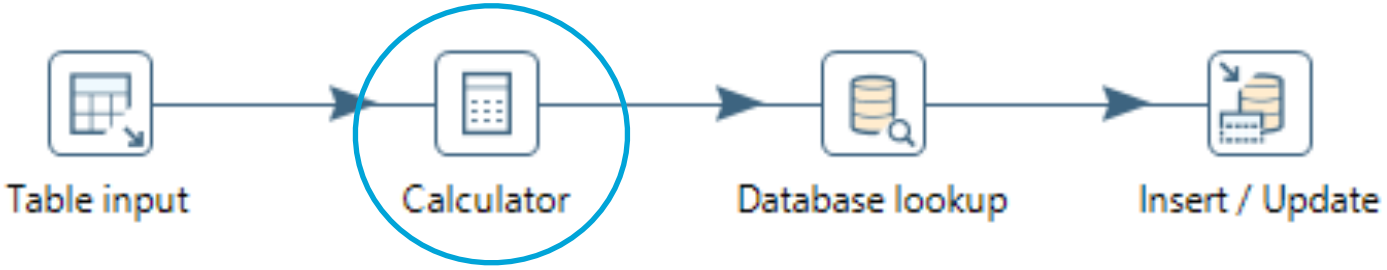

Calculator

Step name: Calculator

Fields:

#	New field	Calculation	Field A	Field B	Field C
1	LINETOTAL	A * B	QUANTITYORDERED	PRICEEACH	

Fact table

- The fact table

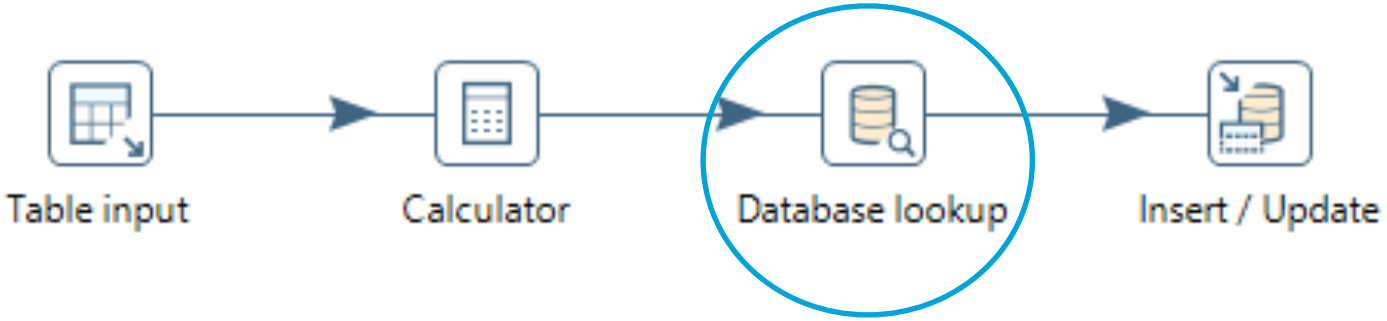

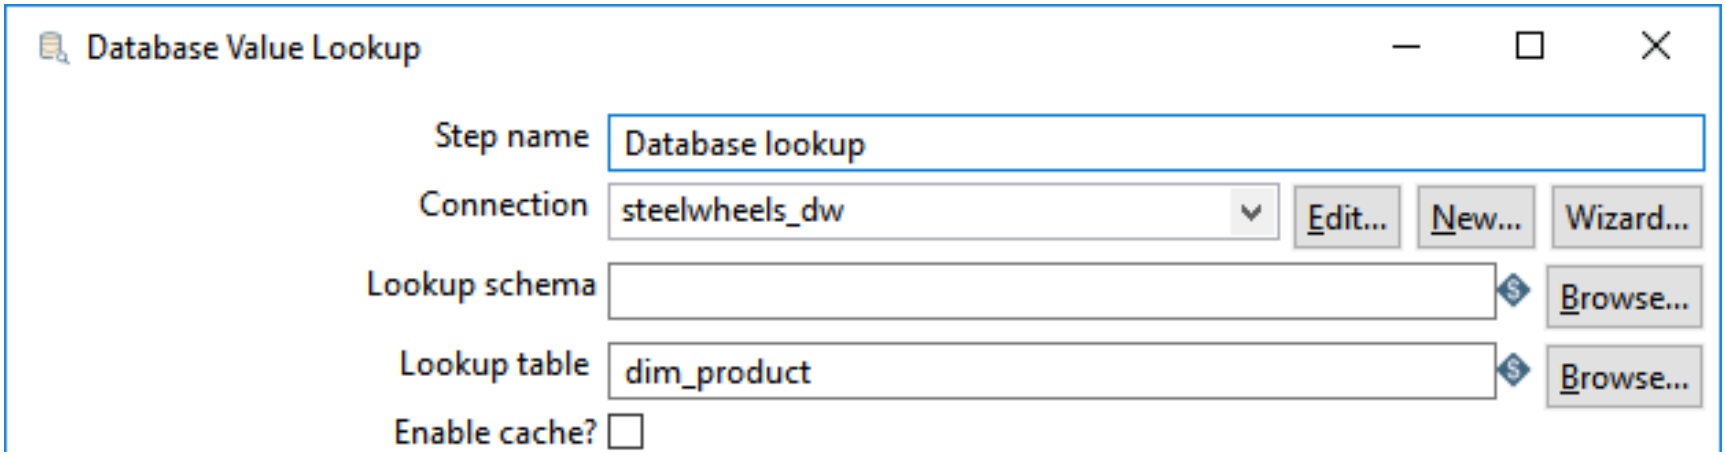

Fact table

- The fact table

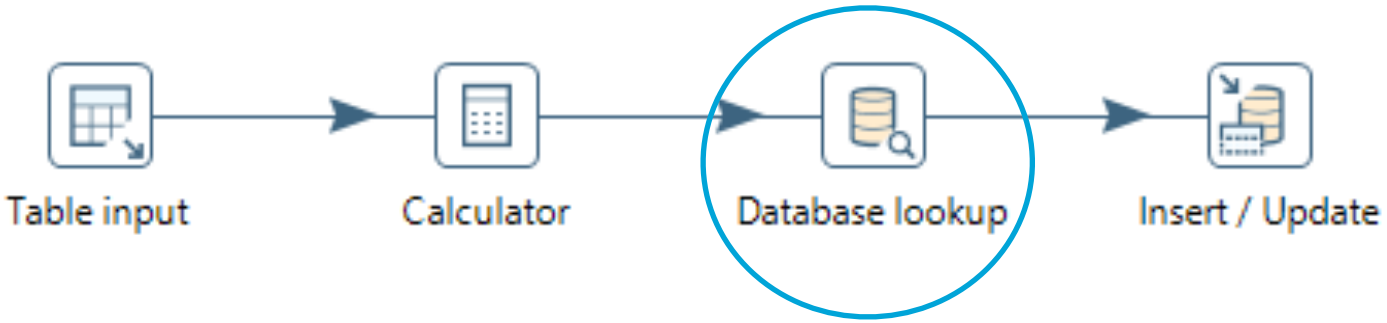

The key(s) to look up the value(s):

#	Table field	Comparator	Field1	Field2
1	PRODUCTCODE	=	PRODUCTCODE	
2	DATE_FROM	<=	ORDERDATE	
3	DATE_TO	>	ORDERDATE	

Fact table

- The fact table

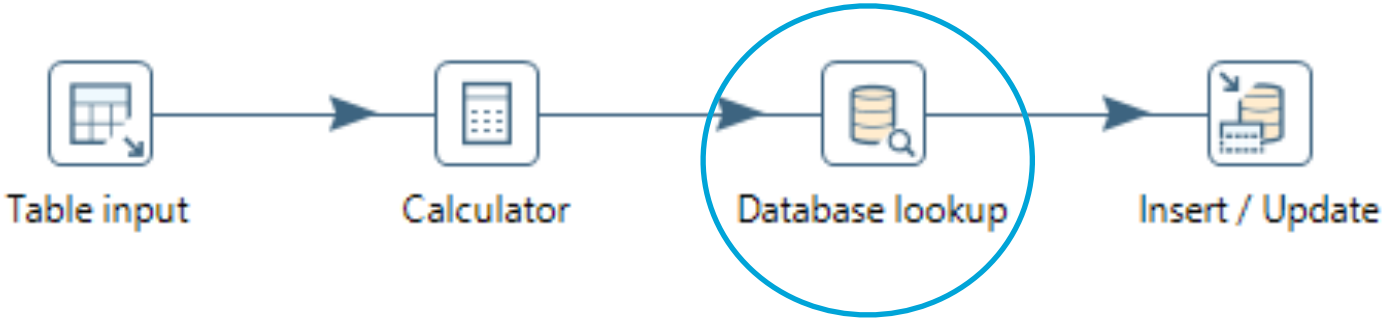

Values to return from the lookup table :

#	Field	New name	Default	Type
1	PRODUCT_ID			Integer

Fact table

- The fact table

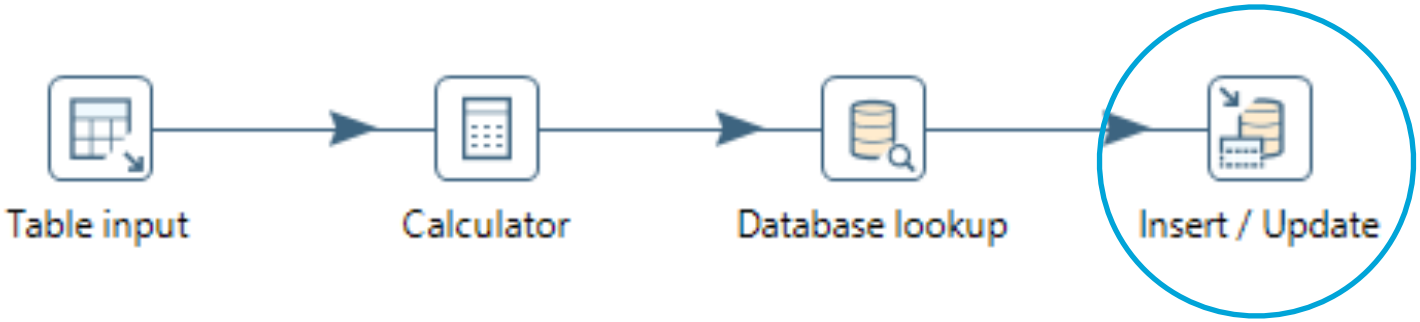

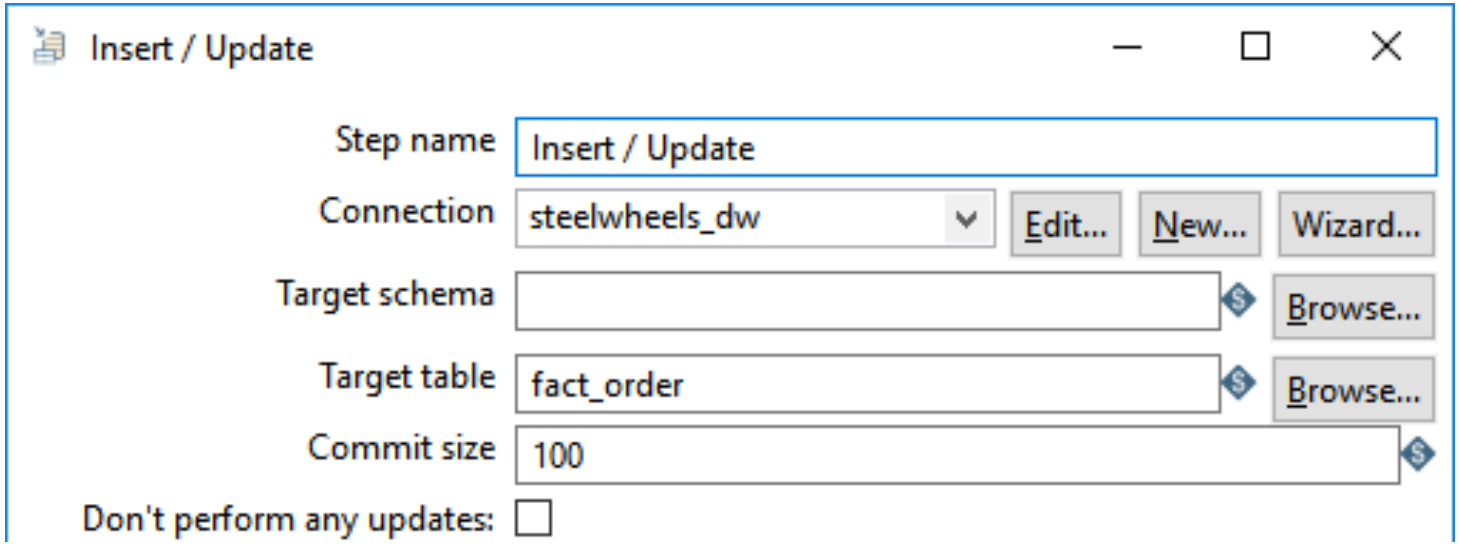

Fact table

- The fact table

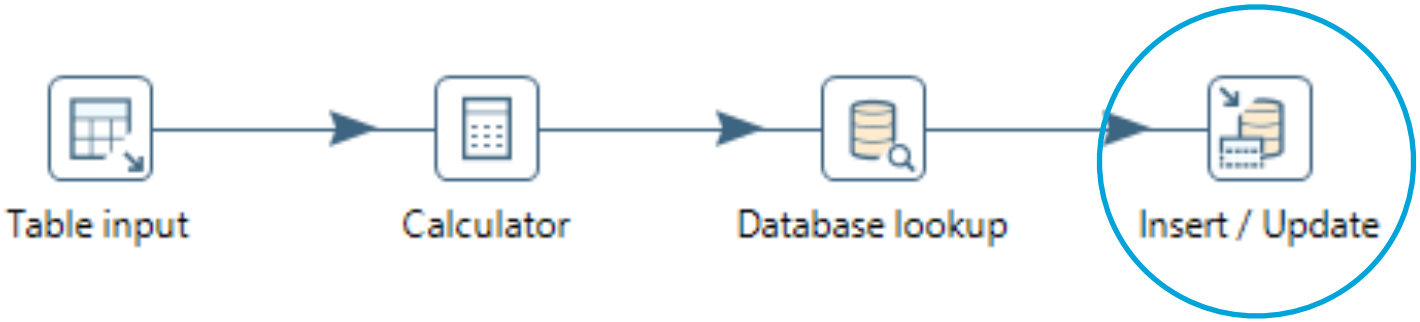

The key(s) to look up the value(s):

#	Table field	Comparator	Stream field1	Str
1	ORDERNUMBER	=	ORDERNUMBER	
2	ORDERLINENUMBER	=	ORDERLINENUMBER	

Get fields

Fact table

- The fact table

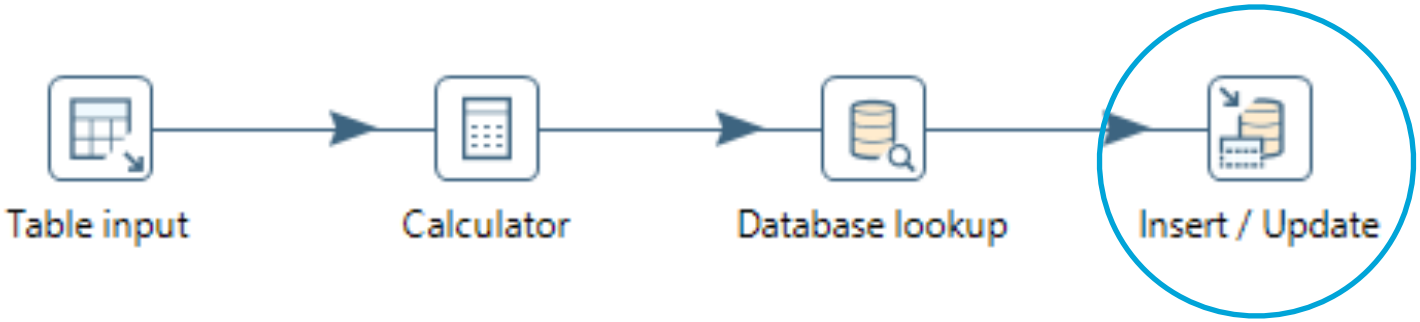

Update fields:

#	Table field	Stream field	Update
1	ORDERNUMBER	ORDERNUMBER	Y
2	ORDERLINENUMBER	ORDERLINENUMBER	Y
3	QUANTITYORDERED	QUANTITYORDERED	Y
4	PRICEEACH	PRICEEACH	Y
5	LINETOTAL	LINETOTAL	Y
6	CUSTOMERNUMBER	CUSTOMERNUMBER	Y
7	PRODUCT_ID	PRODUCT_ID	Y
8	TIME_ID	ORDERDATE	Y

Get update fields

Edit mapping

Fact table

- The fact table

ORDERNUMBER	ORDERLINENUMBER	QUANTITYORDERED	PRICEEACH	LINETOTAL	CUSTOMERNUMBER	PRODUCT_ID	TIME_ID
10100	1	49	34	1666	363	80	2003-01-06 00:00:00
10100	2	50	68	3400	363	27	2003-01-06 00:00:00
10100	3	30	172	5160	363	23	2003-01-06 00:00:00
10100	4	22	87	1914	363	50	2003-01-06 00:00:00
10101	1	26	145	3770	128	33	2003-01-09 00:00:00
10101	2	46	54	2484	128	64	2003-01-09 00:00:00
10101	3	45	31	1395	128	61	2003-01-09 00:00:00
10101	4	25	151	3775	128	29	2003-01-09 00:00:00
10102	1	41	50	2050	181	20	2003-01-10 00:00:00
10102	2	39	123	4797	181	19	2003-01-10 00:00:00
10103	1	36	102	3672	121	65	2003-01-29 00:00:00
10103	2	22	54	1188	121	30	2003-01-29 00:00:00
10103	3	31	104	3224	121	85	2003-01-29 00:00:00
10103	4	42	129	5418	121	6	2003-01-29 00:00:00
10103	5	36	117	4212	121	52	2003-01-29 00:00:00
10103	6	42	106	4452	121	103	2003-01-29 00:00:00
10103	7	45	76	3420	121	90	2003-01-29 00:00:00
10103	8	27	126	3402	121	9	2003-01-29 00:00:00
10103	9	41	47	1927	121	53	2003-01-29 00:00:00
10103	10	35	112	3920	121	17	2003-01-29 00:00:00

Complete ETL process

- Transformations
  - the customer dimension
  - the product dimension
  - the time dimension
  - the fact table

Complete ETL process

- Defining a job
  - A job is a sequence of transformations
    - each transformation runs only upon successful completion of the previous one
  - The complete ETL process for the data warehouse
    - can be run multiple times to update the data warehouse