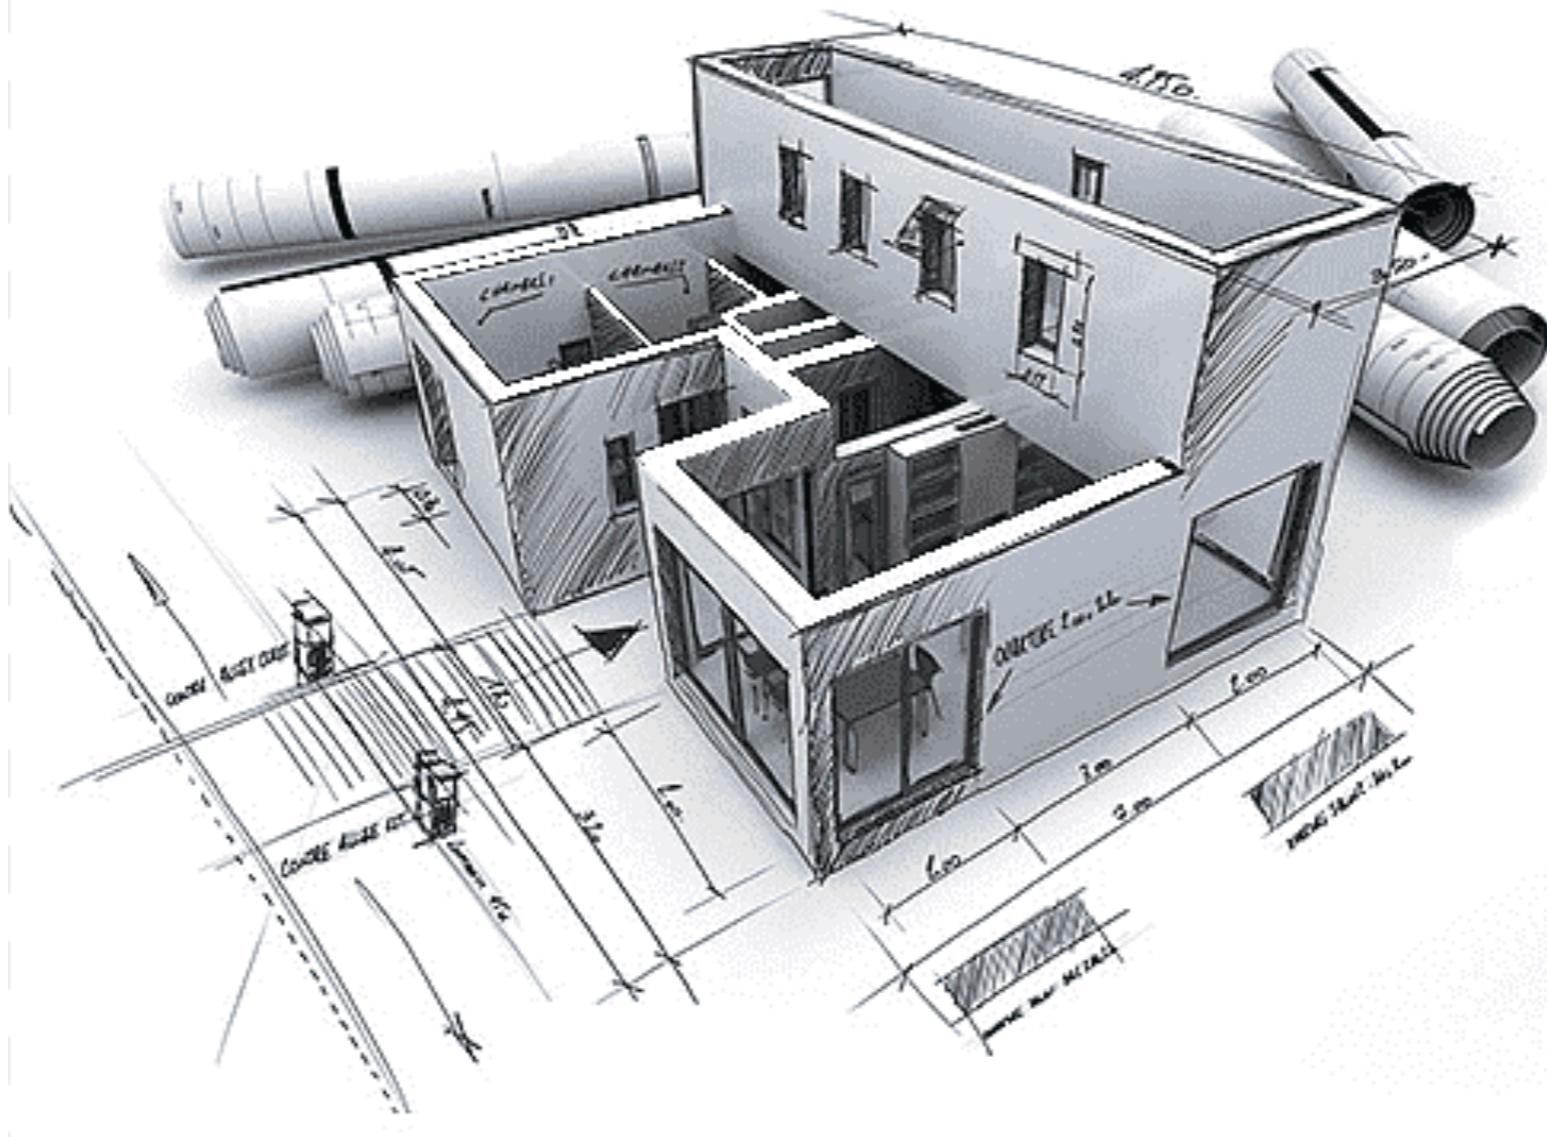


# Planning, Learning and Intelligent Decision Making

## Lecture 2

PADInt 2023



# MC examples (modeling)

# Example: PageRank

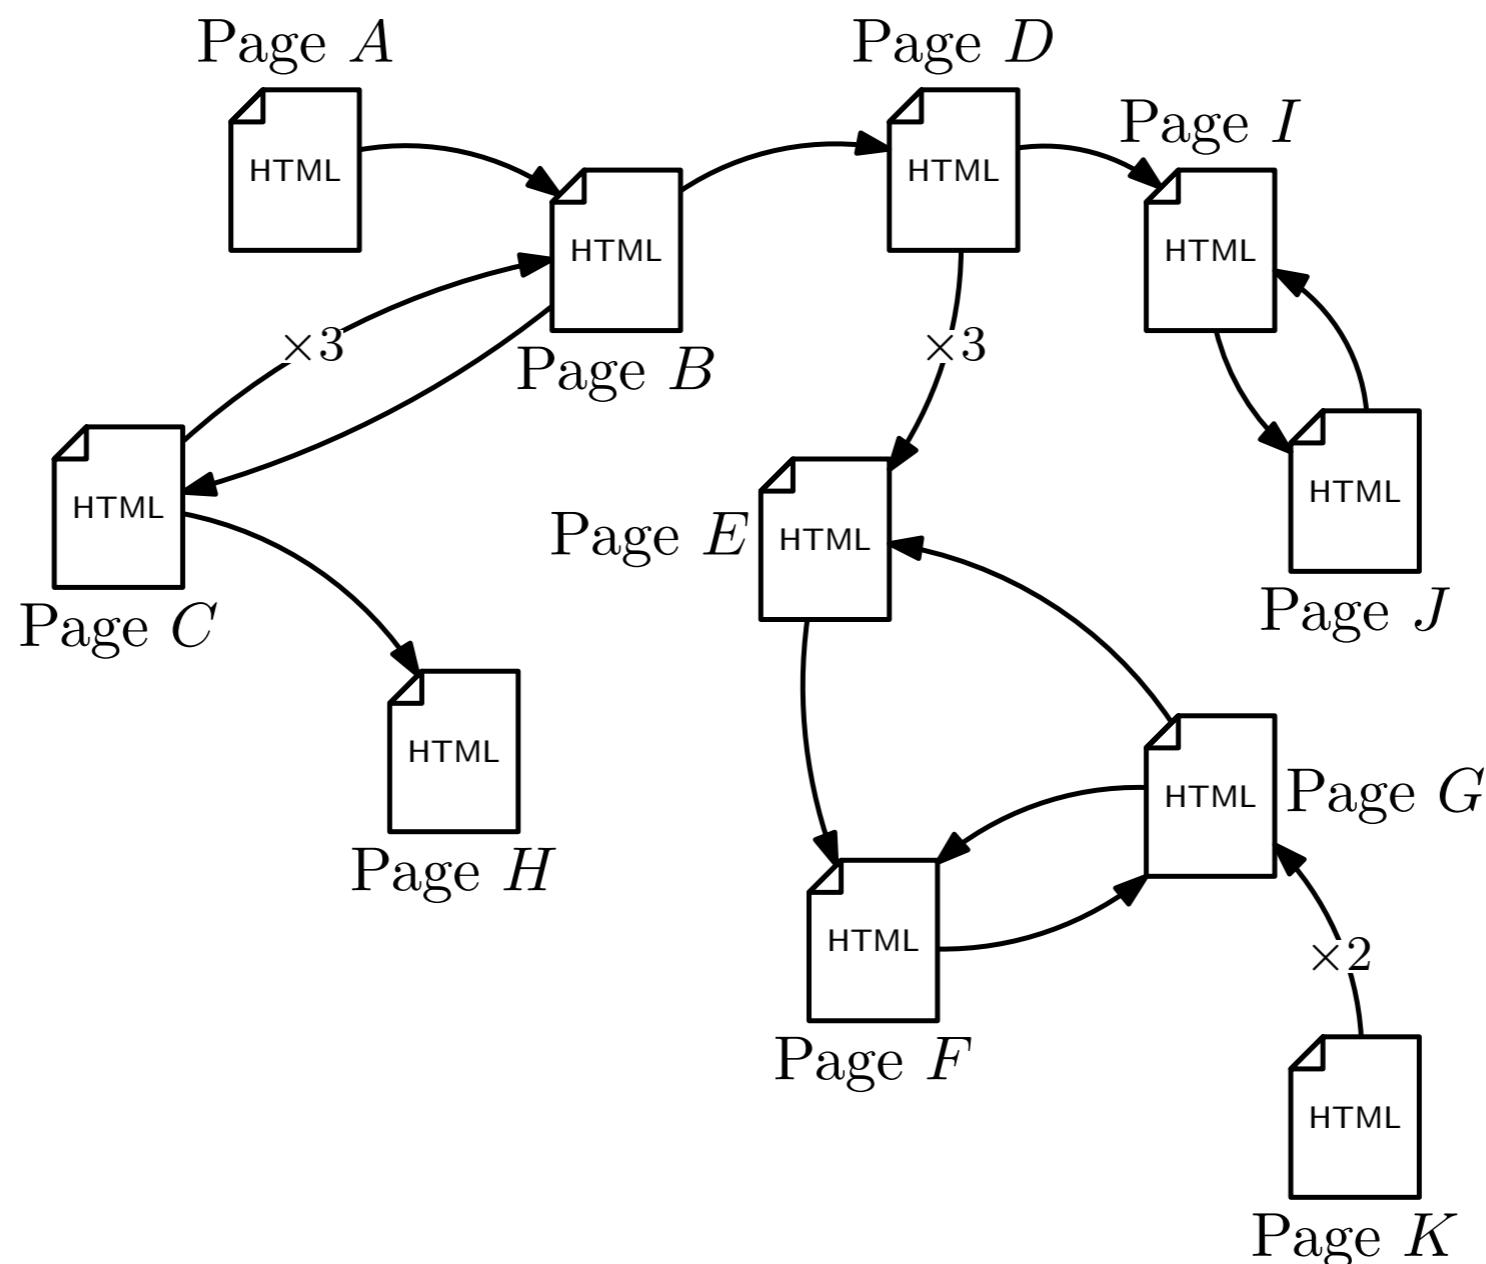
- PageRank is the algorithm used by Google to rank a set of connected documents
- It simulates a “random bot” navigating the web of retrieved documents

# Example: PageRank

- Upon visiting a page, the bot randomly moves to one of the pages linked by the current one
- The rank corresponds to the “amount of time” that the bot spends on each page

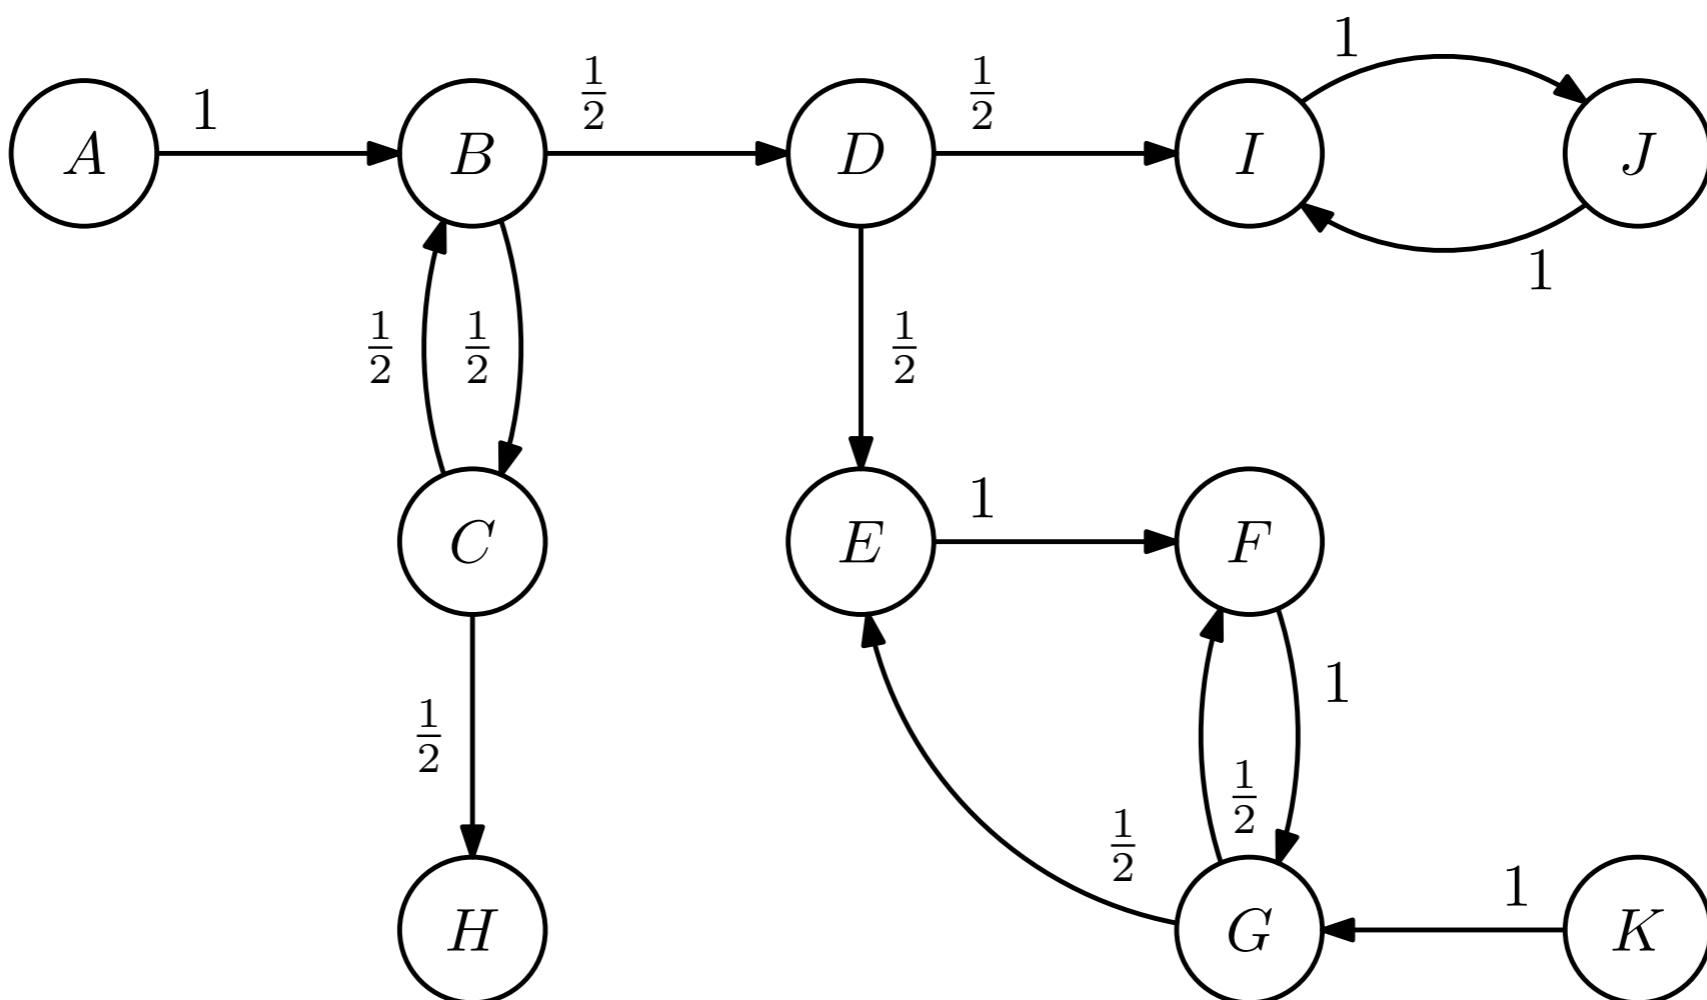
# Example: PageRank

- Given the web of documents:



# Example: PageRank

- States:  $\mathcal{X} = \{A, B, C, D, E, F, G, H, I, J, K\}$
- Transitions:





# Stability of MCs

# What is stability?

- It is often important to understand how a MC behaves in the long run
- Stability concerns the **long-term behavior** of the chain

# What is stability?

- We may want to know:
  - Depending on where the chain starts, can it **reach** any other state?
  - Is the behavior of the chain **cyclic**?
  - How **frequently** does the chain visit each state?

# Irreducibility

- A state  $y$  can be reached from a state  $x$  if

$$P^t(y \mid x) > 0$$

for some  $t$



Positive probability of  
visiting  $y$  after visiting  $x$

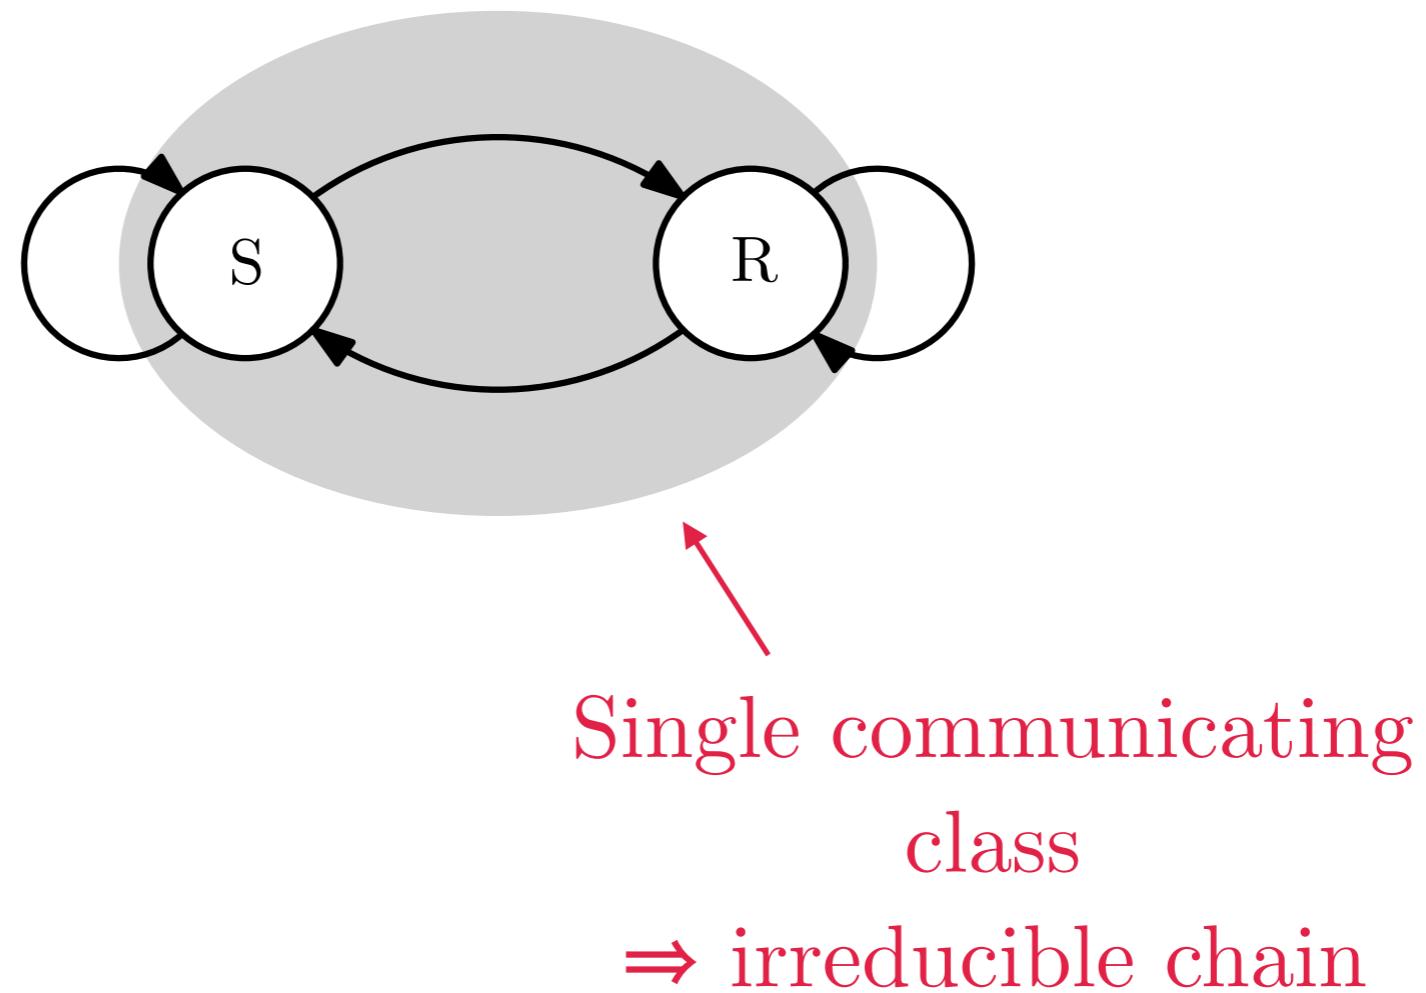
- A chain is **irreducible** if any state  $y$  can be reached from any other state  $x$

# Irreducibility

- We can split the state space of a chain in **communicating classes** (sets of states that are mutually reachable)

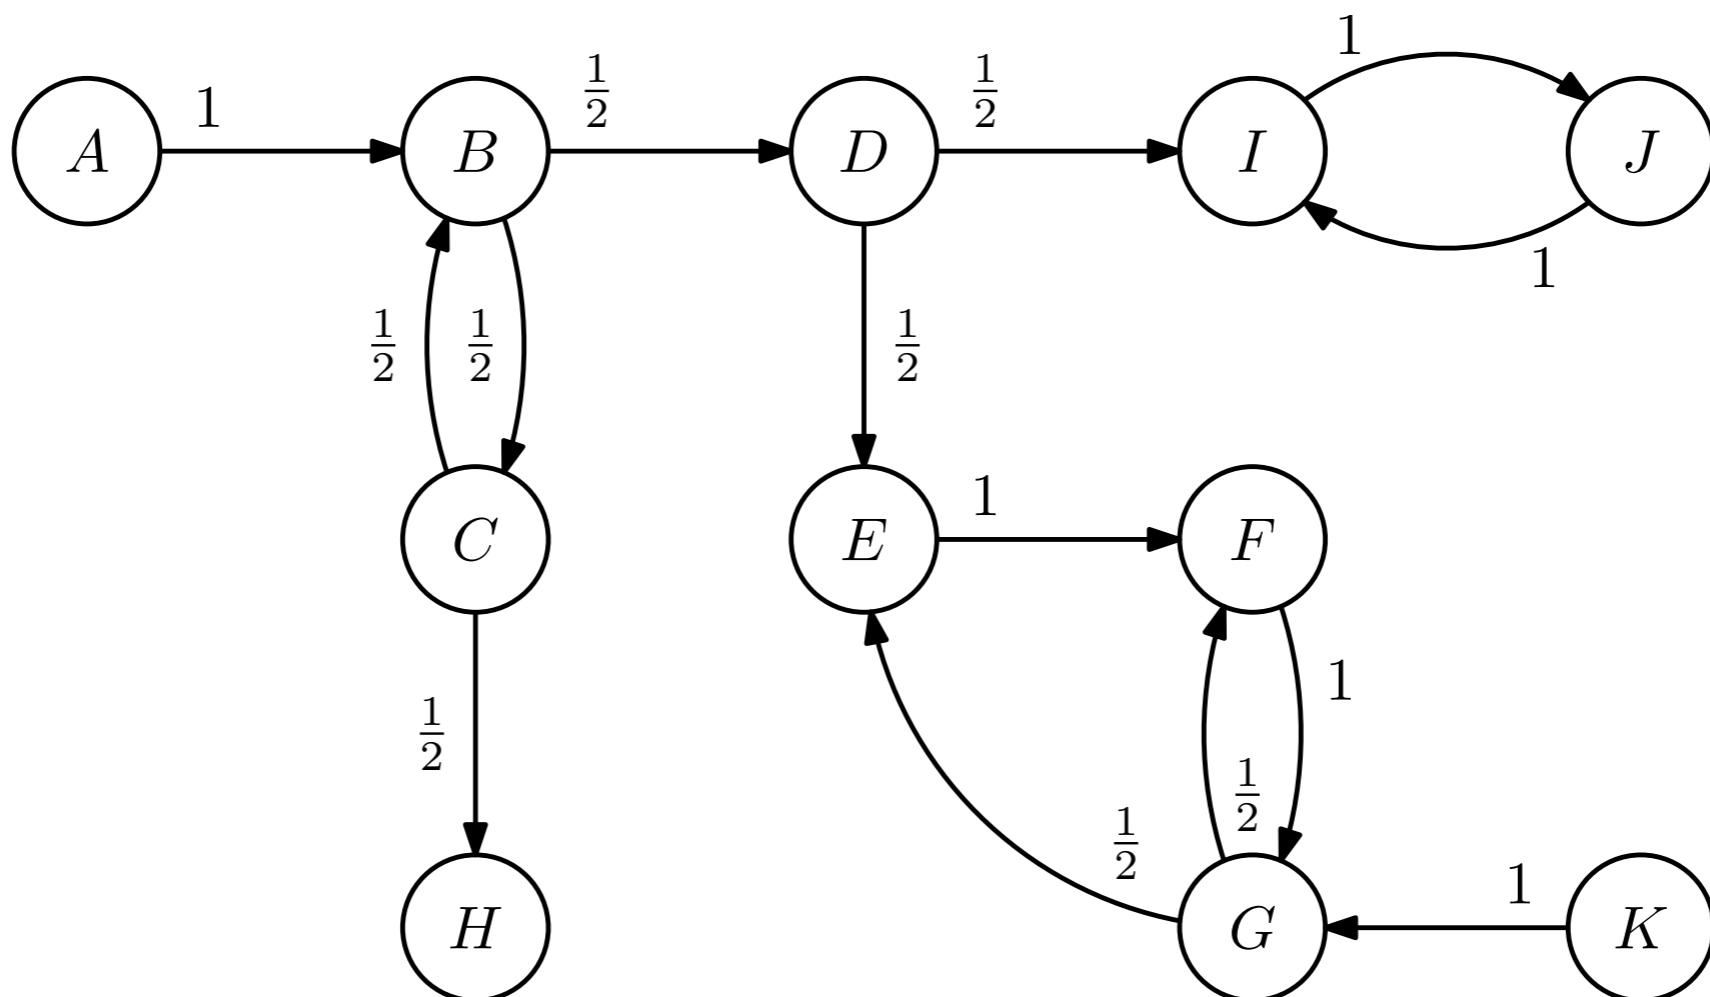
# Irreducibility

- In the weather example:



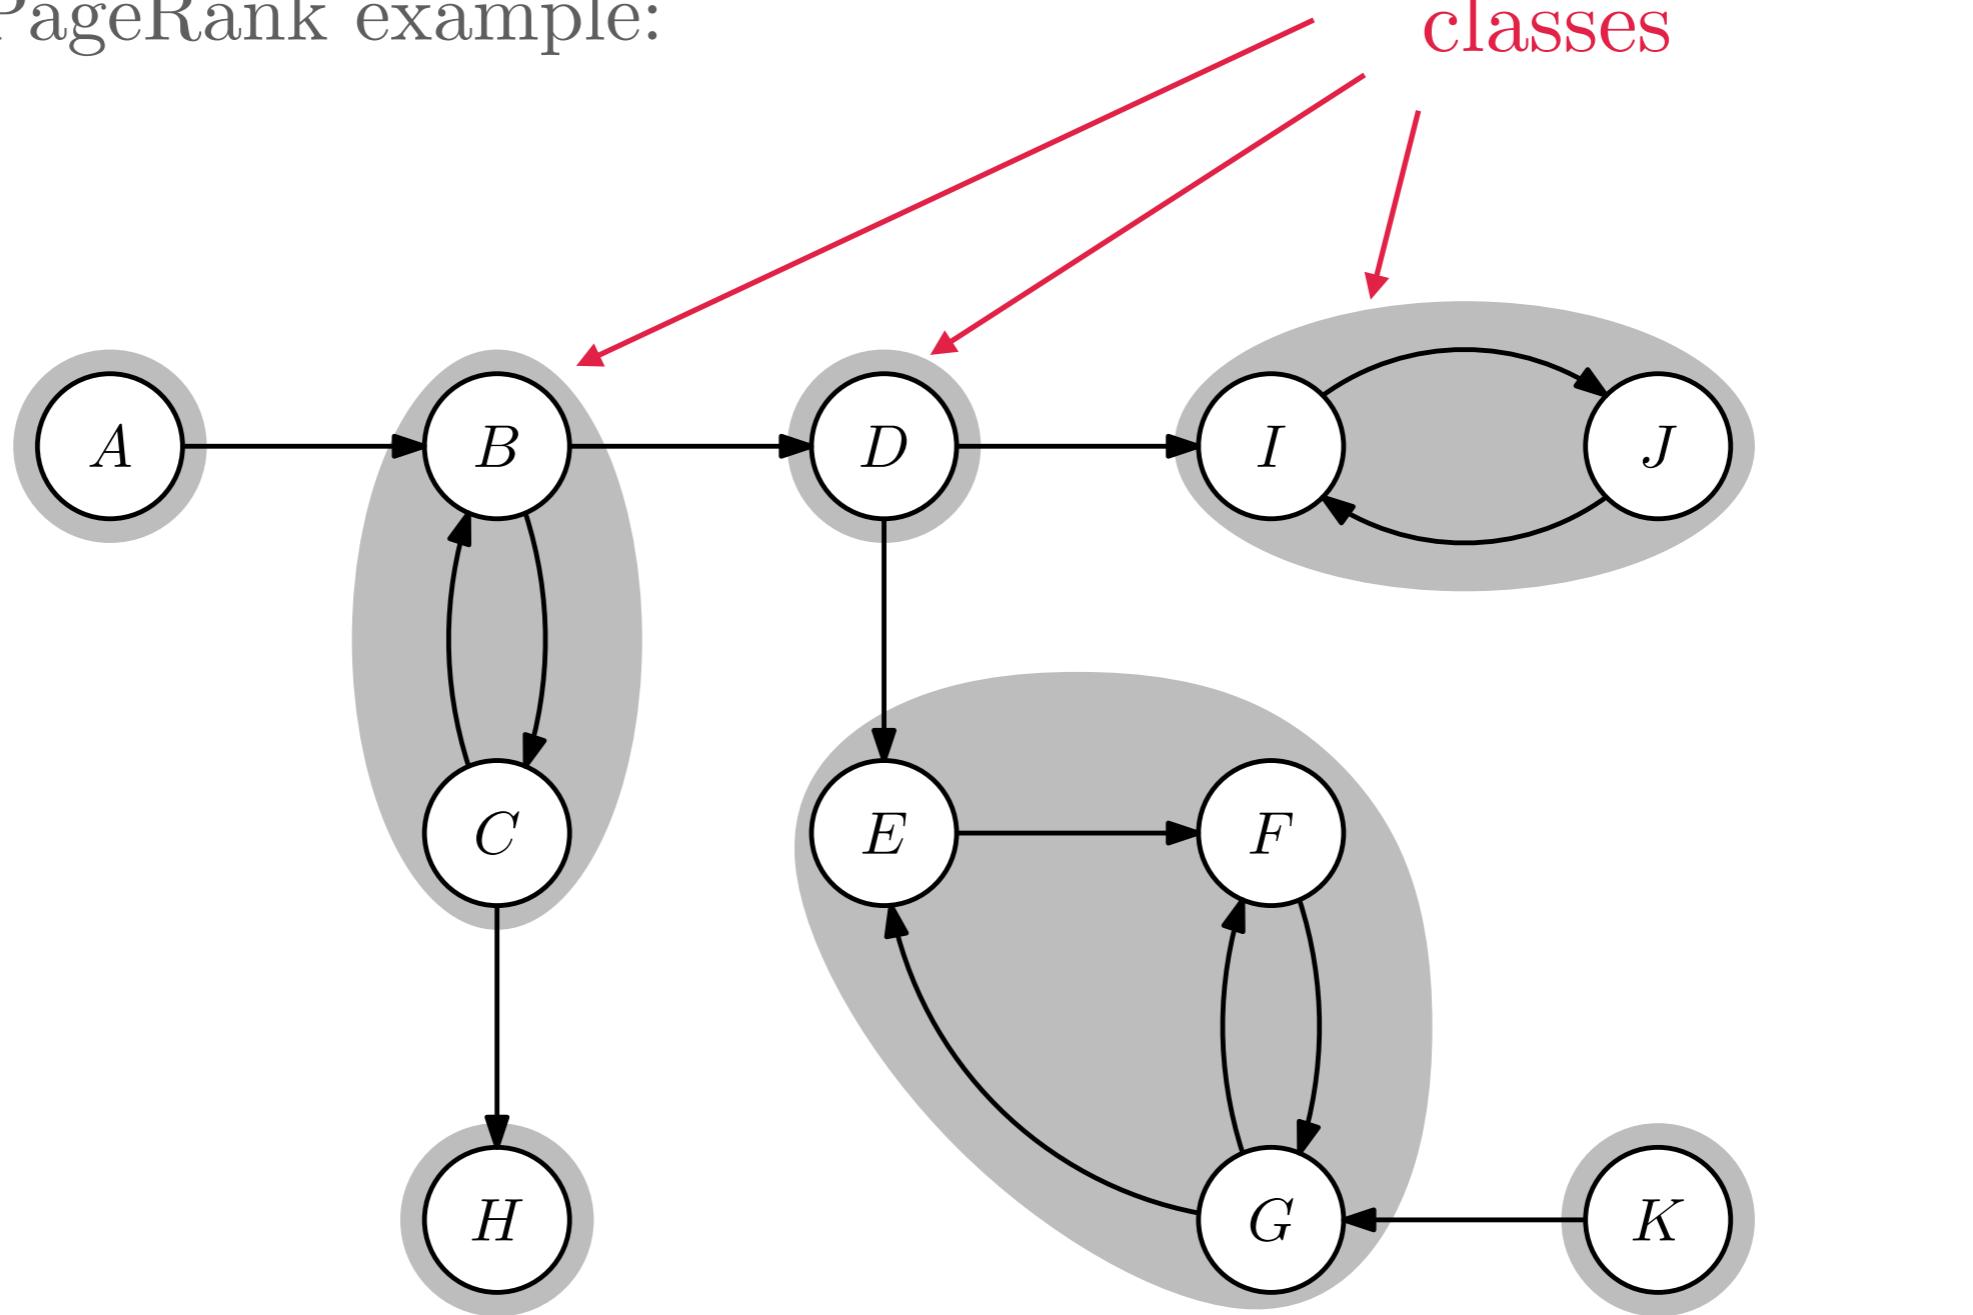
# Irreducibility

- In the PageRank example:



# Irreducibility

- In the PageRank example:



# Aperiodicity

- The period of a state  $x$  is...
  - the **greatest common divider**...
  - of all time steps in which  $x$  can be visited...
  - if the chain departs from  $x$

# Aperiodicity

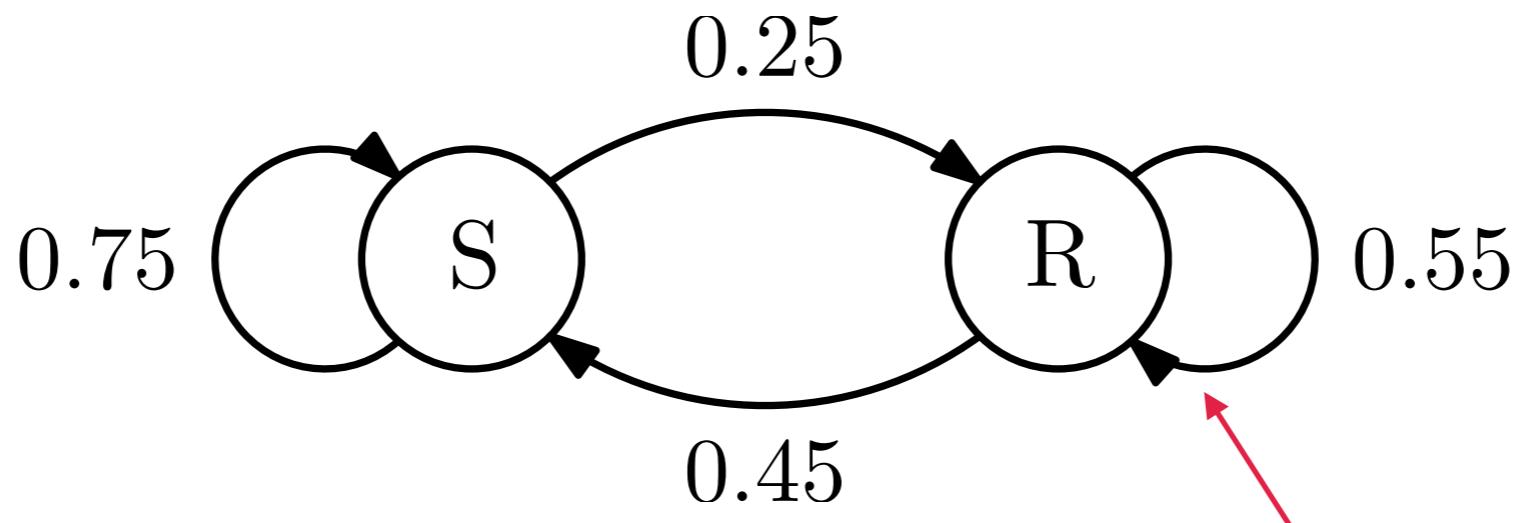
- Formally, the **period of  $x$**  is the number

$$d_x = \gcd \{t \in \mathbb{N} \mid \mathbf{P}^t(x \mid x) > 0, t > 0\}$$

- A state  $x$  is **aperiodic** if  $d_x = 1$
- A chain is **aperiodic** if all states are aperiodic, and **periodic** otherwise

# Aperiodicity

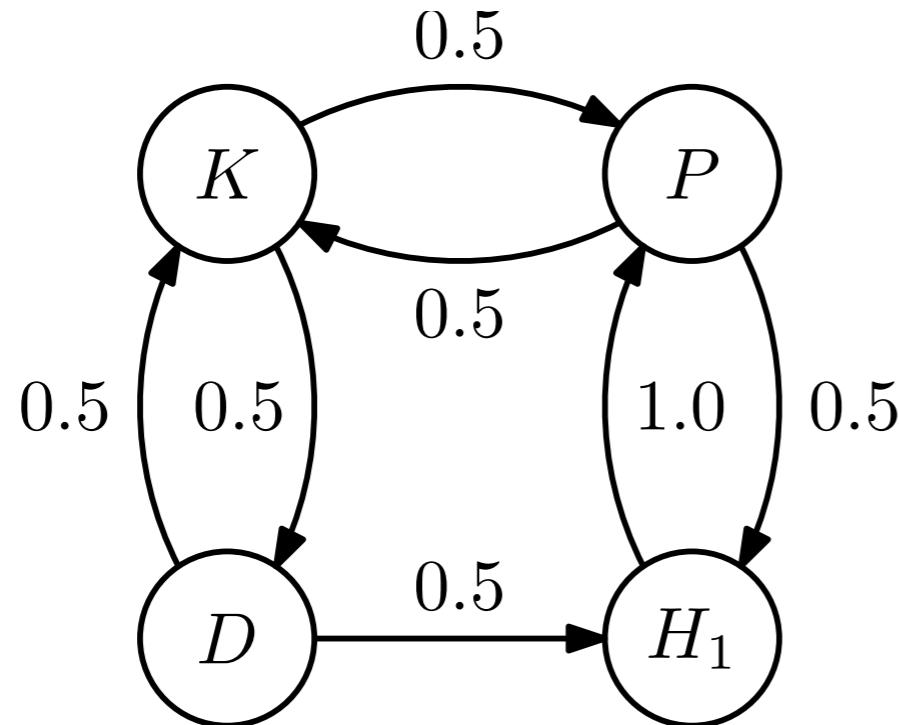
- In the weather example:



$P(R | R) > 0$   
 $\Rightarrow$  aperiodic state

# Aperiodicity

- In the robot example:



If the chain departs from  $P$ , when does it return?



Several possibilities:  
 $t = 2$  (goes to  $H_1$  and returns)  
 $t = 2$  (goes to  $K$  and returns)  
 $t = 4$  (goes to  $K, D$  and returns)

...

The period is  $d = 2$

# Stationary distribution

- Let  $\mu$  be a distribution over  $\mathcal{X}$
- In practical terms,  $\mu$  is a row vector

$$\mu = [\mu(x_1) \quad \mu(x_2) \quad \dots \quad \mu(x_{|\mathcal{X}|})]$$

where

Component  $\mu(x)$  is the probability of  $x$  according to  $\mu$

$$\sum_{x \in \mathcal{X}} \mu(x) = 1$$

# Stationary distribution

- $\mu$  can represent, for example,
  - ... the initial distribution for the chain;
  - ... the predicted distribution after  $t$  steps;
  - ... etc.

# Stationary distribution

- The distribution  $\mu$  is **stationary** if

$$\mu(x) = \sum_{y \in \mathcal{X}} \mu(y) \mathbf{P}(x \mid y)$$

- In other words, if the state at time  $t$  follows the distribution  $\mu$  and  $\mu$  is stationary, then the state at time  $t + 1$  also follows the distribution  $\mu$
- The stationary distribution corresponds to **stable behavior of the chain**

# Key stability results

An irreducible and aperiodic Markov chain **possesses a stationary distribution.**

Positive chain

For an irreducible and aperiodic Markov chain with stationary distribution  $\mu^*$ ,

$$\lim_{t \rightarrow \infty} \mu_0 \mathbf{P}^t = \mu^*$$

for any initial distribution  $\mu_0$ .

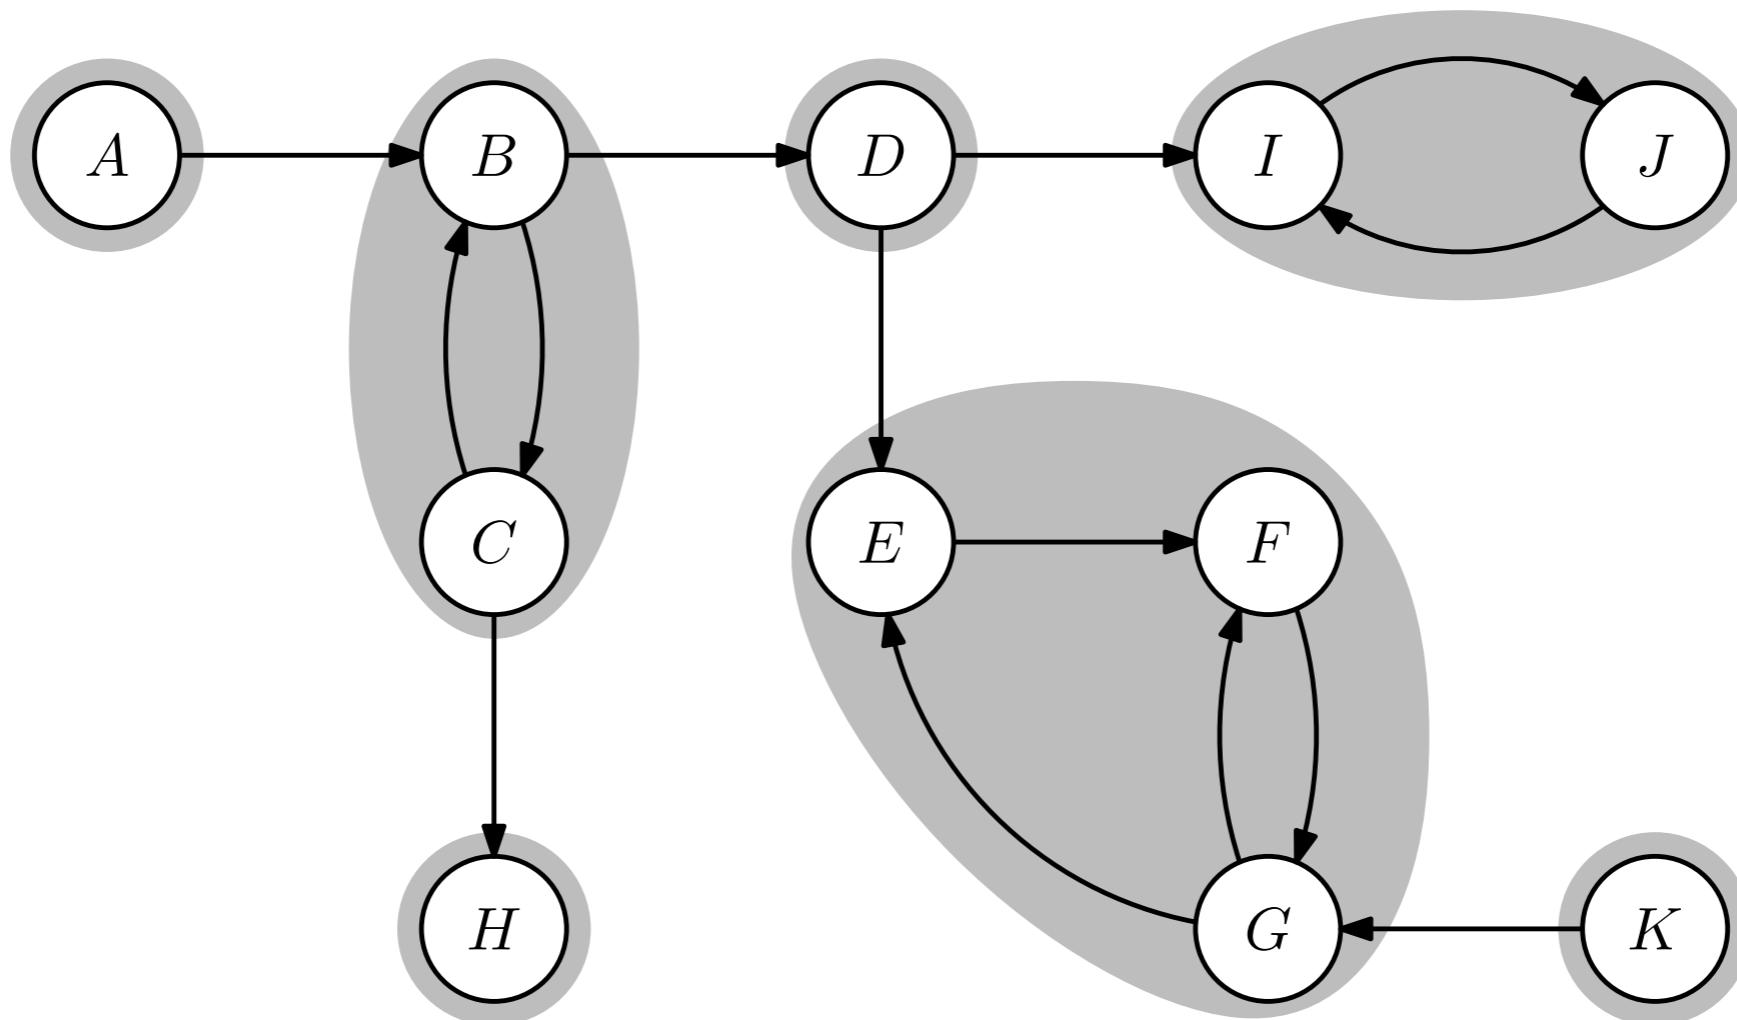
Ergodic chain

# Summarizing...

- A positive chain possesses a stationary distribution  $\mu$ :
  - If  $x_t$  is distributed according to  $\mu$ , then so is  $x_{t+1}$
- An ergodic chain eventually reaches the stationary distribution

# Returning to PageRank

- The chain is **not** irreducible

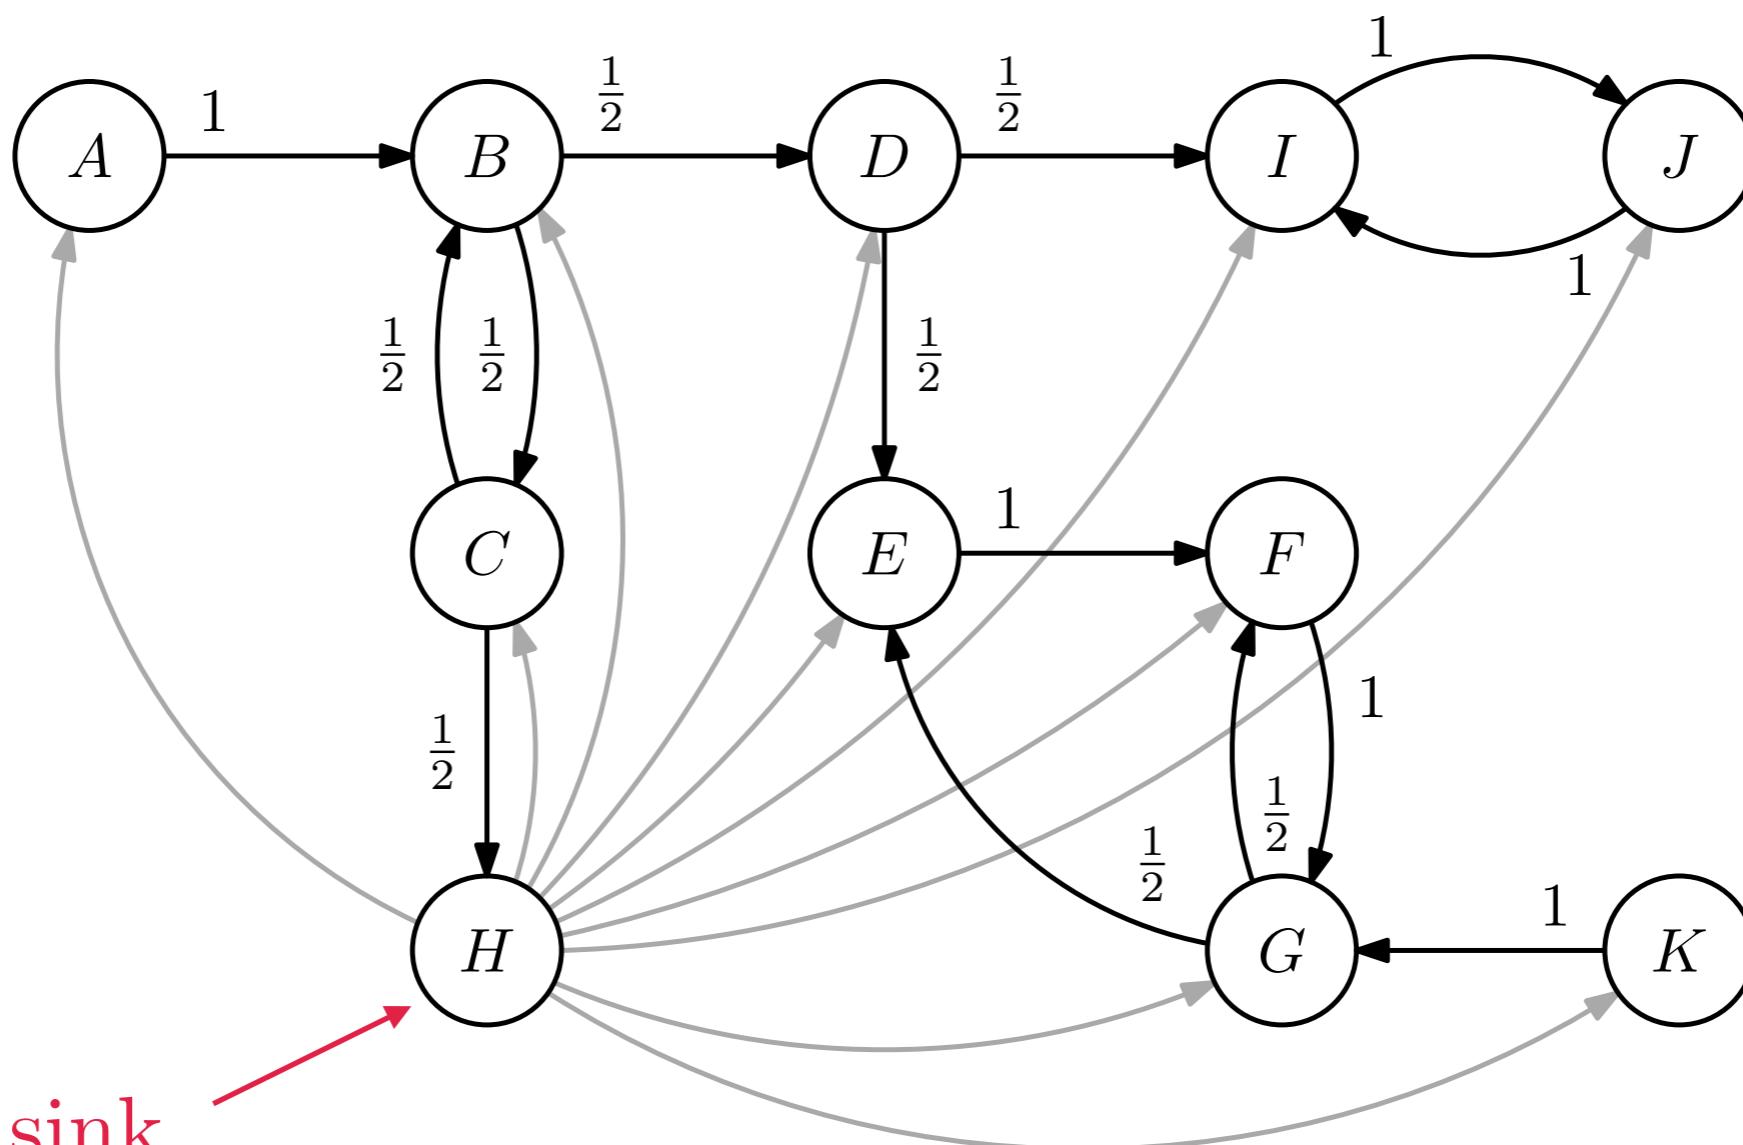


# Returning to PageRank

- PageRank introduces two modifications to the chain:
  - Sinks link to all other nodes
  - There is a probability  $1 - \gamma$  of “teleporting” to an arbitrary node

# Returning to PageRank

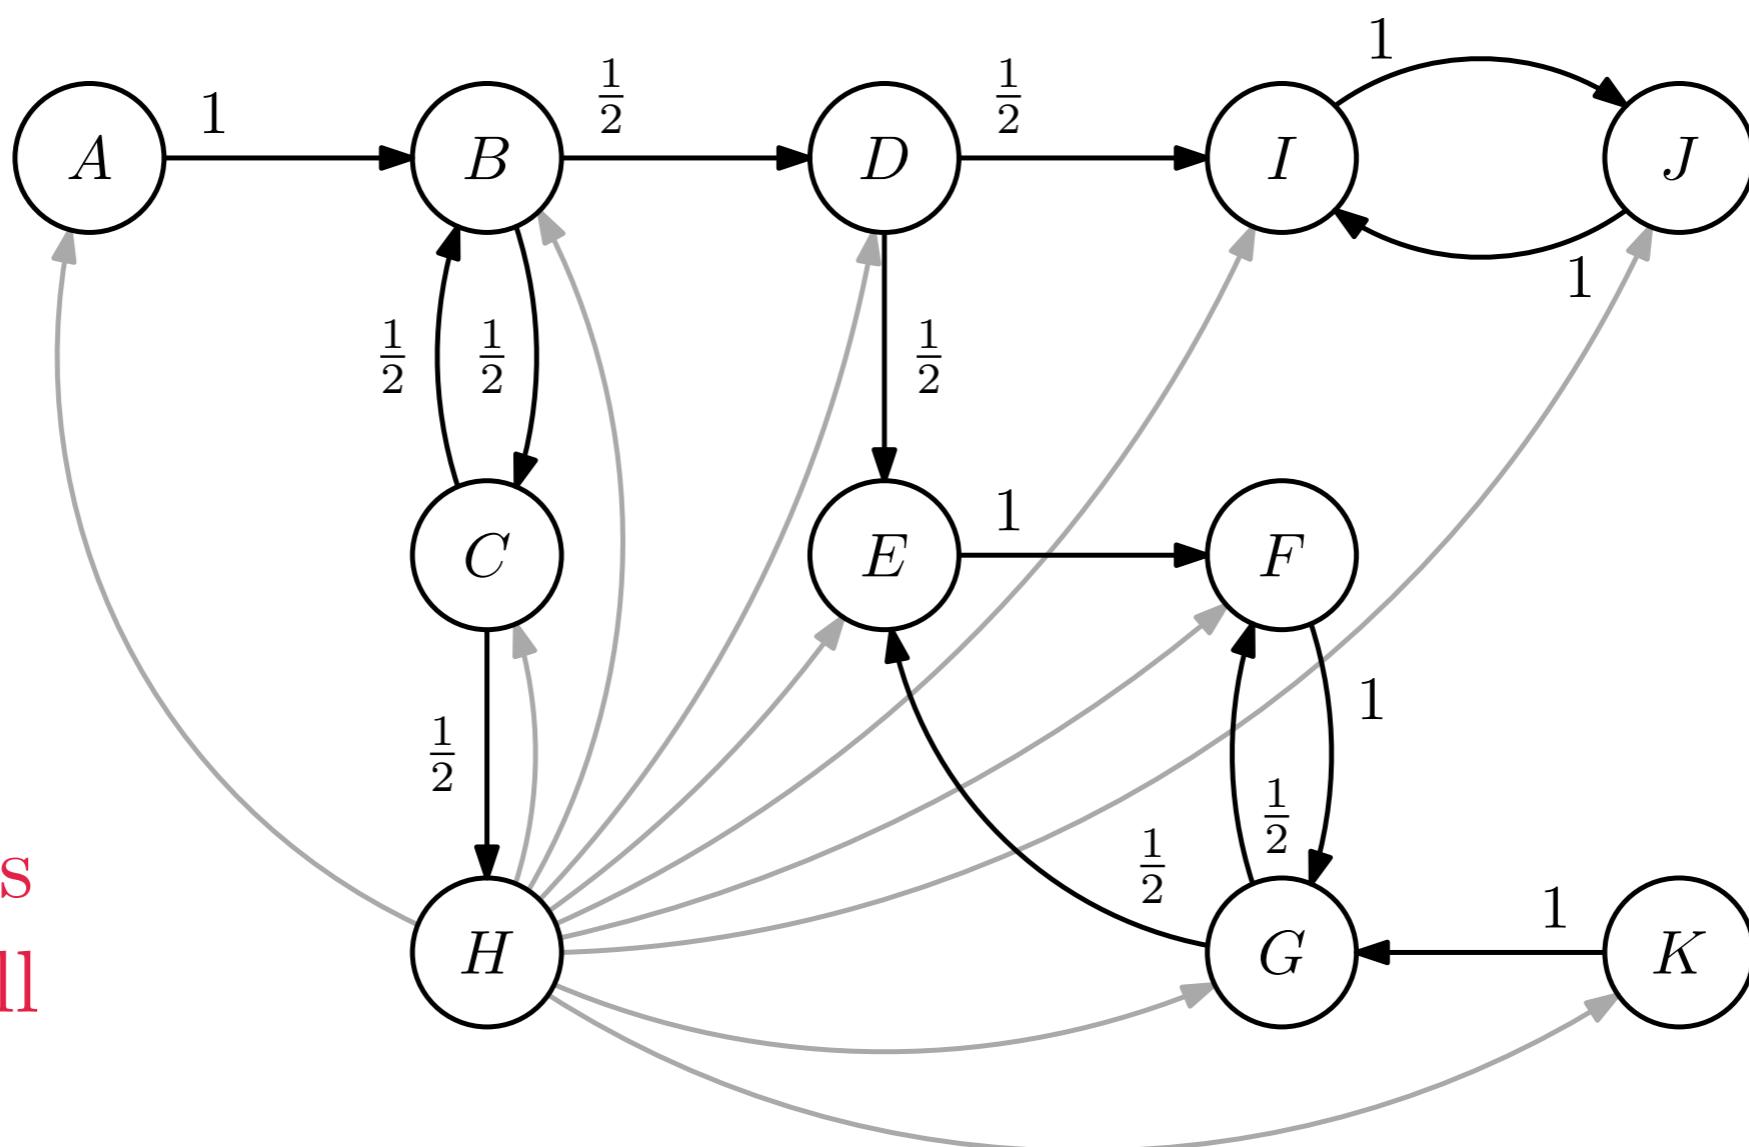
- In our case:



*H is a sink*

# Returning to PageRank

- In our case:



All nodes  
link to all  
nodes

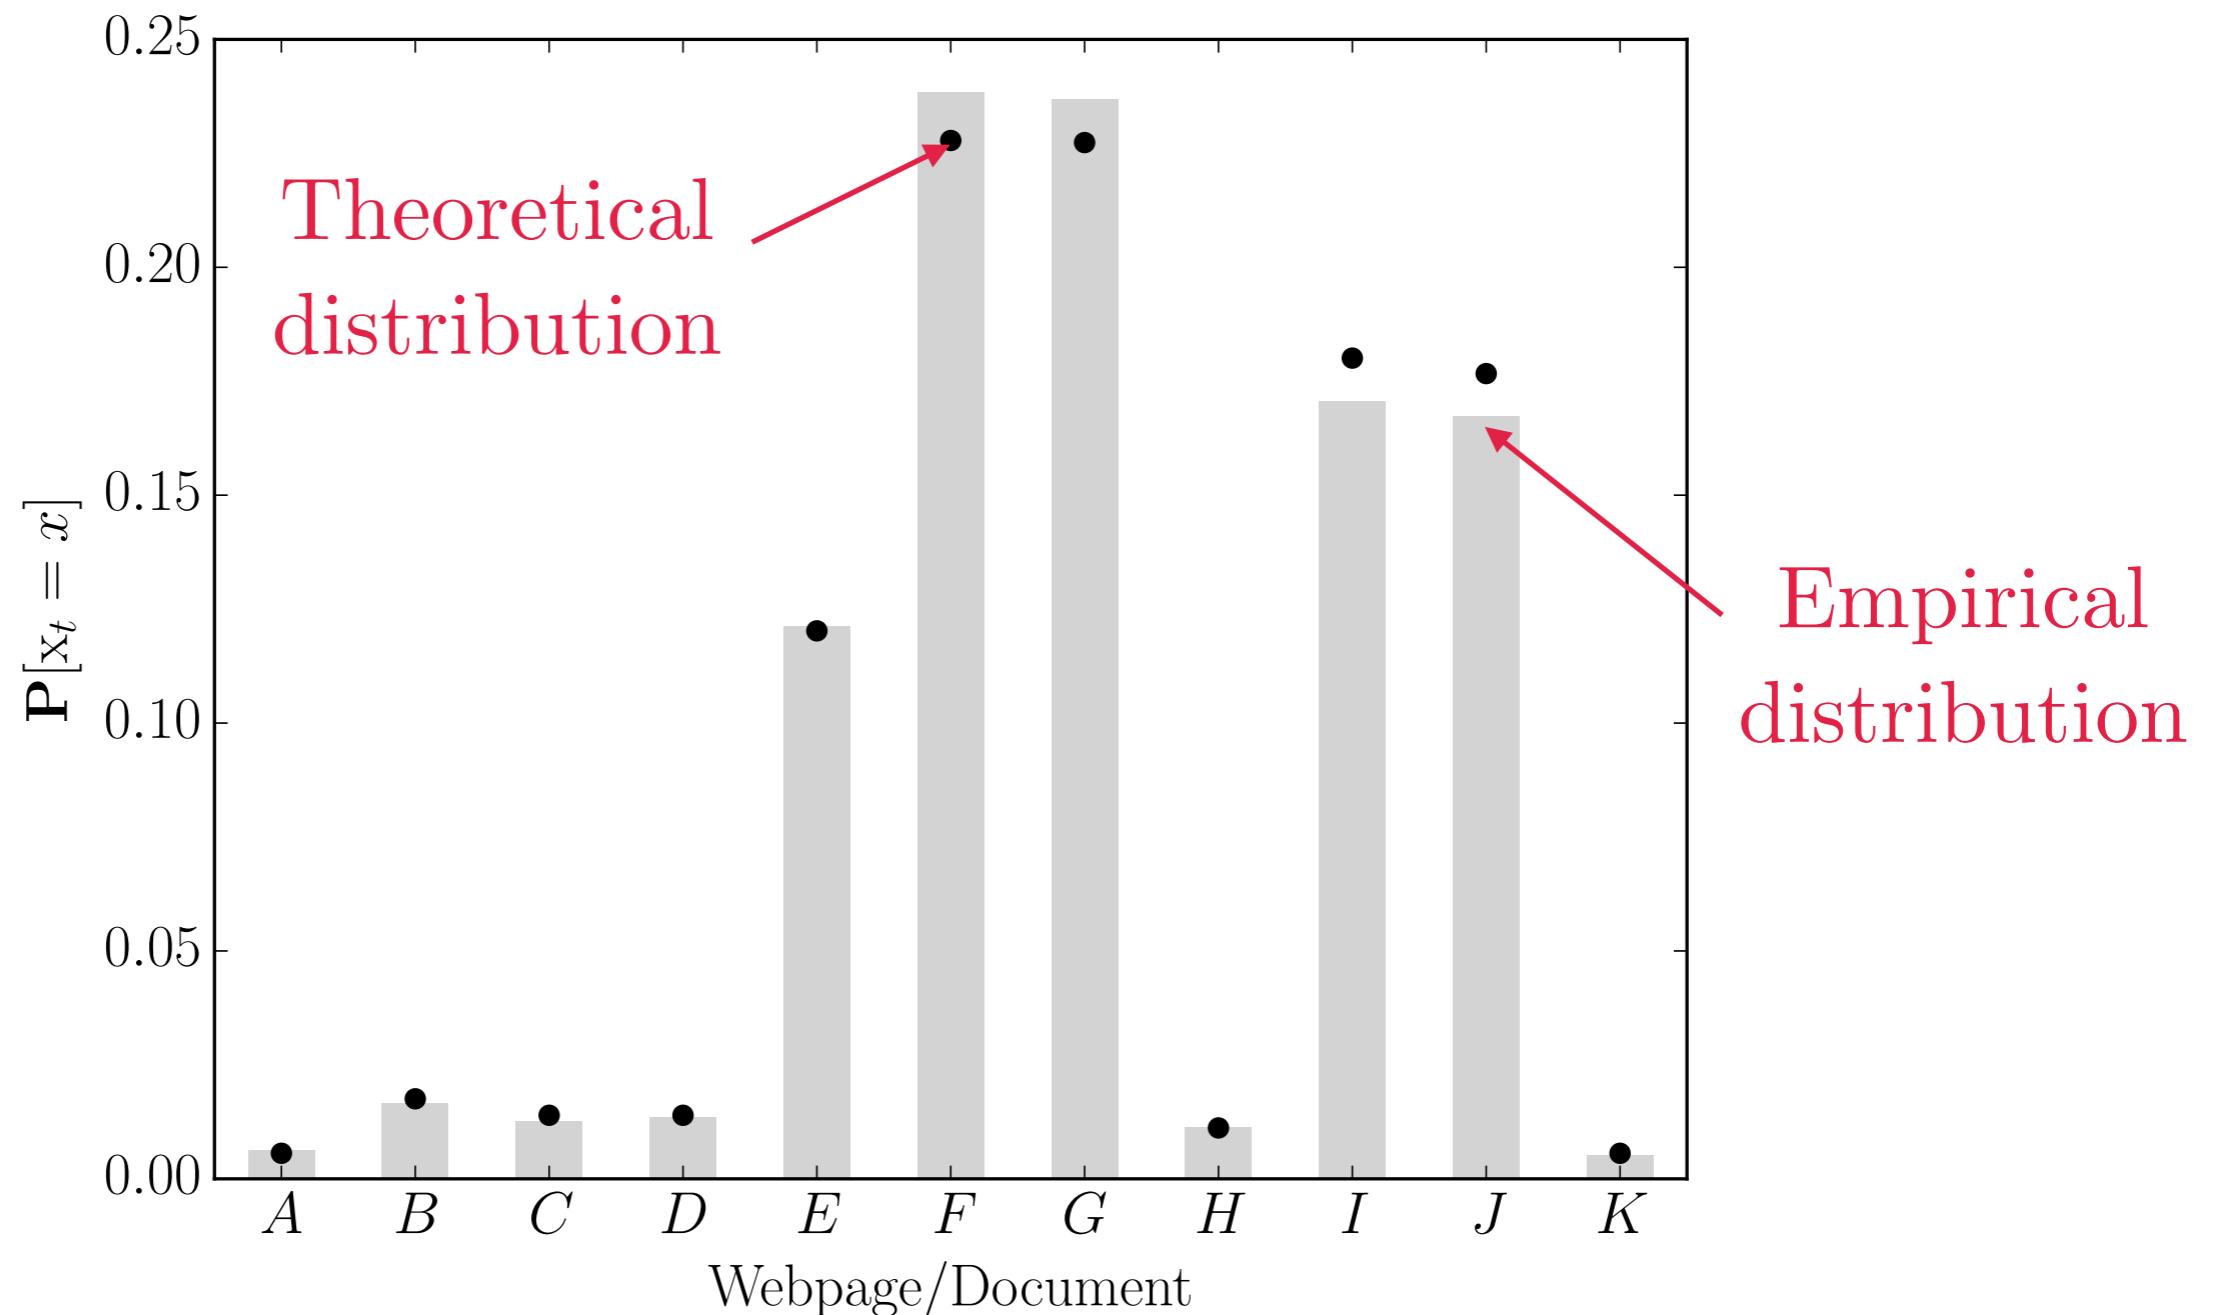
# Returning to PageRank

- Is the chain irreducible?
  - Yes
- Is the chain aperiodic?
  - Yes
- ... then there is a stationary distribution

# Returning to PageRank

- The stationary distribution can be computed analytically (efficiently)
- It provides a convenient criterion for ranking the pages

# Returning to PageRank



# Sequential models (cont.)

The image shows two ancient Greek vases, specifically kylikes, side-by-side against a plain white background. Both vases are made of terracotta and feature a dark brown, almost black, glaze on their necks and shoulders. The main body of the vases is a light tan or beige color. They are decorated with a repeating pattern of concentric circles, which are rendered in a dark brown or black ink. There are four such circles on each vase, arranged in two rows of two. The vase on the left has a single vertical crack near its base. The vase on the right has a small tear at the very bottom edge. Both vases have small, curved handles on either side of their bases.

# The urn problem

# The urn problem

- An oracle has available two urns
  - Each urn is filled with black and white balls
  - Urn 1: 25% white balls; 75% black balls
  - Urn 2: 25% black balls; 75% white balls

# The urn problem

- At each step, the oracle draws a ball from one of the urns
- The ball is put back afterwards

# The urn problem

- The current urn determines which urn the next ball will come from
  - The next ball will come from the same urn with 80% probability
  - The next ball will come from a different urn with 20% probability

# 1. Is this a Markov chain?

- Yes
- The urn-ball in the next step depends on the urn-ball in the current step

## 2. What are the states?

- All possible urn-ball pairs:
  - $\mathcal{X} = \{(1, w), (1, b), (2, w), (2, b)\}$

### 3. Transition probabilities

- E.g., suppose that the initial urn-ball was  $(1, w)$

- What is the next state?

- $(1, w)$  with probability 0.2

80% probability

- $(1, b)$  with probability 0.6

- $(2, w)$  with probability 0.15

20% probability

- $(2, b)$  with probability 0.05

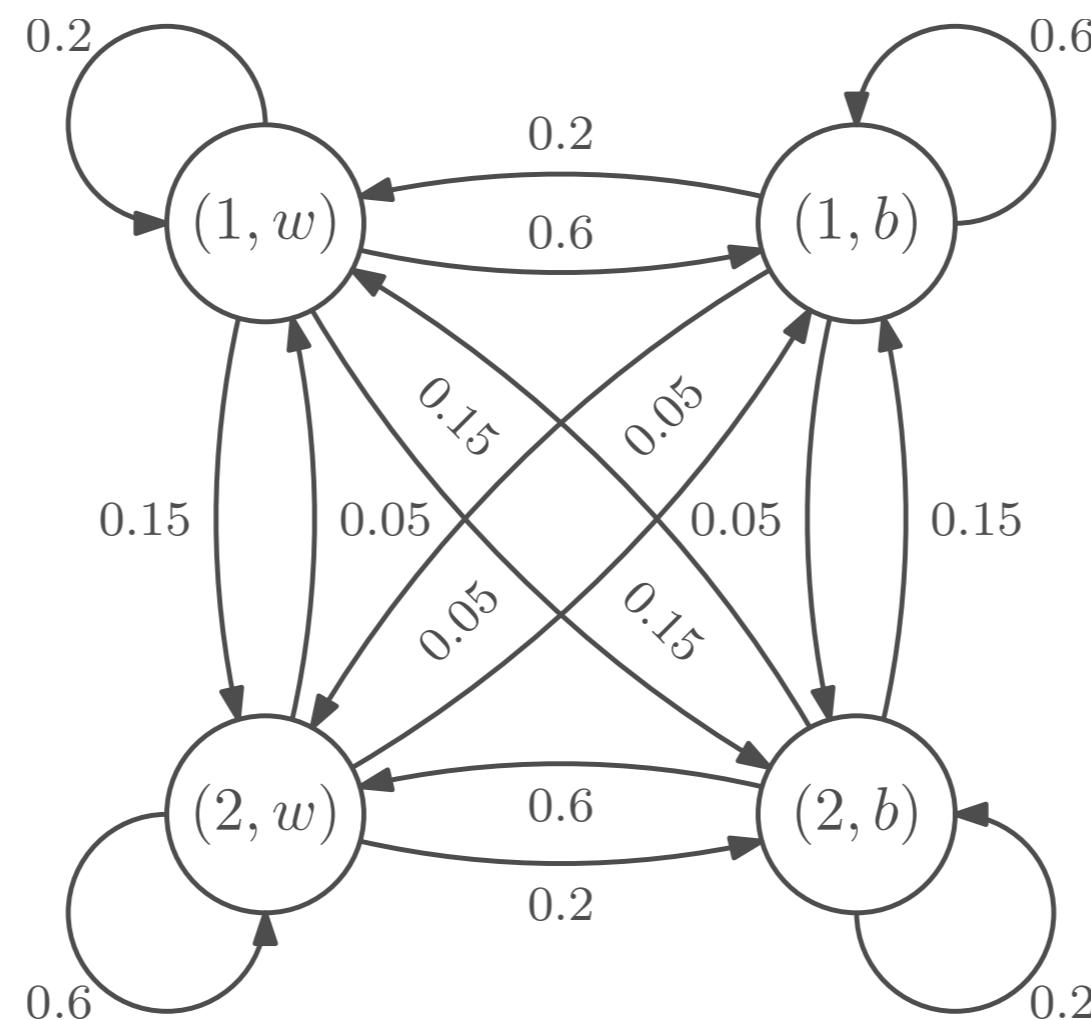
### 3. Transition probabilities?

- We can represent the transition probabilities as a matrix:

|        | (1, w) | (1, b) | (2, w) | (2, b) |
|--------|--------|--------|--------|--------|
| (1, w) | 0.2    | 0.6    | 0.15   | 0.05   |
| (1, b) | 0.2    | 0.6    | 0.15   | 0.05   |
| (2, w) | 0.05   | 0.15   | 0.6    | 0.2    |
| (2, b) | 0.05   | 0.15   | 0.6    | 0.2    |

### 3. Transition probabilities?

- We can also represent the game with a transition diagram:



# Can we make predictions?

- E.g., suppose that the initial urn-ball was  $(1, b)$
- What is the next state?



|      |      |      |      |
|------|------|------|------|
| 0.2  | 0.6  | 0.15 | 0.05 |
| 0.2  | 0.6  | 0.15 | 0.05 |
| 0.05 | 0.15 | 0.6  | 0.2  |
| 0.05 | 0.15 | 0.6  | 0.2  |

# Can we make predictions?

- E.g., suppose that the initial urn-ball was  $(1, b)$
- What is the state at time  $t = 3$ ?

$$\mathbf{P}^3 = \begin{bmatrix} 0.15 & 0.46 & 0.29 & 0.10 \\ 0.15 & 0.46 & 0.29 & 0.10 \\ 0.10 & 0.29 & 0.46 & 0.15 \\ 0.10 & 0.29 & 0.46 & 0.15 \end{bmatrix}$$

A red arrow points to the first row of the matrix  $\mathbf{P}^3$ .

# Can we make predictions?

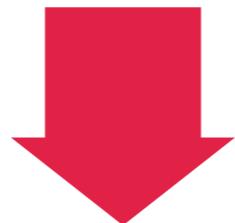
- E.g., suppose that the initial urn-ball was  $(1, b)$
- What is the state at time  $t = 3$ ?
  - $(1, w)$  with probability 0.15
  - $(1, b)$  with probability 0.46
  - $(2, w)$  with probability 0.29
  - $(2, b)$  with probability 0.10

# And what if...?

# The urn problem (revisited)

- At each step, the oracle draws a ball from one of the urns
- The ball is put back afterwards
- The oracle reveals only the color of the ball

Partial observability



The state cannot be  
fully observed

# 1. Is this a Markov chain?

- No!
- The information available at each moment (ball color) is not enough to predict next ball color
- Why?
- The next distribution depends on the urn!

# Can we make predictions?

- Suppose that we know that:
  - At time  $t = 0$ , the ball was white and drawn from urn 1
  - At time  $t = 1$ , the ball was black
  - What is the state at time  $t = 2$ ?

# Can we make predictions?

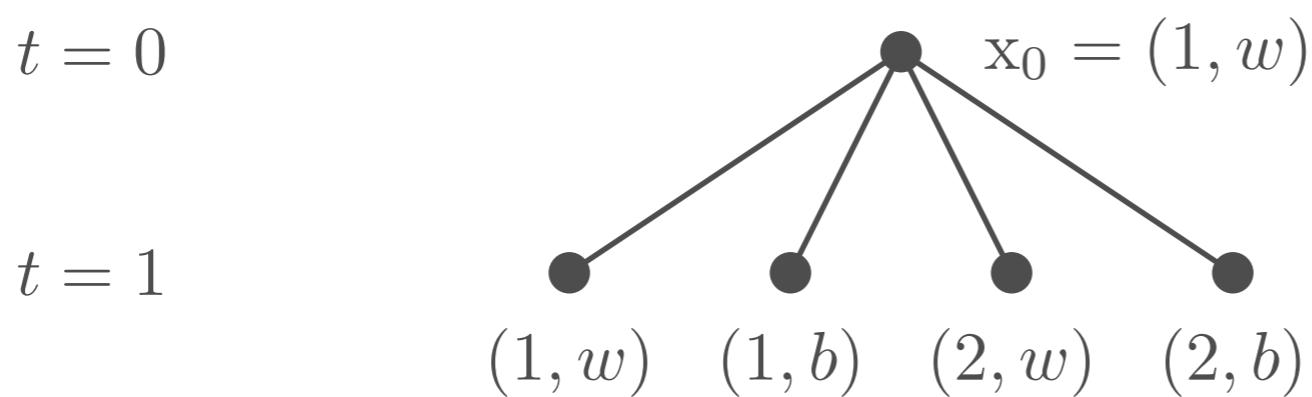
- We have:

$$t = 0$$

- $x_0 = (1, w)$

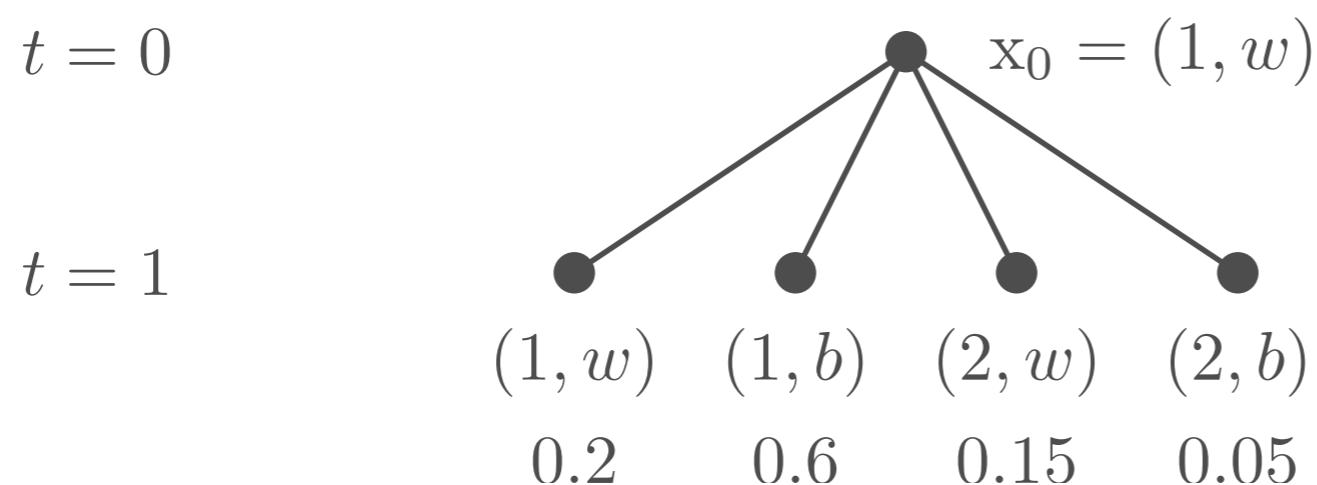
# Can we make predictions?

- We have:



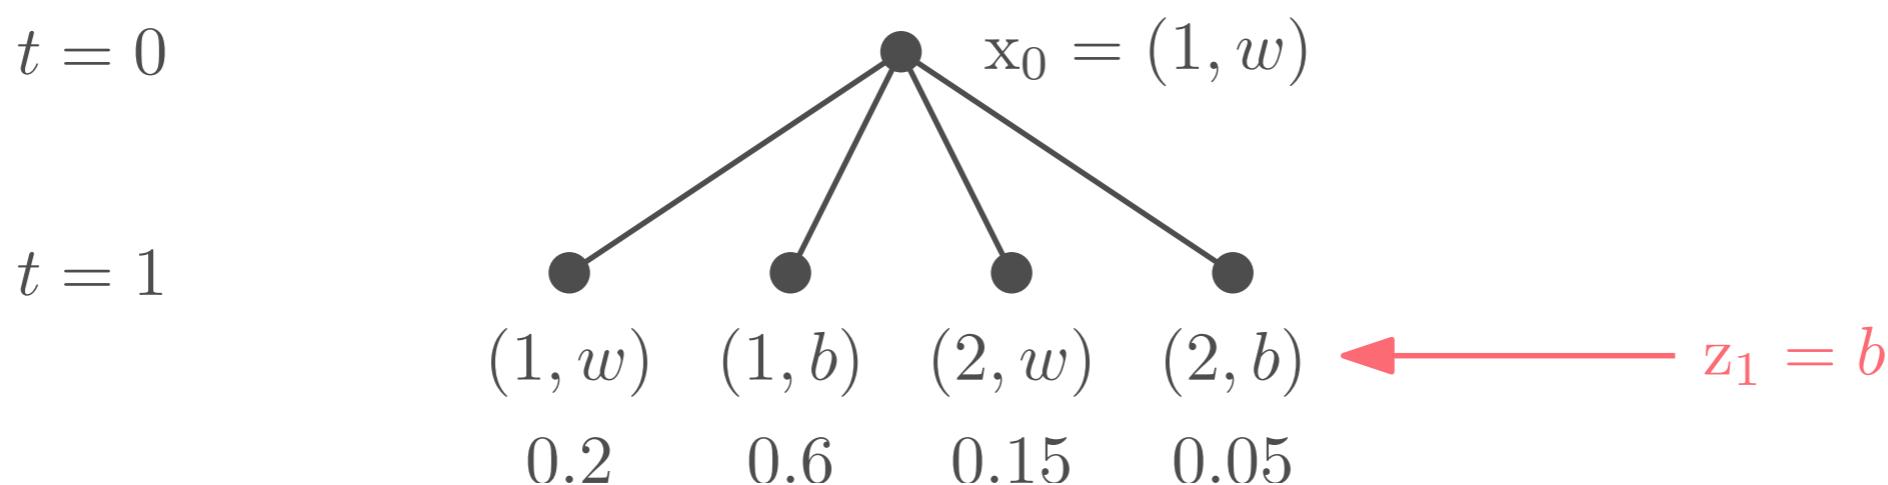
# Can we make predictions?

- We have:



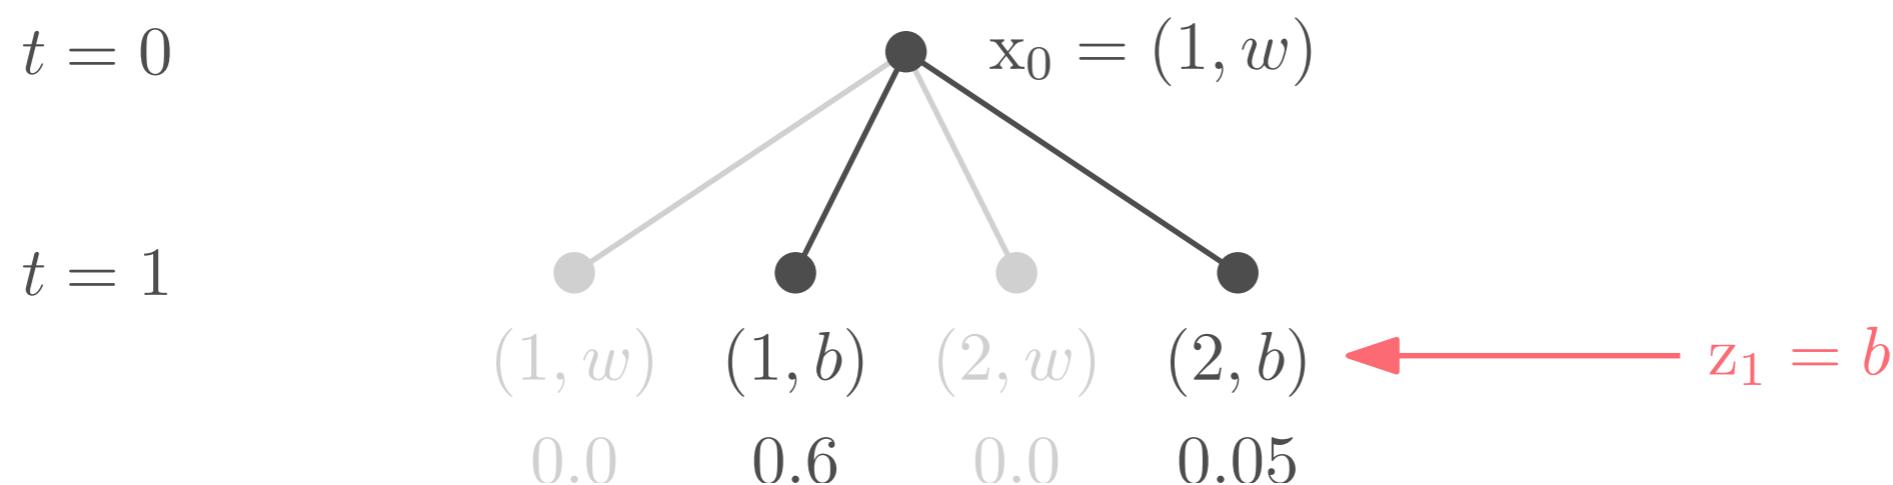
# Can we make predictions?

- We have:



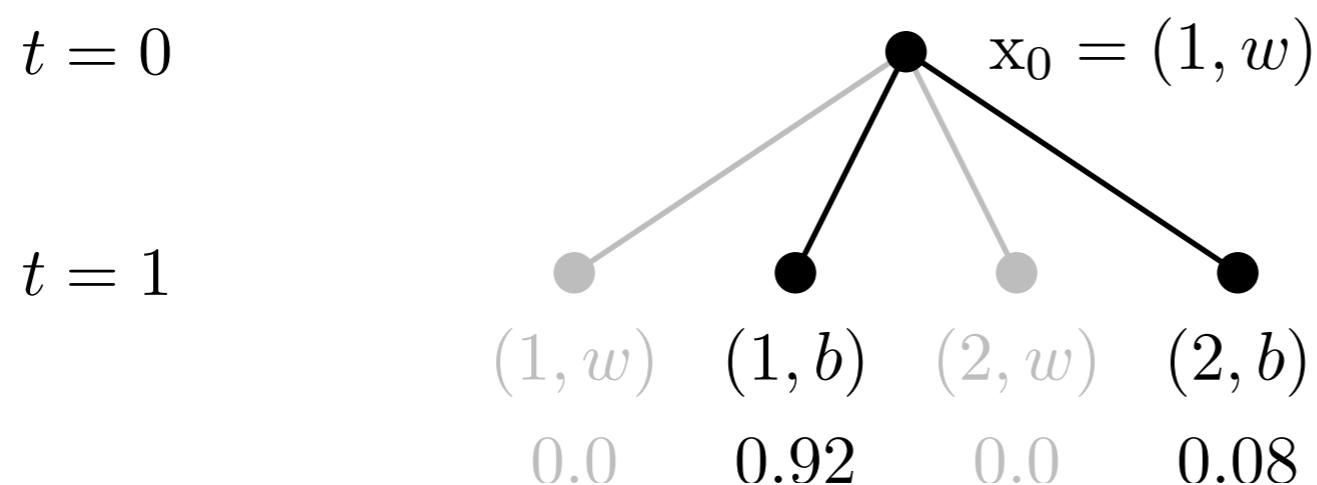
# Can we make predictions?

- We have:



# Can we make predictions?

- We have:



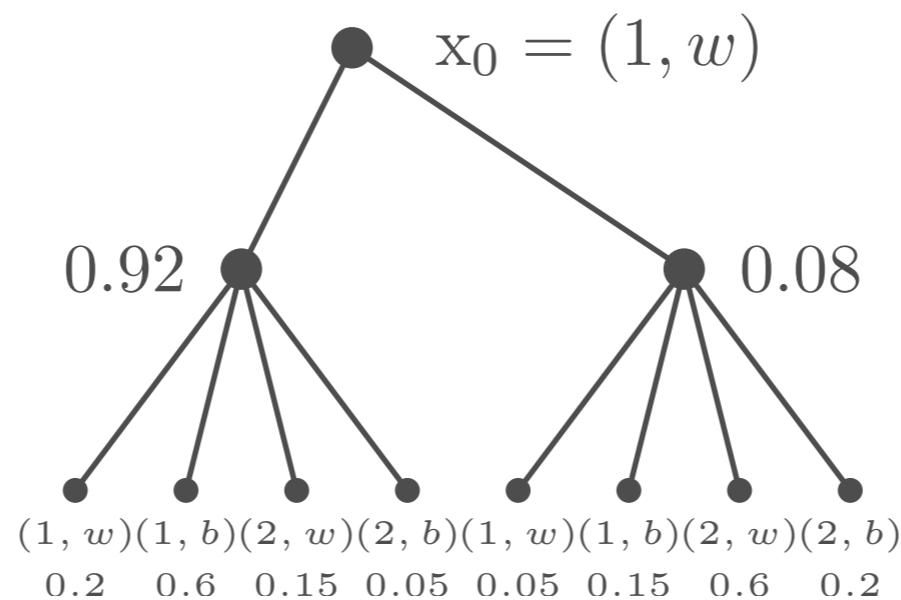
# Can we make predictions?

- We have:

$t = 0$

$t = 1$

$t = 2$



What is the probability of  $(1, w)$ ?

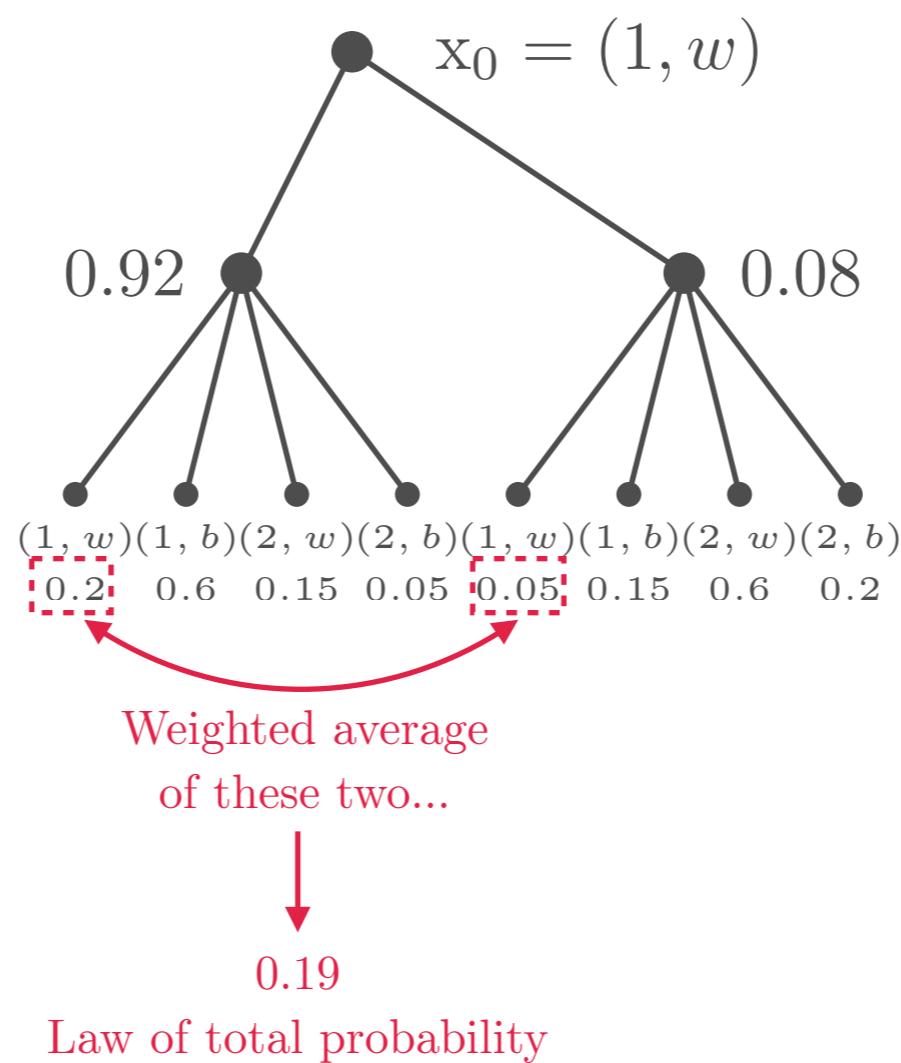
# Can we make predictions?

- We have:

$t = 0$

$t = 1$

$t = 2$



# Can we make predictions?

- In formal terms:

Observation



$$\mathbb{P}[x_2 = (1, w) \mid x_0 = (1, w), z_1 = b]$$

# Can we make predictions?

- In formal terms:

$$\mathbb{P}[x_2 = (1, w) \mid x_0 = (1, w), z_1 = b]$$

$$= \mathbb{P}[x_2 = (1, w) \mid x_1 = x, x_0 = (1, w), z_1 = b]$$



Intermediate state

# Can we make predictions?

- In formal terms:

$$\mathbb{P}[x_2 = (1, w) \mid x_0 = (1, w), z_1 = b]$$

$$= \mathbb{P}[x_2 = (1, w) \mid x_1 = x, x_0 = (1, w), z_1 = b] \mathbb{P}[x_1 = x \mid x_0 = (1, w), z_1 = b]$$



Probability of  
intermediate state

# Can we make predictions?

- In formal terms:

$$\mathbb{P}[x_2 = (1, w) \mid x_0 = (1, w), z_1 = b]$$

$$= \sum_{x \in \mathcal{X}} \mathbb{P}[x_2 = (1, w) \mid x_1 = x, x_0 = (1, w), z_1 = b] \mathbb{P}[x_1 = x \mid x_0 = (1, w), z_1 = b]$$

Weighted mean  
of probabilities  
(total probability law)

Bayes theorem

# Can we make predictions?

- In formal terms:

$$\mathbb{P}[x_2 = (1, w) \mid x_0 = (1, w), z_1 = b]$$

$$= \sum_{x \in \mathcal{X}} \mathbb{P}[x_2 = (1, w) \mid x_1 = x, x_0 = (1, w), z_1 = b] \frac{\mathbb{P}[z_1 = b \mid x_1 = x, x_0 = (1, w)] \mathbb{P}[x_1 = x \mid x_0 = (1, w)]}{\mathbb{P}[z_1 = b \mid x_0 = (1, w)]}$$

We need only  $x_1$   
to predict  $x_2$

The observation  
depends only on  $x_1$

Transition  
probabilities

Normalization  
factor

# Can we make predictions?

- In formal terms:

$$\mathbb{P}[x_2 = (1, w) \mid x_0 = (1, w), z_1 = b]$$

$$= \frac{1}{\rho} \sum_{x \in \mathcal{X}} \mathbf{P}((1, w) \mid x) \mathbb{P}[z_1 = b \mid x_1 = x] \mathbf{P}(x \mid (1, w))$$



If  $x_1 = (1, w)$  or  $x_1 = (2, w)$   
this probability is 0,  
otherwise it's 1

# Can we make predictions?

- In formal terms:

$$\mathbb{P}[x_2 = (1, w) \mid x_0 = (1, w), z_1 = b]$$

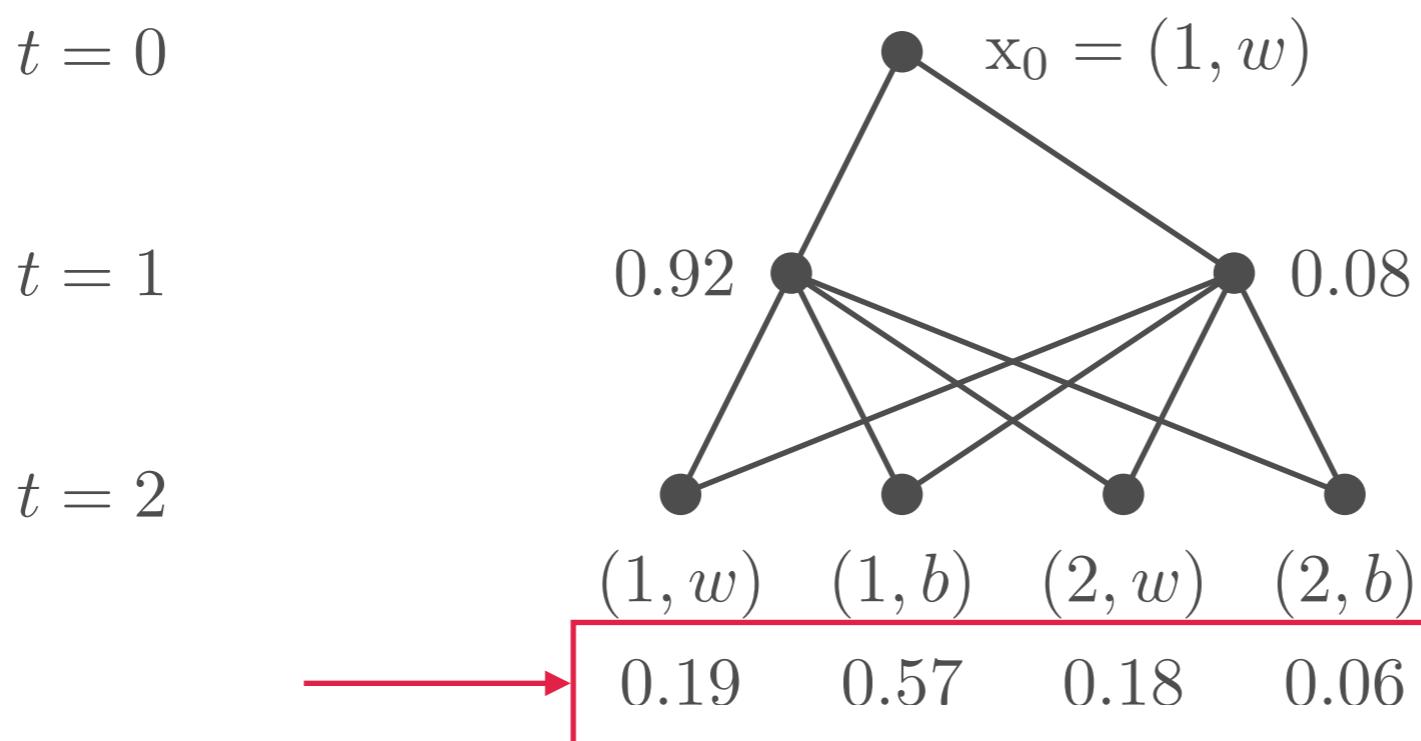
$$= \frac{1}{\rho} (\mathbf{P}((1, w) \mid (1, b)) \mathbf{P}((1, b) \mid (1, w)) + \mathbf{P}((1, w) \mid (2, b)) \mathbf{P}((2, b) \mid (1, w)))$$

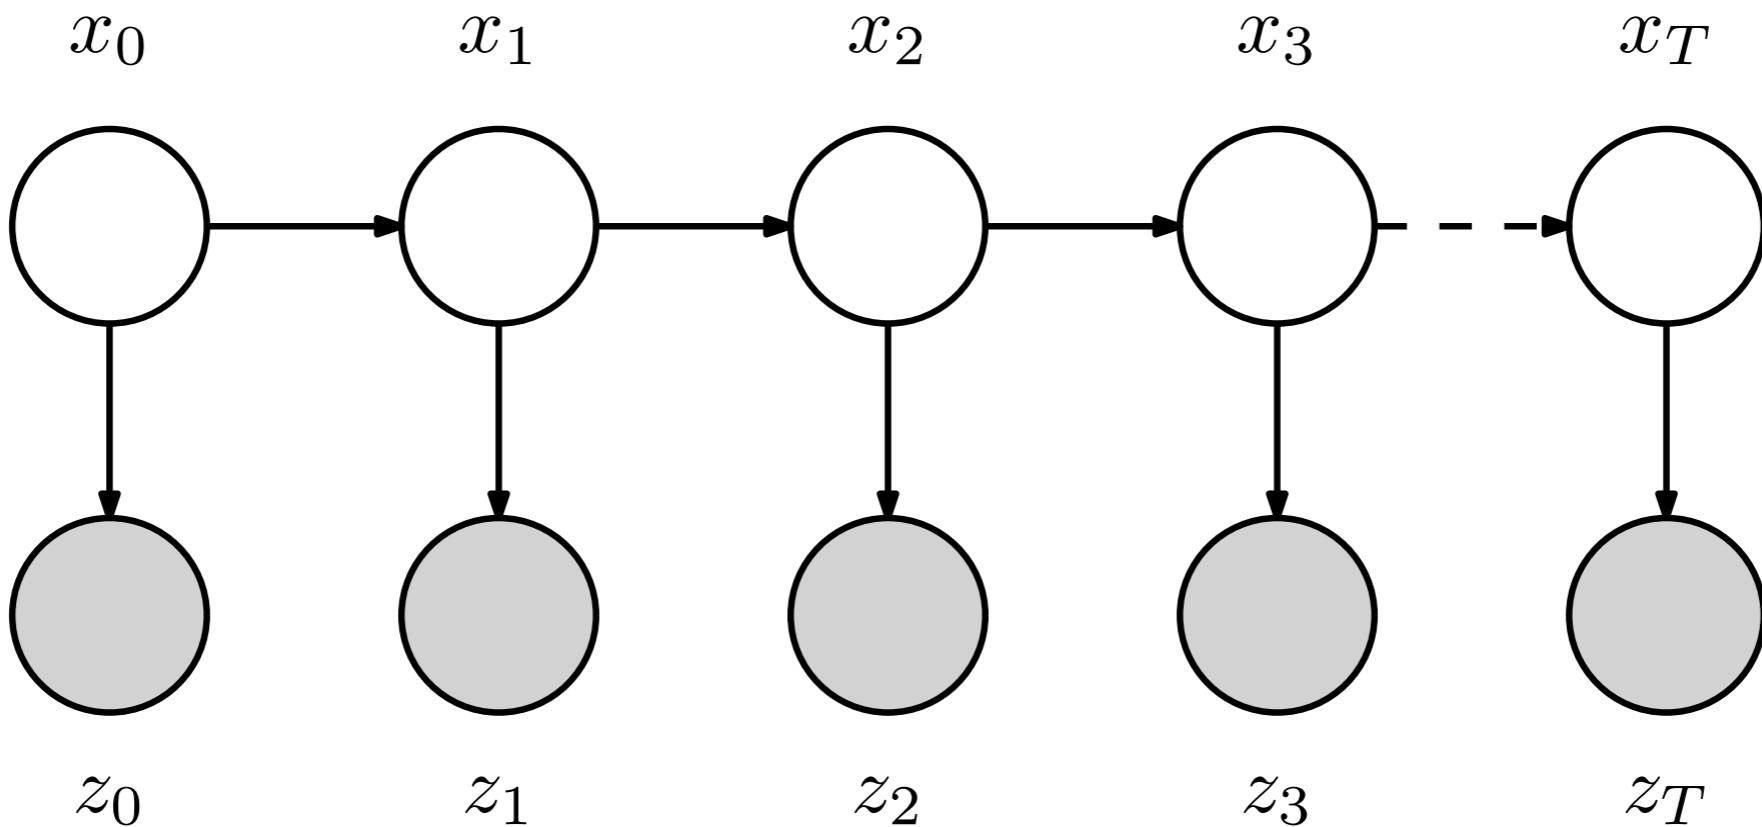
- More generally:

$$\mathbb{P}[x_2 = x \mid x_0 = (1, w), z_1 = b] = \frac{1}{\rho} (\mathbf{P}(x \mid (1, b)) \mathbf{P}((1, b) \mid (1, w)) + \mathbf{P}(x \mid (2, b)) \mathbf{P}((2, b) \mid (1, w)))$$

# Can we make predictions?

- We have:

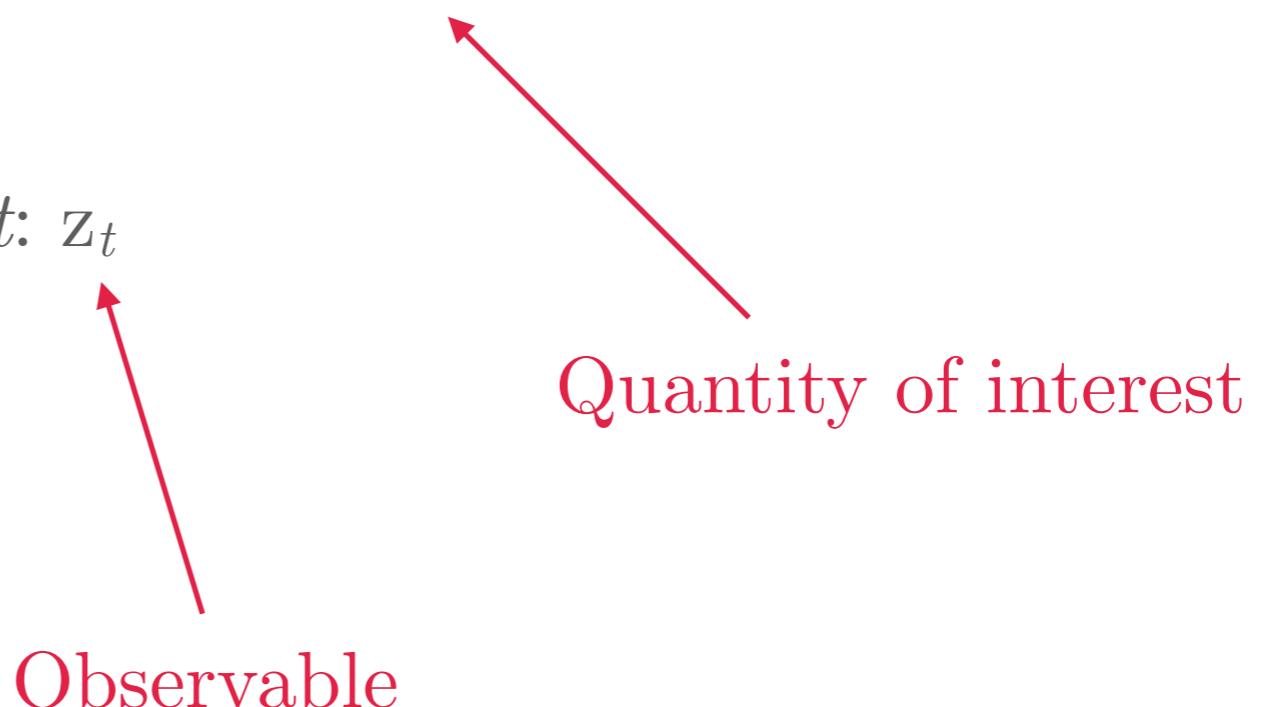




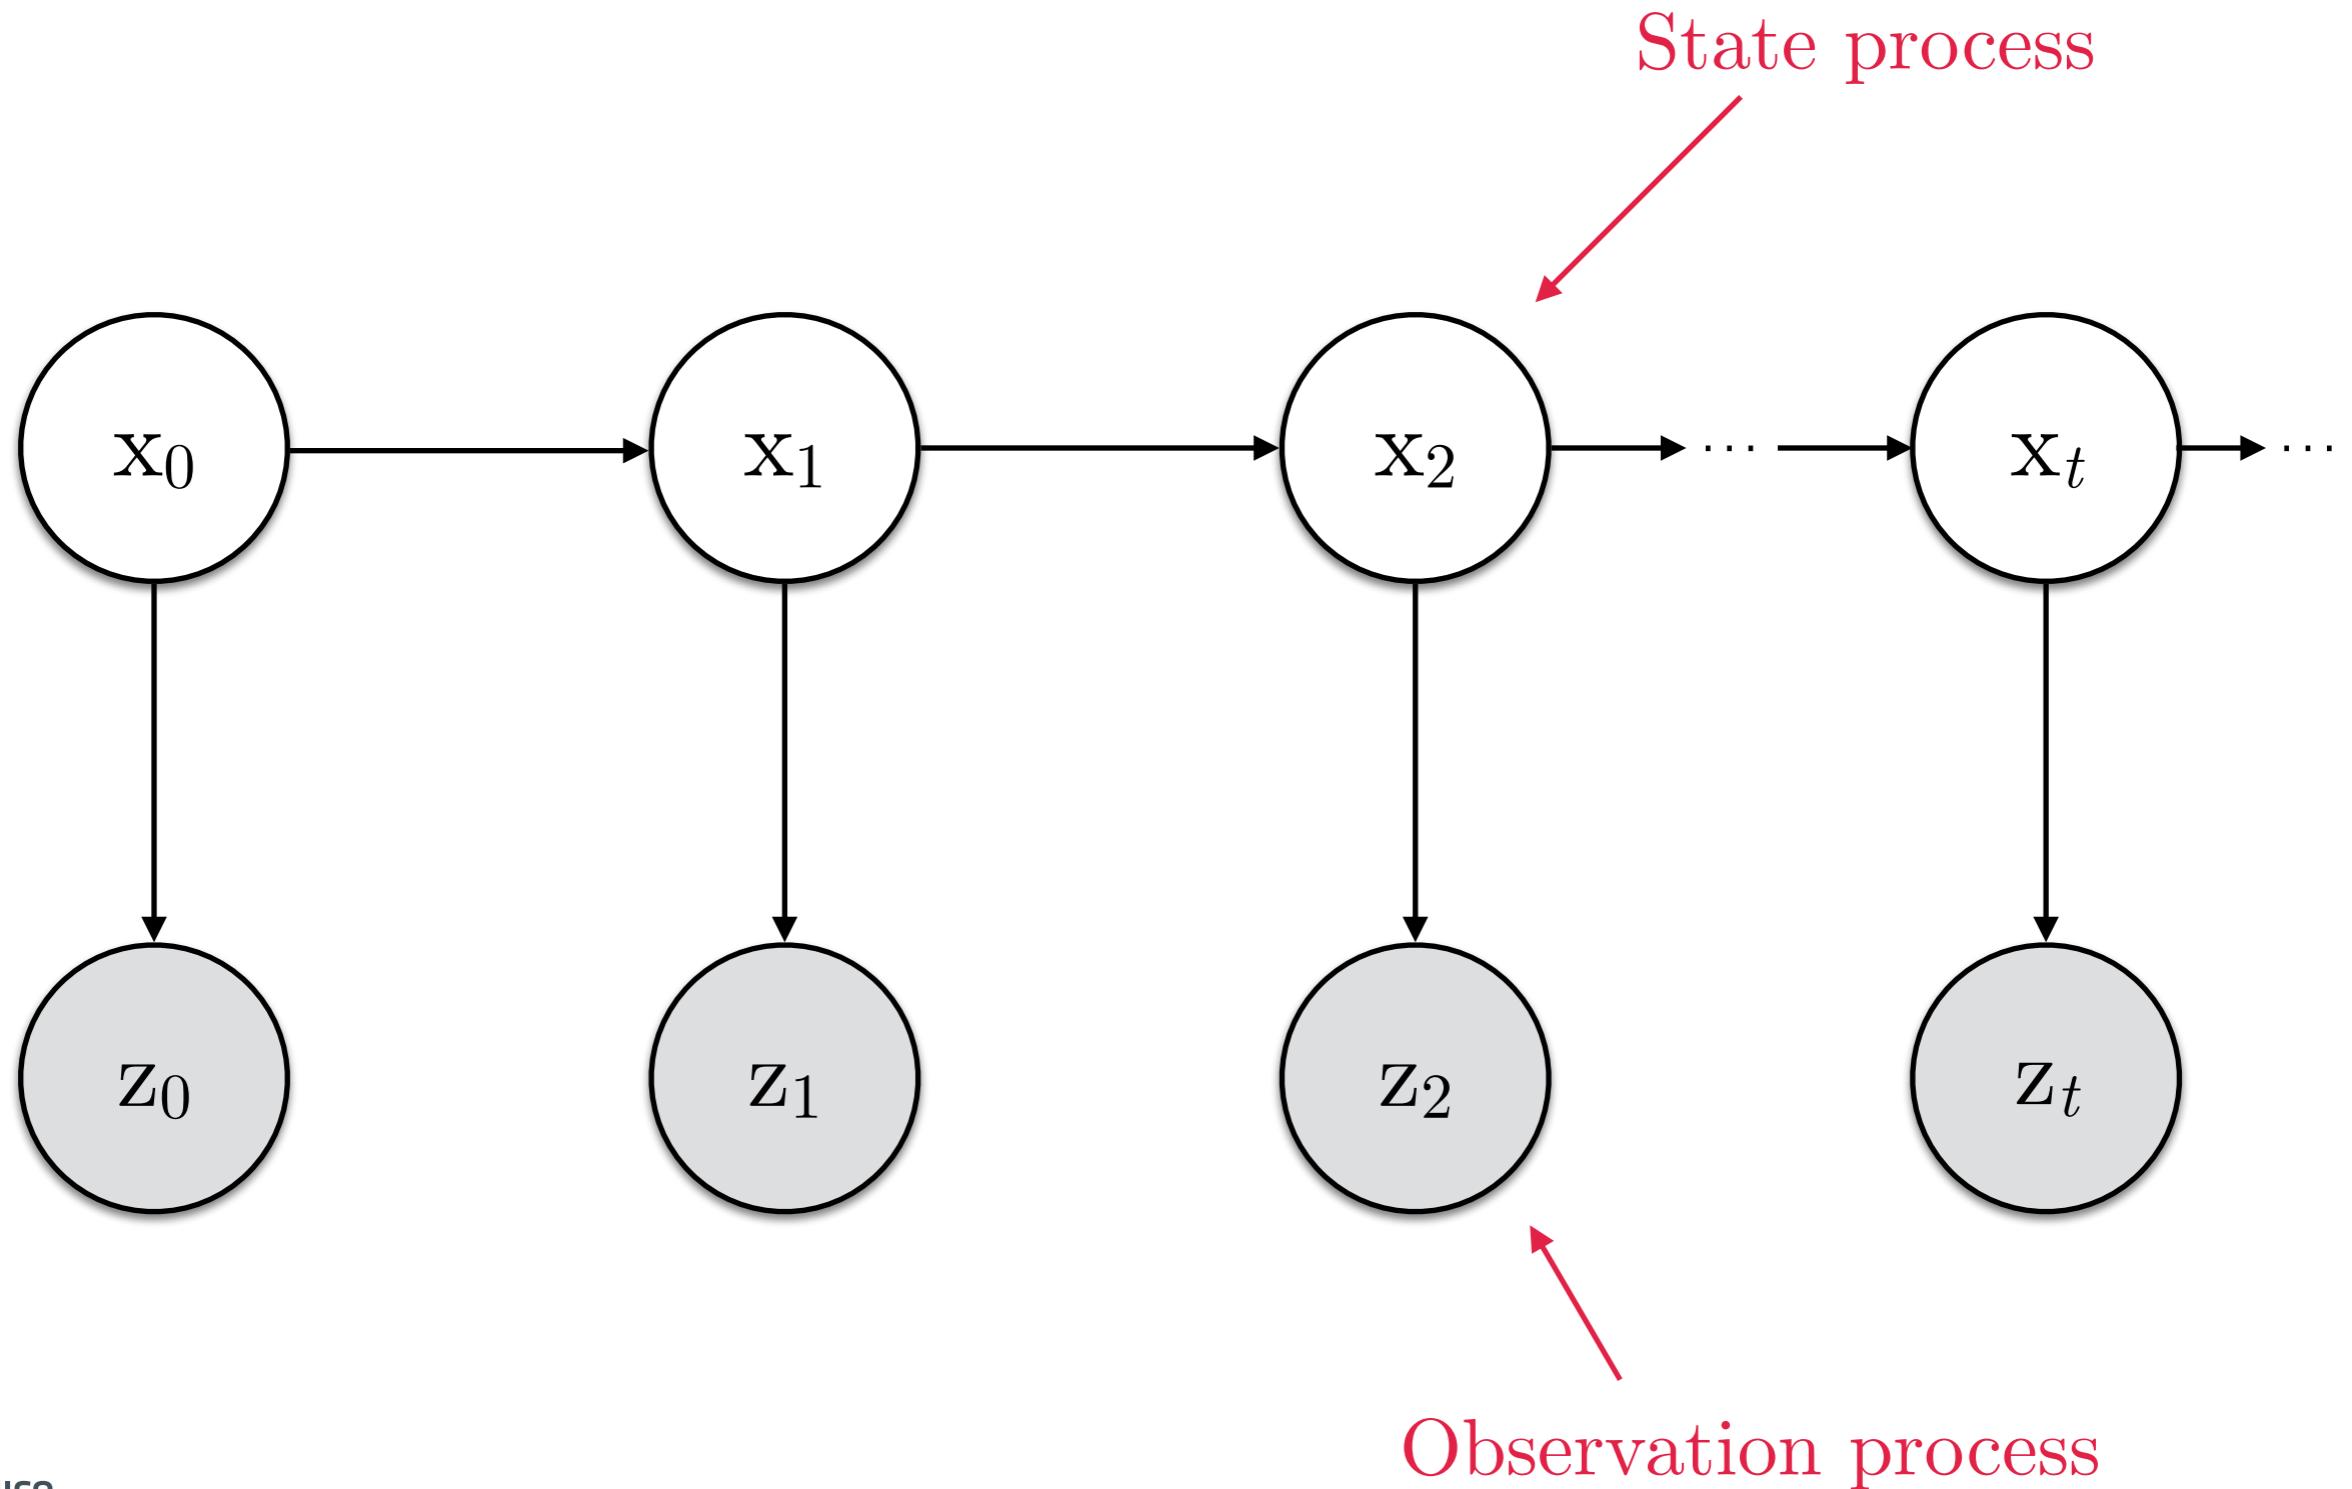
# Hidden Markov models

# Hidden Markov model

- Model for **sequential process** with partial observability
- Process evolves in **discrete time steps** (as in Markov chains)
- State of the process at time step  $t$ :  $x_t$
- Observation at time step  $t$ :  $z_t$



# Hidden Markov model



# Hidden Markov model

## Markov state

The state at instant  $t$  is enough to predict the state at instant  $t + 1$ :

$$\mathbb{P} [\mathbf{x}_{t+1} = y \mid \mathbf{x}_{0:t} = \mathbf{x}_{0:t}, \mathbf{z}_{0:t} = \mathbf{z}_{0:t}] = \mathbb{P} [\mathbf{x}_{t+1} = y \mid \mathbf{x}_t = x_t]$$



Depends only on the  
last state

# Hidden Markov model

## State-dependent observations

The state at instant  $t$  is enough to predict the observation at instant  $t$ :

$$\mathbb{P} [z_t = z \mid \mathbf{x}_{0:t} = \mathbf{x}_{0:t}, \mathbf{z}_{0:t-1} = \mathbf{z}_{0:t-1}] = \mathbb{P} [z_t = z \mid x_t = x_t]$$



Depends only on the  
last state

# Hidden Markov model

## Markov state

The state at instant  $t$  is enough to predict the state at instant  $t + 1$ :

$$\mathbb{P} [\mathbf{x}_{t+1} = y \mid \mathbf{x}_{0:t} = \mathbf{x}_{0:t}, \mathbf{z}_{0:t} = \mathbf{z}_{0:t}] = \mathbb{P} [\mathbf{x}_{t+1} = y \mid \mathbf{x}_t = x_t]$$

## State-dependent observations

The state at instant  $t$  is enough to predict the observation at instant  $t$ :

$$\mathbb{P} [\mathbf{z}_t = z \mid \mathbf{x}_{0:t} = \mathbf{x}_{0:t}, \mathbf{z}_{0:t-1} = \mathbf{z}_{0:t-1}] = \mathbb{P} [\mathbf{z}_t = z \mid \mathbf{x}_t = x_t]$$

# Hidden Markov model

- Other assumptions (for most of this course):
  - There is only a **finite number** of possible states
  - $\mathcal{X}$  is the set of possible states (**state space**)
  - There is only a **finite number** of observations
  - $\mathcal{Z}$  is the set of possible observations (**observation space**)

# Hidden Markov model

- Other assumptions:
  - The transition probabilities  $\mathbb{P} [x_{t+1} = y \mid x_t = x]$  do not depend on  $t$
  - The observation probabilities  $\mathbb{P} [z_t = z \mid x_t = x]$  do not depend on  $t$

# Transition probability matrix

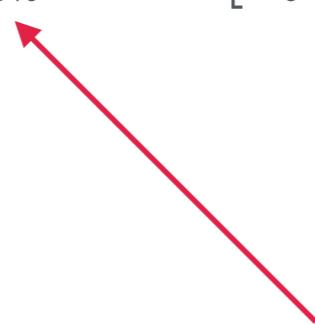
- As with Markov chains, we store the transition probabilities in a **matrix  $P$**

$$[P]_{xy} = \mathbb{P} [x_{t+1} = y \mid x_t = x]$$

# Observation probability matrix

- Similarly, we store the observation probabilities in a **matrix  $O$**

$$[O]_{xz} = \mathbb{P} [z_t = z \mid x_t = x]$$



Number in row  $x$  column  $z$  is  
the probability of observing  
 $z$  in  $x$

# Example

- The urn problem

|        | (1, w) | (1, b) | (2, w) | (2, b) |
|--------|--------|--------|--------|--------|
| (1, w) | 0.2    | 0.6    | 0.15   | 0.05   |
| (1, b) | 0.2    | 0.6    | 0.15   | 0.05   |
| (2, w) | 0.05   | 0.15   | 0.6    | 0.2    |
| (2, b) | 0.05   | 0.15   | 0.6    | 0.2    |

|        | w | b |
|--------|---|---|
| (1, w) |   |   |
| (1, b) |   |   |
| (2, w) |   |   |
| (2, b) |   |   |

# Example

- The urn problem

|        | (1, w) | (1, b) | (2, w) | (2, b) |
|--------|--------|--------|--------|--------|
| (1, w) | 0.2    | 0.6    | 0.15   | 0.05   |
| (1, b) | 0.2    | 0.6    | 0.15   | 0.05   |
| (2, w) | 0.05   | 0.15   | 0.6    | 0.2    |
| (2, b) | 0.05   | 0.15   | 0.6    | 0.2    |

|        | w | b |
|--------|---|---|
| (1, w) | 1 | 0 |
| (1, b) |   |   |
| (2, w) |   |   |
| (2, b) |   |   |

# Example

- The urn problem

|        | (1, w) | (1, b) | (2, w) | (2, b) |
|--------|--------|--------|--------|--------|
| (1, w) | 0.2    | 0.6    | 0.15   | 0.05   |
| (1, b) | 0.2    | 0.6    | 0.15   | 0.05   |
| (2, w) | 0.05   | 0.15   | 0.6    | 0.2    |
| (2, b) | 0.05   | 0.15   | 0.6    | 0.2    |

|        | w | b |
|--------|---|---|
| (1, w) | 1 | 0 |
| (1, b) | 0 | 1 |
| (2, w) |   |   |
| (2, b) |   |   |

# Example

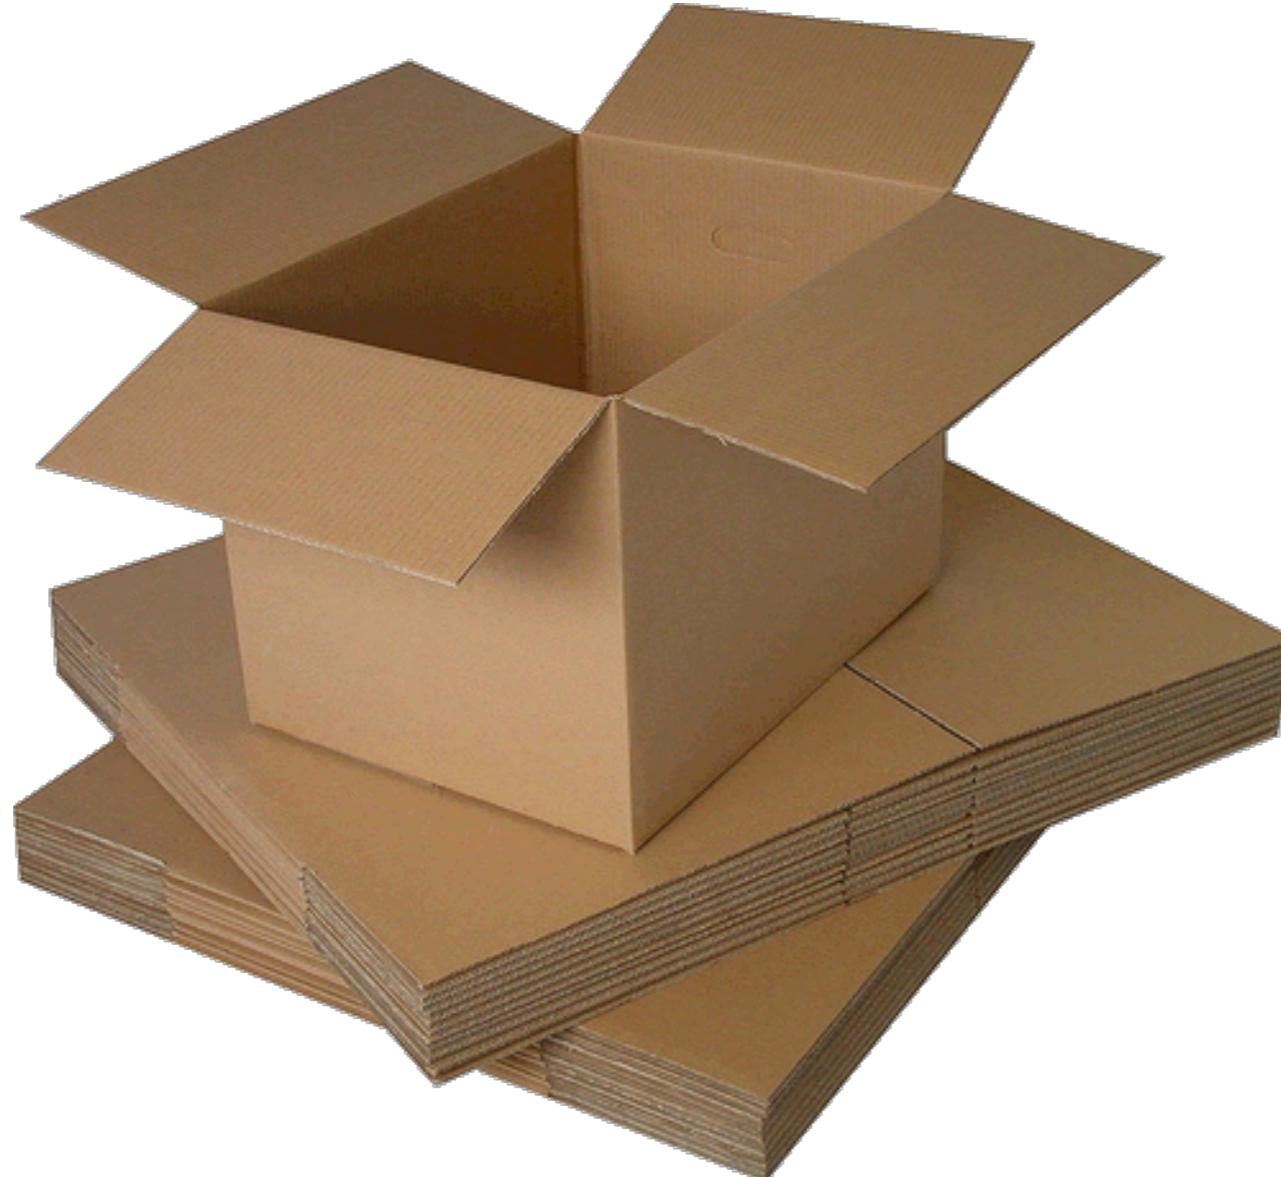
- The urn problem

|        | (1, w) | (1, b) | (2, w) | (2, b) |
|--------|--------|--------|--------|--------|
| (1, w) | 0.2    | 0.6    | 0.15   | 0.05   |
| (1, b) | 0.2    | 0.6    | 0.15   | 0.05   |
| (2, w) | 0.05   | 0.15   | 0.6    | 0.2    |
| (2, b) | 0.05   | 0.15   | 0.6    | 0.2    |

|        | w | b |
|--------|---|---|
| (1, w) | 1 | 0 |
| (1, b) | 0 | 1 |
| (2, w) | 1 | 0 |
| (2, b) | 0 | 1 |

# Summarizing...

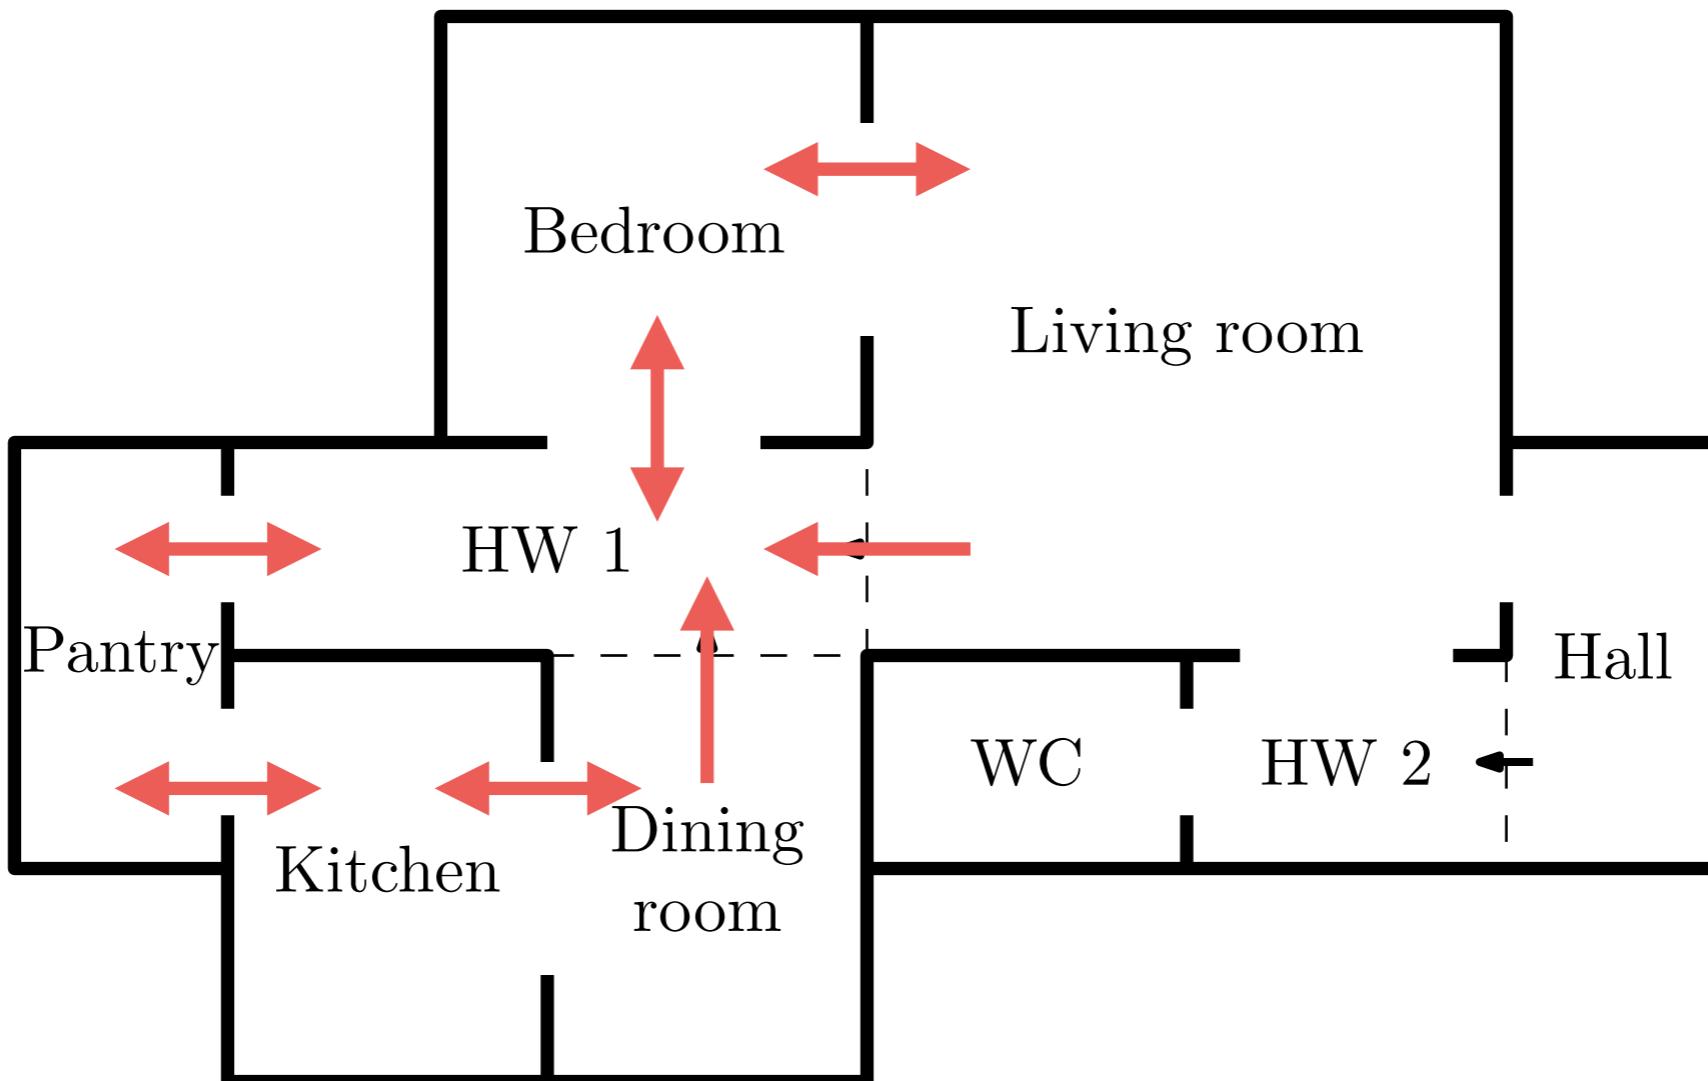
- A HMM can be represented compactly as a tuple  $(\mathcal{X}, \mathcal{Z}, \mathbf{P}, \mathbf{O})$ 
  - $\mathcal{X}$  is the set of possible states
  - $\mathcal{Z}$  is the set of possible observations
  - $\mathbf{P}$  is the transition probability matrix
  - $\mathbf{O}$  is the observation probability matrix



# Examples

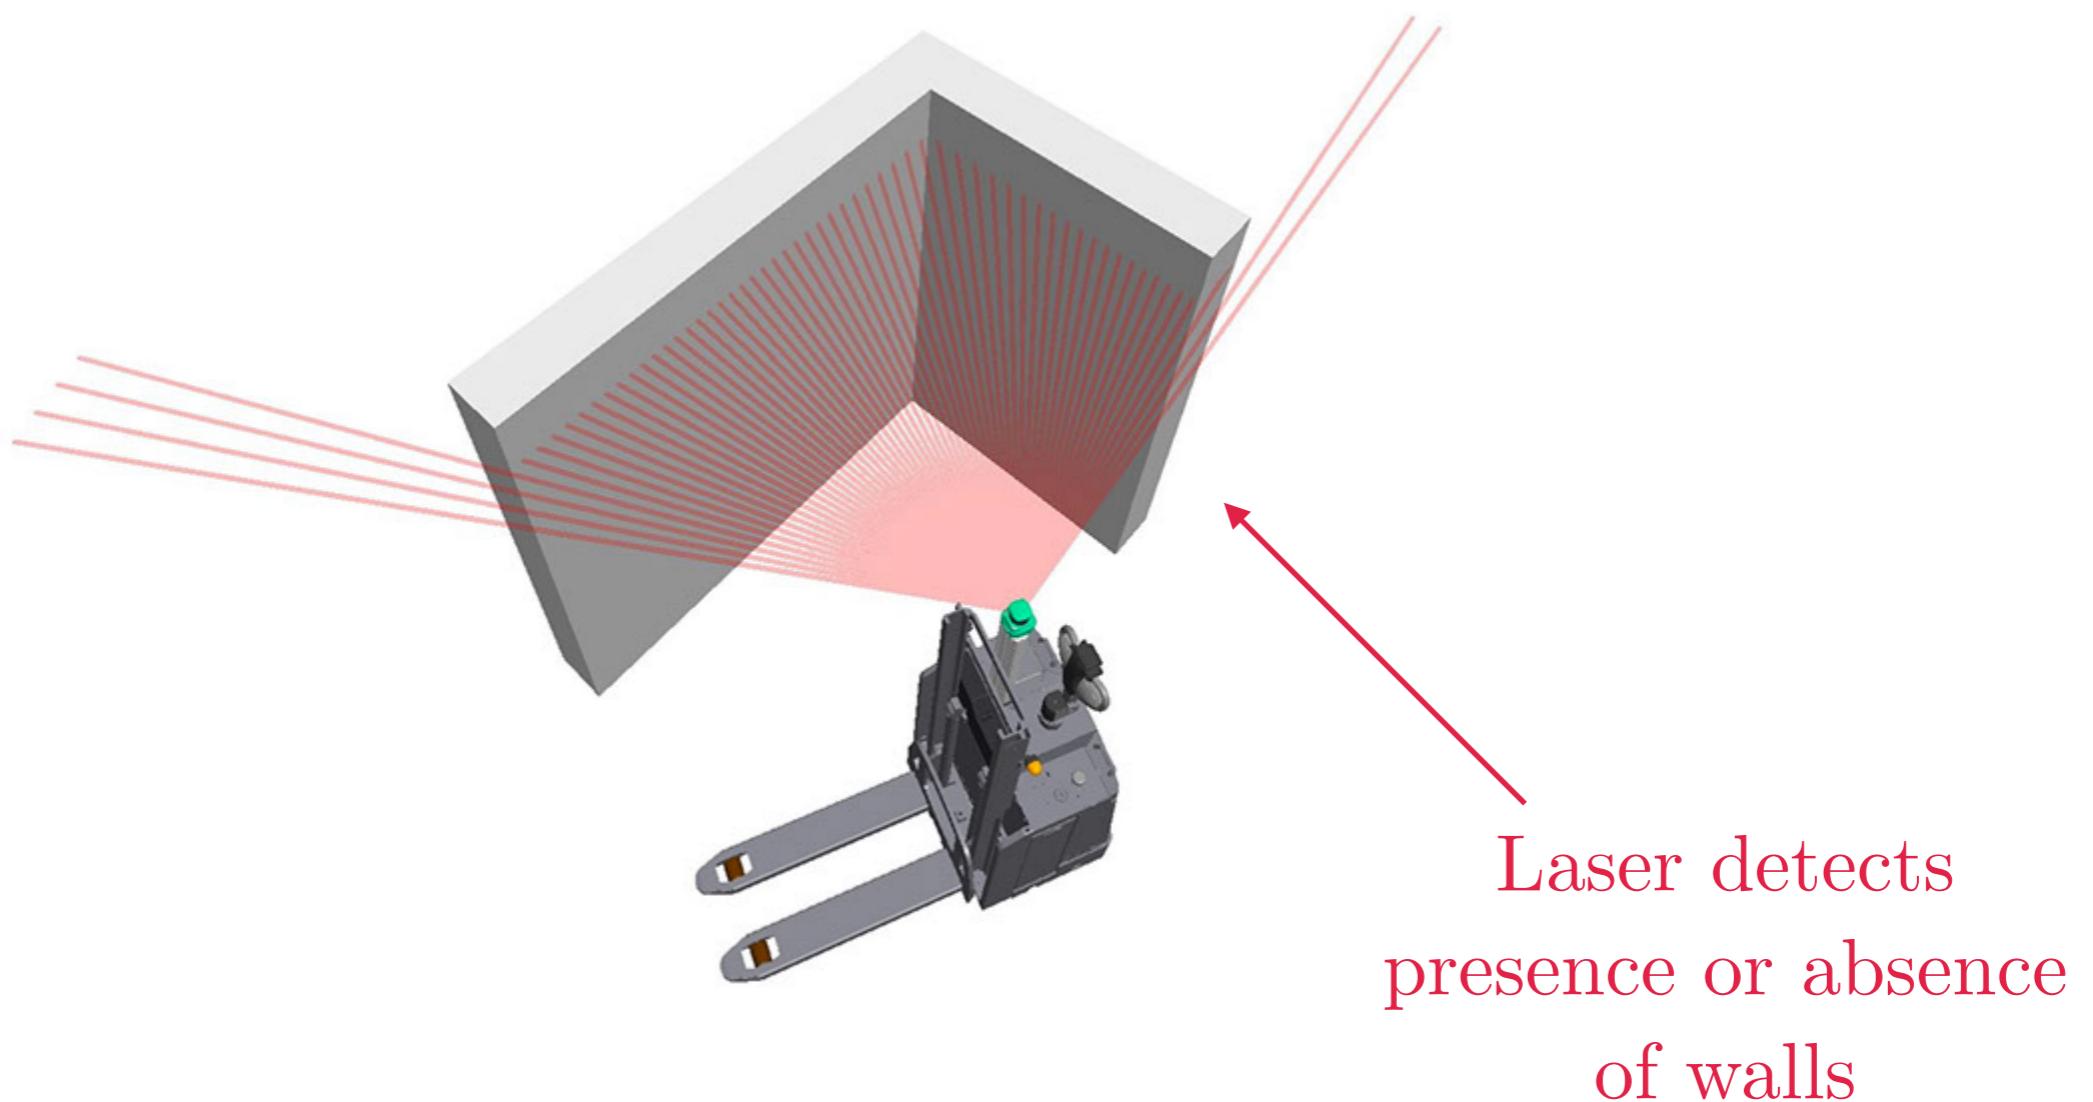
# 1. Household robot

- Remember?



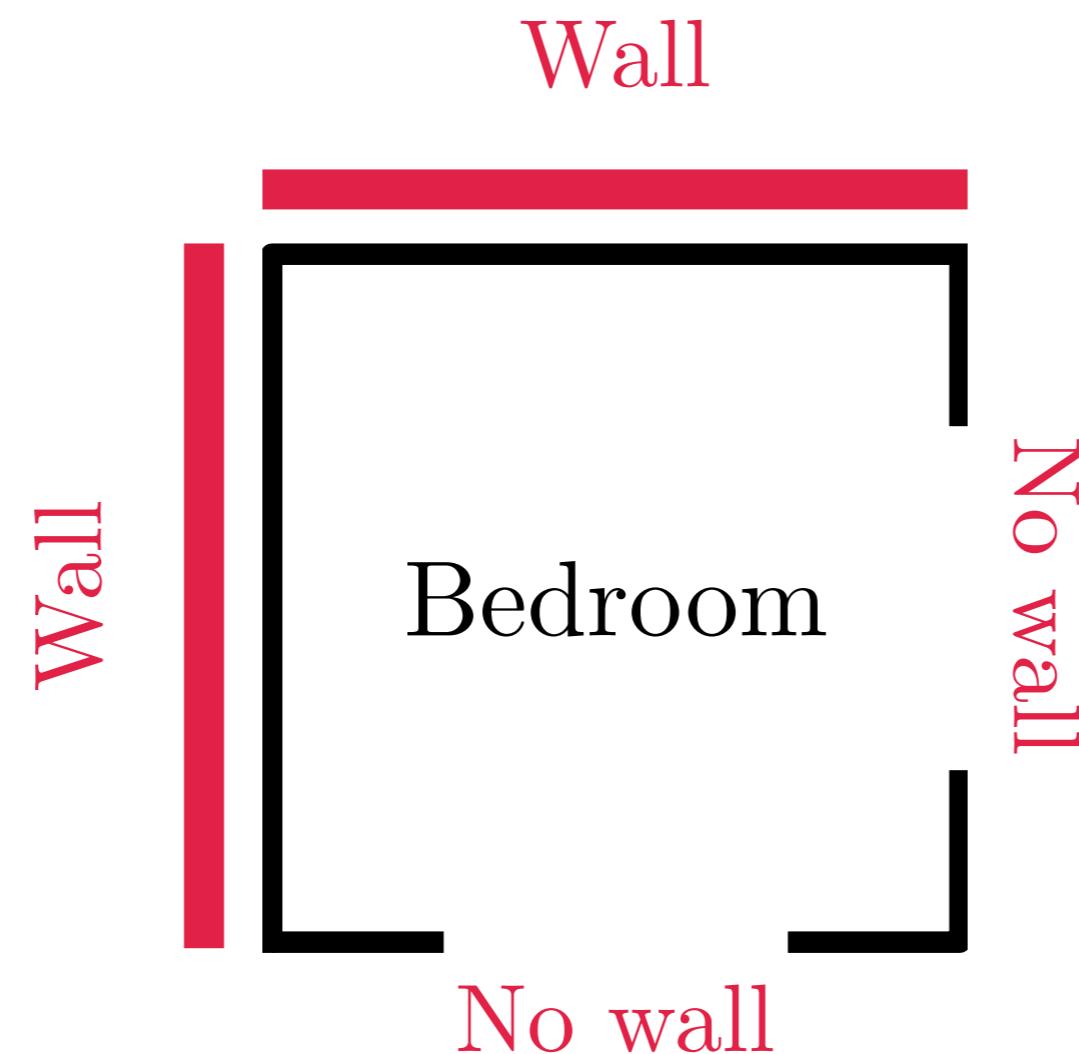
# 1. Household robot

- Robot navigates using a laser



# 1. Household robot

- For example:



# 1. Household robot

- However, laser is not perfect
  - It fails to detect existing walls with 5% probability
  - It detects non-existing walls with 10% probability (in some situations with 20% probability)
  - Detection of a wall independent of adjacent walls

# Is this an HMM?

- State verifies the Markov property?
  - Yes - the position of the robot at time  $t + 1$  depends only of position of the robot at time  $t$
- Observation depends only on state?
  - Yes - the wall detections depends only on robot position

**It is an HMM!**

# What are the states?

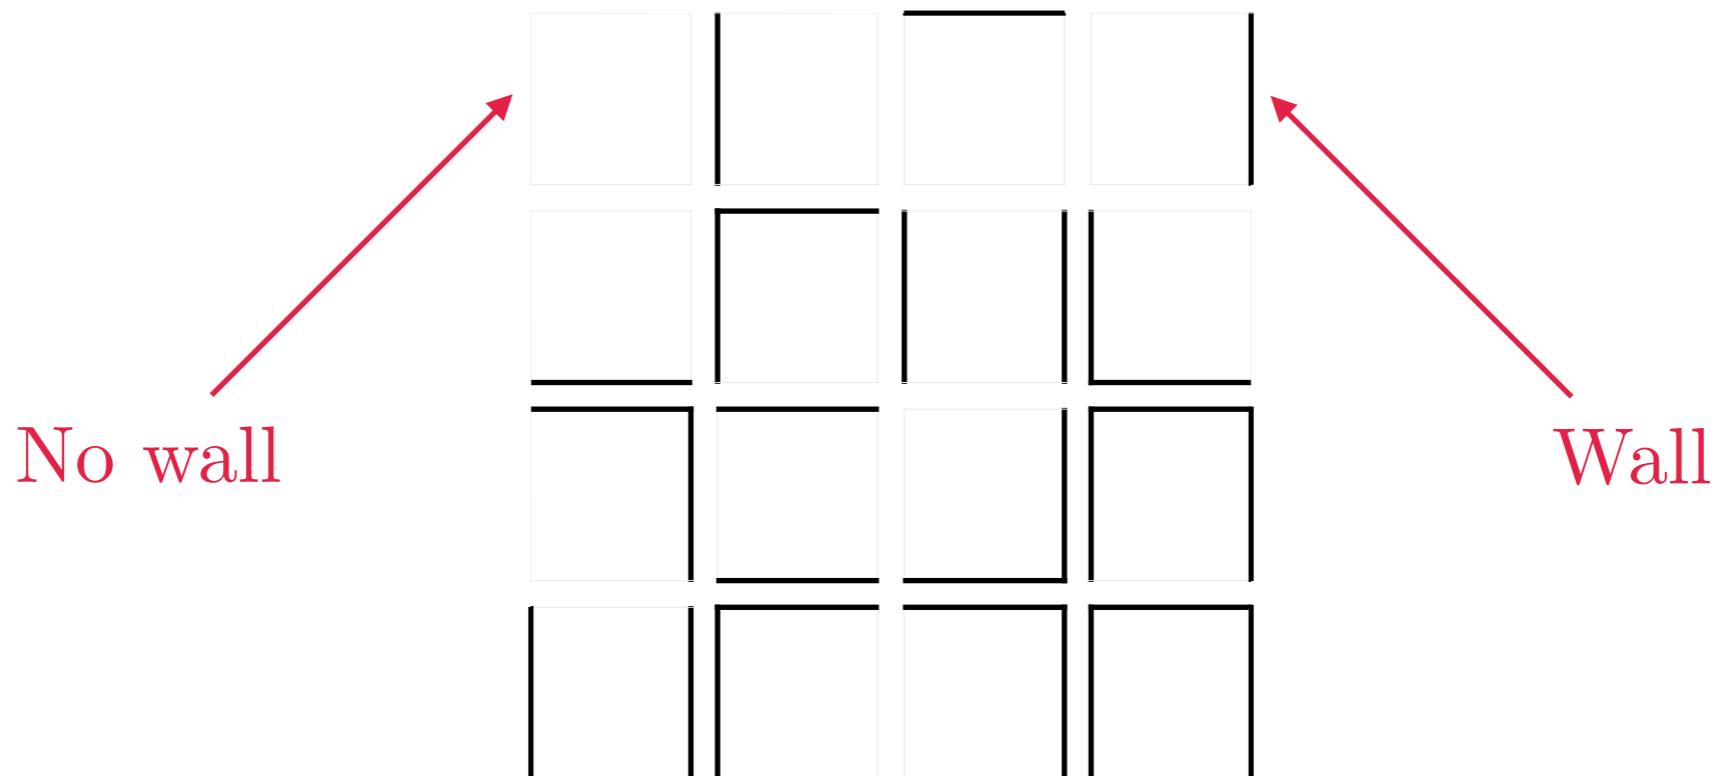
- Possible positions of the robot:
  - $\mathcal{X} = \{K, P, D, H_1, B, L\}$

# What are the observations?

- What does the robot “see”?

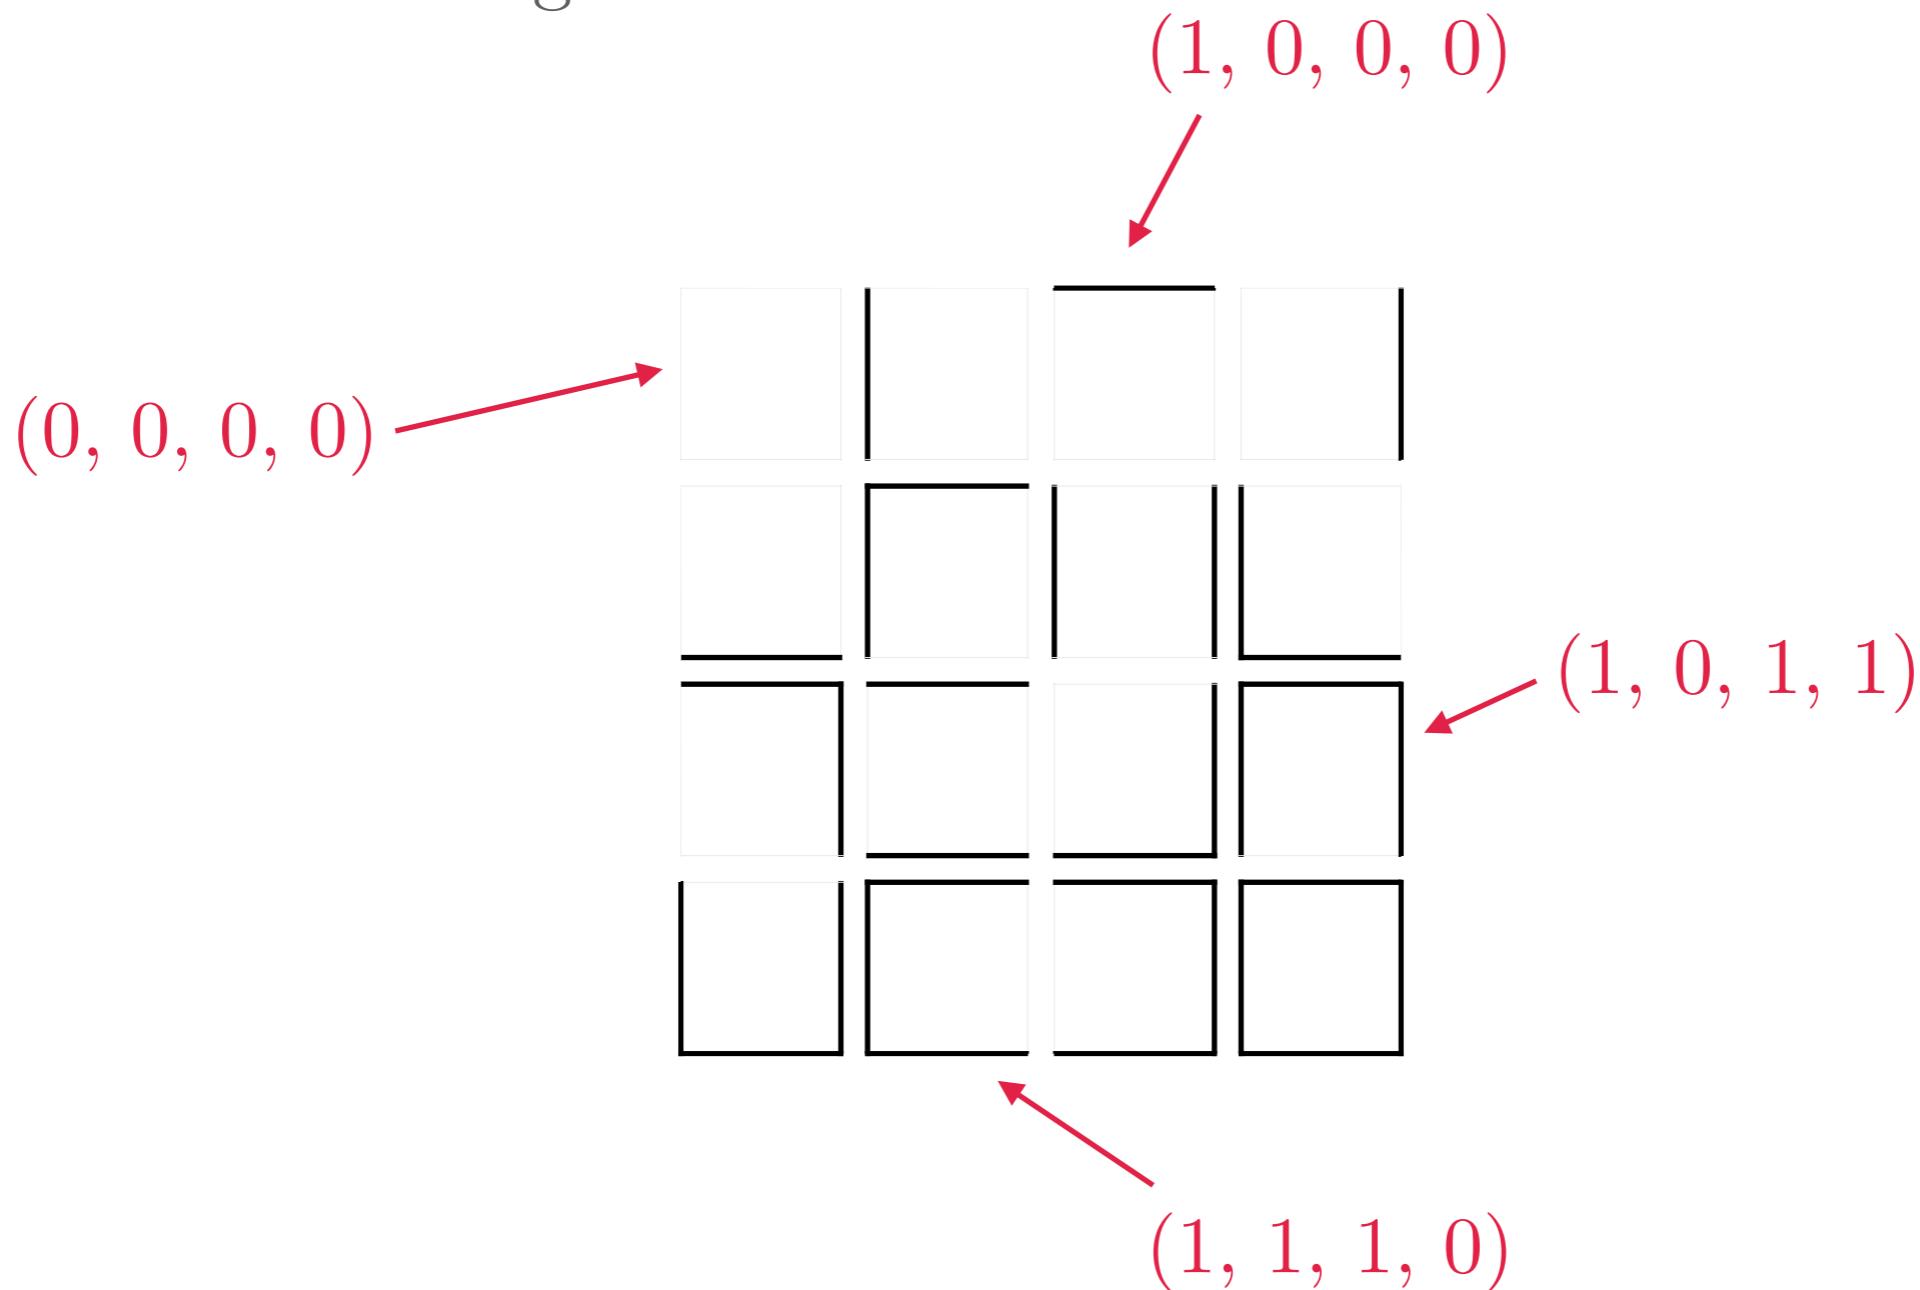
# What are the observations?

- Possible wall configurations:



# What are the observations?

- Possible wall configurations:



# What are the observations?

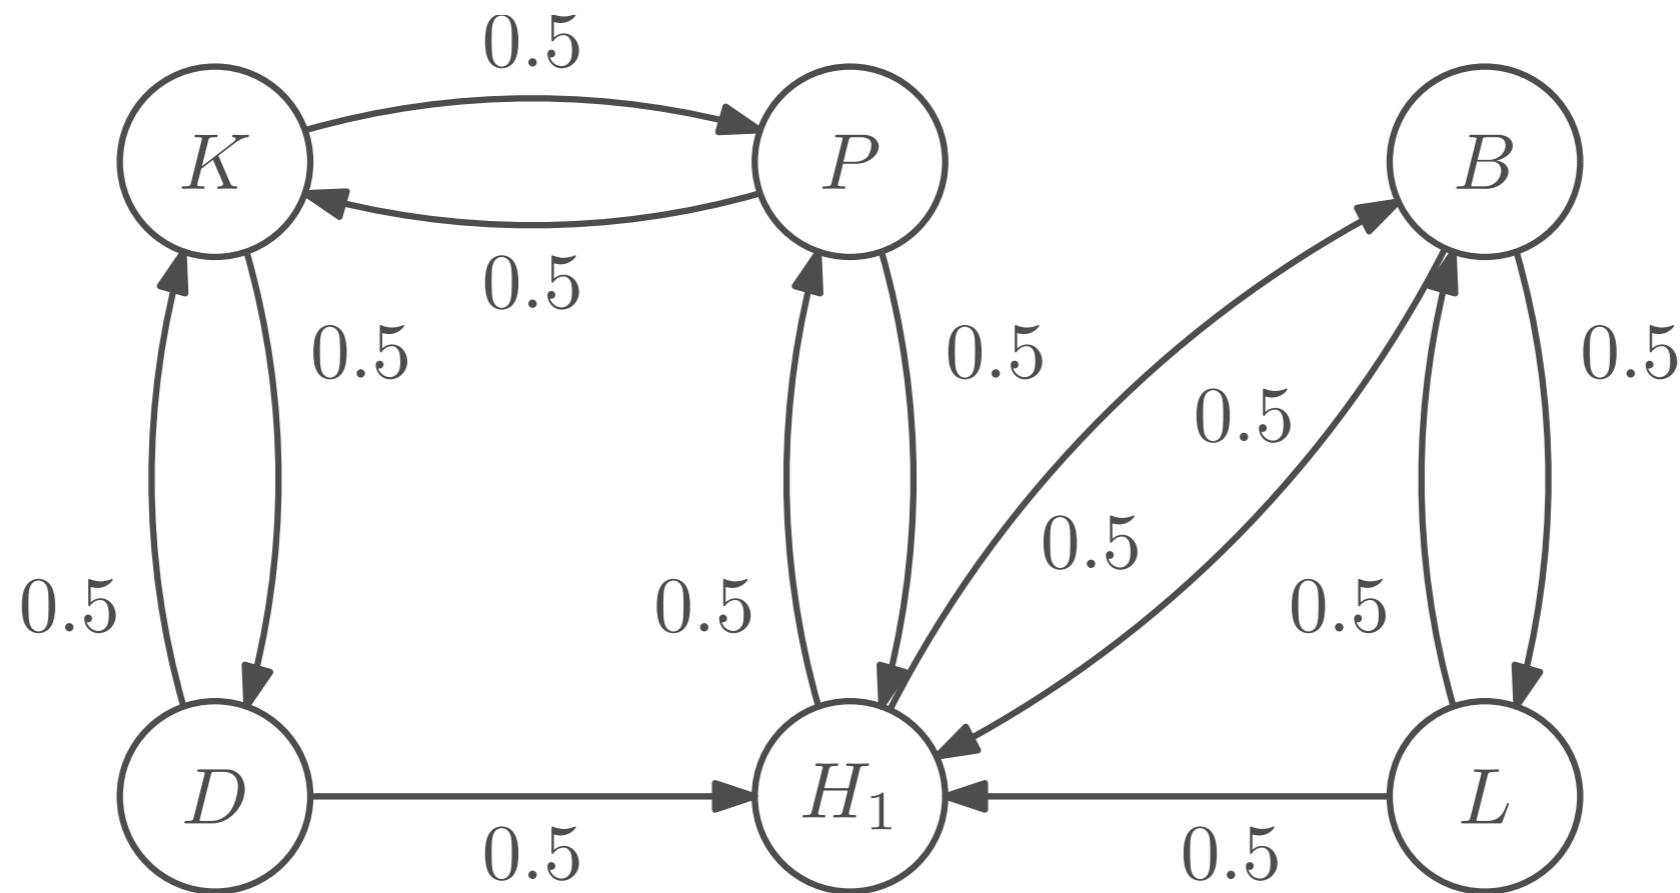
- Possible wall configurations:
- $\mathcal{Z} = \{(0, 0, 0, 0), (0, 0, 0, 1), (0, 0, 1, 0), (0, 0, 1, 1), (0, 1, 0, 0), (0, 1, 0, 1), (0, 1, 1, 0), (0, 1, 1, 1), (1, 0, 0, 0), (1, 0, 0, 1), (1, 0, 1, 0), (1, 0, 1, 1), (1, 1, 0, 0), (1, 1, 0, 1), (1, 1, 1, 0), (1, 1, 1, 1)\}$



(U, D, L, R)

# Transition probabilities

- We can represent the underlying state transitions using a transition diagram:



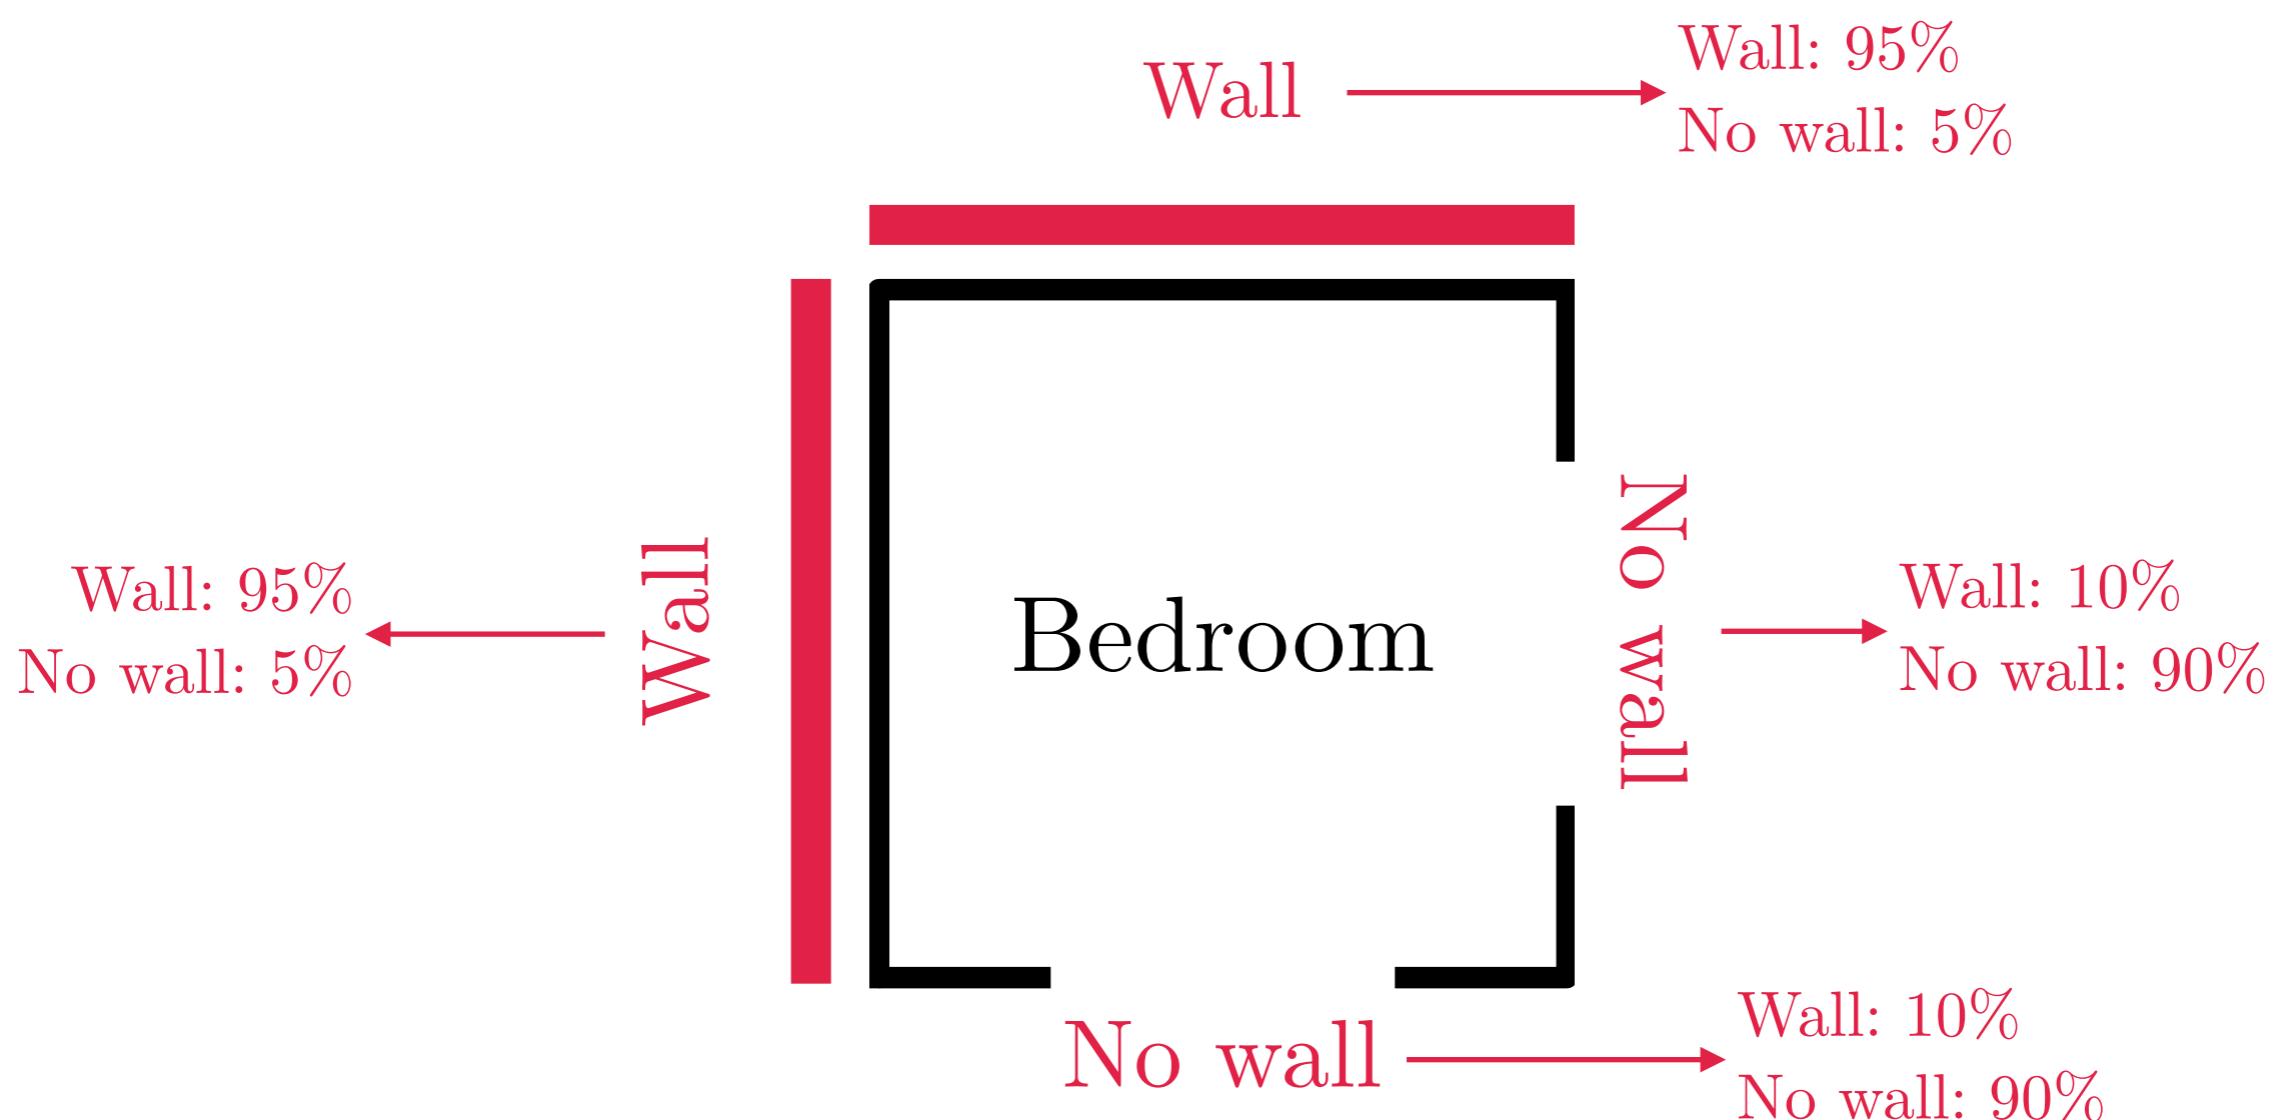
# Transition probabilities

- ... or as a matrix:

$$P = \begin{bmatrix} 0.0 & 0.5 & 0.5 & 0.0 & 0.0 & 0.0 \\ 0.5 & 0.0 & 0.0 & 0.5 & 0.0 & 0.0 \\ 0.5 & 0.0 & 0.0 & 0.5 & 0.0 & 0.0 \\ 0.0 & 0.5 & 0.0 & 0.0 & 0.5 & 0.0 \\ 0.0 & 0.0 & 0.0 & 0.5 & 0.0 & 0.5 \\ 0.0 & 0.0 & 0.0 & 0.5 & 0.5 & 0.0 \end{bmatrix}$$

# Observation probabilities

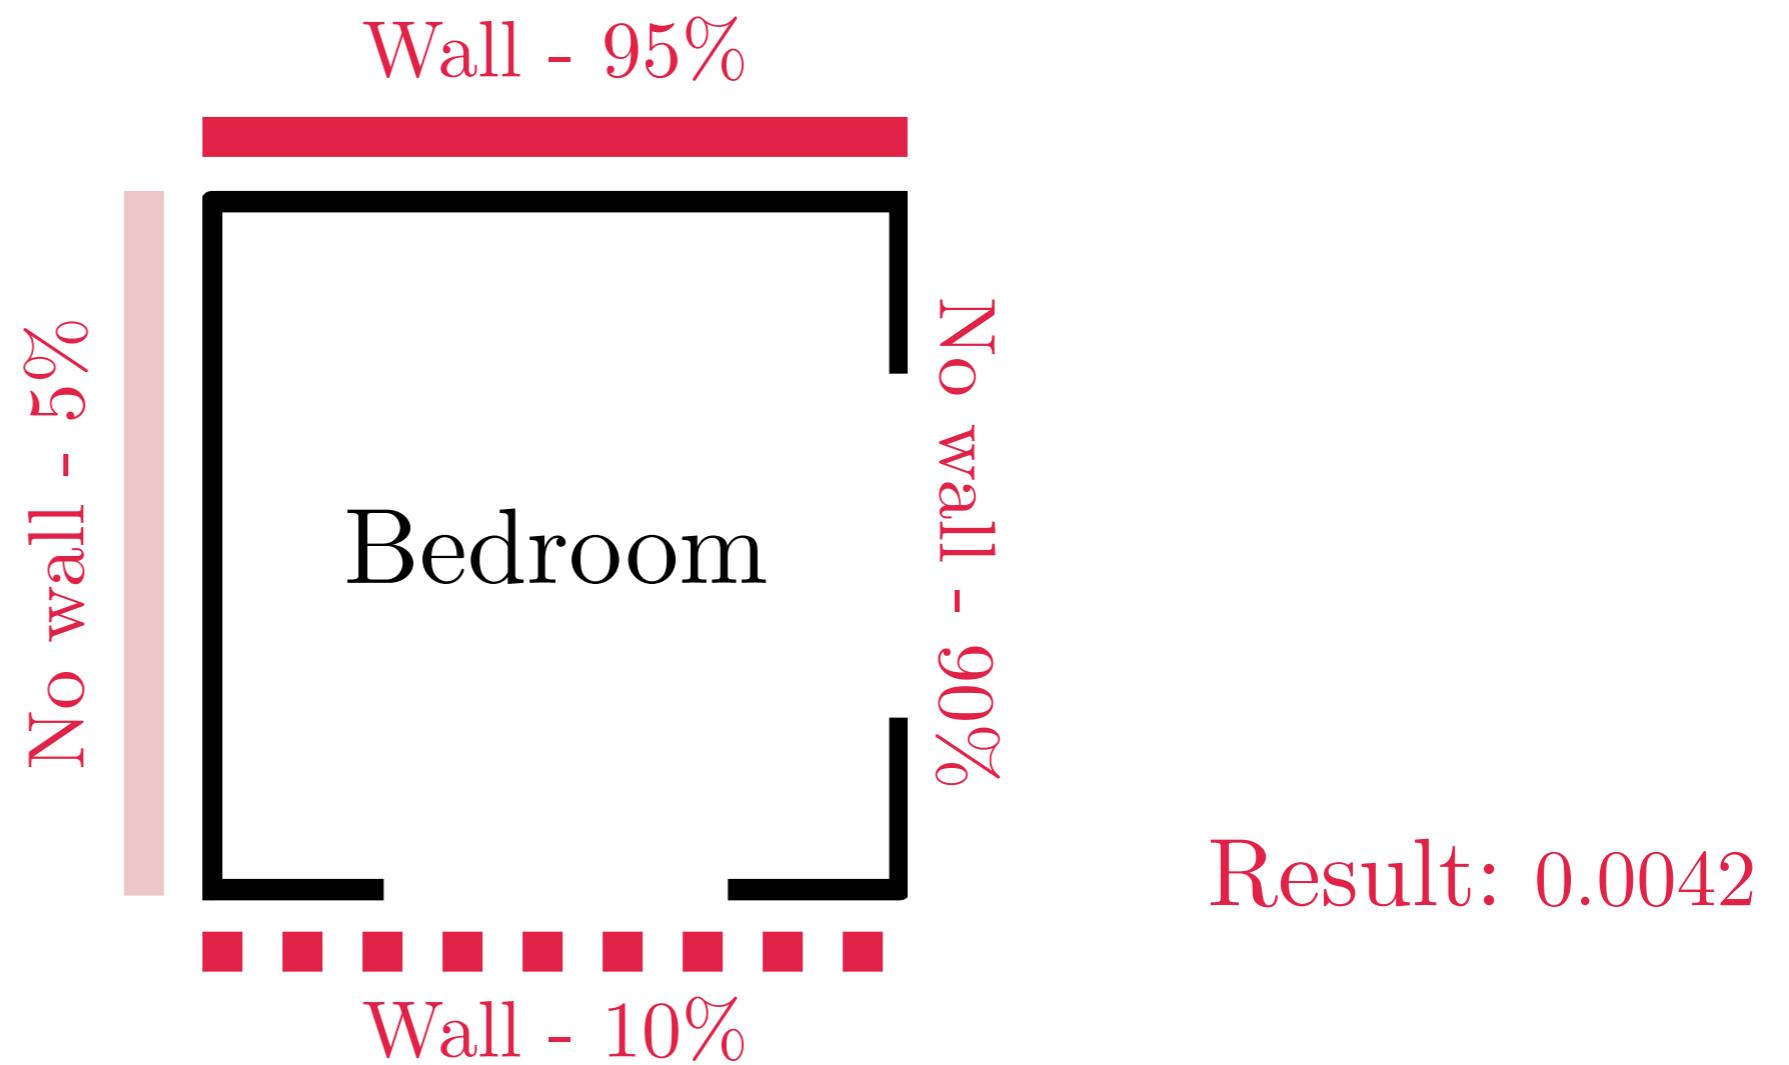
- The observation probabilities can be computed for each wall independently:



# Observation probabilities

- For example:

$$\mathbb{P}[z_t = (\underline{1}, \underline{1}, 0, 0) \mid x_t = B]$$



# Observation probabilities

- This yields a large matrix:

$$O = \begin{bmatrix} 0.0018 & 0.0004 & 0.0342 & \dots & 0.0009 & 0.0722 & 0.0180 \\ 0.0001 & 0.0021 & 0.0021 & \dots & 0.0045 & 0.0045 & 0.0857 \\ 0.0020 & 0.0002 & 0.0384 & \dots & 0.0005 & 0.0812 & 0.0090 \\ 0.5184 & 0.0576 & 0.1296 & \dots & 0.0016 & 0.0036 & 0.0004 \\ 0.0020 & 0.0385 & 0.0002 & \dots & 0.0812 & 0.0004 & 0.0090 \\ 0.0324 & 0.0036 & 0.0081 & \dots & 0.0076 & 0.0171 & 0.0019 \end{bmatrix}.$$

# HMM problems

# HMM problems

- Estimation:
  - Estimate states from sequences of observations
- ~~Inference:~~
  - Infer an HMM model from data

# Estimation

- Filtering:  
Forward algorithm
  - Given a sequence of observations, estimate the final state
- Smoothing:  
Viterbi algorithm
  - Given a sequence of observations, estimate the sequence of states
- Prediction:
  - Given a sequence of observations, predict future states



# Filtering

# Filtering

- We are given a sequence of observations  $\mathbf{z}_{0:T}$
- We want to estimate

What is the  
final state?



$$\mathbb{P}_{\mu_0} [\mathbf{x}_T = \mathbf{x} \mid \mathbf{z}_{0:t} = \mathbf{z}_{0:T}]$$

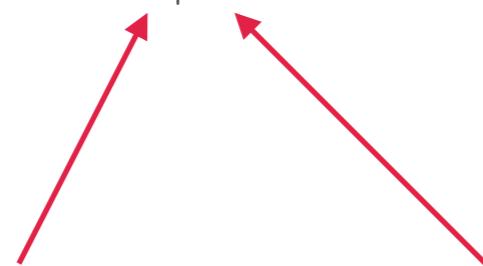
where  $\mu_0$  is the initial distribution, i.e.,

$$\mu_0(x) = \mathbb{P} [\mathbf{x}_0 = \mathbf{x}]$$

# Filtering

- For simplicity, write

$$\mu_{T|0:T}(x) = \mathbb{P}_{\mu_0} [\mathbf{x}_T = x \mid \mathbf{z}_{0:t} = \mathbf{z}_{0:T}]$$

  
Distribution at time  $T$       Given observations  $0:T$

# Forward mapping

## Forward mapping

Given a sequence of observations  $\mathbf{z}_{0:t}$ , the forward mapping

$\alpha_t : \mathcal{X} \rightarrow \mathbb{R}$  is defined for each  $t$  as

$$\alpha_t(x) = \mathbb{P}_{\mu_0} [\mathbf{x}_t = x, \mathbf{z}_{0:t} = \mathbf{z}_{0:t}]$$

Joint  
probability



How likely is it that I end up in  $x$  having observed  $\mathbf{z}_{0:t}$ ?

# So what?

- Forward mapping has several useful properties
  1. We can compute  $\mu_{T|0:T}$  from  $\alpha_T$ :

$$\mu_{T|0:T}(x) = \frac{\alpha_T(x)}{\sum_{y \in \mathcal{X}} \alpha_T(y)}$$

It's what we  
want (up to a  
factor)!

# So what?

- Forward mapping has several useful properties
  1. We can compute  $\mu_{T|0:T}$  from  $\alpha_T$
  2. The forward mapping can be computed recursively:

$$\alpha_T(x) = \mathbf{O}(z_T \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{T-1}(y)$$

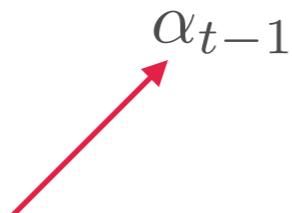
# Vectorizing the forward update

- Let's look at the update

$$\alpha_t(x) = \mathbf{O}(z_t \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{t-1}(y)$$

using vectors:

$$[\square \square \square \cdots \square]$$



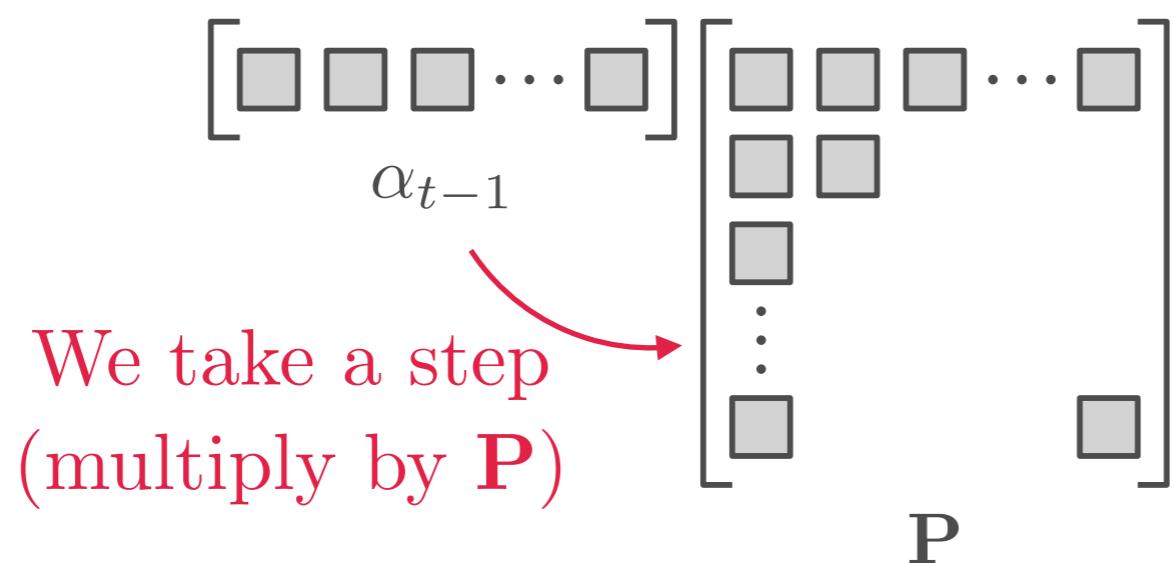
We start with our  
vector at time  $t - 1$   
(assume row vector)

# Vectorizing the forward update

- Let's look at the update

$$\alpha_t(x) = \mathbf{O}(z_t \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{t-1}(y)$$

using vectors:



# Vectorizing the forward update

- Let's look at the update

$$\alpha_t(x) = \mathbf{O}(z_t \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{t-1}(y)$$

using vectors:

$$\begin{bmatrix} \square & \square & \square & \cdots & \square \end{bmatrix}$$
$$\alpha_{t-1} \mathbf{P}$$

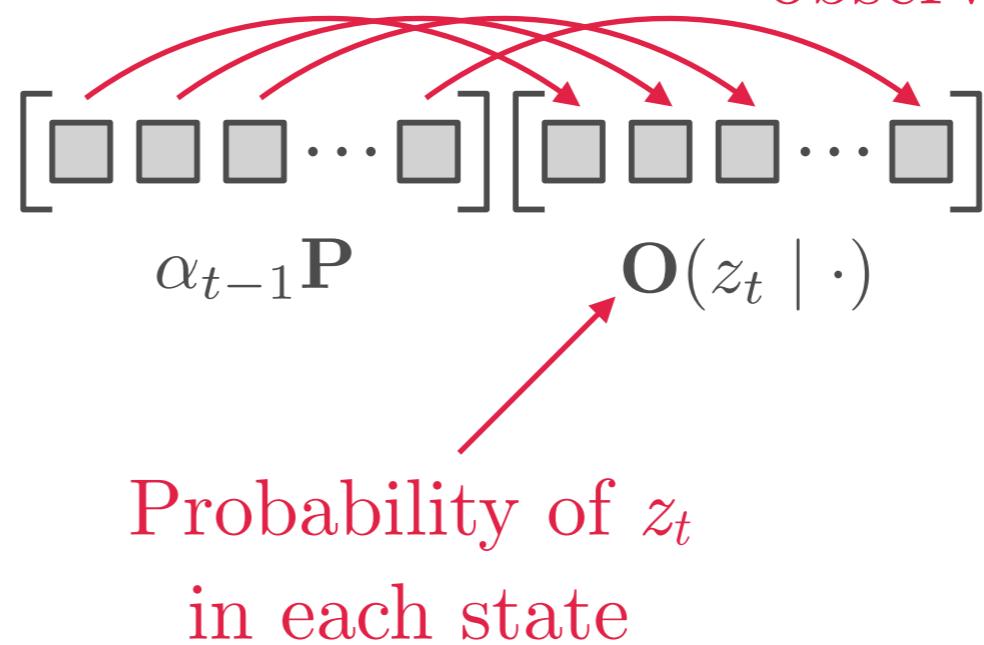
# Vectorizing the forward update

- Let's look at the update

$$\alpha_t(x) = \mathbf{O}(z_t | x) \sum_{y \in \mathcal{X}} \mathbf{P}(x | y) \alpha_{t-1}(y)$$

using vectors:

We multiply each value  
by the corresponding  
observation probability

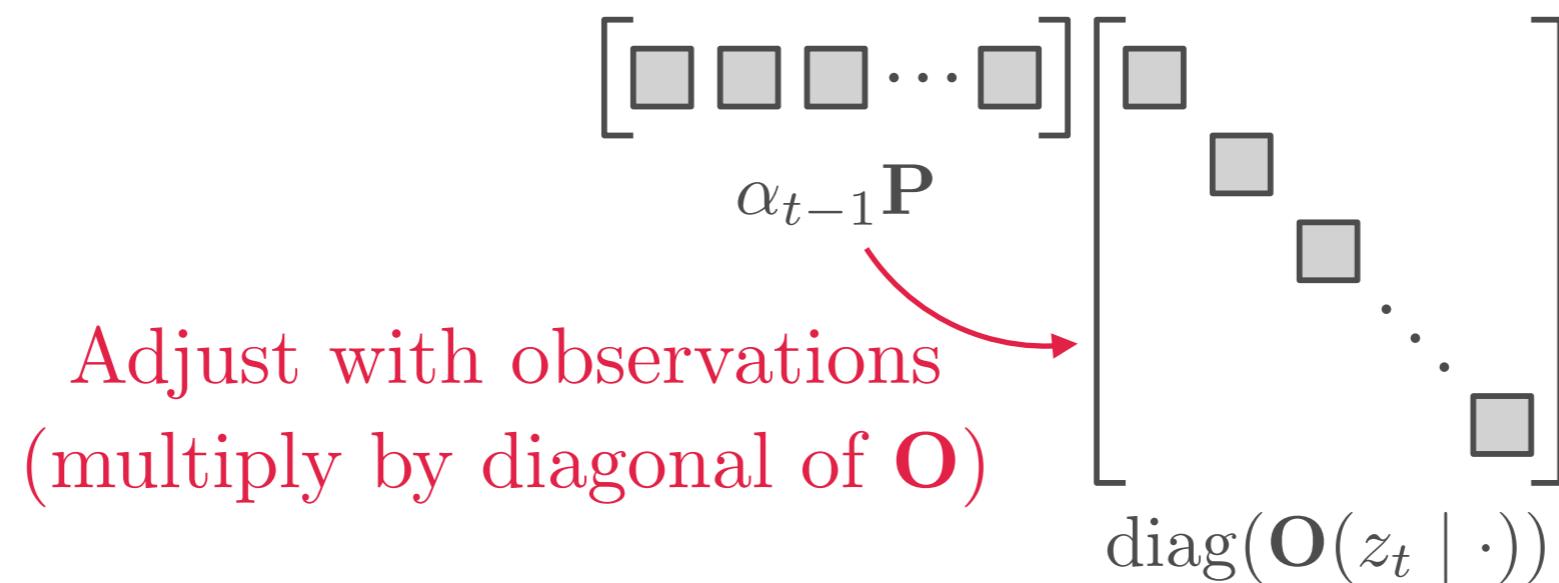


# Vectorizing the forward update

- Let's look at the update

$$\alpha_t(x) = \mathbf{O}(z_t \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{t-1}(y)$$

using vectors:

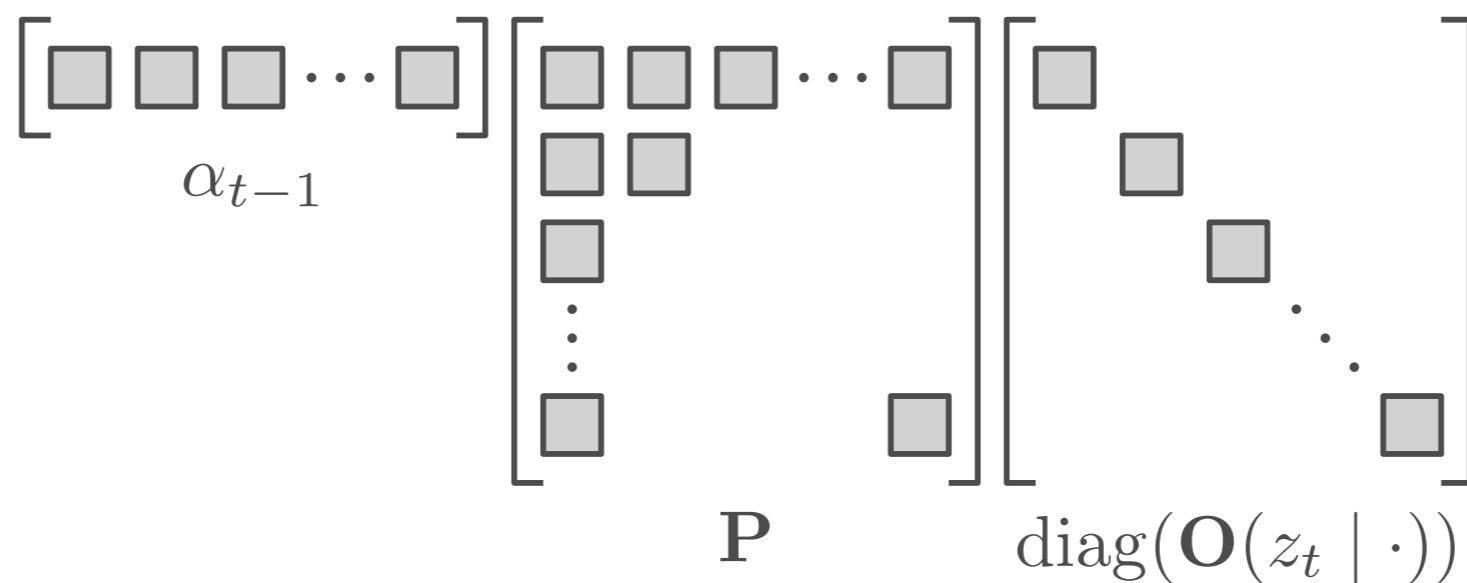


# Vectorizing the forward update

- Let's look at the update

$$\alpha_t(x) = \mathbf{O}(z_t \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{t-1}(y)$$

using vectors:

$$\begin{matrix} [\square \square \square \cdots \square] \\ \alpha_{t-1} \end{matrix} \begin{matrix} [\square \square \square \cdots \square] \\ \mathbf{P} \end{matrix} \begin{matrix} [\square \quad \quad \quad \cdots \quad \quad \quad \square] \\ \text{diag}(\mathbf{O}(z_t \mid \cdot)) \end{matrix}$$


# Vectorizing the forward update

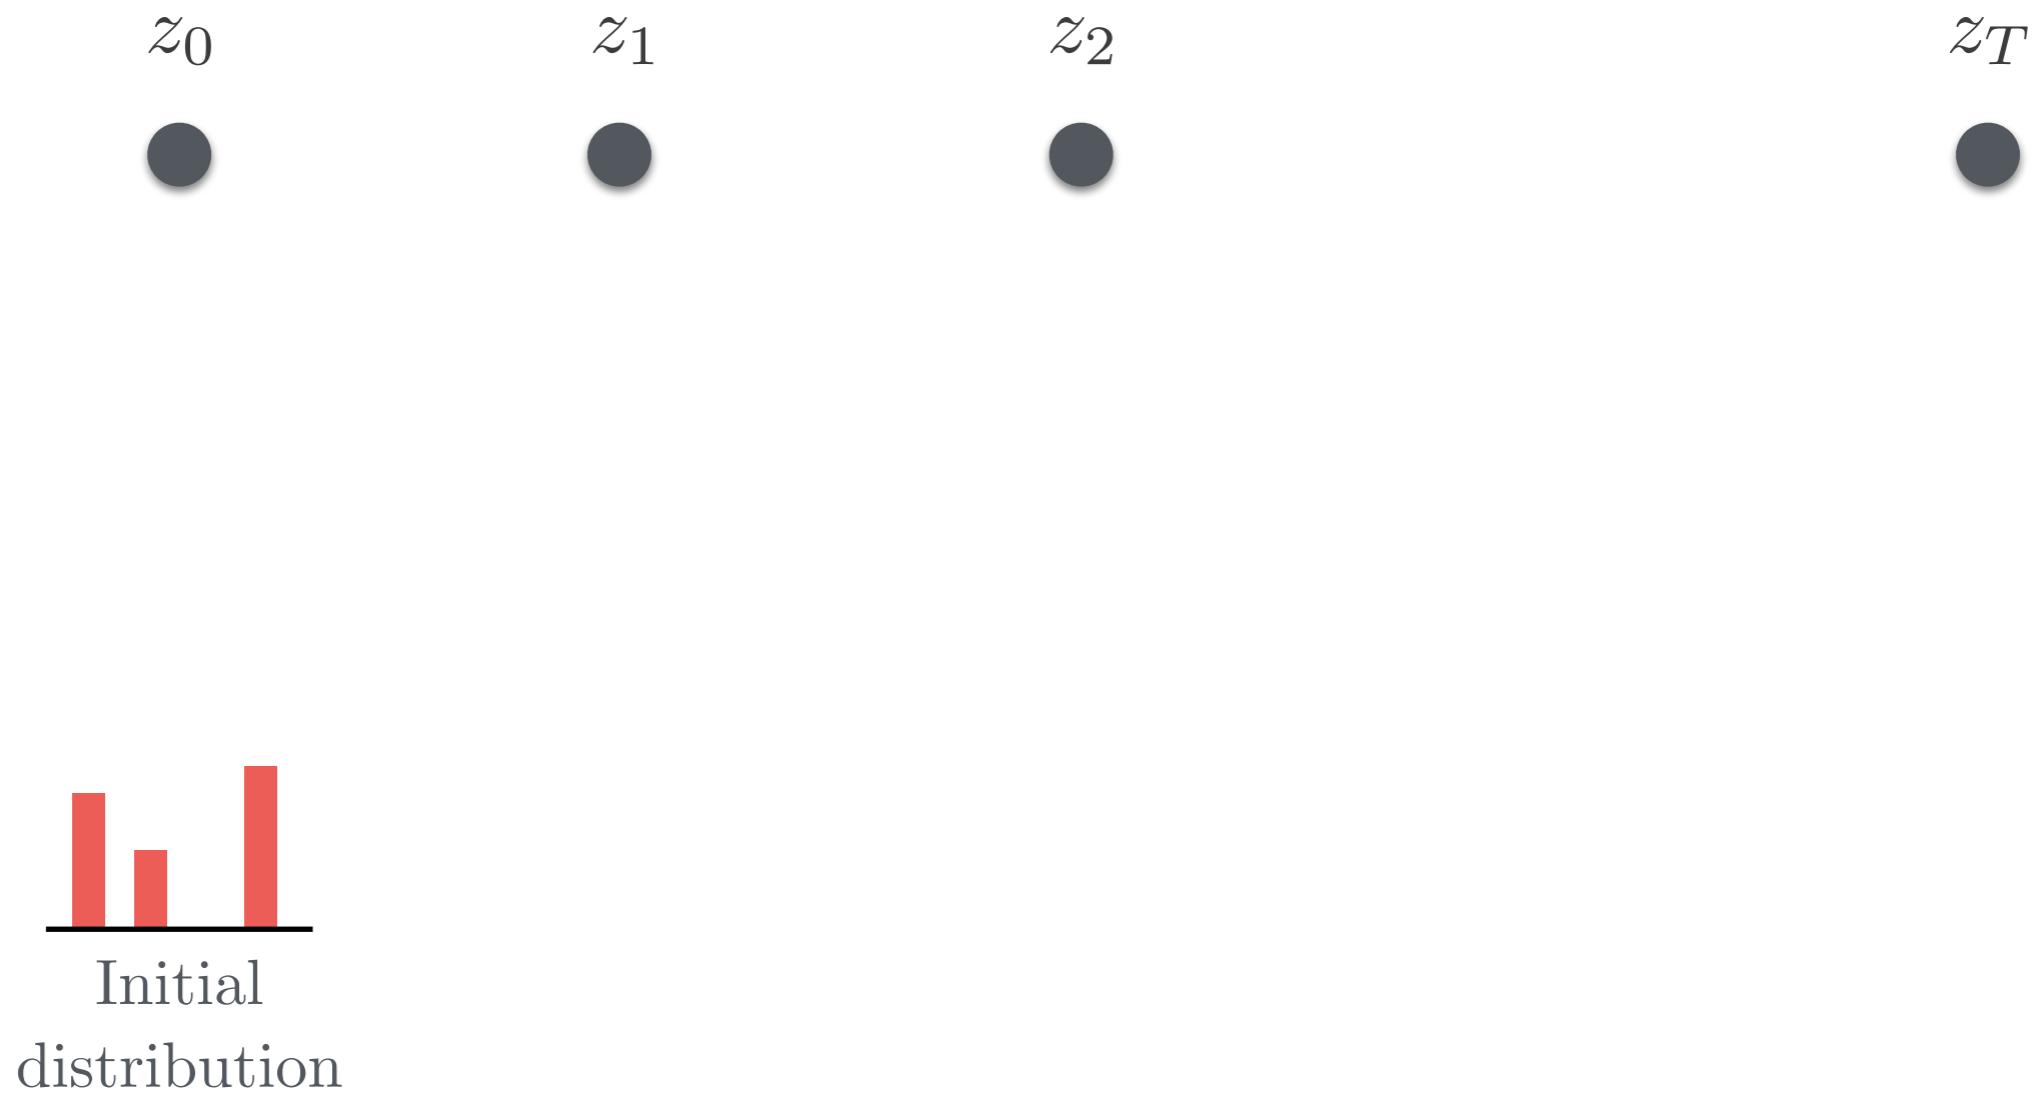
- Let's look at the update

$$\alpha_t(x) = \mathbf{O}(z_t \mid x) \sum_{y \in \mathcal{X}} \mathbf{P}(x \mid y) \alpha_{t-1}(y)$$

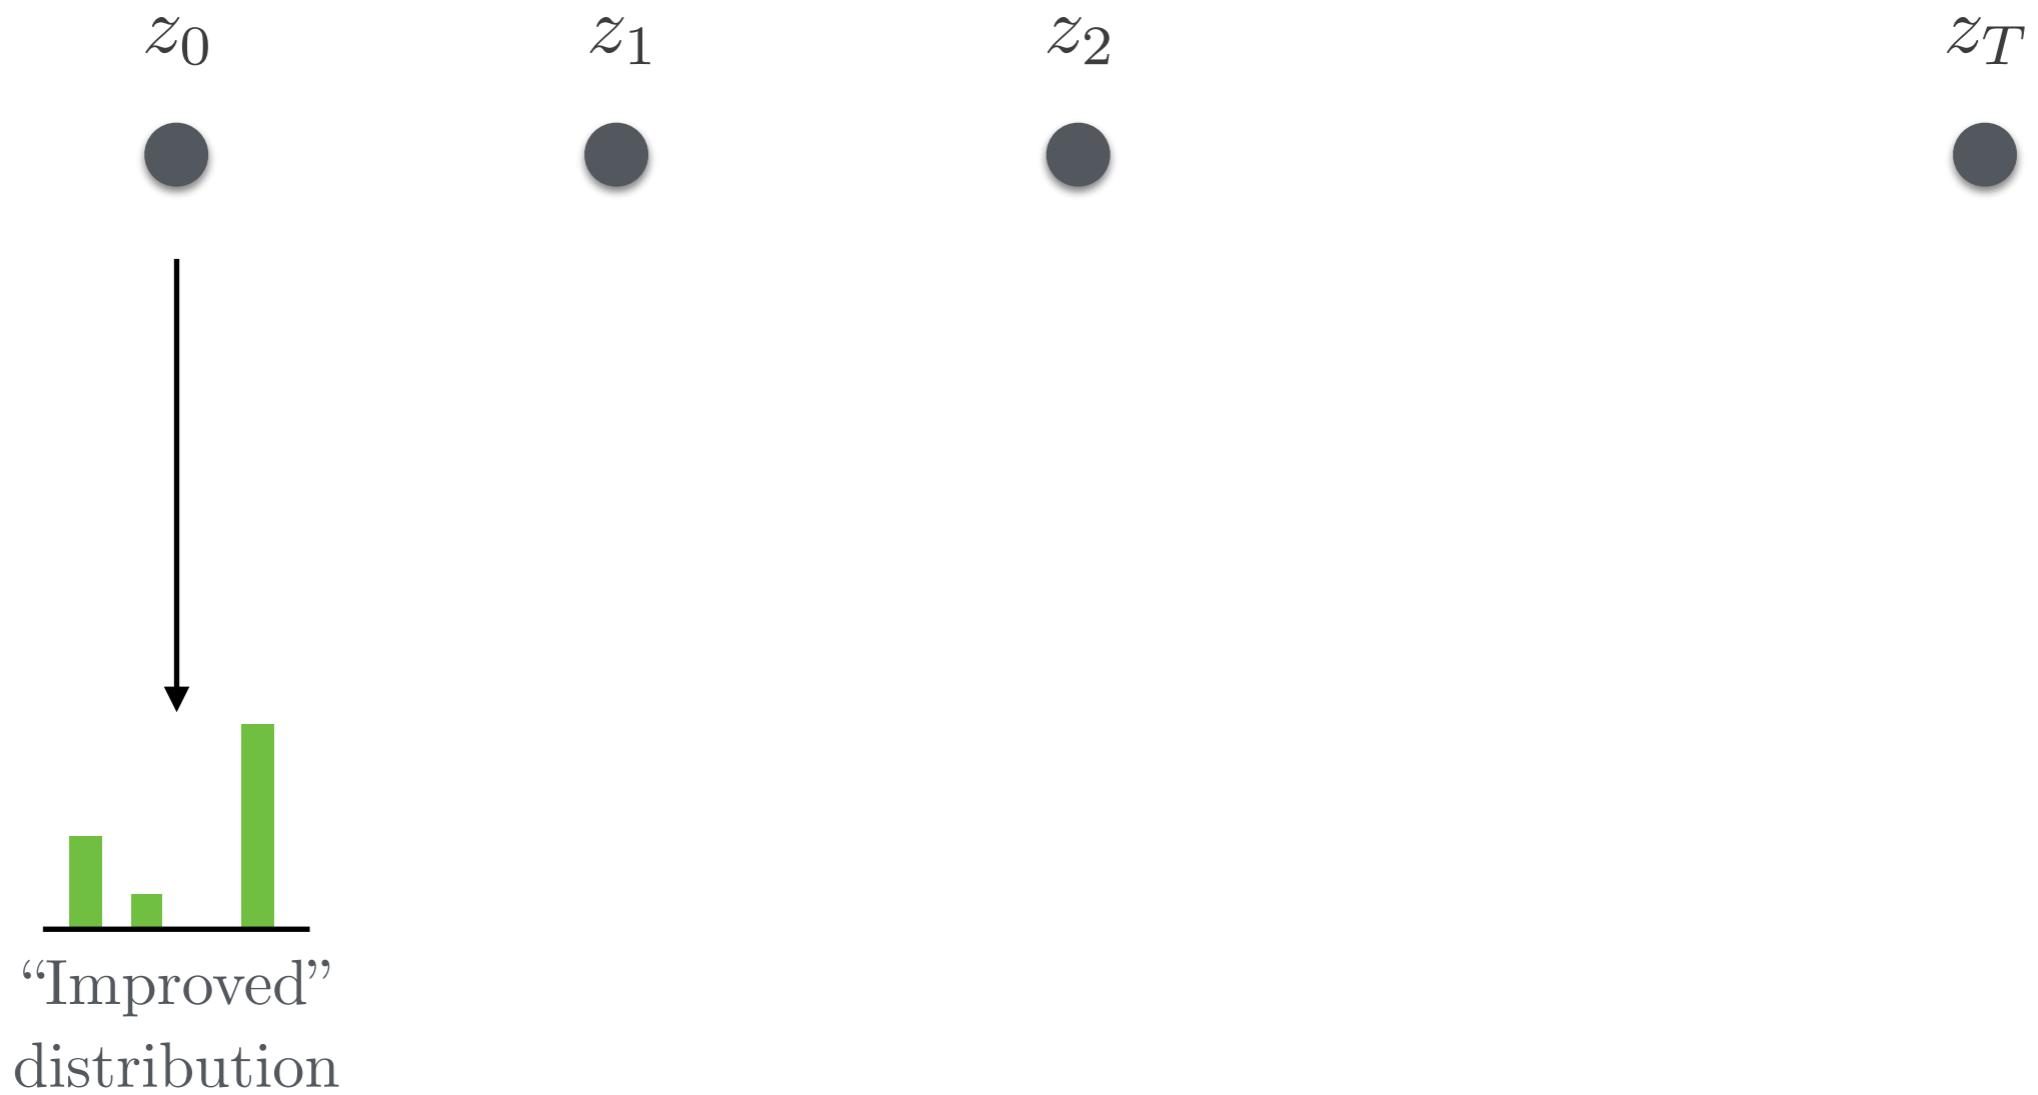
using a **column** vector:

$$\begin{bmatrix} \text{diag}(\mathbf{O}(z_t \mid \cdot)) \\ \vdots \\ \text{diag}(\mathbf{O}(z_t \mid \cdot)) \end{bmatrix} \mathbf{P}^\top \begin{bmatrix} \alpha_{t-1} \\ \vdots \\ \alpha_{t-1} \end{bmatrix}$$

# Forward algorithm

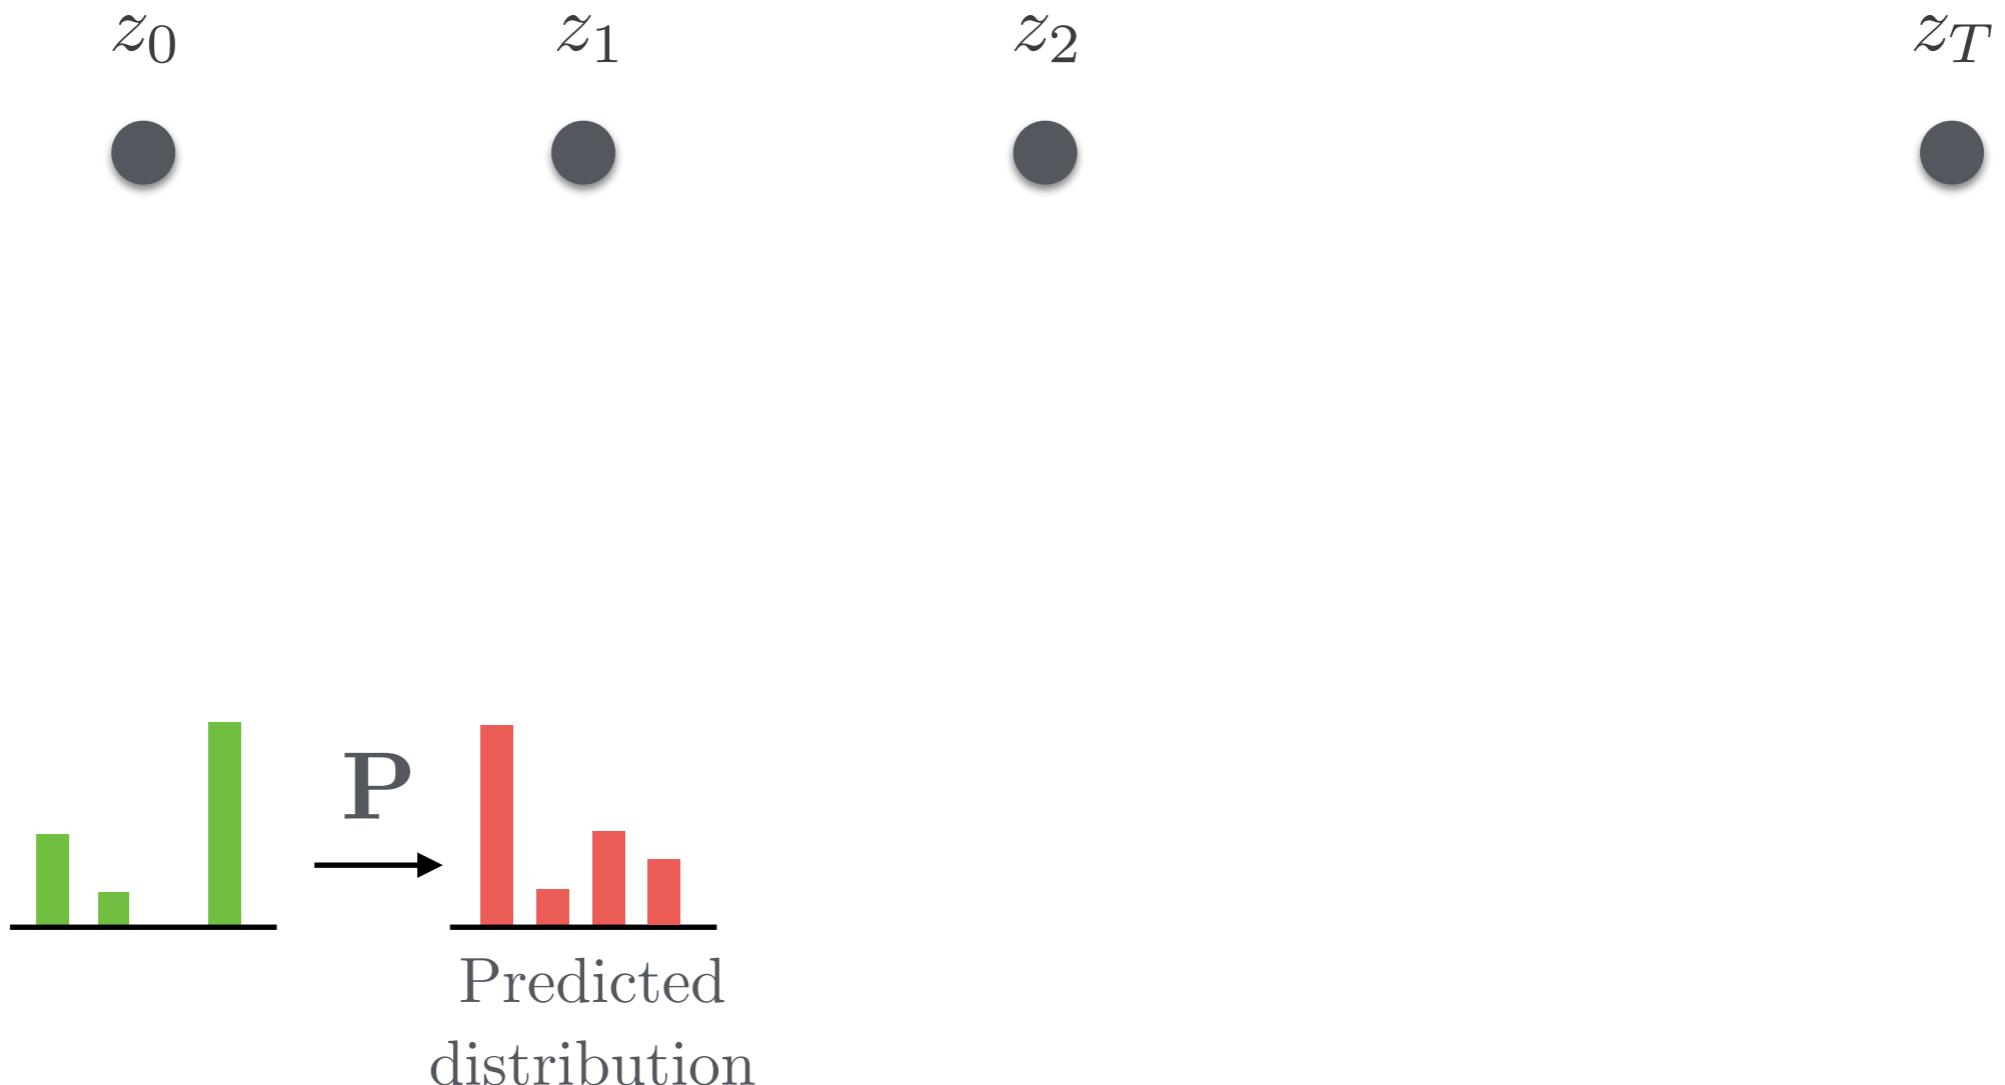


# Forward algorithm



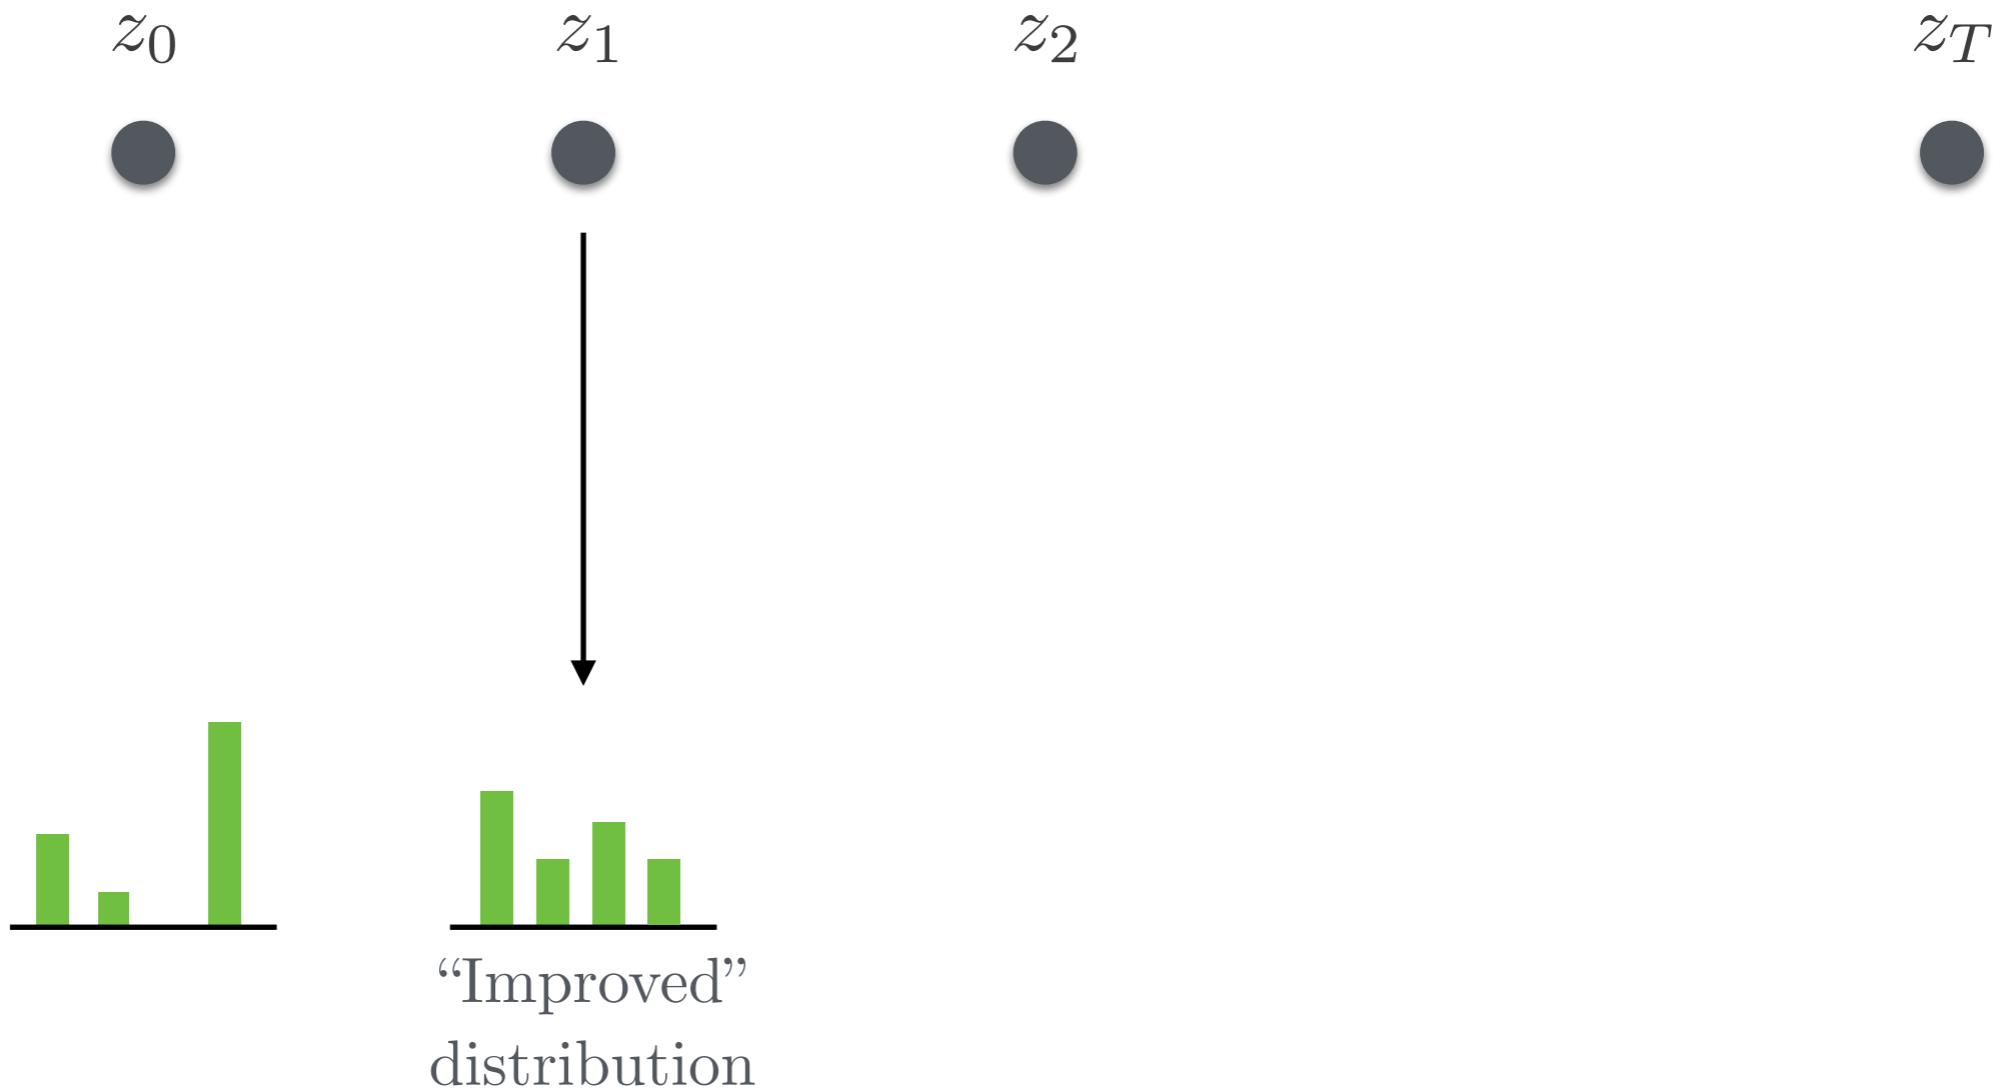
$$\text{diag}(\mathbf{O}(z_0 \mid \cdot) \boldsymbol{\mu}_0^\top$$

# Forward algorithm



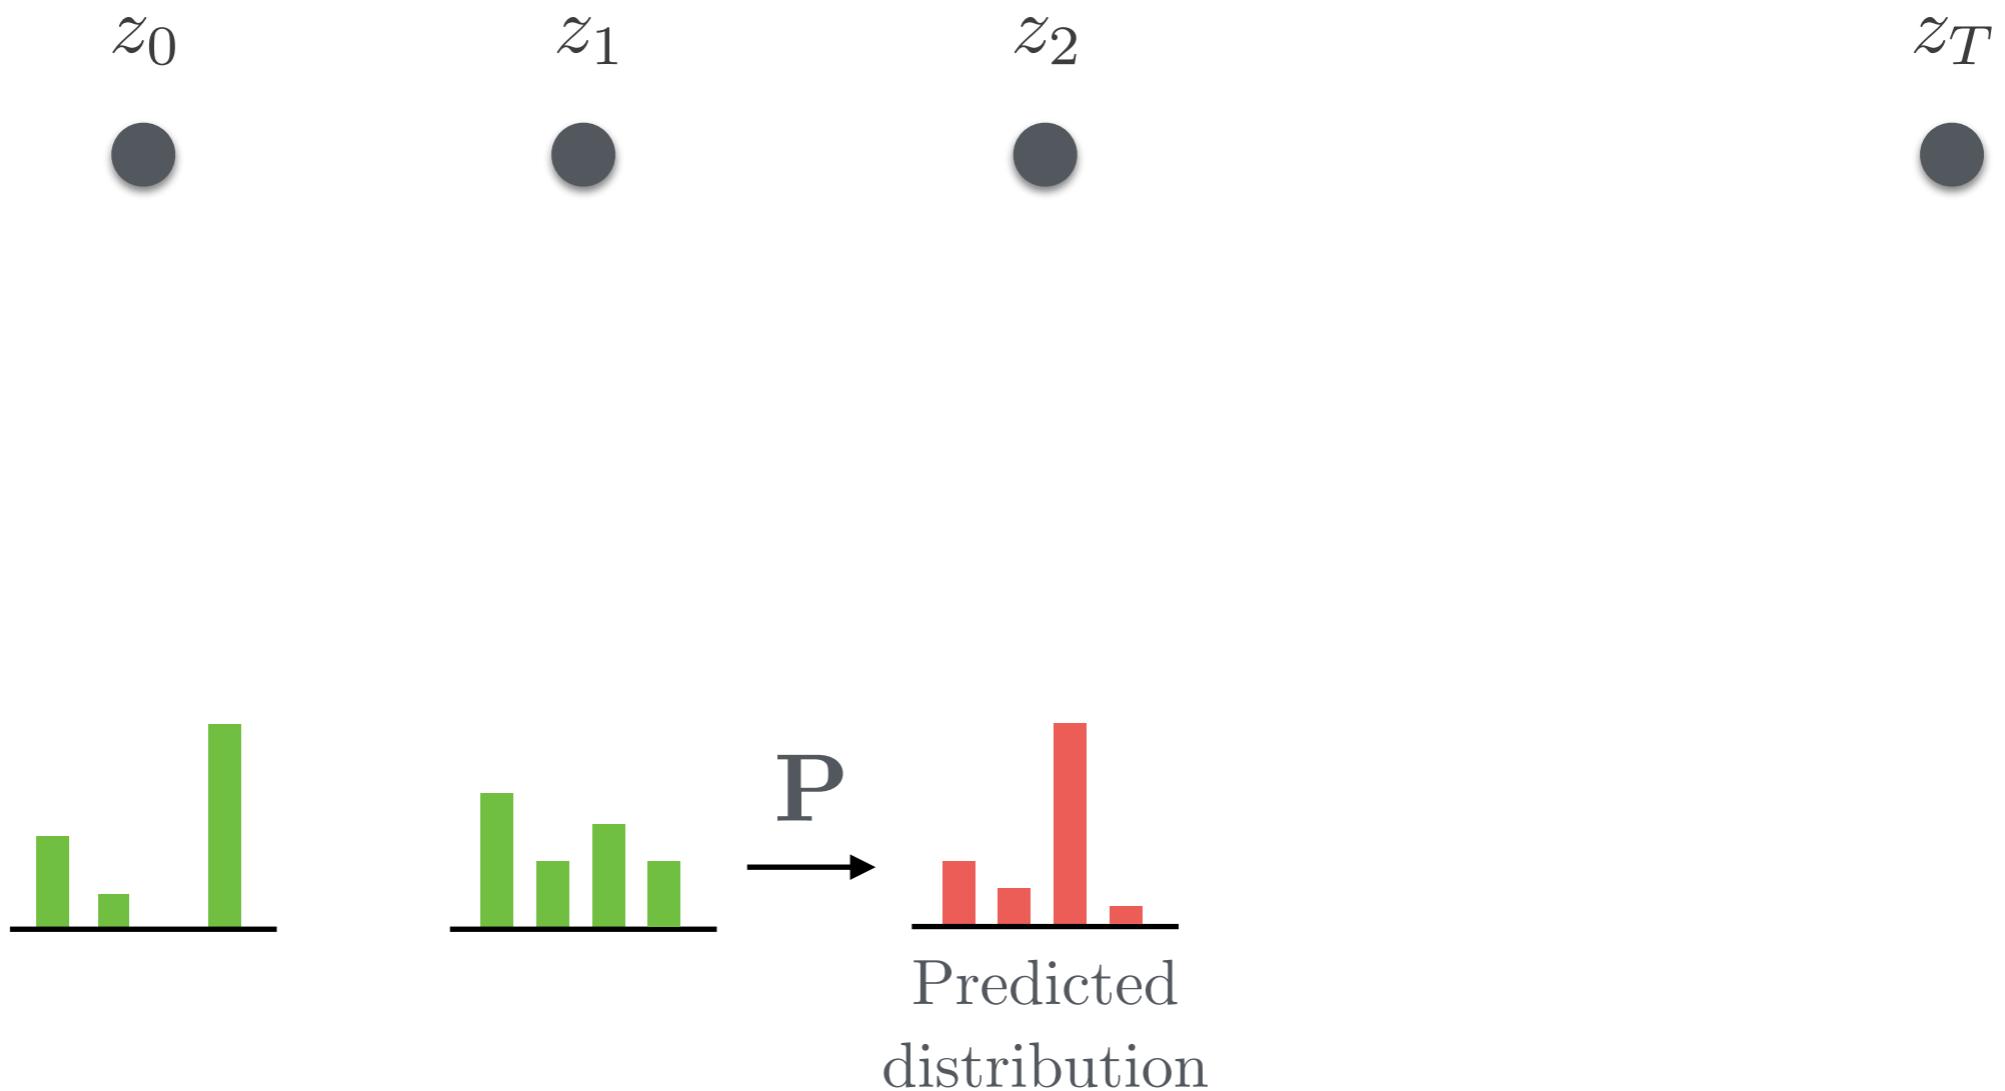
$$\mathbf{P}^\top \boldsymbol{\alpha}_0$$

# Forward algorithm



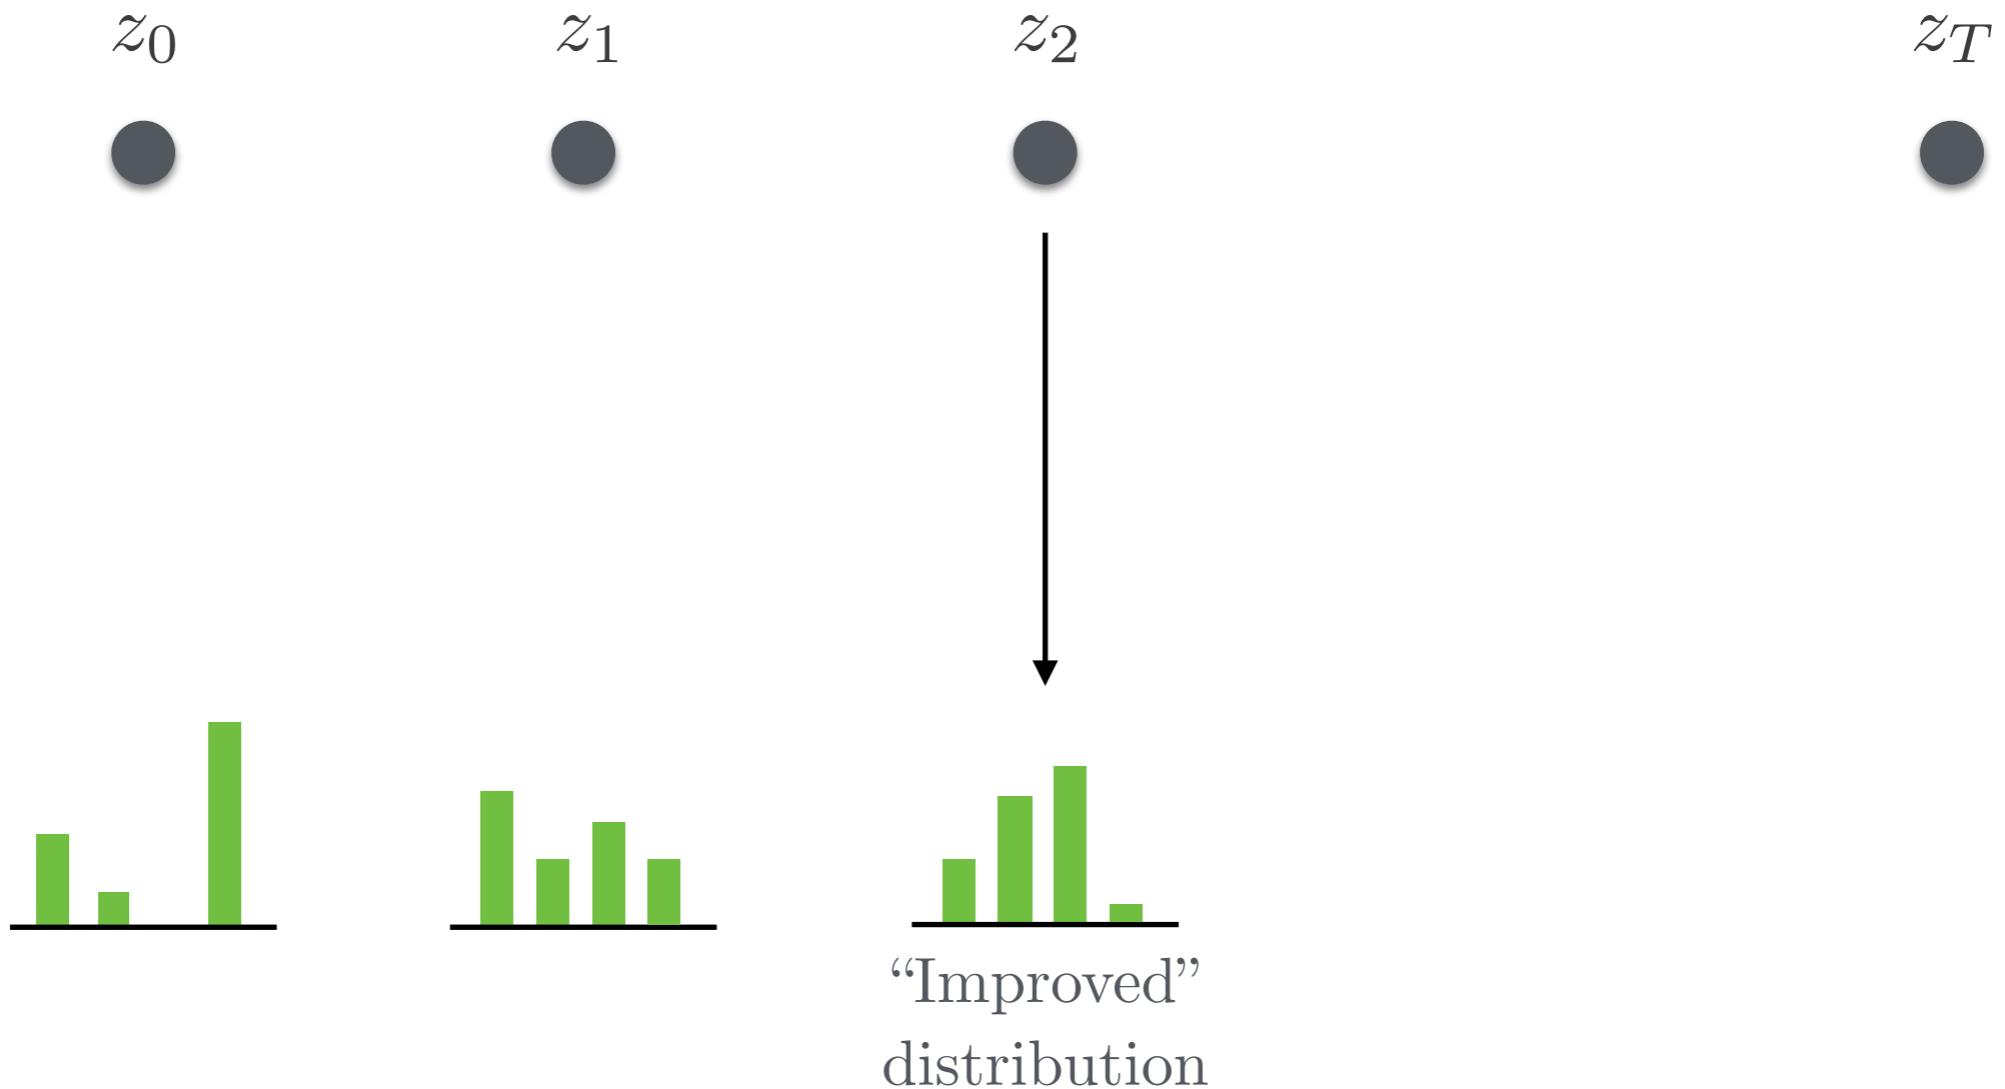
$$\text{diag}(\mathbf{O}(z_1 \mid \cdot)) \mathbf{P}^\top \boldsymbol{\alpha}_0$$

# Forward algorithm



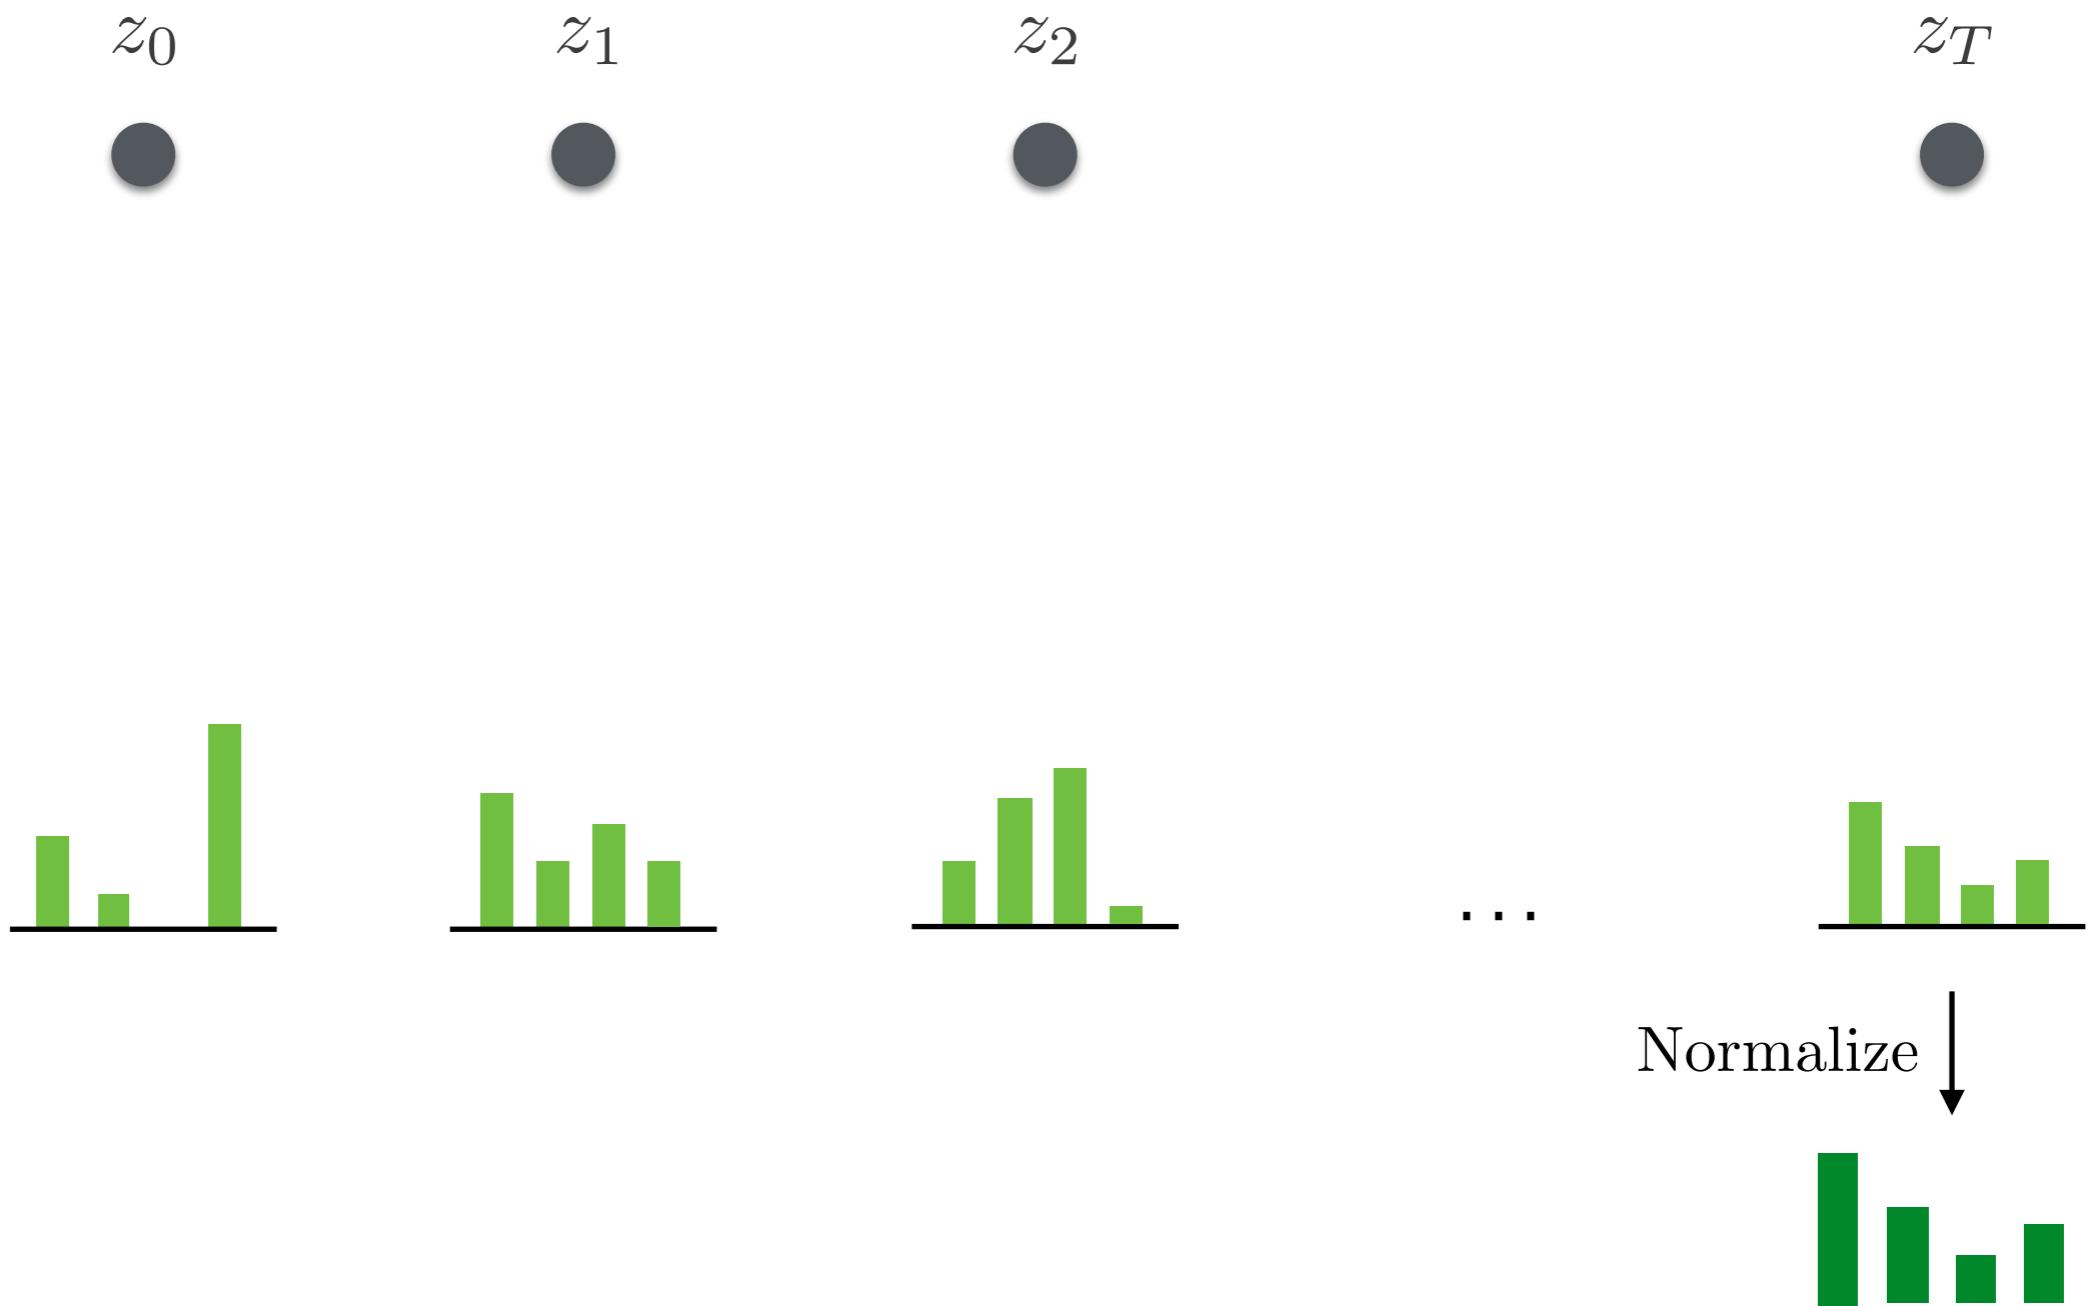
$$\mathbf{P}^\top \boldsymbol{\alpha}_1$$

# Forward algorithm



$$\text{diag}(\mathbf{O}(z_2 \mid \cdot)) \mathbf{P}^\top \boldsymbol{\alpha}_1$$

# Forward algorithm



# Forward algorithm

Given observation sequence  $z_{0:T}$

1. Multiply initial distribution by  $\mathbf{O}(z_0 | :)$

Improve  
initial  
distribution



# Forward algorithm

Given observation sequence  $z_{0:T}$

1. Multiply initial distribution by  $\mathbf{O}(z_0 | :)$
  2. At each time step:
    - a. Multiply current distribution by  $\mathbf{P}$
- Predict  
1-step  
move
- 

# Forward algorithm

Given observation sequence  $z_{0:T}$

1. Multiply initial distribution by  $\mathbf{O}(z_0 | :)$
  2. At each time step:
    - a. Multiply current distribution by  $\mathbf{P}$
    - b. Multiply by  $\mathbf{O}(z_t | :)$
- Check  
prediction with  
observation
- 

# Forward algorithm

Given observation sequence  $z_{0:T}$

1. Multiply initial distribution by  $\mathbf{O}(z_0 | :)$
2. At each time step:
  - a. Multiply current distribution by  $\mathbf{P}$
  - b. Multiply by  $\mathbf{O}(z_t | :)$
3. Normalize

# Forward algorithm

**Require:** Observation sequence  $z_{0:T}$

1. Initialize  $\alpha_0 = \text{diag}(\mathbf{O}(z_0 | \cdot))\mu_0^\top$
2. for  $t = 1, \dots, T$  do

If we have an initial observation

$$\alpha_t = \text{diag}(\mathbf{O}(z_t | \cdot))\mathbf{P}^\top \alpha_{t-1}$$

3. end for

Here's our update!

4. return  $\mu_{T|0:T} = \alpha_T / \text{sum}(\alpha_T)$

# Example: The urn problem

- Suppose that

$$\mu_0 = [0.125 \quad 0.375 \quad 0.375 \quad 0.125]$$

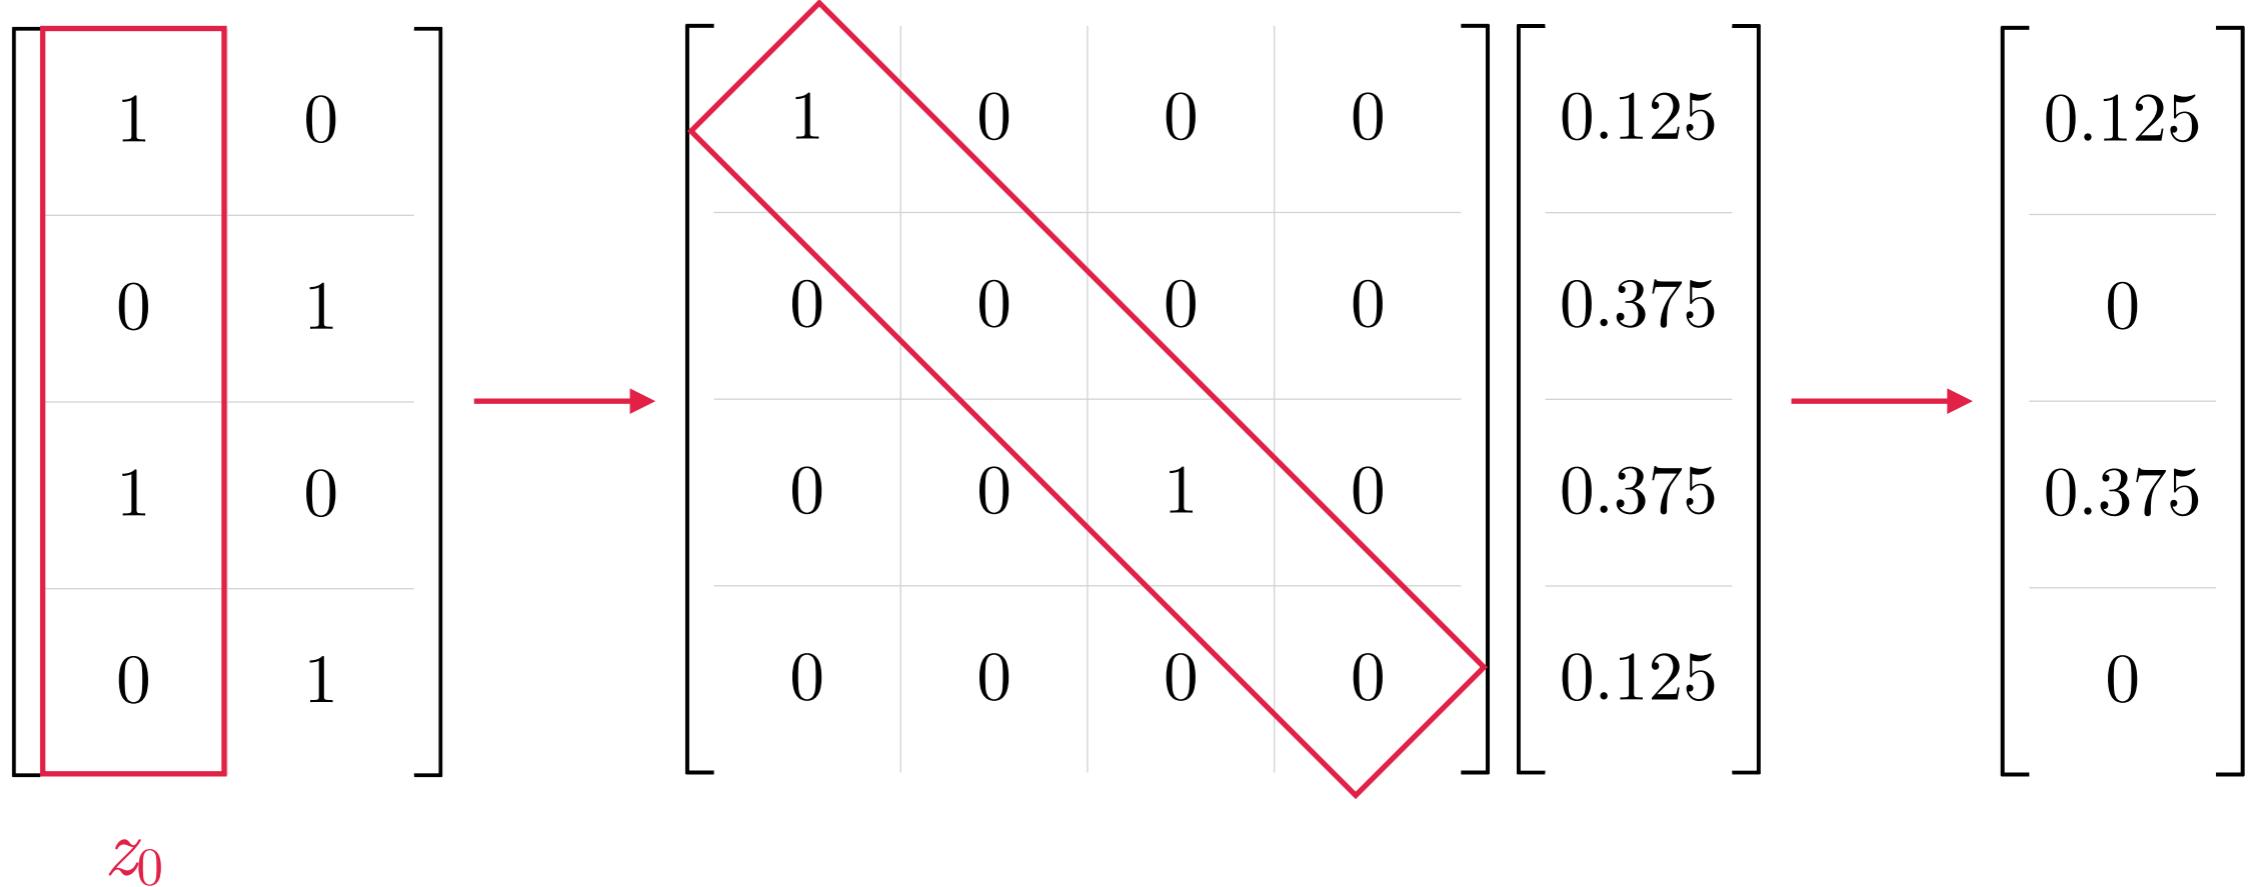
- We observe the sequence of observations

$$z_{0:2} = \{w, w, b\}$$

- What is the state at time  $t = 2$ ?

# Step 1: Initialize $a_0$

- $\alpha_0 = \text{diag}(\mathbf{O}(z_0 | \cdot) \boldsymbol{\mu}_0^\top)$



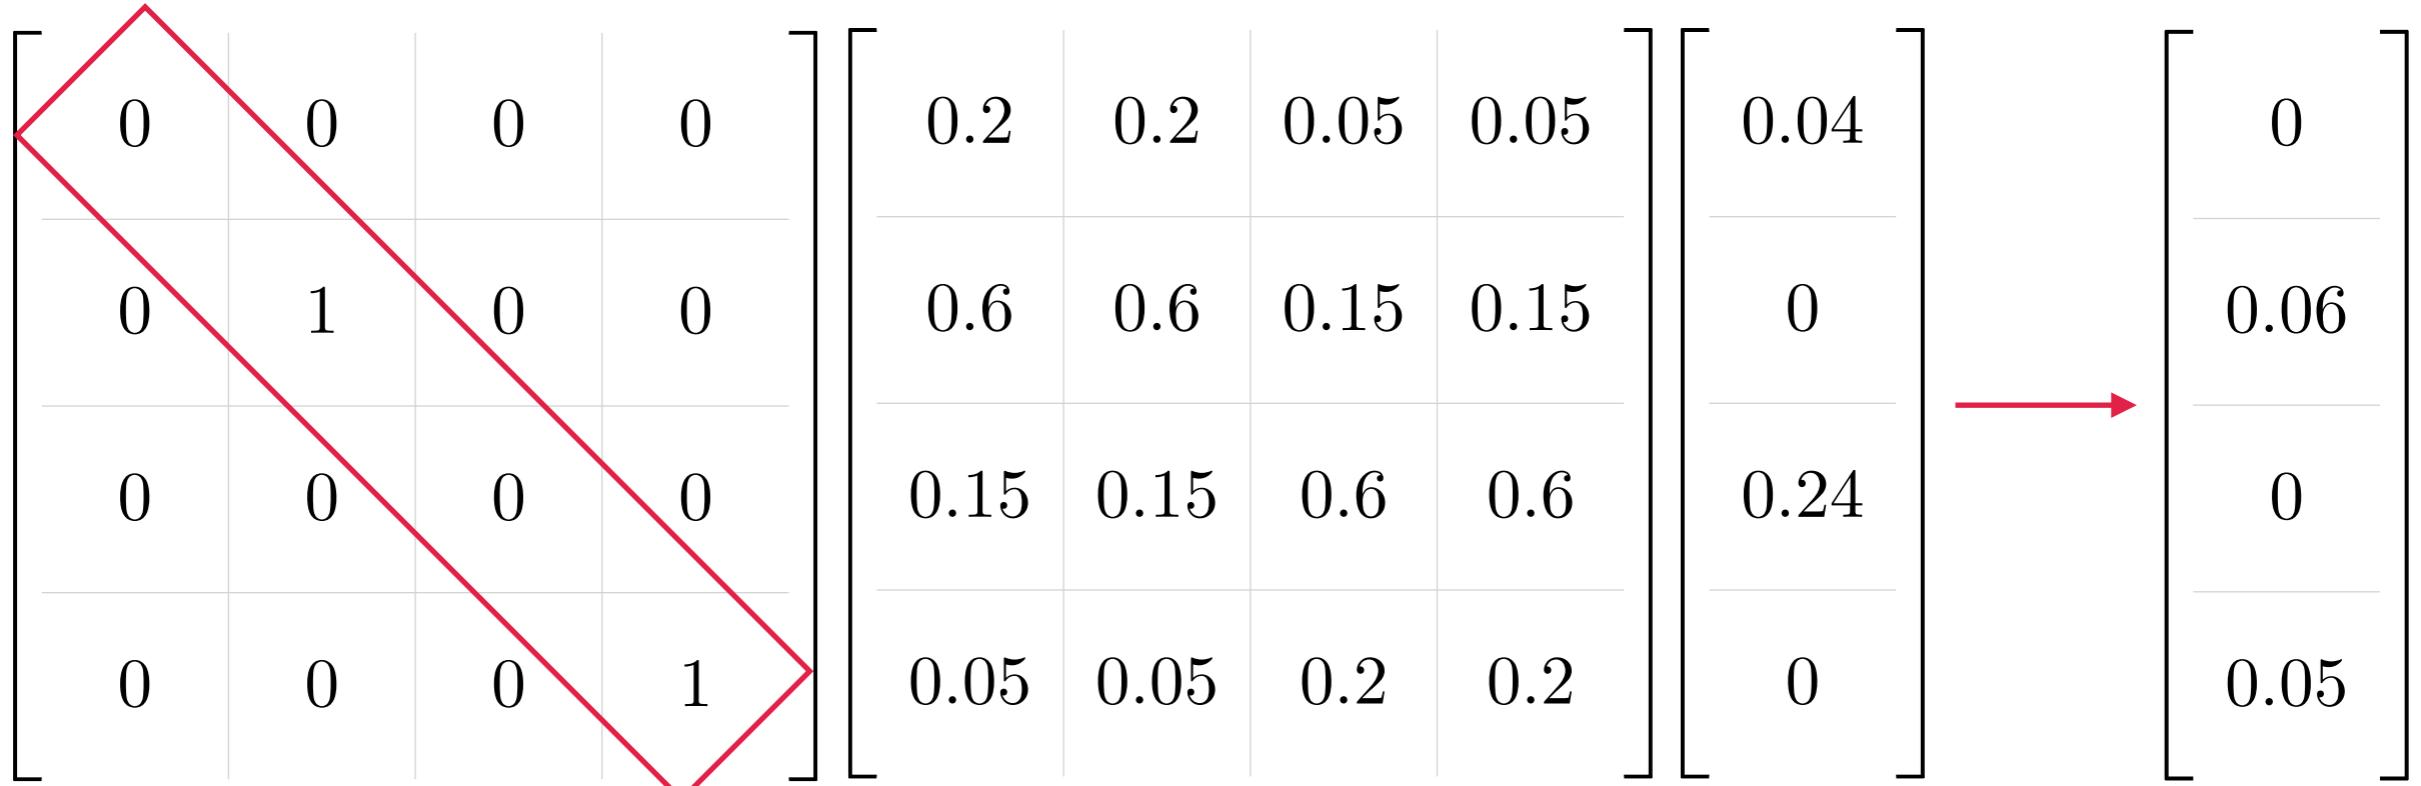
## Step 2: Compute $\alpha_1$

- $\alpha_1 = \text{diag}(\mathbf{O}(z_1 \mid \cdot))\mathbf{P}^\top \alpha_0$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0.2 & 0.2 & 0.05 & 0.05 \\ 0.6 & 0.6 & 0.15 & 0.15 \\ 0.15 & 0.15 & 0.6 & 0.6 \\ 0.05 & 0.05 & 0.2 & 0.2 \end{bmatrix} \begin{bmatrix} 0.125 \\ 0 \\ 0.375 \\ 0 \end{bmatrix} \xrightarrow{\hspace{1cm}} \begin{bmatrix} 0.04 \\ 0 \\ 0.24 \\ 0 \end{bmatrix}$$

# Step 3: Compute $\alpha_2$

- $\alpha_2 = \text{diag}(\mathbf{O}(z_2 \mid \cdot))\mathbf{P}^\top \alpha_1$



# Final step: Compute $\mu_2 | 0:2$

- We finally get:

$$\begin{aligned}\mu_{2|0:2} &= \alpha_2 / \text{sum}(\alpha_2) \\ &= [0 \quad 0.552 \quad 0 \quad 0.448]\end{aligned}$$