

# **Towards a dependable consensus: Broadcast primitives**

Highly dependable systems

Lecture 3

Lecturers: Miguel Matos and Rodrigo Miragaia Rodrigues

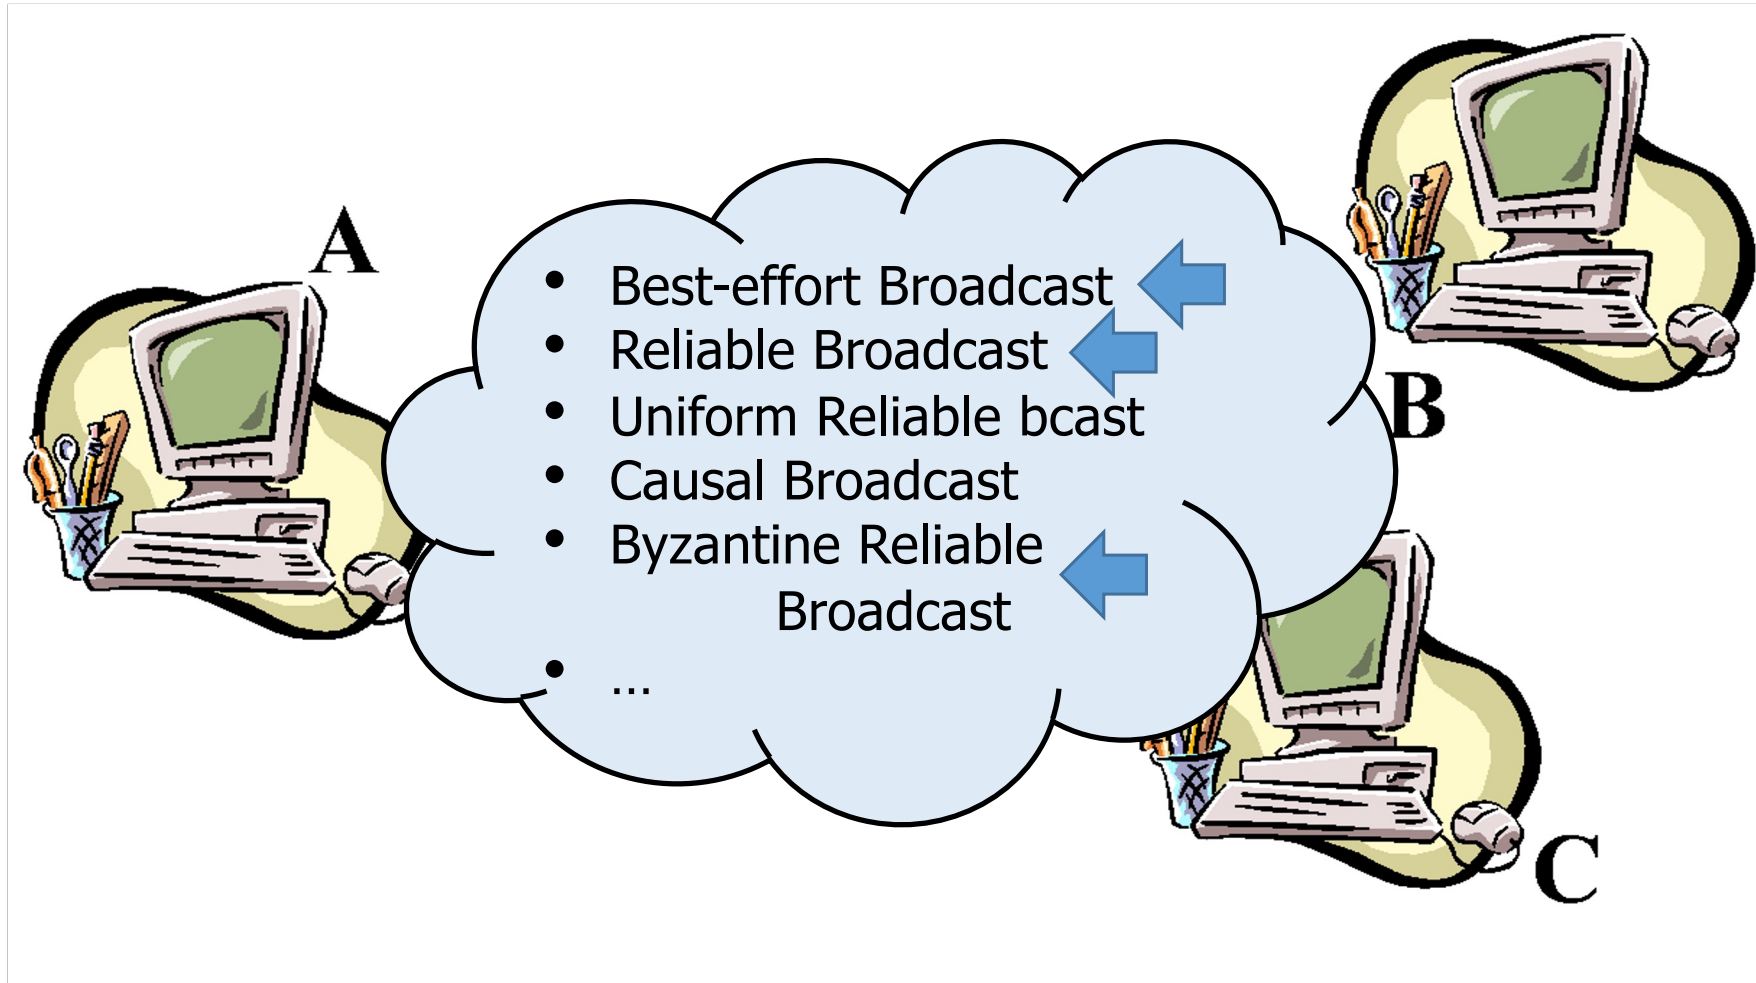
# Last lecture: safety and liveness

- Given a trace (sequence of outputs) of a distributed system
  - A safety property obeys:
    - If a finite trace does not obey the property, no extension of that trace obeys that property
  - A liveness property obeys:
    - A finite trace that does not obey the property can be extended so that the liveness property is upheld
- Any specification can be expressed in terms of liveness and safety properties

# Last lecture - Byzantine Leader Election

- Properties
  - **Eventual succession:** if more than  $f$  correct processes that trust some process  $p$  complain about  $p$ , then every correct process eventually trusts a different process than  $p$
  - **Putsch resistance:** A correct process does not trust a new leader unless at least one correct process has complained against the previous leader
  - **Eventual agreement:** there is a time after which no two correct processes trust different processes
- Eventually every correct process trusts some process that appears to perform its task in the higher-level algorithm.

# Broadcast Abstractions




# Intuition

- Broadcast is useful per se, for instance, in applications where some processes subscribe to events published by other processes (e.g., stocks)
- But also, a crucial building blocks in other protocols, namely consensus
- The receivers might require some **reliability** guarantees from the broadcast service (we say sometimes quality of service QoS) that the underlying network does not provide

# Overview

We shall consider three forms of reliability for a broadcast primitive

- **(1) Best-effort broadcast** 
- **(2) Reliable broadcast in the crash fault model**
- **(3) Reliable broadcast in the arbitrary fault model**
- We provide **specifications** and then **algorithms**

# Best-effort Broadcast (beb)

- **Events**

Request: <bebBroadcast, m>

Indication: <bebDeliver, src, m>

- **Properties: BEB1, BEB2, BEB3**

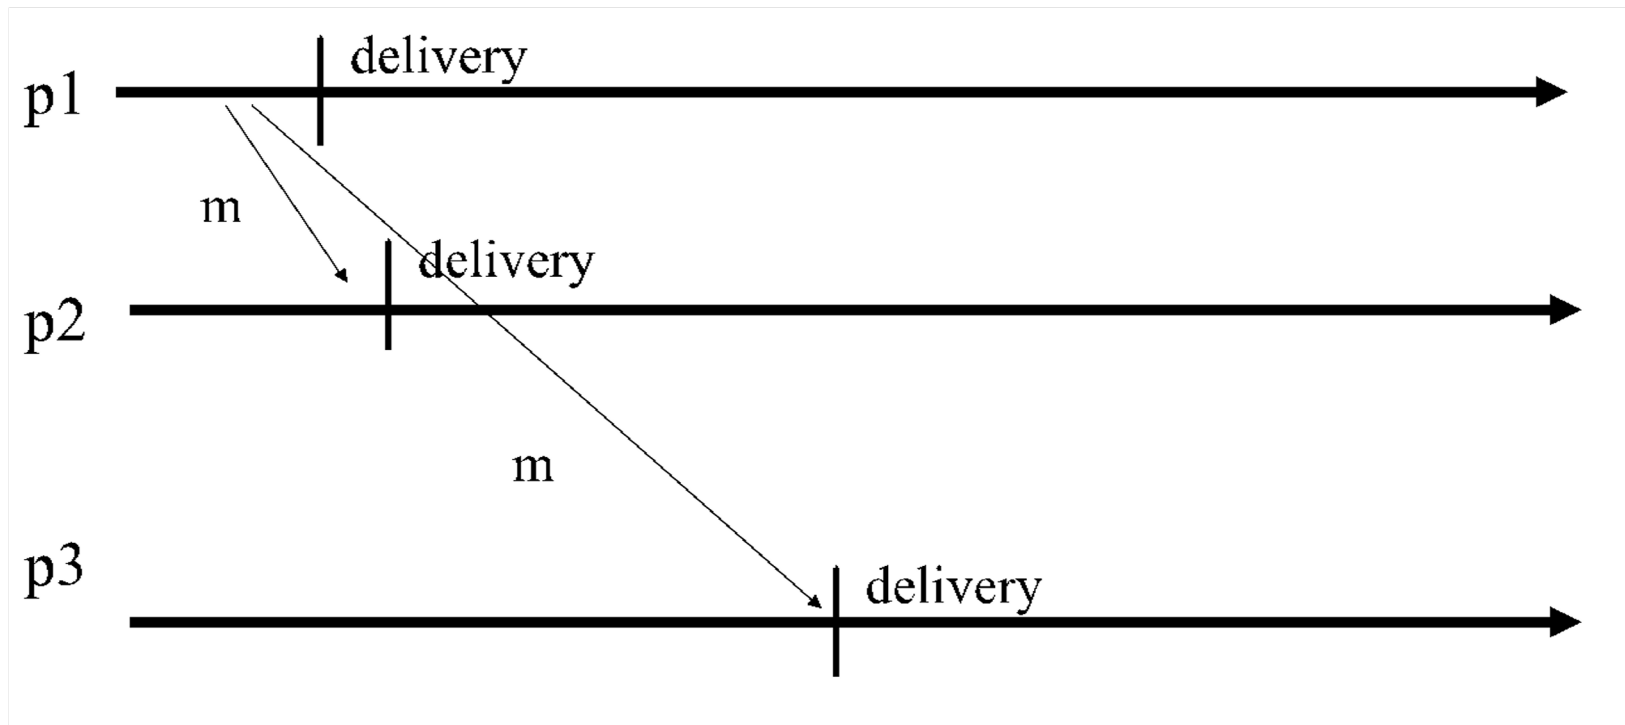
# Best-effort broadcast (beb)

## ■ Properties

- **BEB1. Validity:** If  $p_i$  and  $p_j$  are correct, then every message broadcast by  $p_i$  is eventually delivered by  $p_j$
- **BEB2. No duplication:** No message is delivered more than once
- **BEB3. No creation:** No message is delivered unless it was broadcast




# Best-effort Broadcast



Meets the specification


# Best-effort Broadcast



Meets the specification

# Overview

We shall consider three forms of reliability for a broadcast primitive

- **(1) Best-effort broadcast** 
- **(2) Reliable broadcast in the crash fault model**
- **(3) Reliable broadcast in the arbitrary fault model**
- We provide **specifications** and then **algorithms**

# Algorithm (beb)

- **Implements:** BestEffortBroadcast (beb).
- **Uses:** PerfectLinks (pp2p).
- **upon event**  $\langle \text{bebBroadcast}, m \rangle$  **do**  
    **for all**  $p_i \in S$  **do**  
        **trigger**  $\langle \text{pp2pSend}, p_i, m \rangle$ ;
- **upon event**  $\langle \text{pp2pDeliver}, p_i, m \rangle$  **do**  
    **trigger**  $\langle \text{bebDeliver}, p_i, m \rangle$ ;


# Algorithm (beb)

- **Proof (sketch)**

- **BEB1. Validity:** By the validity property of perfect links and the very facts that (1) the sender sends the message to all and (2) every correct process that pp2pDelivers a message bebDelivers it
- **BEB2. No duplication:** By the no duplication property of perfect links
- **BEB3. No creation:** By the no creation property of perfect links

# Overview

We shall consider three forms of reliability for a broadcast primitive

- **(1) Best-effort broadcast**
- **(2) Reliable broadcast in the crash fault model** 
- **(3) Reliable broadcast in the arbitrary fault model**
- We provide **specifications** and then **algorithms**

# Reliable Broadcast (rb)

- **Events**

- Request:  $\langle \text{rbBroadcast}, m \rangle$
- Indication:  $\langle \text{rbDeliver}, \text{src}, m \rangle$

- **Properties: RB1, RB2, RB3, RB4**

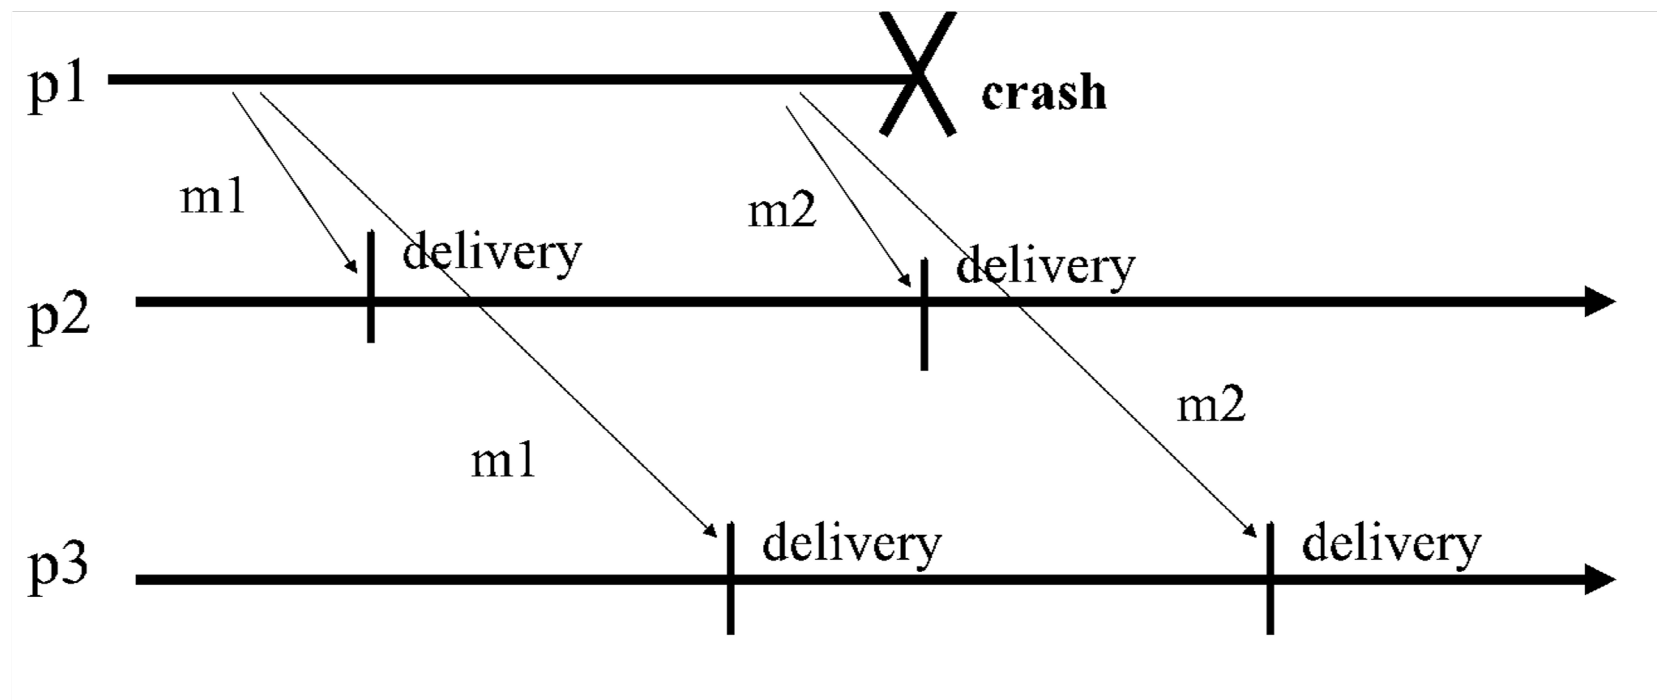
# Reliable broadcast

- **Properties**

- **RB1 = BEB1.**
- **RB2 = BEB2.**
- **RB3 = BEB3.**
- **RB4. Agreement:** For any message  $m$ , if a correct process delivers  $m$ , then every correct process delivers  $m$

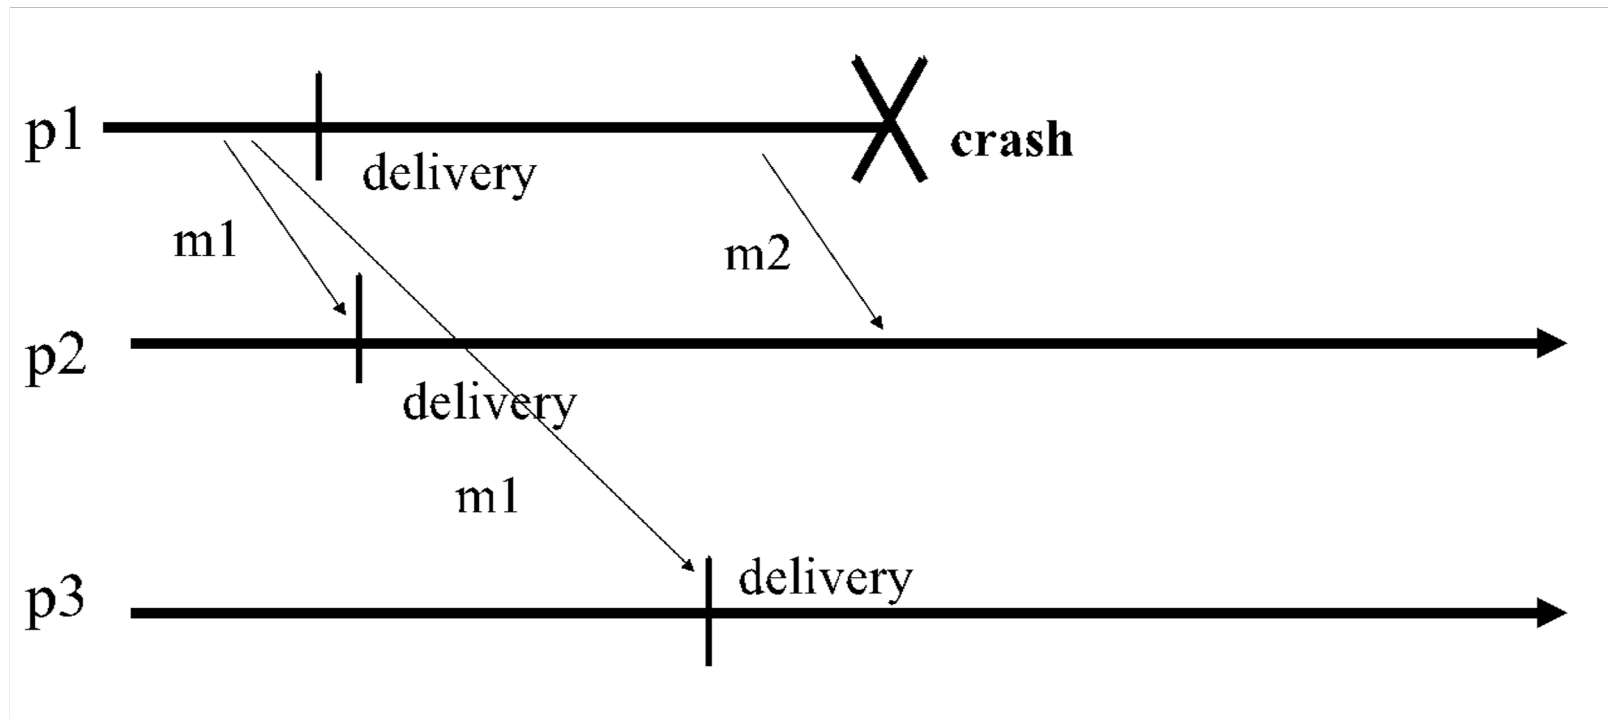


# Reliable Broadcast



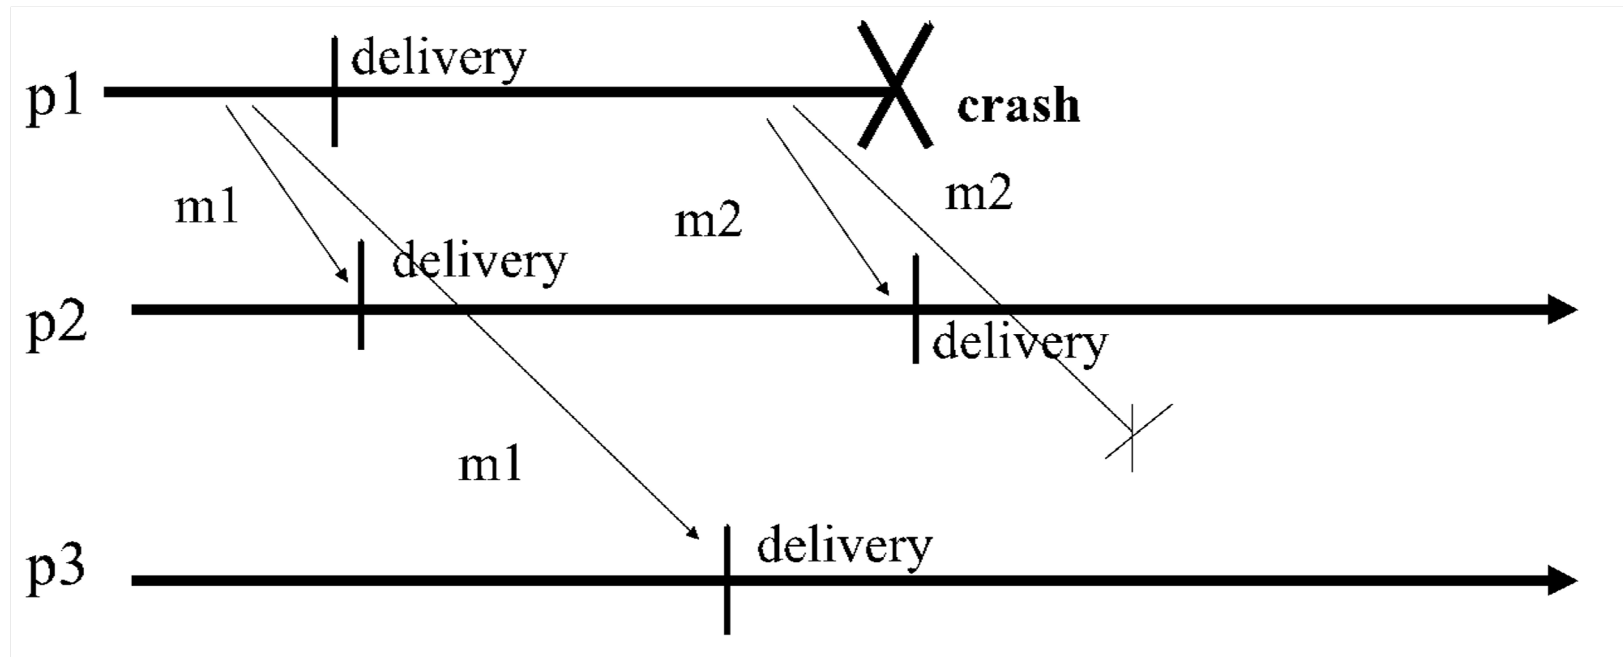
Meets the specification

# Reliable Broadcast



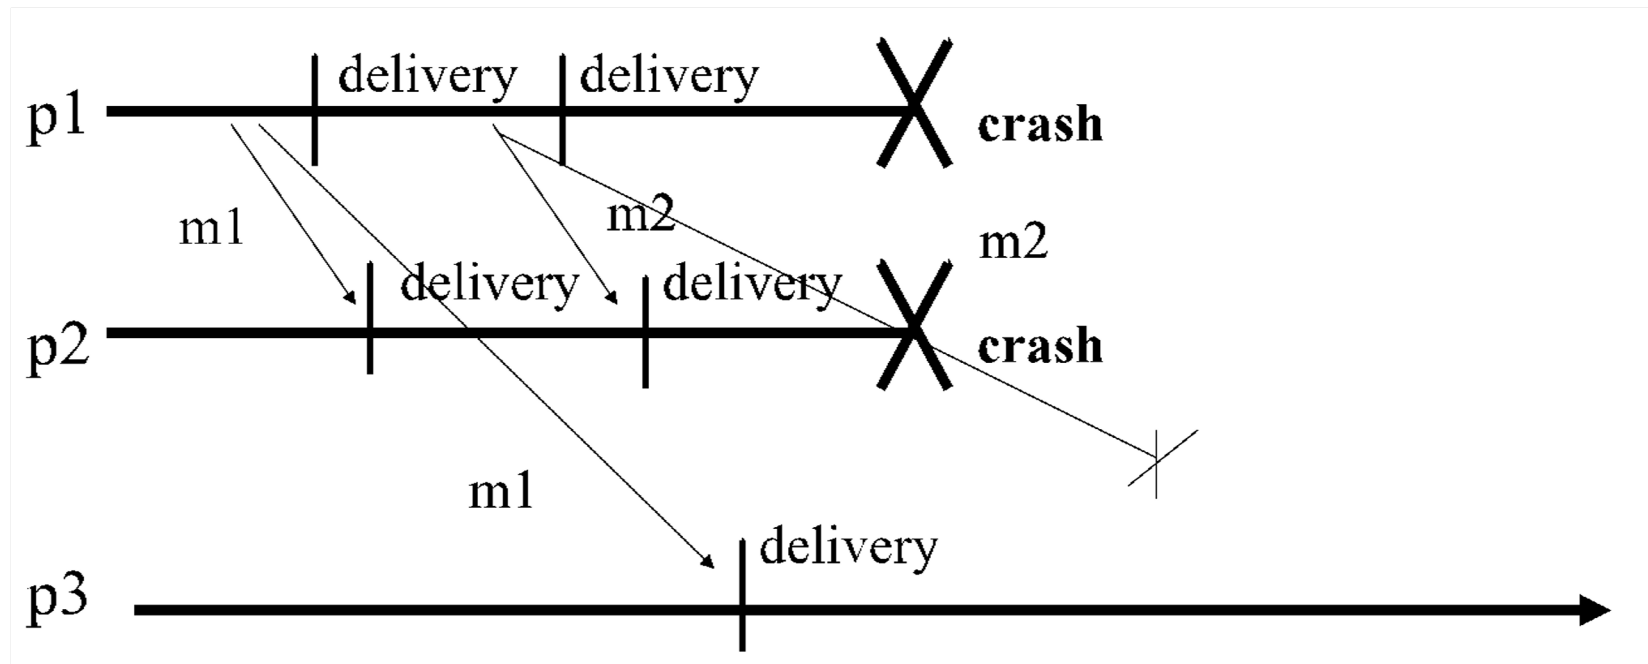
Meets the specification

# Reliable Broadcast



Does not meet the specification



# Reliable Broadcast



Meets the specification

# Overview

We shall consider three forms of reliability for a broadcast primitive

- **(1) Best-effort broadcast**
- **(2) Reliable broadcast in the crash fault model** 
- **(3) Reliable broadcast in the arbitrary fault model**
- We provide **specifications** and then **algorithms** 

# Eager Reliable Broadcast

## Implements:

ReliableBroadcast, **instance** *rb*.

## Uses:

BestEffortBroadcast, **instance** *beb*.

**upon event**  $\langle rb, Init \rangle$  **do**

*delivered* :=  $\emptyset$ ;

**upon event**  $\langle rb, Broadcast \mid m \rangle$  **do**

**trigger**  $\langle beb, Broadcast \mid [DATA, self, m] \rangle$ ;

**upon event**  $\langle beb, Deliver \mid p, [DATA, s, m] \rangle$  **do**

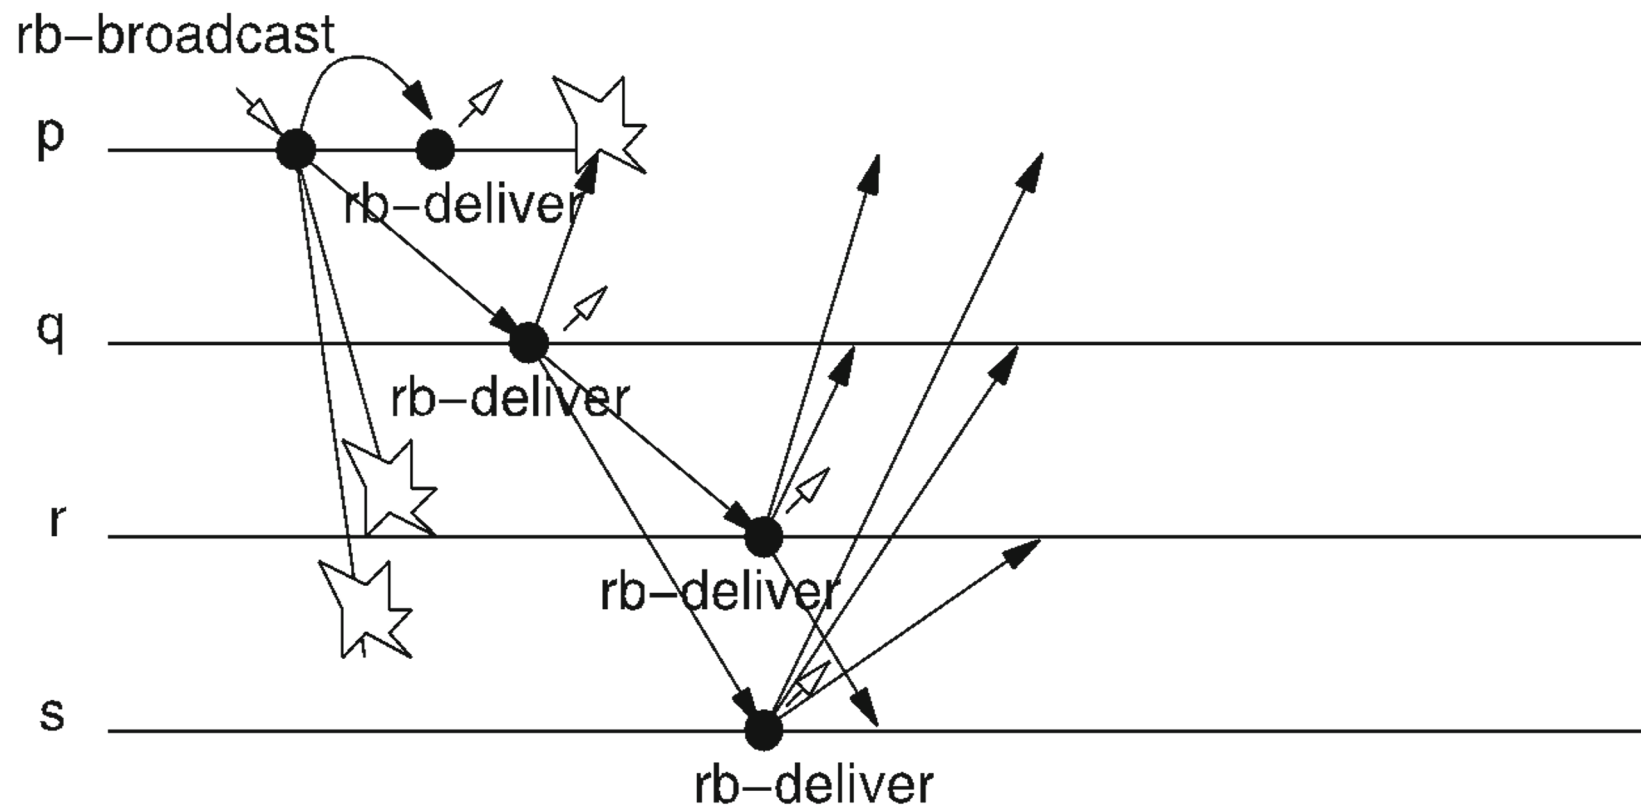
**if**  $m \notin delivered$  **then**

*delivered* := *delivered*  $\cup \{m\}$ ;

**trigger**  $\langle rb, Deliver \mid s, m \rangle$ ;

**trigger**  $\langle beb, Broadcast \mid [DATA, s, m] \rangle$ ;

# Eager Reliable Broadcast



# Proof (sketch)

- **RB1. RB2. RB3:** as for the 1st algorithm
- **RB4. Agreement:**
  - every correct process immediately relays every message it *rb*-delivers
  - the *validity* property of the underlying best-effort broadcast primitive ensures that all other correct process will eventually deliver the message.

Note: Does not use a failure detector  $\Leftrightarrow$   
Does not make any synchrony assumptions



# Overview

We shall consider three forms of reliability for a broadcast primitive

- **(1) Best-effort broadcast**
- **(2) Reliable broadcast in the crash fault model**
- **(3) Reliable broadcast in the arbitrary fault model**

- We provide **specifications** and then **algorithms**



# Reliable broadcast in the byzantine model

- A byzantine process may, e.g.:
  - broadcast messages different from the ones sent by applications,
  - forge messages so that they look as originated by other processes
    - violation of the No creation property
  - send different messages to different processes
    - violation of the *agreement* property

# Reliable broadcast in the byzantine model

- Authenticated Perfect Links are a useful tool:
  - disallows forging messages as originated from other processes
  - still does not solve the problems above
- Digital signatures can also be helpful:
  - allow any process to verify the authenticity of a message
  - but faulty processes may sign “illegal” msgs
    - e.g., sending two different messages, after having signed them

# Byzantine Consistent and Reliable Broadcasts

- We consider two variants of reliable broadcasts in the byzantine model:
  - Byzantine Consistent Broadcast
  - Byzantine Reliable Broadcast

# Byzantine Consistent and Reliable Broadcasts

## ■ Byzantine Consistent Broadcast

- If the sender  $s$  is correct then every correct process should later deliver  $m$ .
- If  $s$  is faulty, then every correct process delivers the same message, if it delivers one at all.
  - i.e., correct processes are not guaranteed to deliver the same set of messages if sender is faulty

## ■ Byzantine Reliable Broadcast

- Ensures that if a correct process delivers  $m$ , then every correct process delivers  $m$ 
  - independently of whether the sender is faulty

# Byzantine Consistent Broadcast: specification

- **BCB1: *Validity*:** If a correct process  $p$  broadcasts a msg  $m$ , then every correct process eventually delivers  $m$ .
- **BCB2: *No duplication*:** Every correct process delivers at most one message.
- **BCB3: *Integrity*:** If some correct process delivers a message  $m$  with sender  $p$  and process  $p$  is correct, then  $m$  was previously broadcast by  $p$ .
- **BCB4: *Consistency*:** If some correct process delivers a message  $m$  and another correct process delivers a message  $m'$ , then  $m = m'$ .

Note: "Correct" now means non-Byzantine-faulty

# Overview

We shall consider three forms of reliability for a broadcast primitive

- (1) **Best-effort broadcast**
- (2) **Reliable broadcast in the crash fault model**
- (3) **Reliable broadcast in the arbitrary fault model** ←

- We provide **specifications** and then **algorithms**




# Byzantine Consistent Broadcast: Algorithms

- Byzantine Consistent Broadcast
  - If the sender  $s$  is correct then every correct process should later deliver  $m$ .
  - If  $s$  is faulty, then every correct process delivers the same message, if it delivers one at all.
    - i.e., correct processes are not guaranteed to deliver the same set of messages
- Two algorithms:
  - Authenticated Echo Broadcast
    - exchanges a quadratic number of messages
  - Signed Echo Broadcast
    - linear no. of messages but uses digital signatures (costly)
  - Leverage **Byzantine quorums**




# Quorums and Byzantine Quorums

- Quorums are sets of subsets with a specific intersection property:
  - they always overlap in at least one correct process
  - example:
    - majority of processes, assuming  $f < N/2$  **crash failure model**
    - intersection of 2 majority quorums is necessarily a correct process with crash faults
    - no longer true with byzantine faults!



# Derivation of Byzantine quorum sizes

- Safety: every pair of quorums must intersect in at least one correct process



# Derivation of Byzantine quorum sizes

- $Q + Q - N \geq f+1$
- To simplify, we take the optimal size (no need to intersect more than strictly necessary)
- $2Q - N = f+1$

# Derivation of Byzantine quorum sizes

- Liveness: a quorum always needs to be available
- Challenge:  $f$  processes may never reply (crashed or deliberately silent)
- Thus quorums can't be larger than  $N - f$
- $N - f \geq Q$
- Again, we turn this into an equality (don't use more processes than strictly necessary)
- $N - f = Q$

# Solving a system of two equations

- Note that  $f$  is a system parameter (constant)


$$\begin{cases} 2Q - N = f+1 \\ N - f = Q \end{cases}$$

- Thus, solving these equations for  $N$  and  $Q$  yields:

$$N = 3f+1$$

$$Q = 2f+1$$

# Byzantine Quorums: examples



# Acknowledgements

- Rachid Guerraoui, EPFL